

Apache Spark – Apache HBase Connector

Feature Rich and Efficient Access to HBase through Spark SQL

Weiqing Yang

Jun, 2017

Who am I

- Currently focused on Apache Spark and Hadoop, etc
- Contribute to Apache Spark, HBase, Ambari, Hadoop
- Software Engineer at Hortonworks

Agenda

Motivation

Overview

Architecture & Implementation

Usage

Motivation

- ◆ Limited Spark Support in HBase Upstream
 - RDD level
 - But Spark Is Moving to DataFrame/Dataset
- ◆ Existing Connectors in DataFrame Level
 - Complicated Design
 - Embedding Optimization Plan inside Catalyst Engine
 - Stability Impact with Coprocessor
 - Serialized RDD Lineage to HBase
 - Heavy Maintenance Overhead

Overview

Apache Spark– Apache HBase Connector (SHC)

- ◆ Combine Spark and HBase
 - Spark Catalyst Engine for Query Plan and Optimization
 - HBase as Fast Access KV Store
 - Implement Standard External Data Source with Build-in Filter, Maintain Easily
- ◆ Full Fledged DataFrame Support
 - Spark SQL
 - Integrated Language Query
- ◆ High Performance
 - Partition Pruning, Data Locality, Column Pruning, Predicate Pushdown
 - Use Spark UnhandledFilters API
 - Cache Spark HBase Connections

Data Coder & Data Schema

Support Different Data Coders

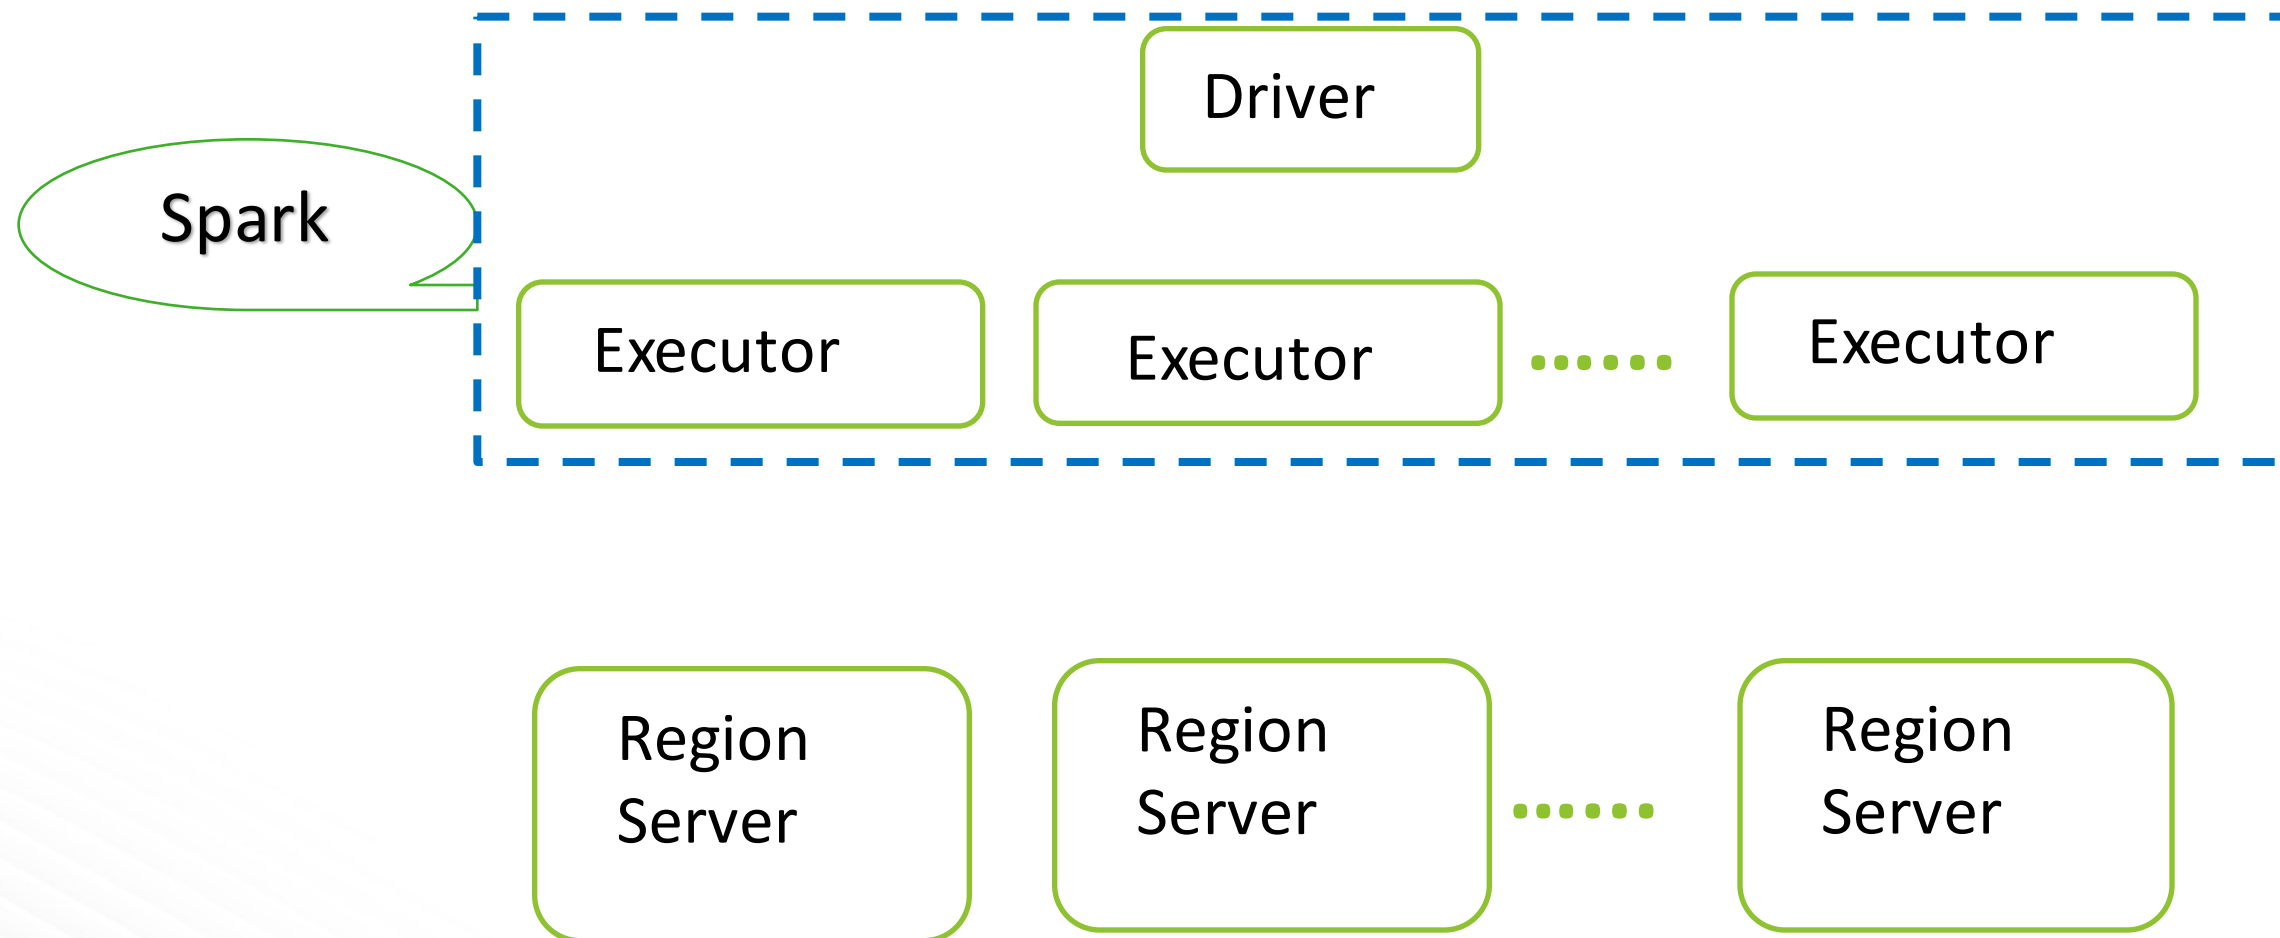
- PrimitiveType: Native Support Java Primitive Types
- Avro: Native Support Avro Encoding/Decoding
- Phoenix: Phoenix Encoding/Decoding
- Plug-In Data Coder
- Can Run on the Top of Existing HBase Tables

Support Composite Key

- `def cat = s"""{
 | "table":{"namespace":"default","name":"shcExampleTable","tableCoder":"Phoenix"},
 | "rowkey":"key1:key2",
 | "columns":{
 | "col00":{"cf":"rowkey","col":"key1","type":"string"},
 | "col01":{"cf":"rowkey","col":"key2","type":"int"},
 ...
 }
}"""`

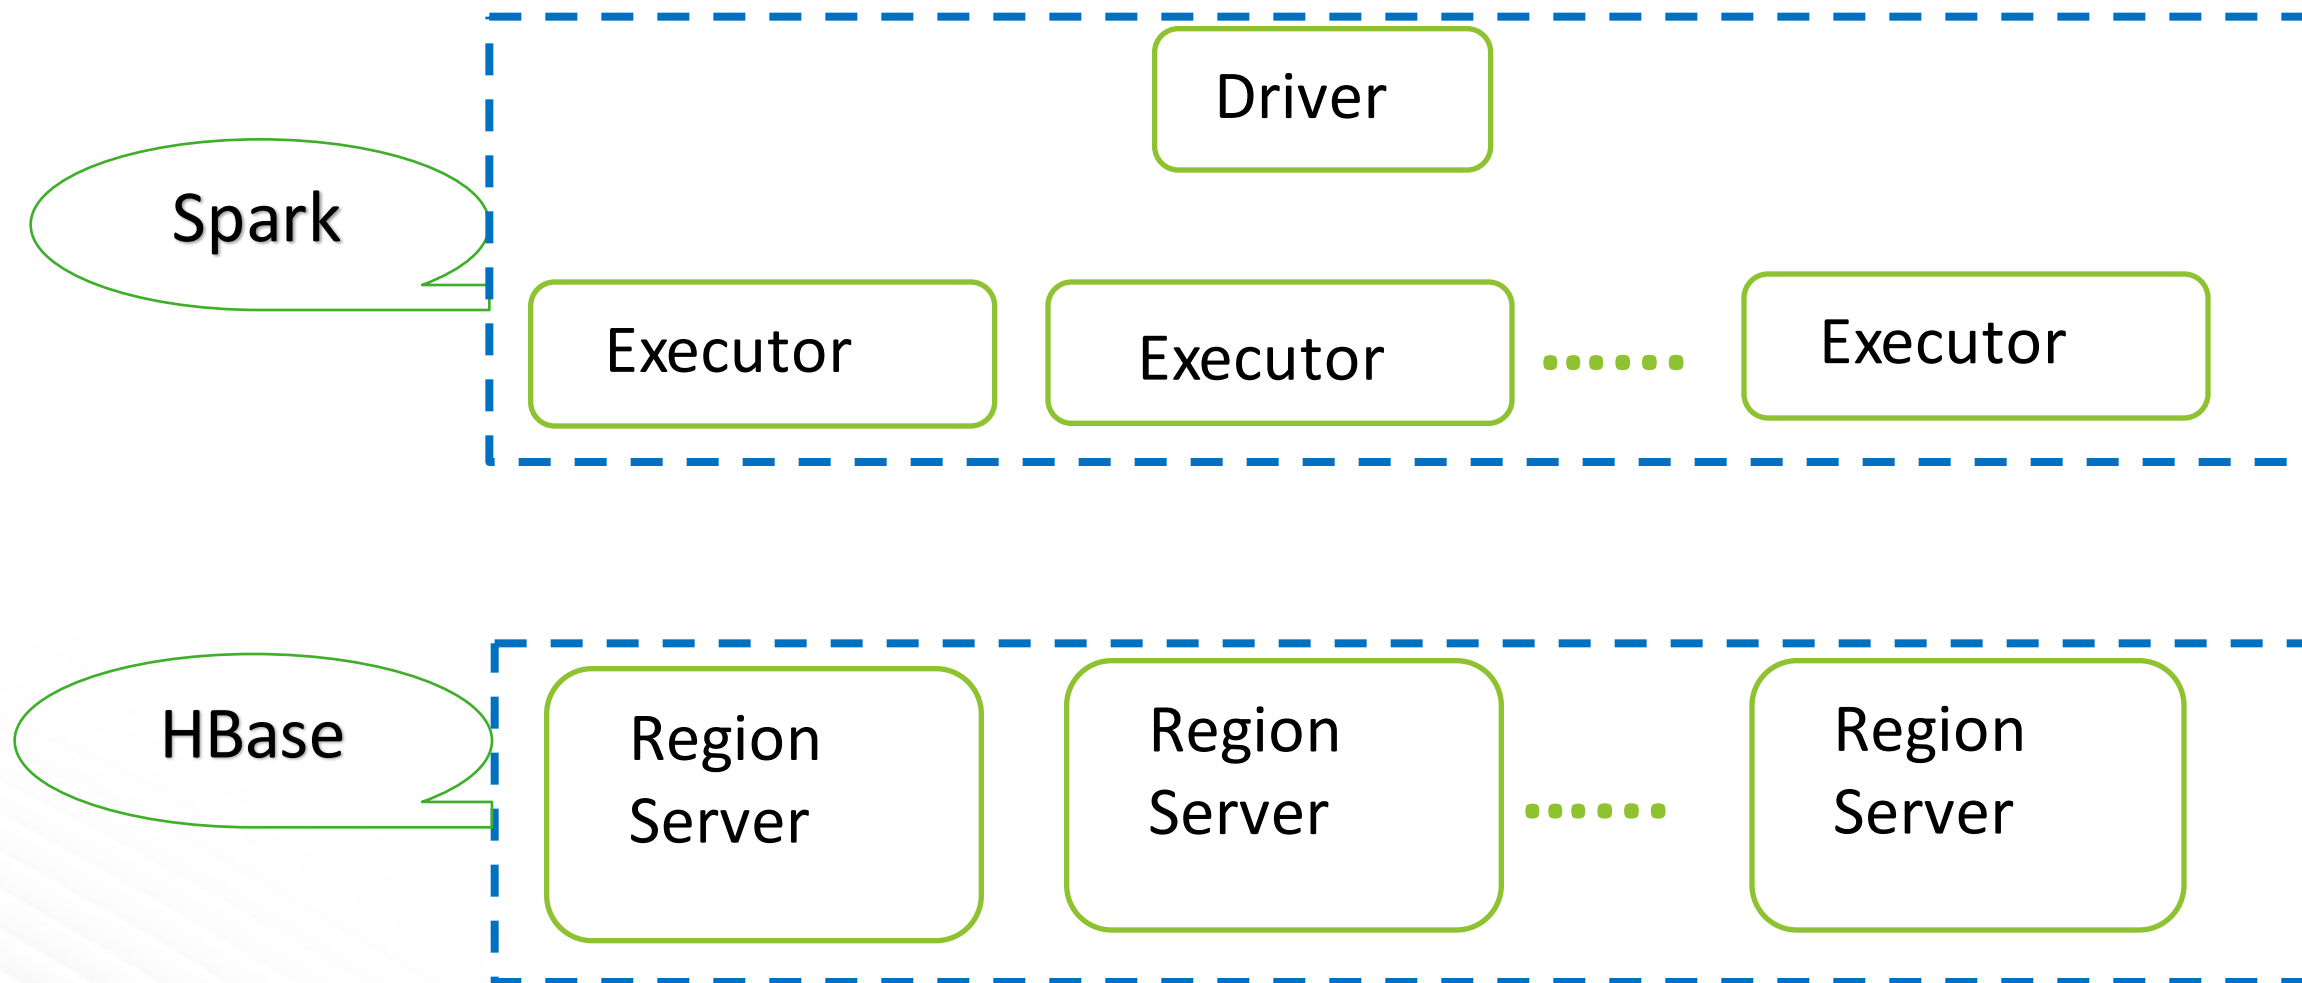
Architecture & Implementation

Architecture




Picture 1. SHC architecture

Architecture



Picture 1. SHC architecture

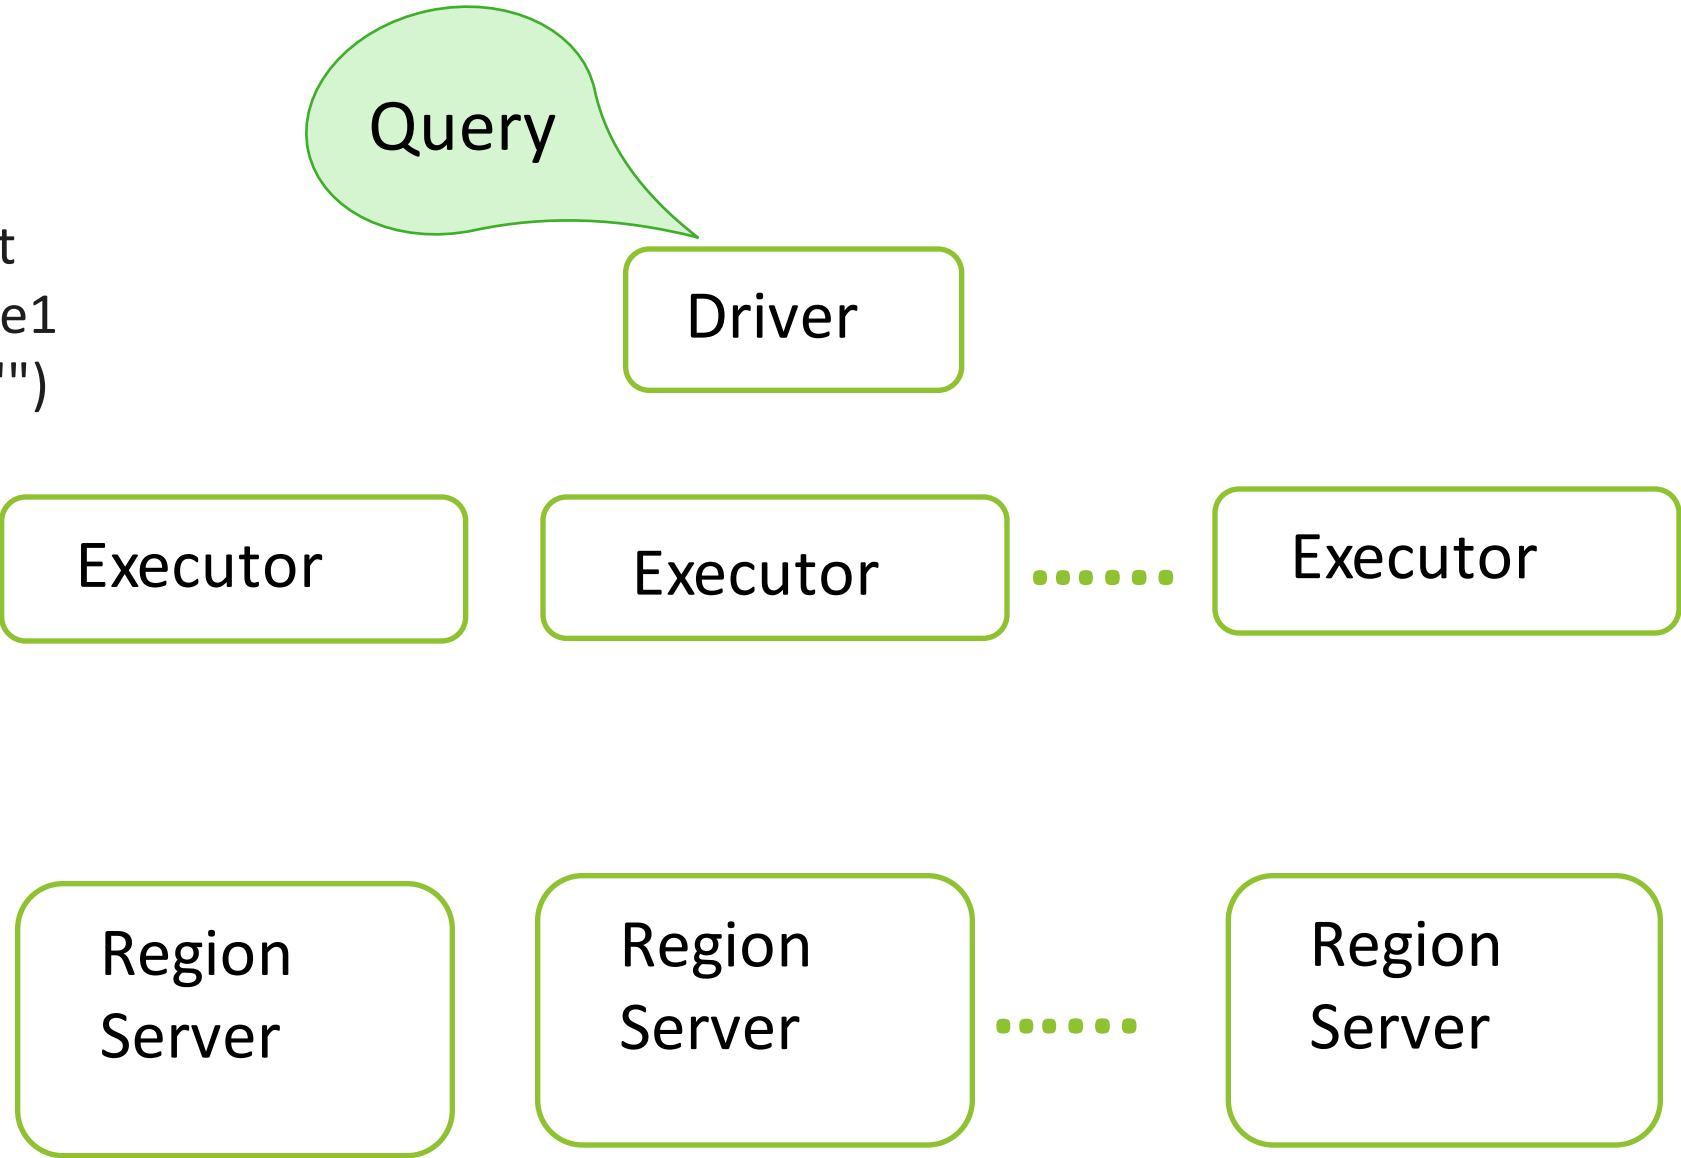
Architecture



Picture 1. SHC architecture

Architecture

```
sqlContext.sql("select  
count(col1) from table1  
where key < 'row050'")
```

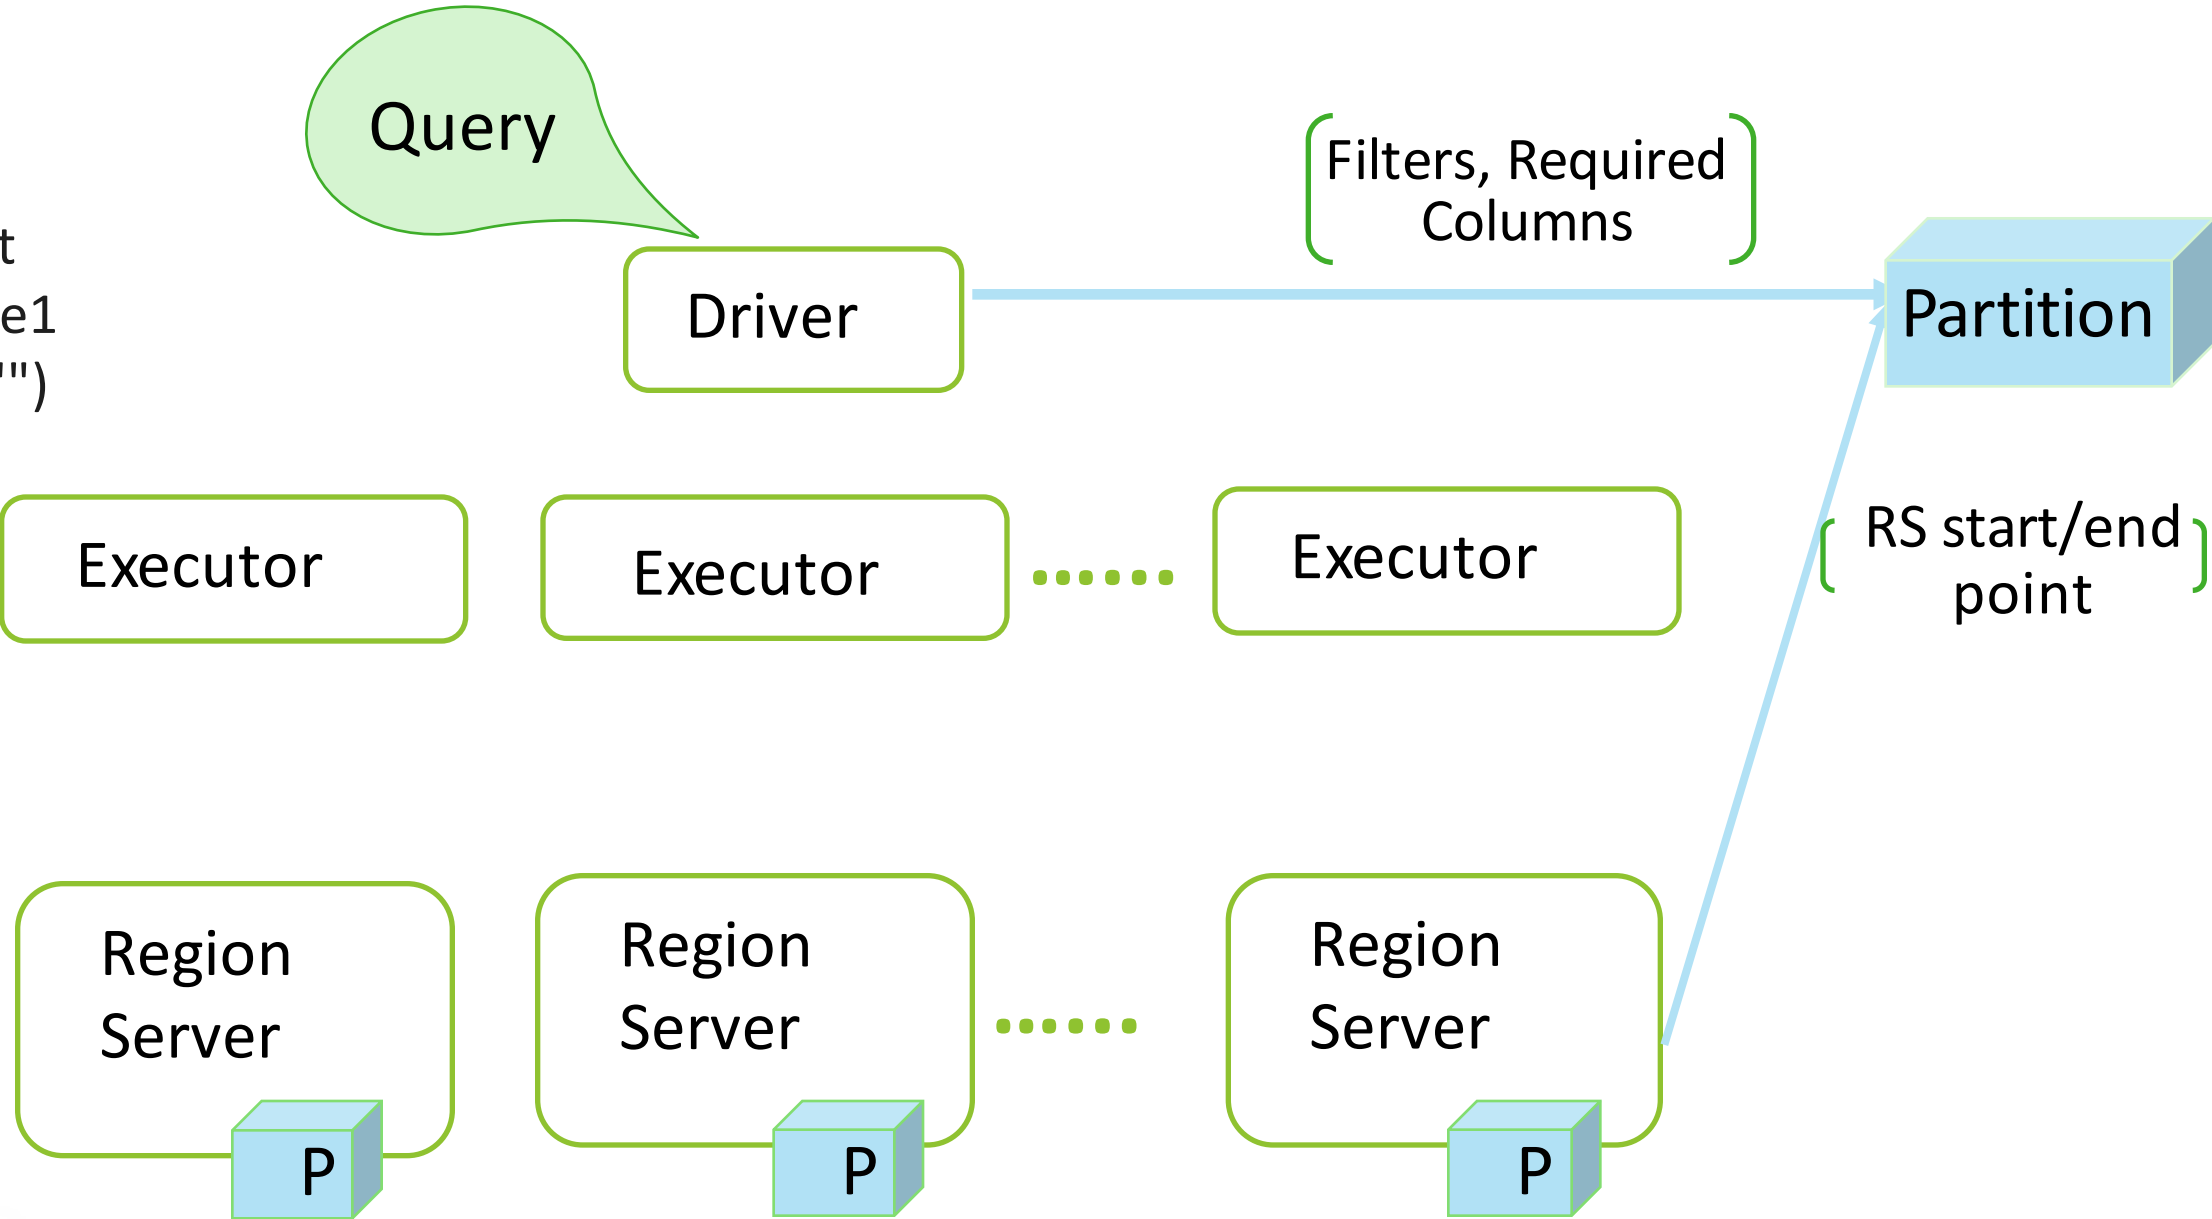


Picture 1. SHC architecture



Architecture

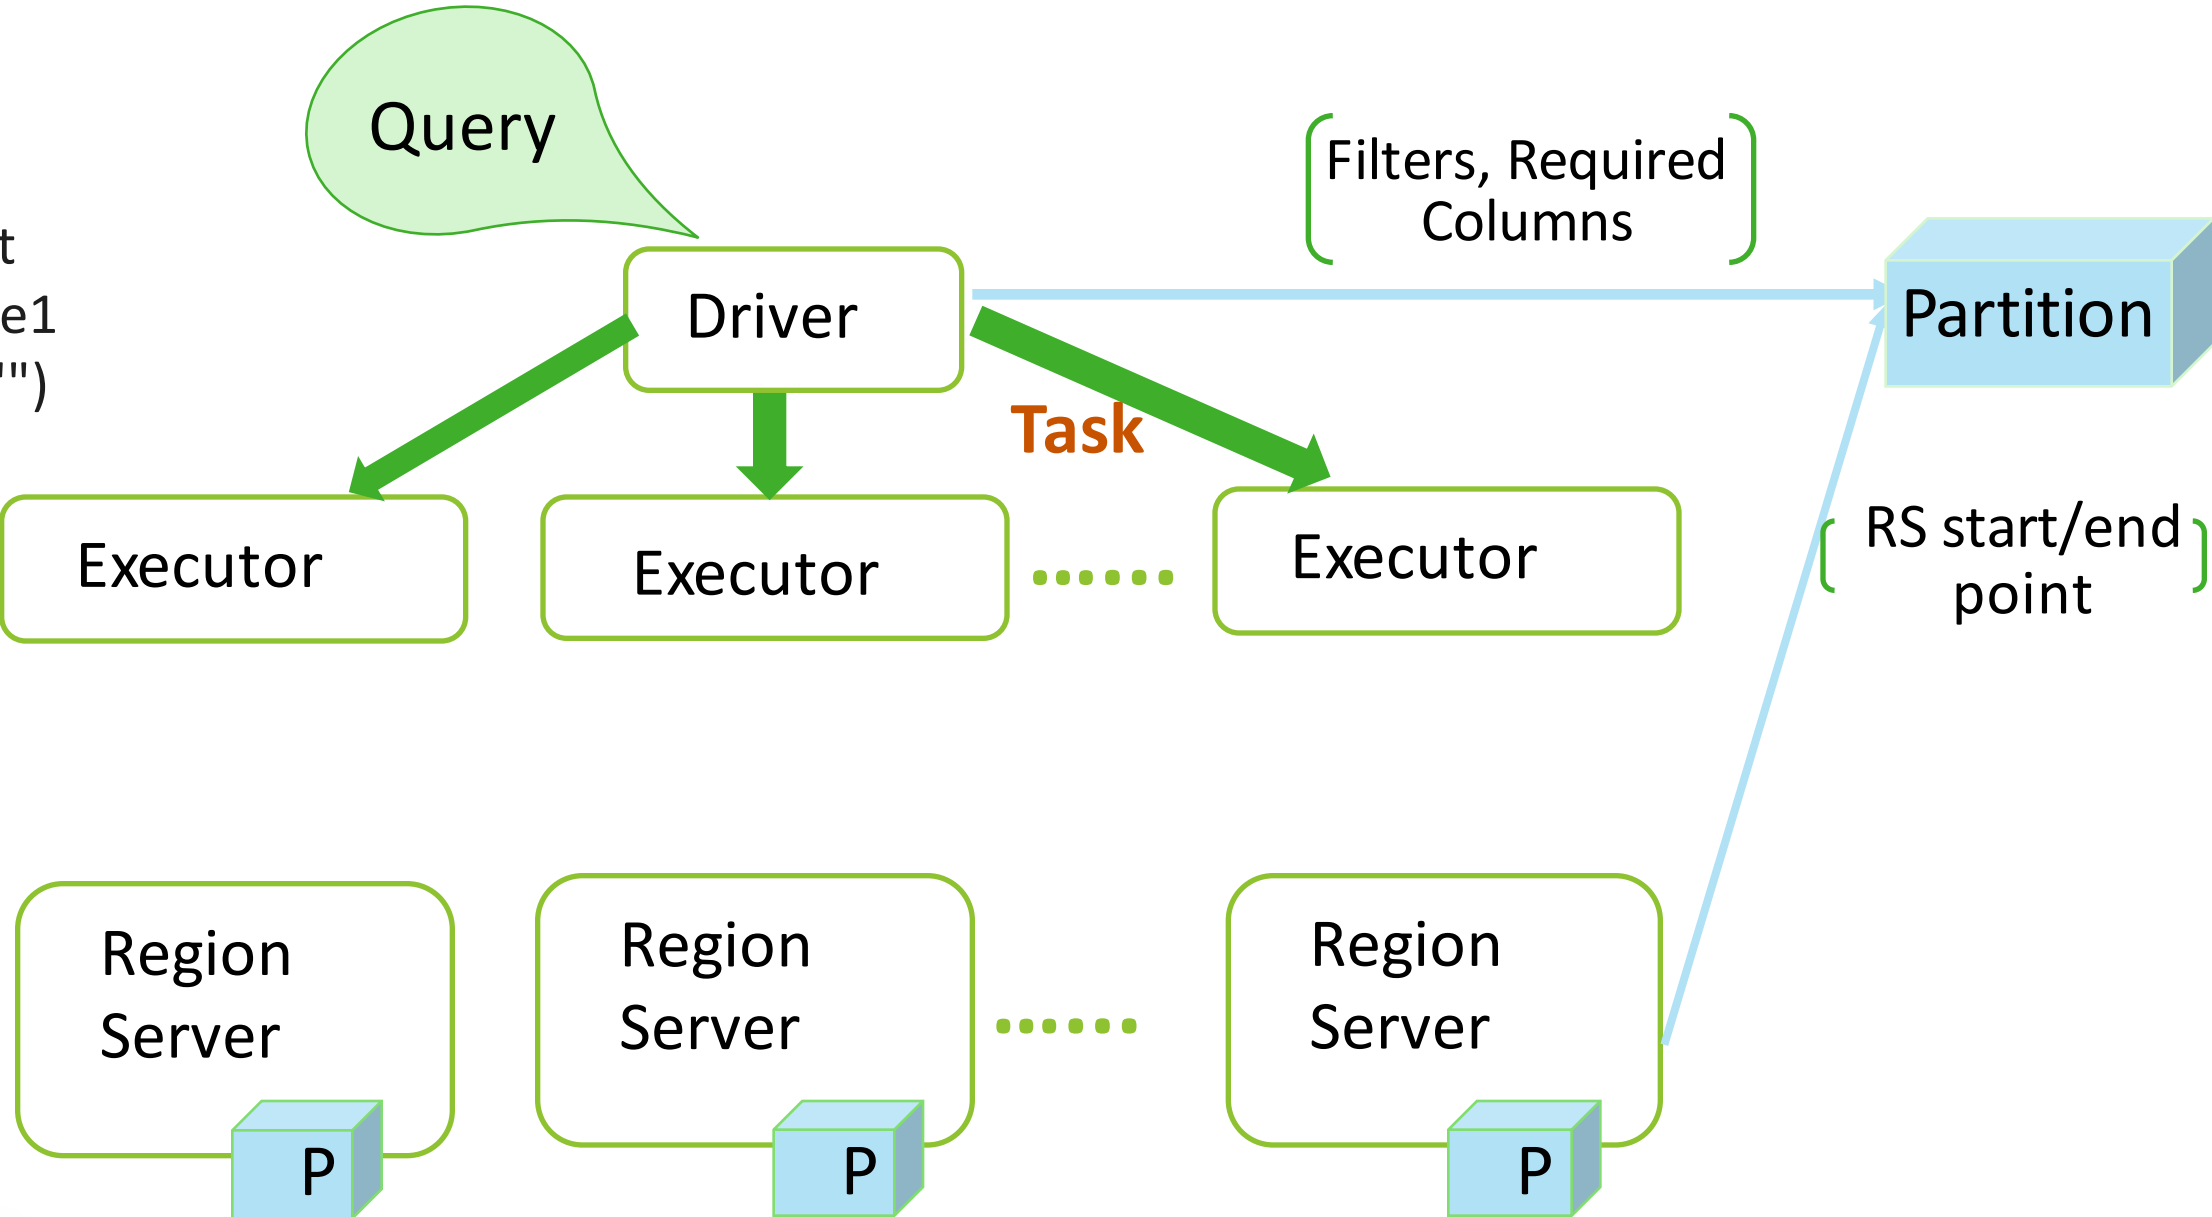
```
sqlContext.sql("select  
count(col1) from table1  
where key < 'row050'")
```



Picture 1. SHC architecture

Architecture

```
sqlContext.sql("select  
count(col1) from table1  
where key < 'row050'")
```

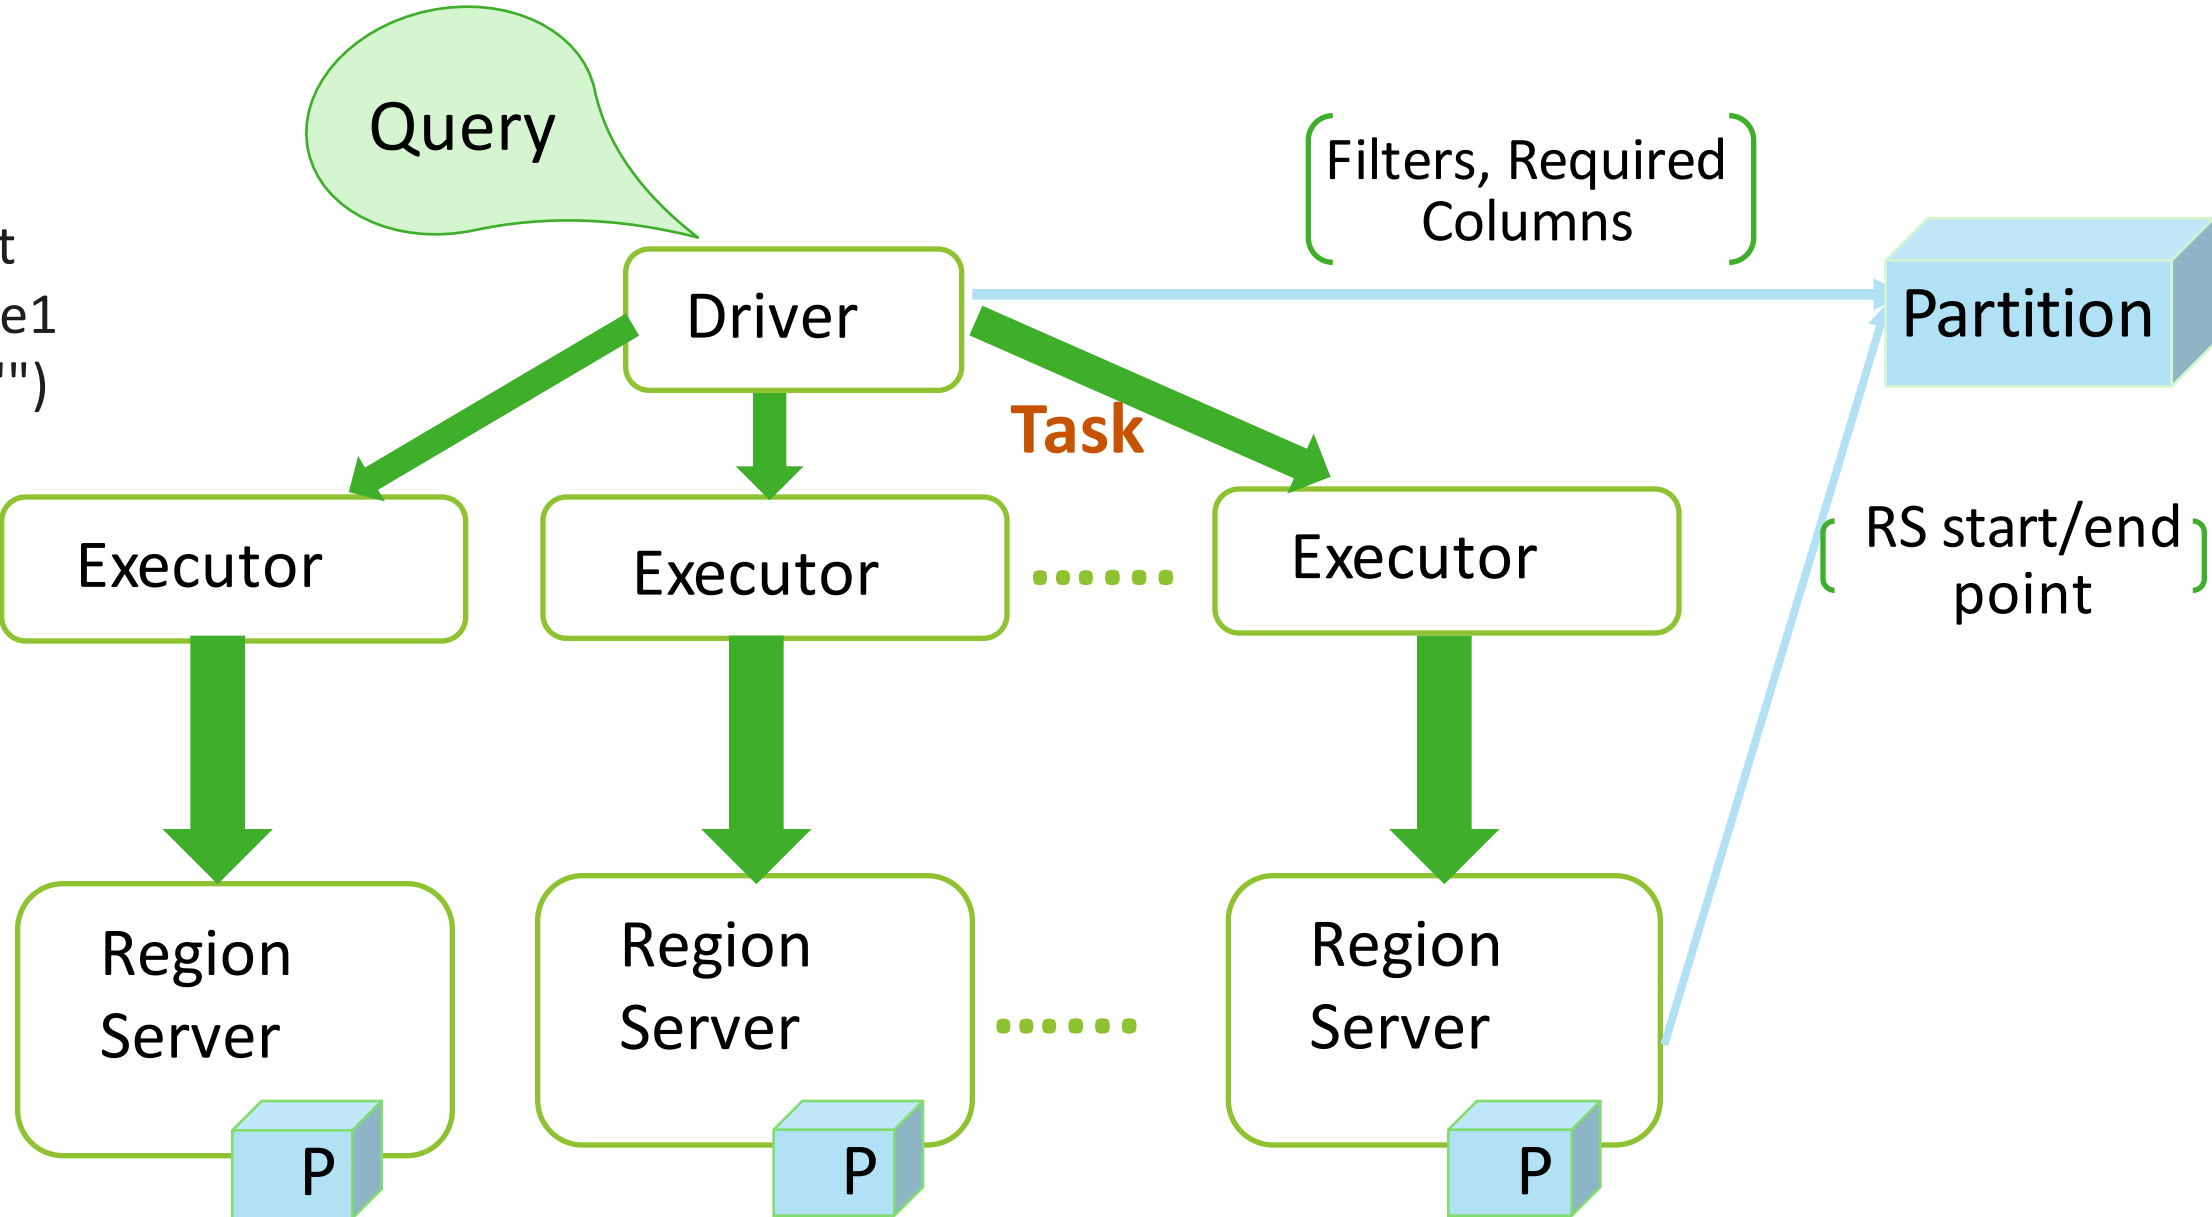


Picture 1. SHC architecture



Architecture

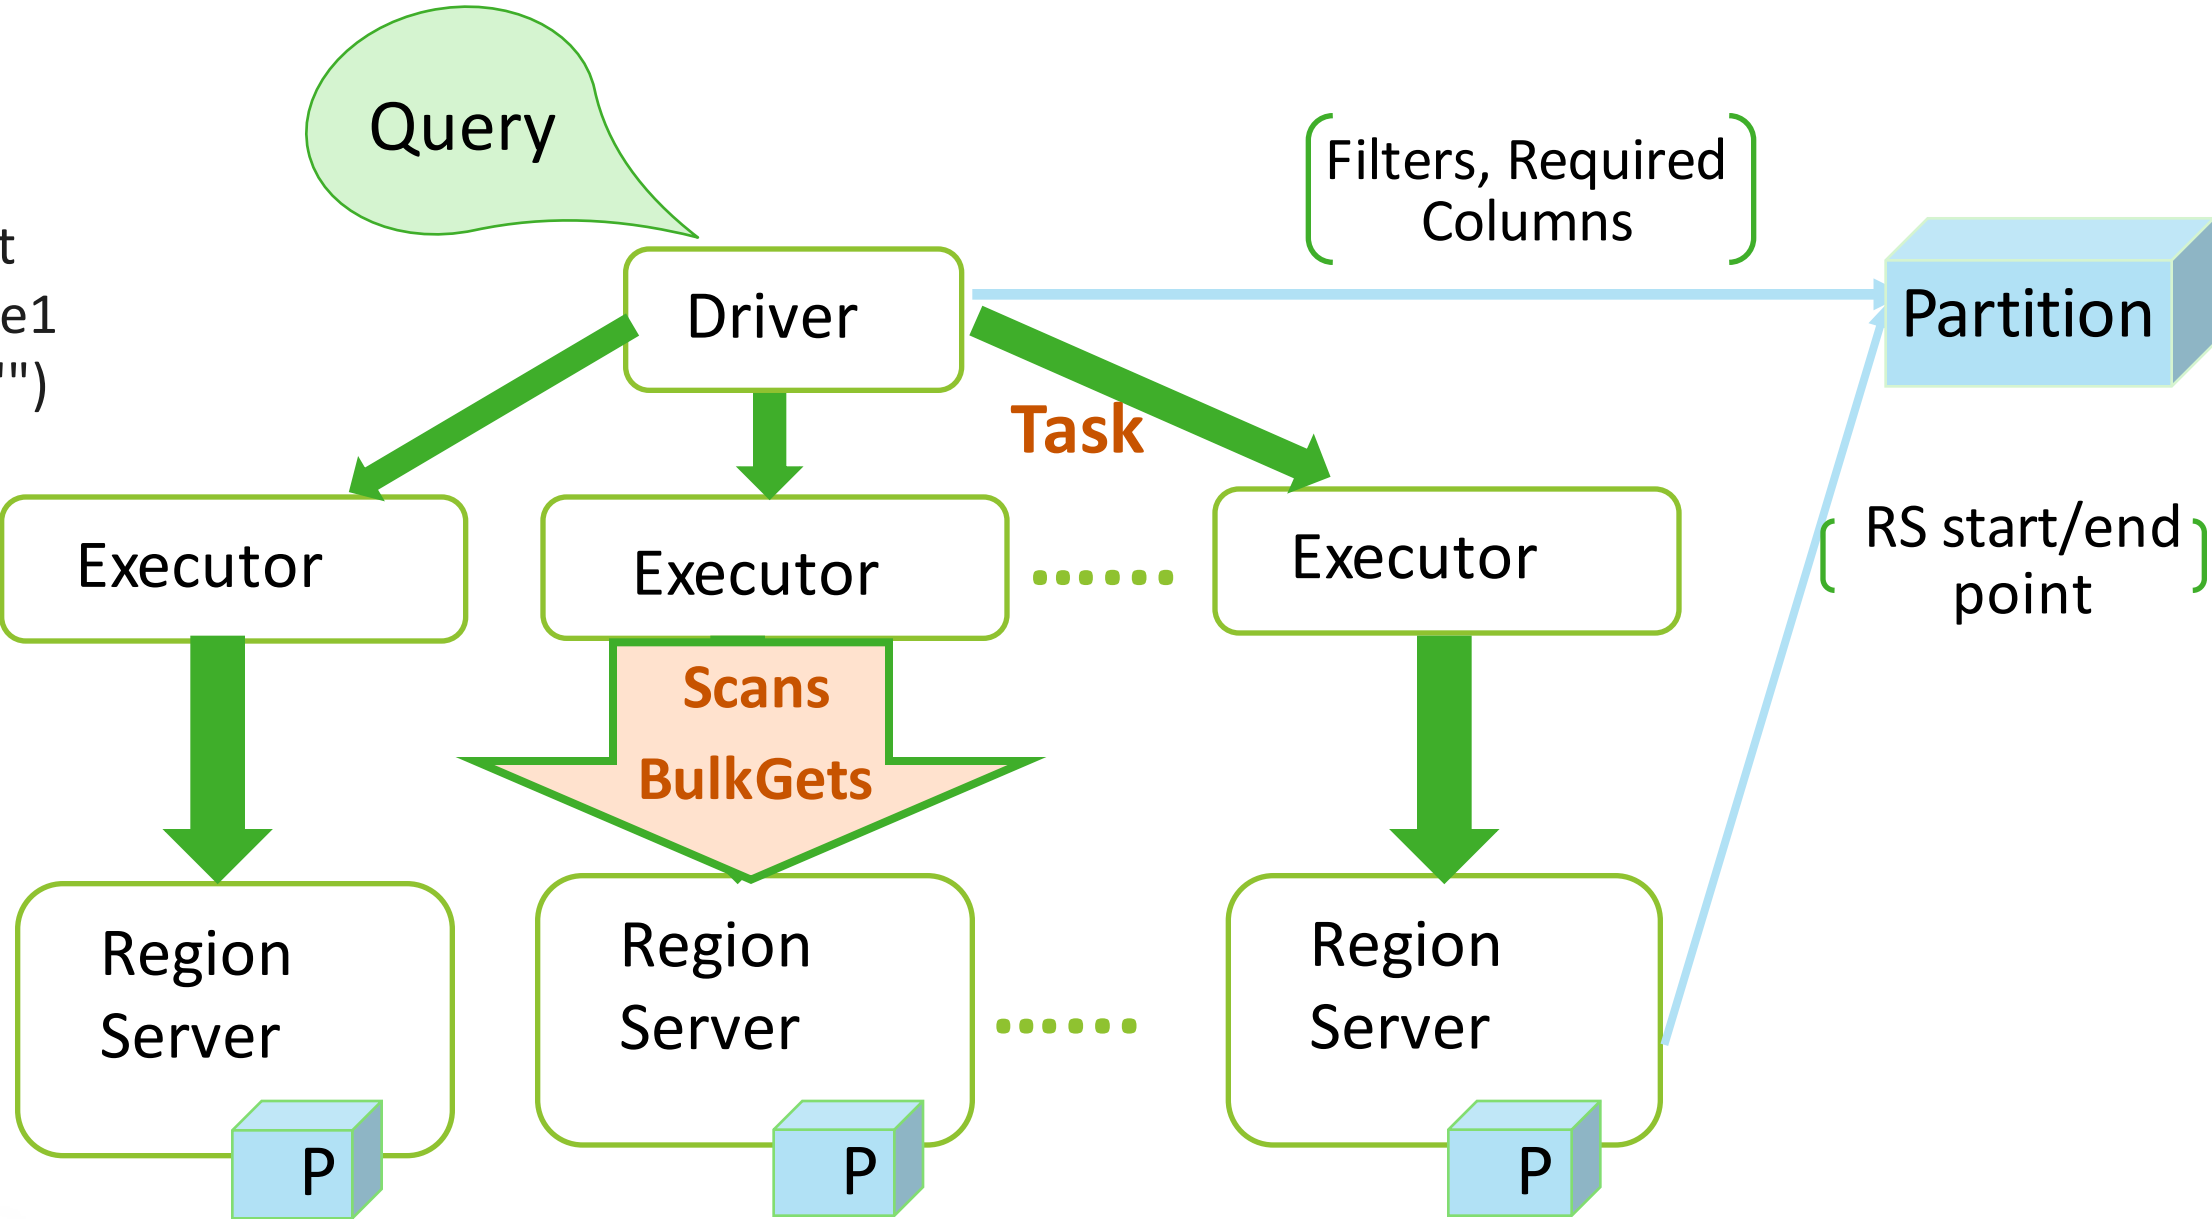
```
sqlContext.sql("select  
count(col1) from table1  
where key < 'row050'")
```



Picture 1. SHC architecture

Architecture

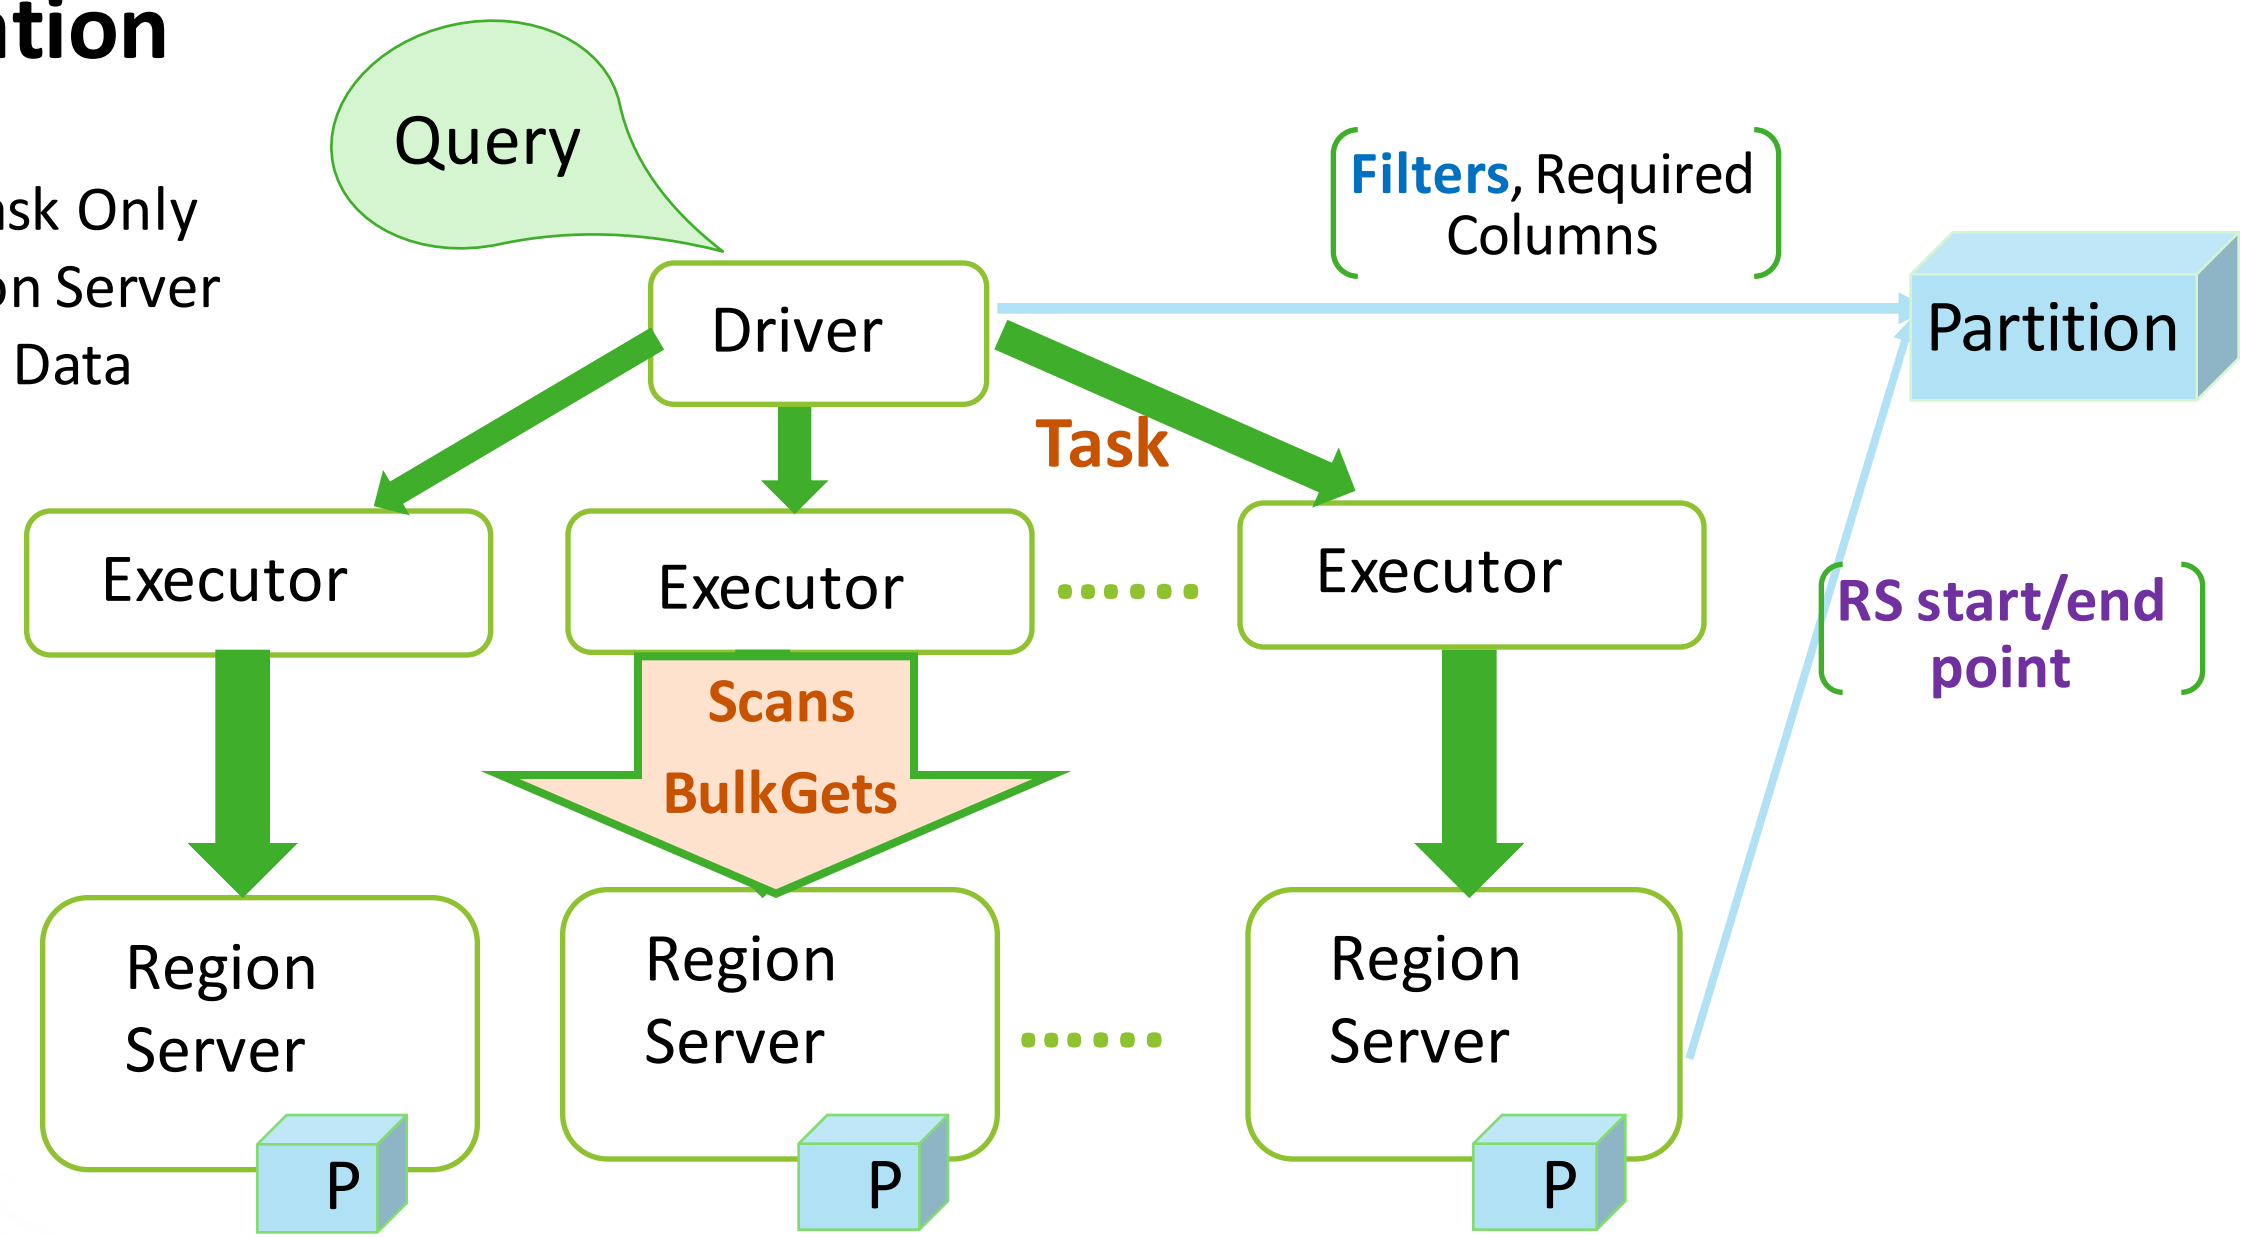
```
sqlContext.sql("select  
count(col1) from table1  
where key < 'row050'")
```



Picture 1. SHC architecture

Implementation

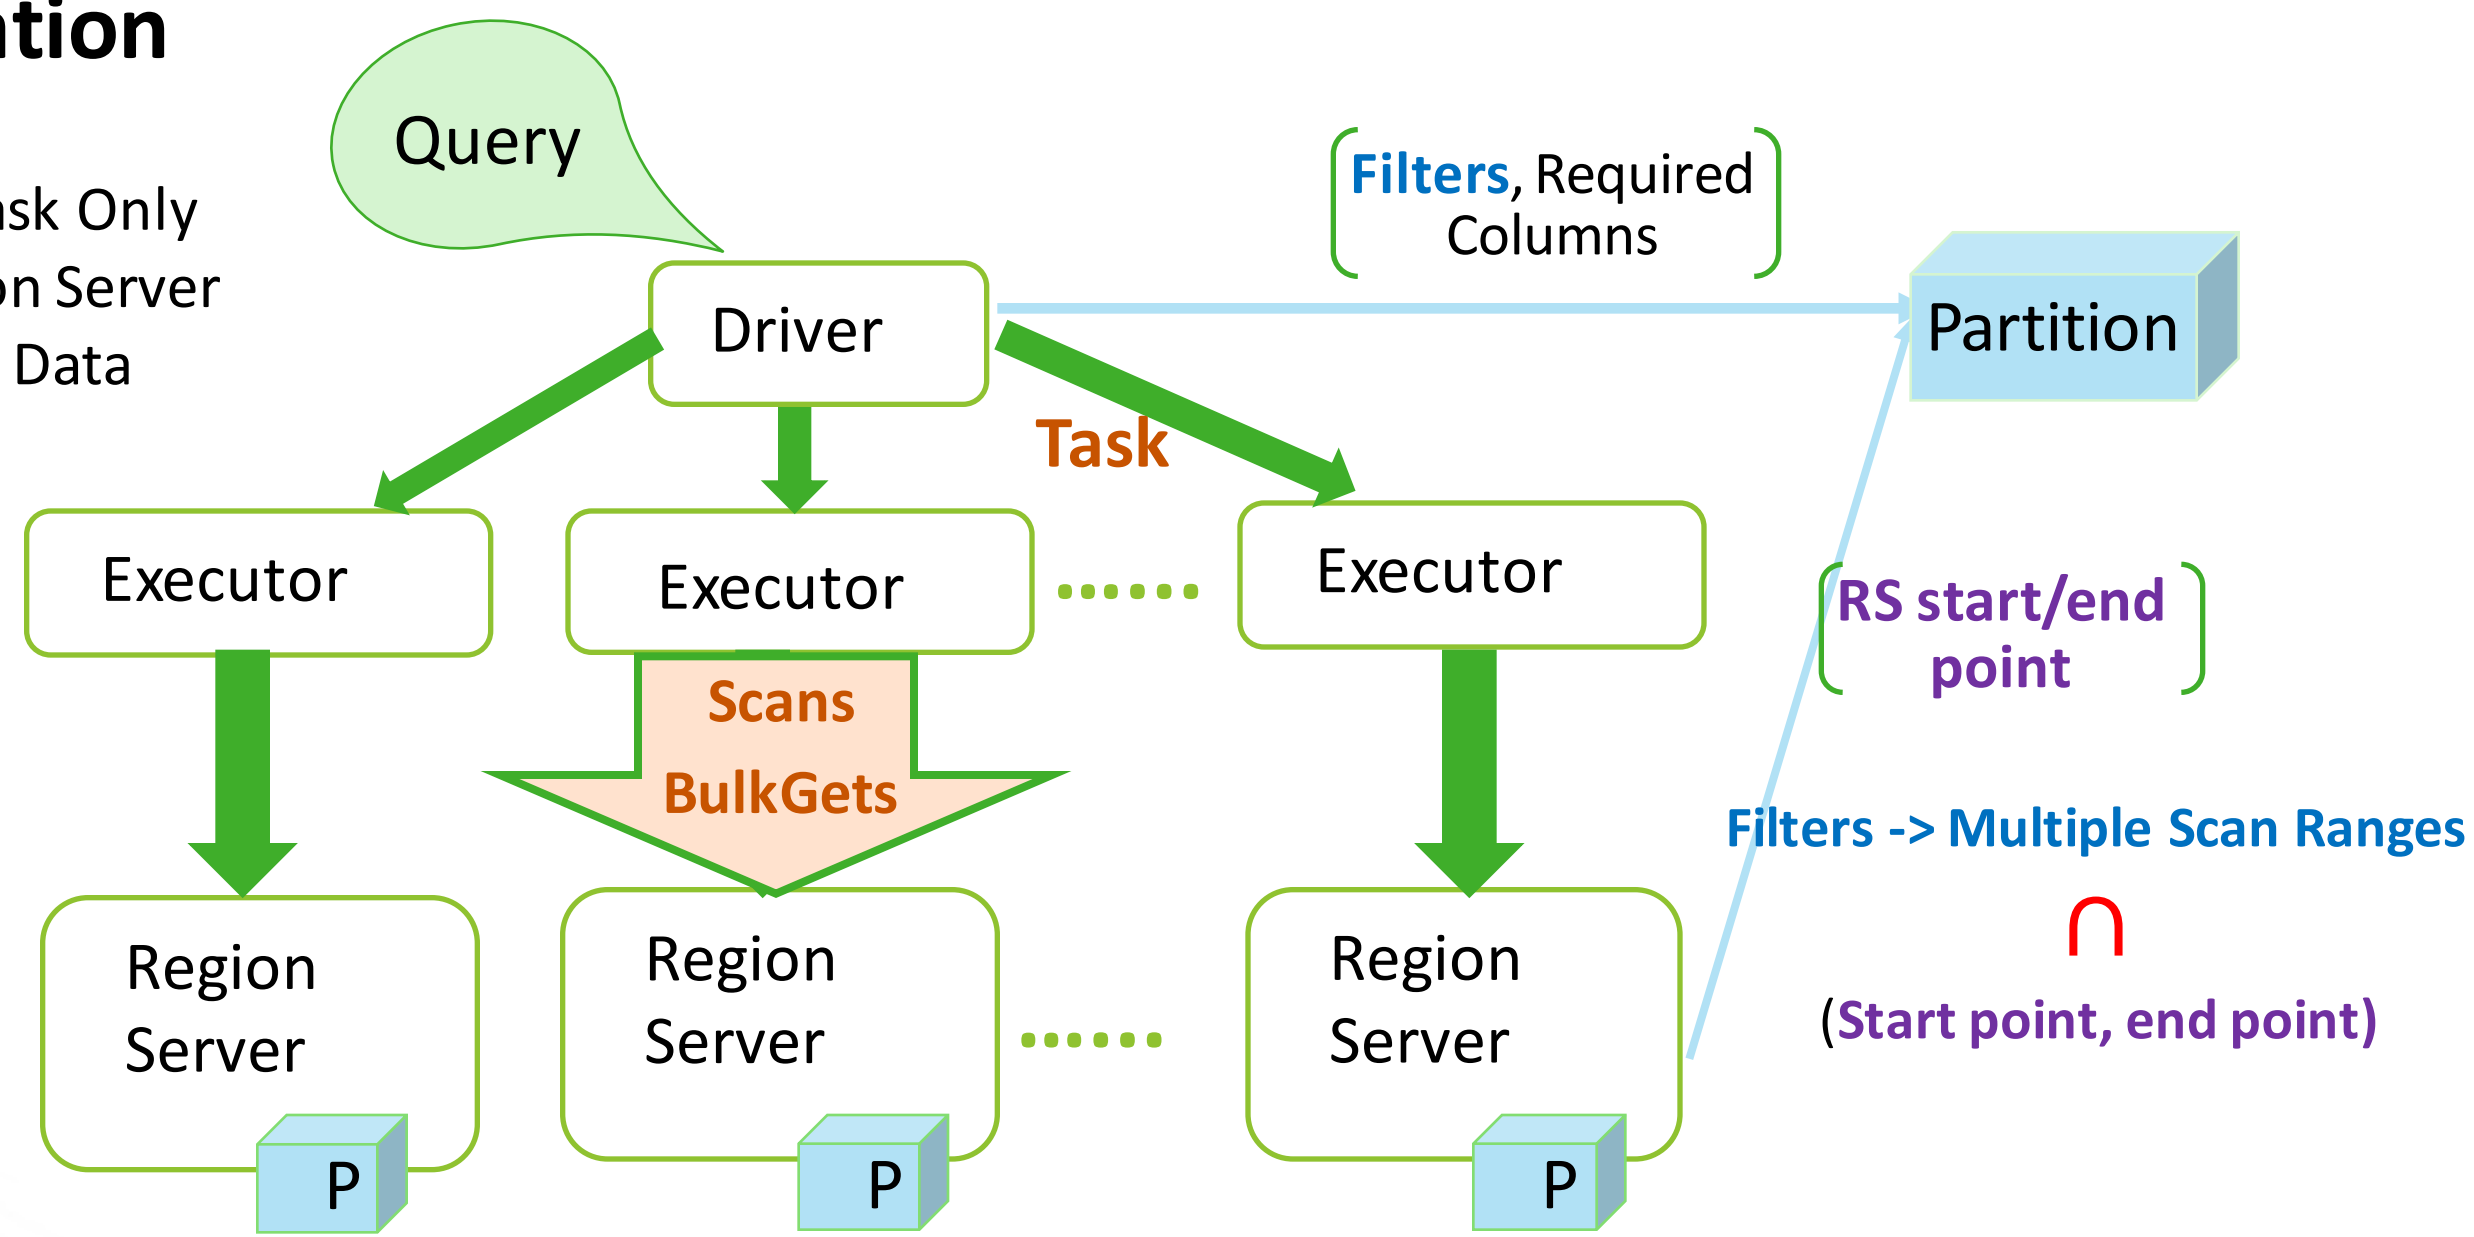
Partition Pruning: Task Only Performed in Region Server Holding Requested Data



Picture 1. SHC architecture

Implementation

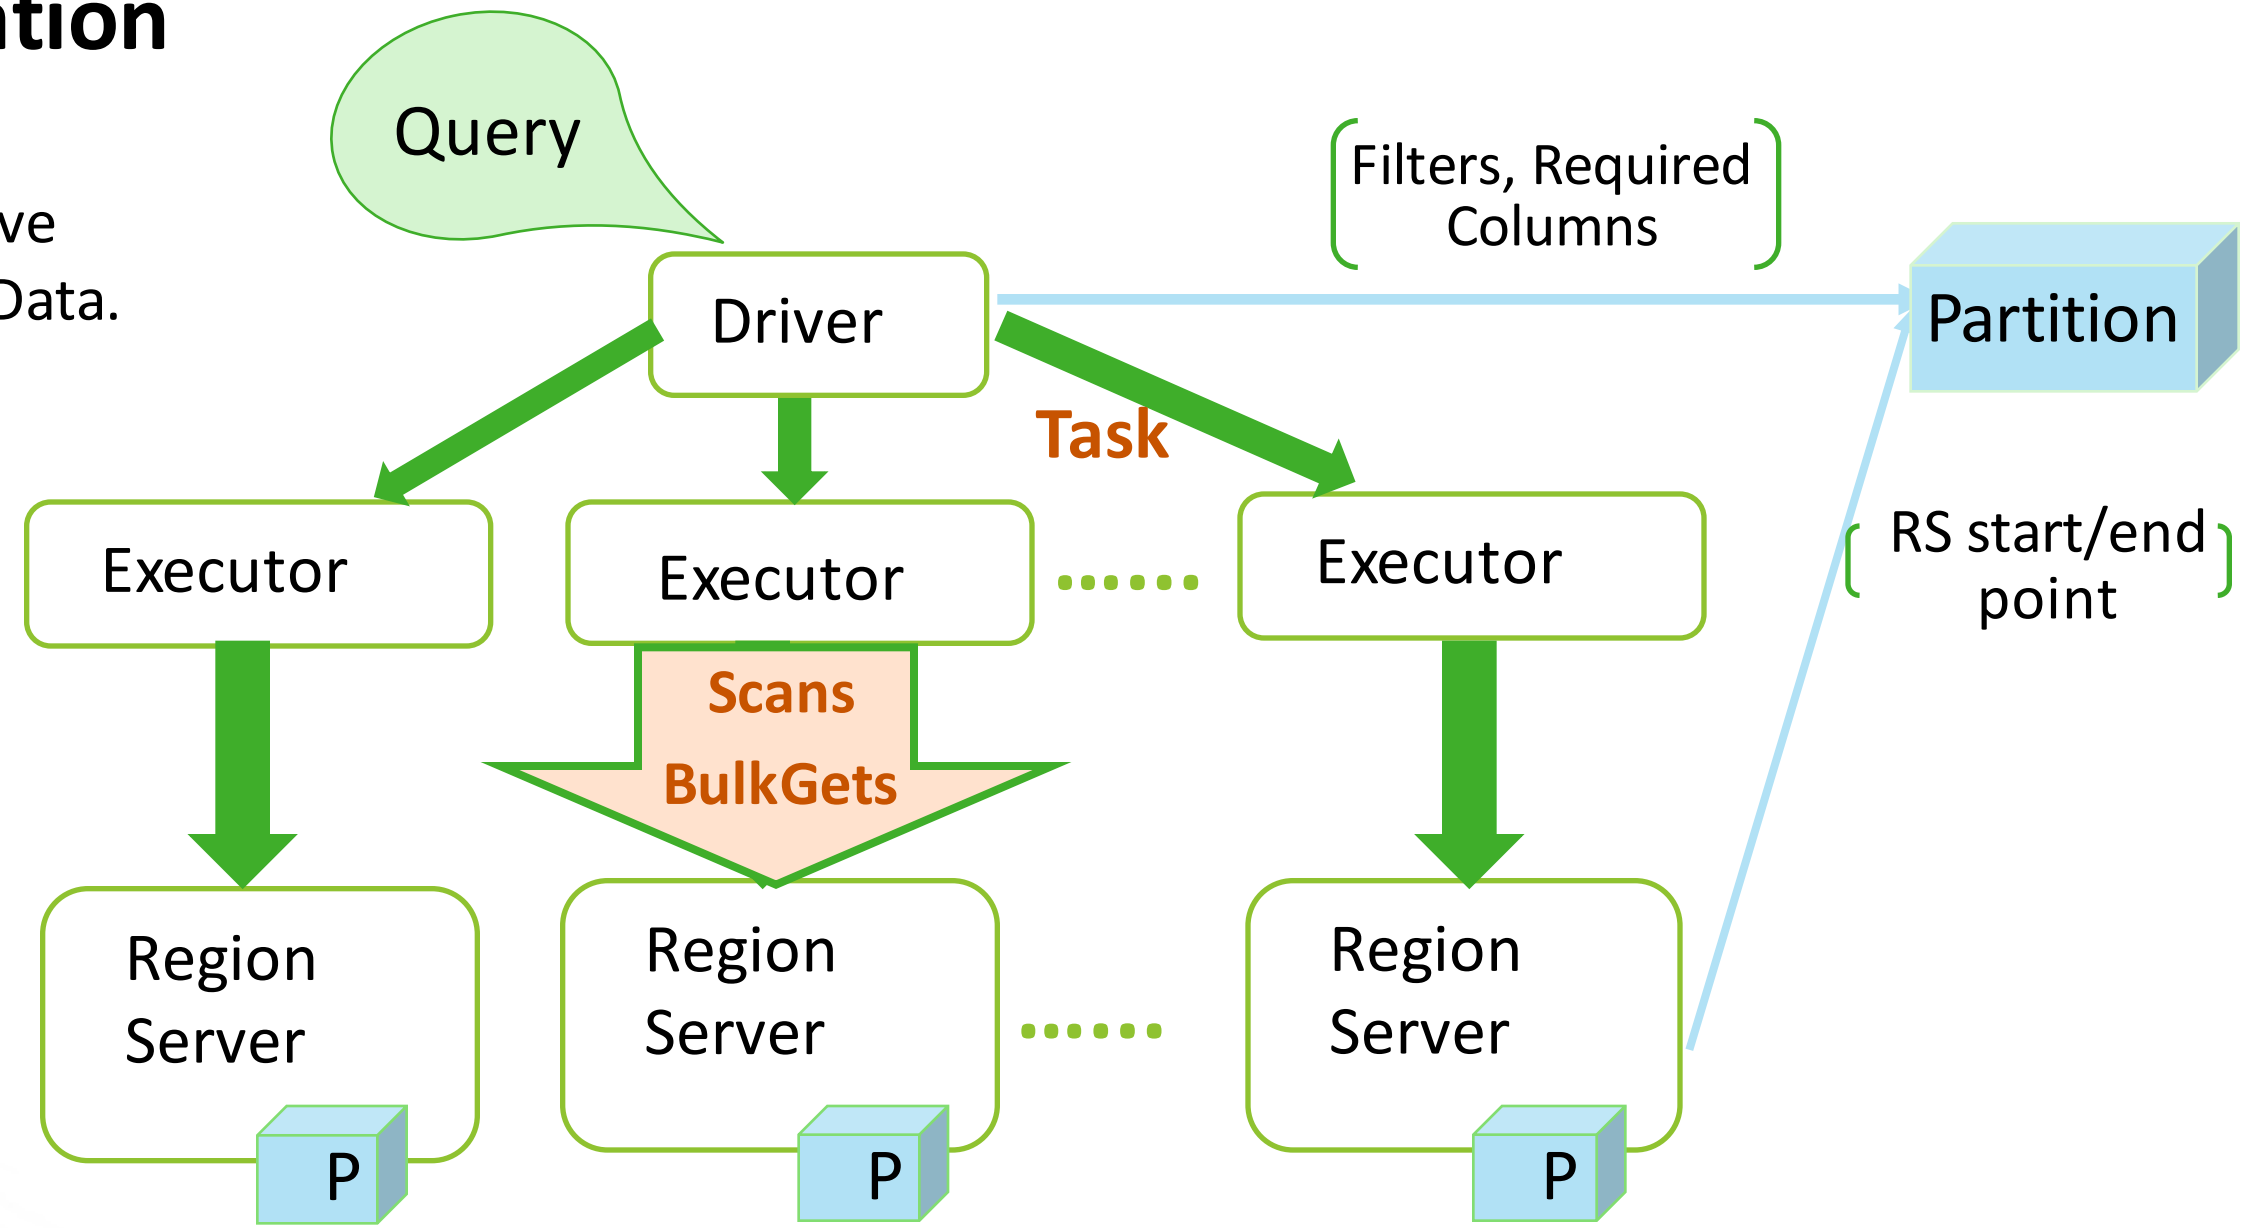
Partition Pruning: Task Only Performed in Region Server Holding Requested Data



Picture 1. SHC architecture

Implementation

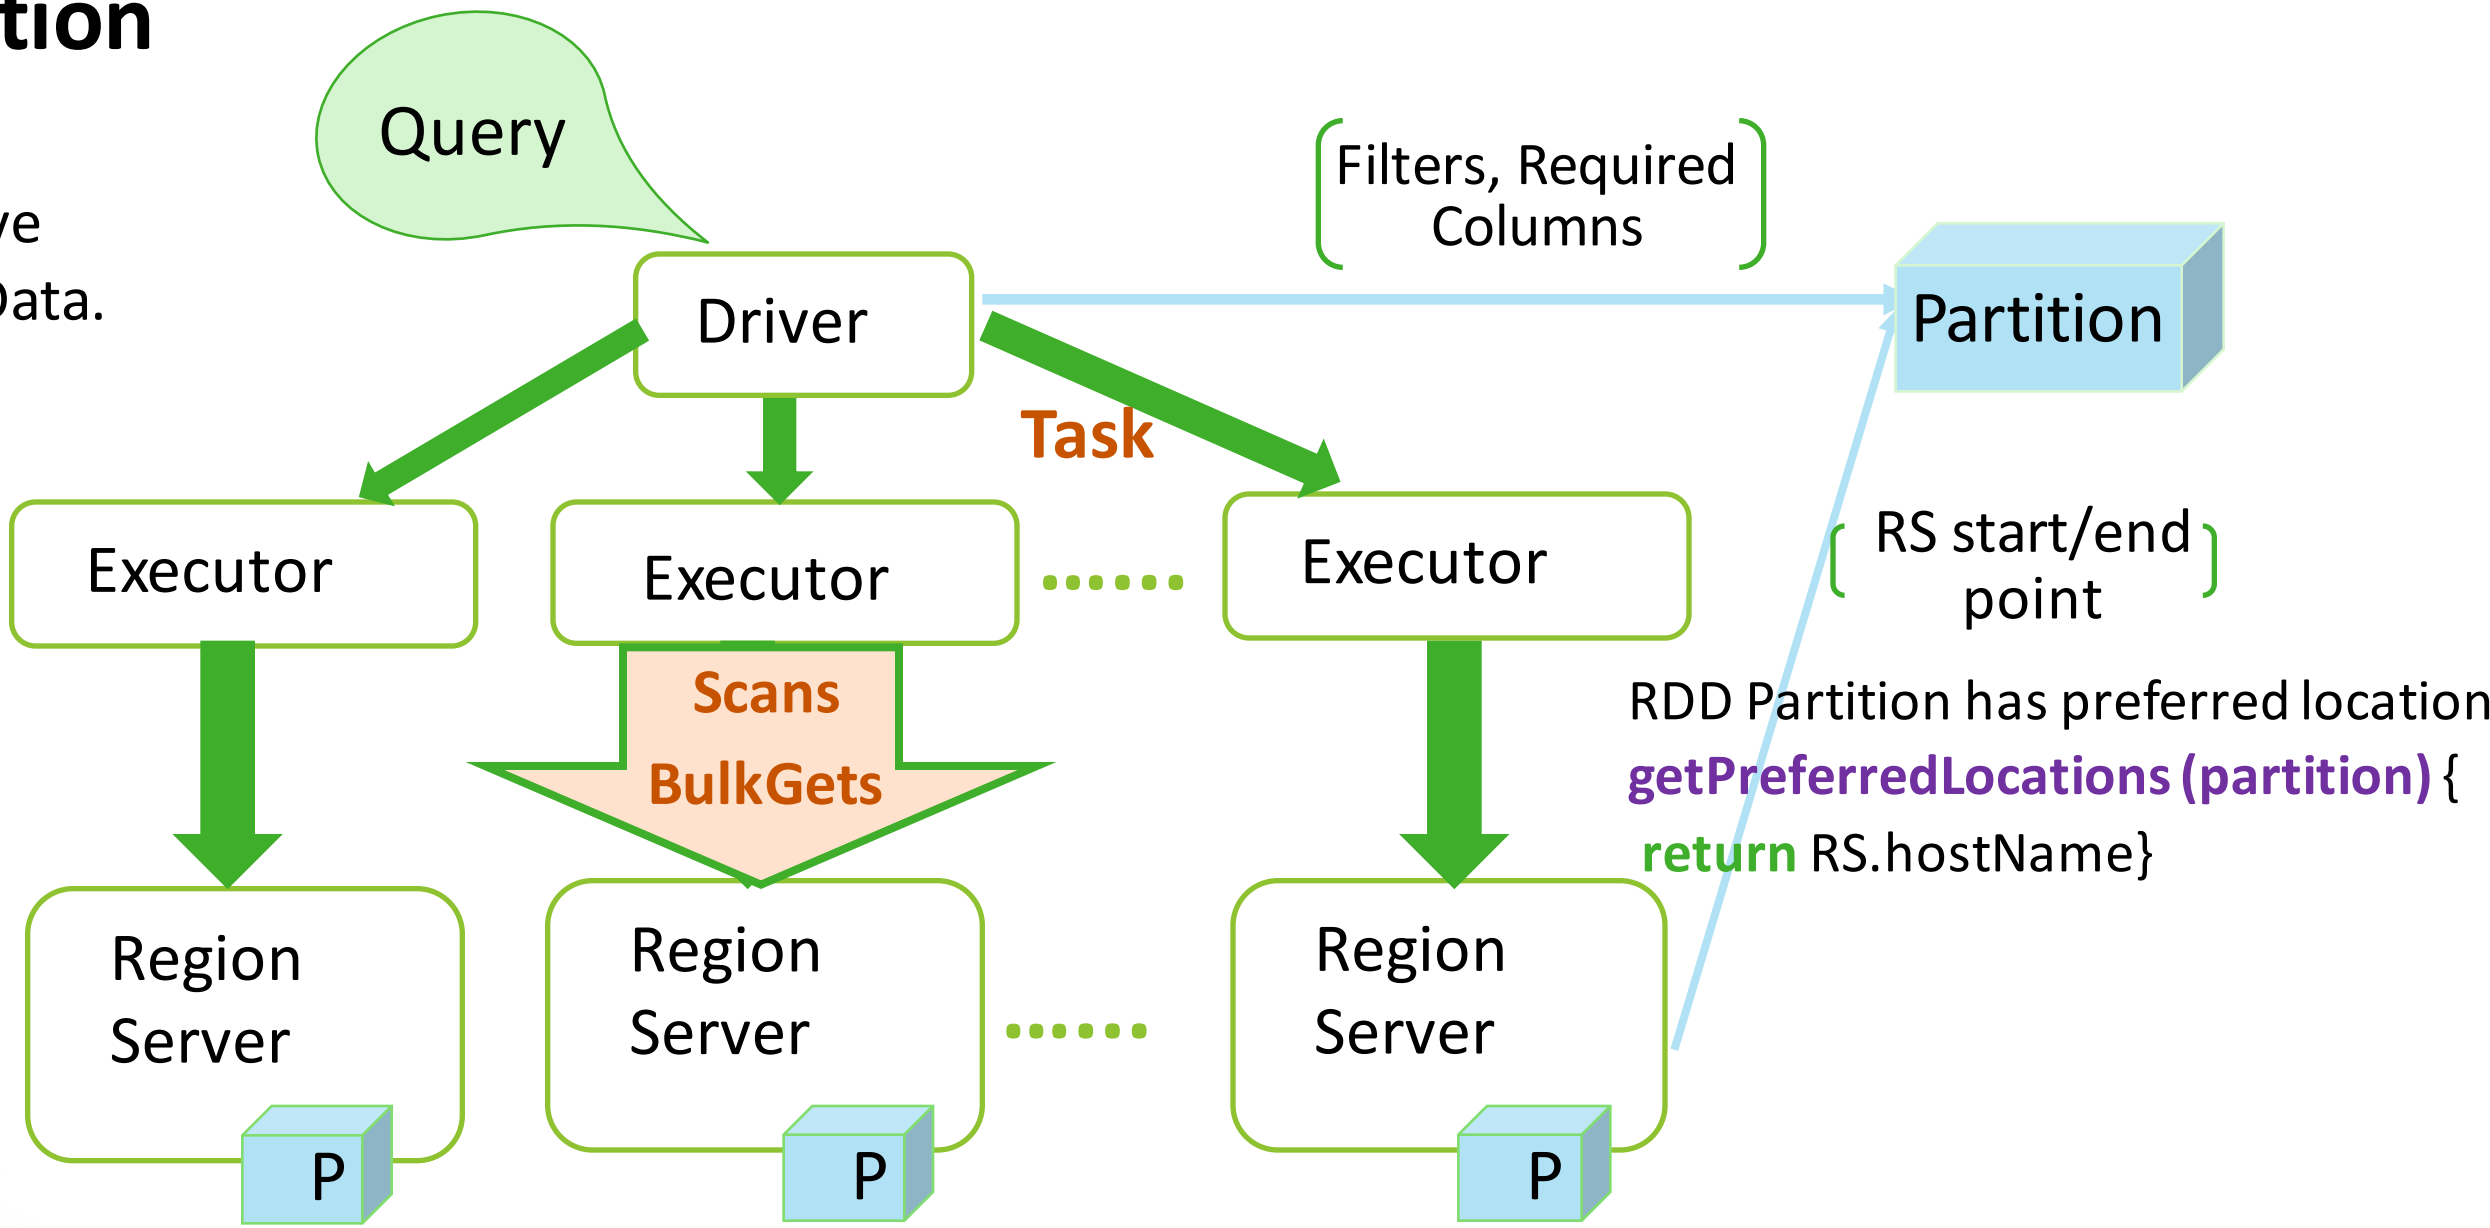
Data Locality: Move
Computation to Data.



Picture 1. SHC architecture

Implementation

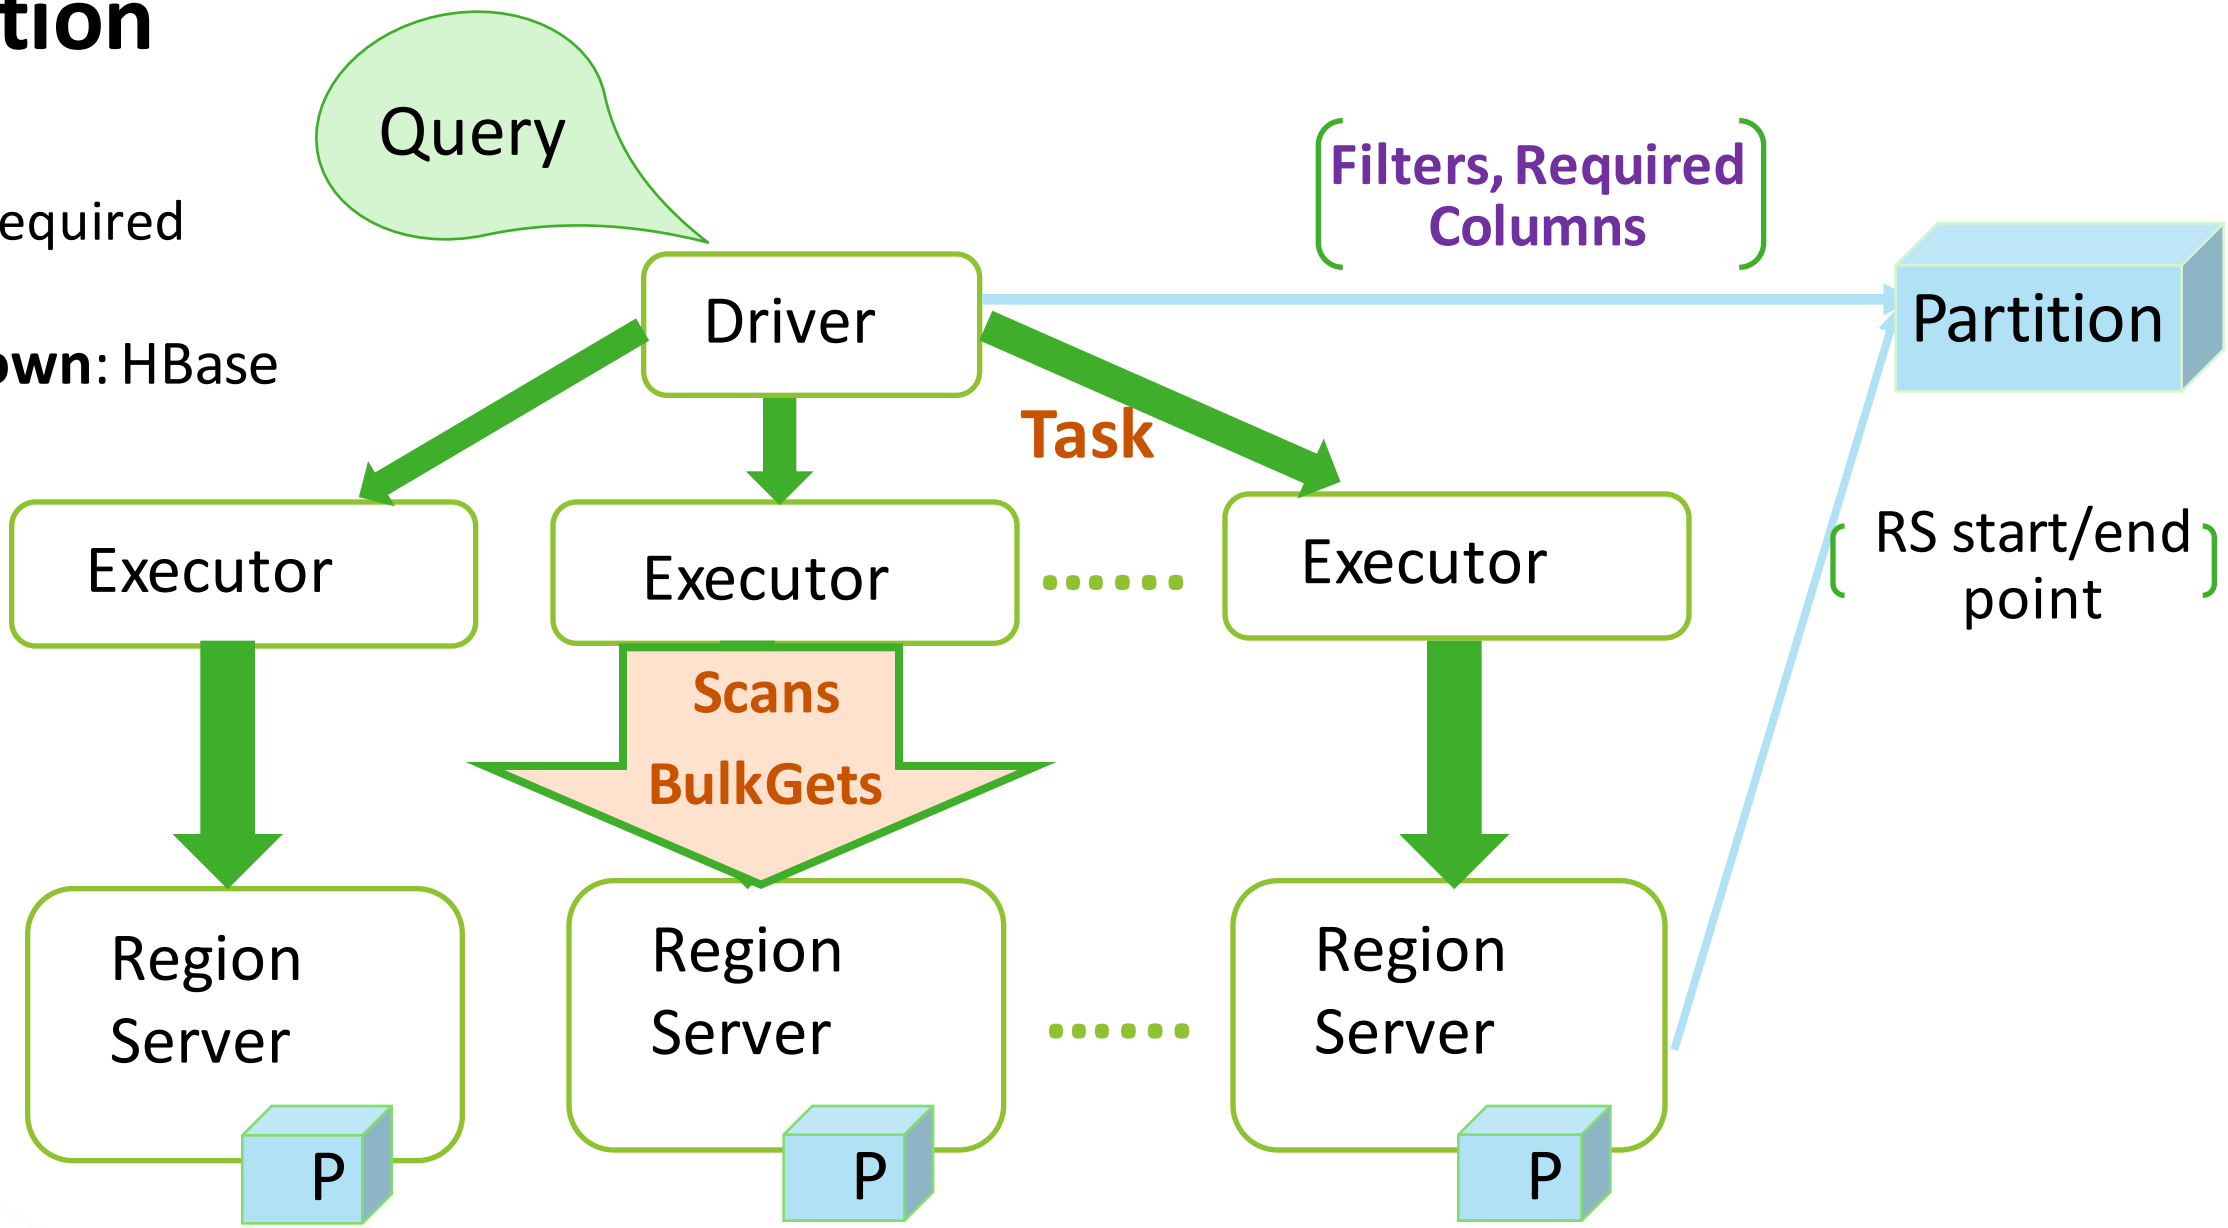
Data Locality: Move
Computation to Data.



Picture 1. SHC architecture

Implementation

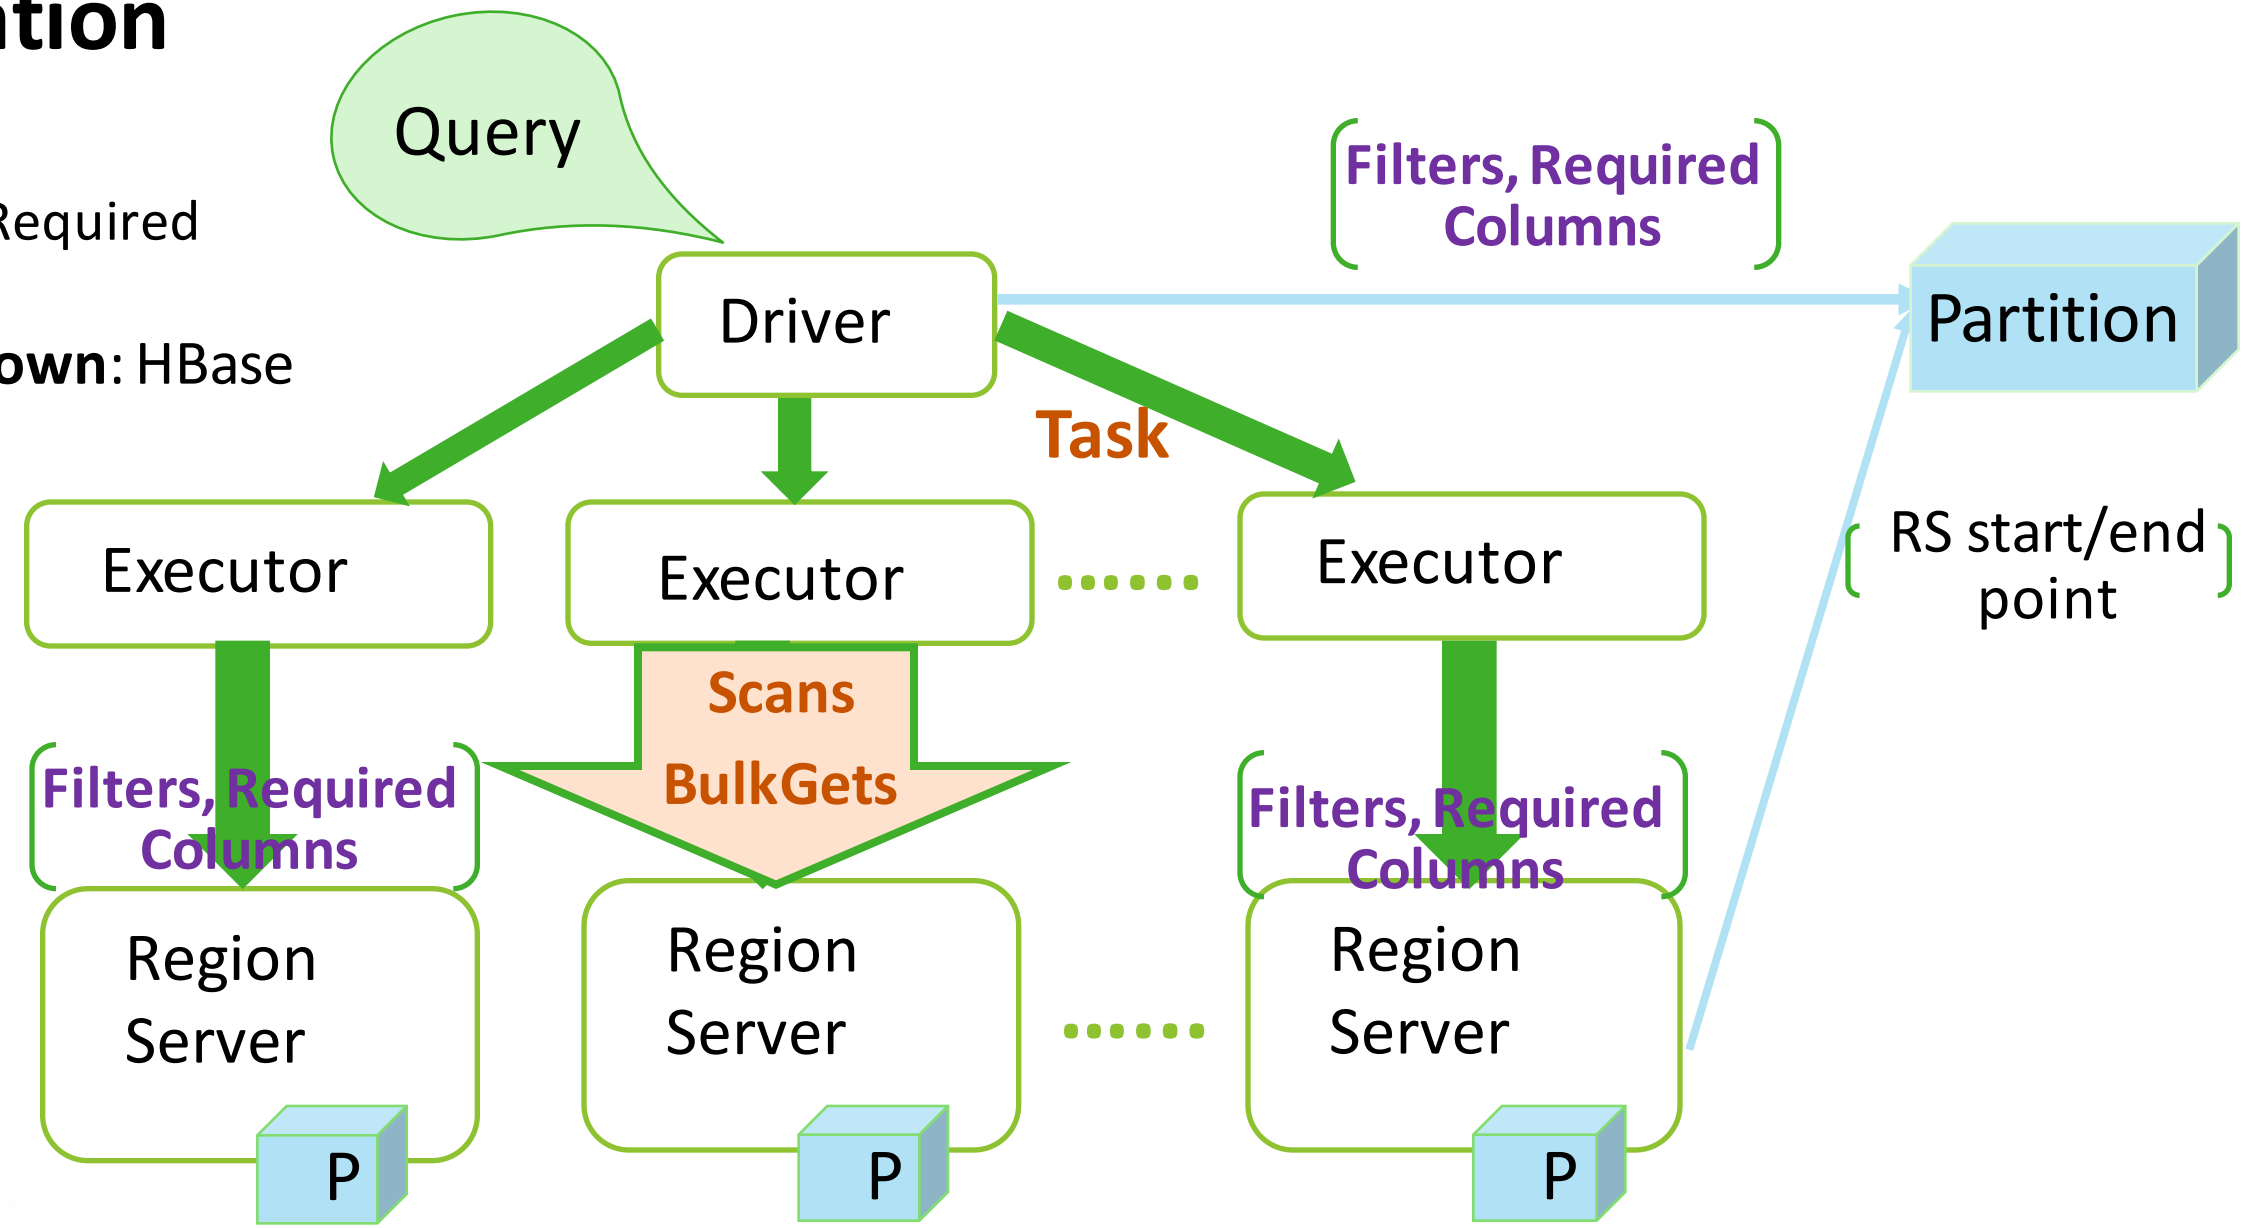
Column Pruning: Required Column
Predicate Pushdown: HBase built-in Filters



Picture 1. SHC architecture

Implementation

Column Pruning: Required Column
Predicate Pushdown: HBase built-in Filters

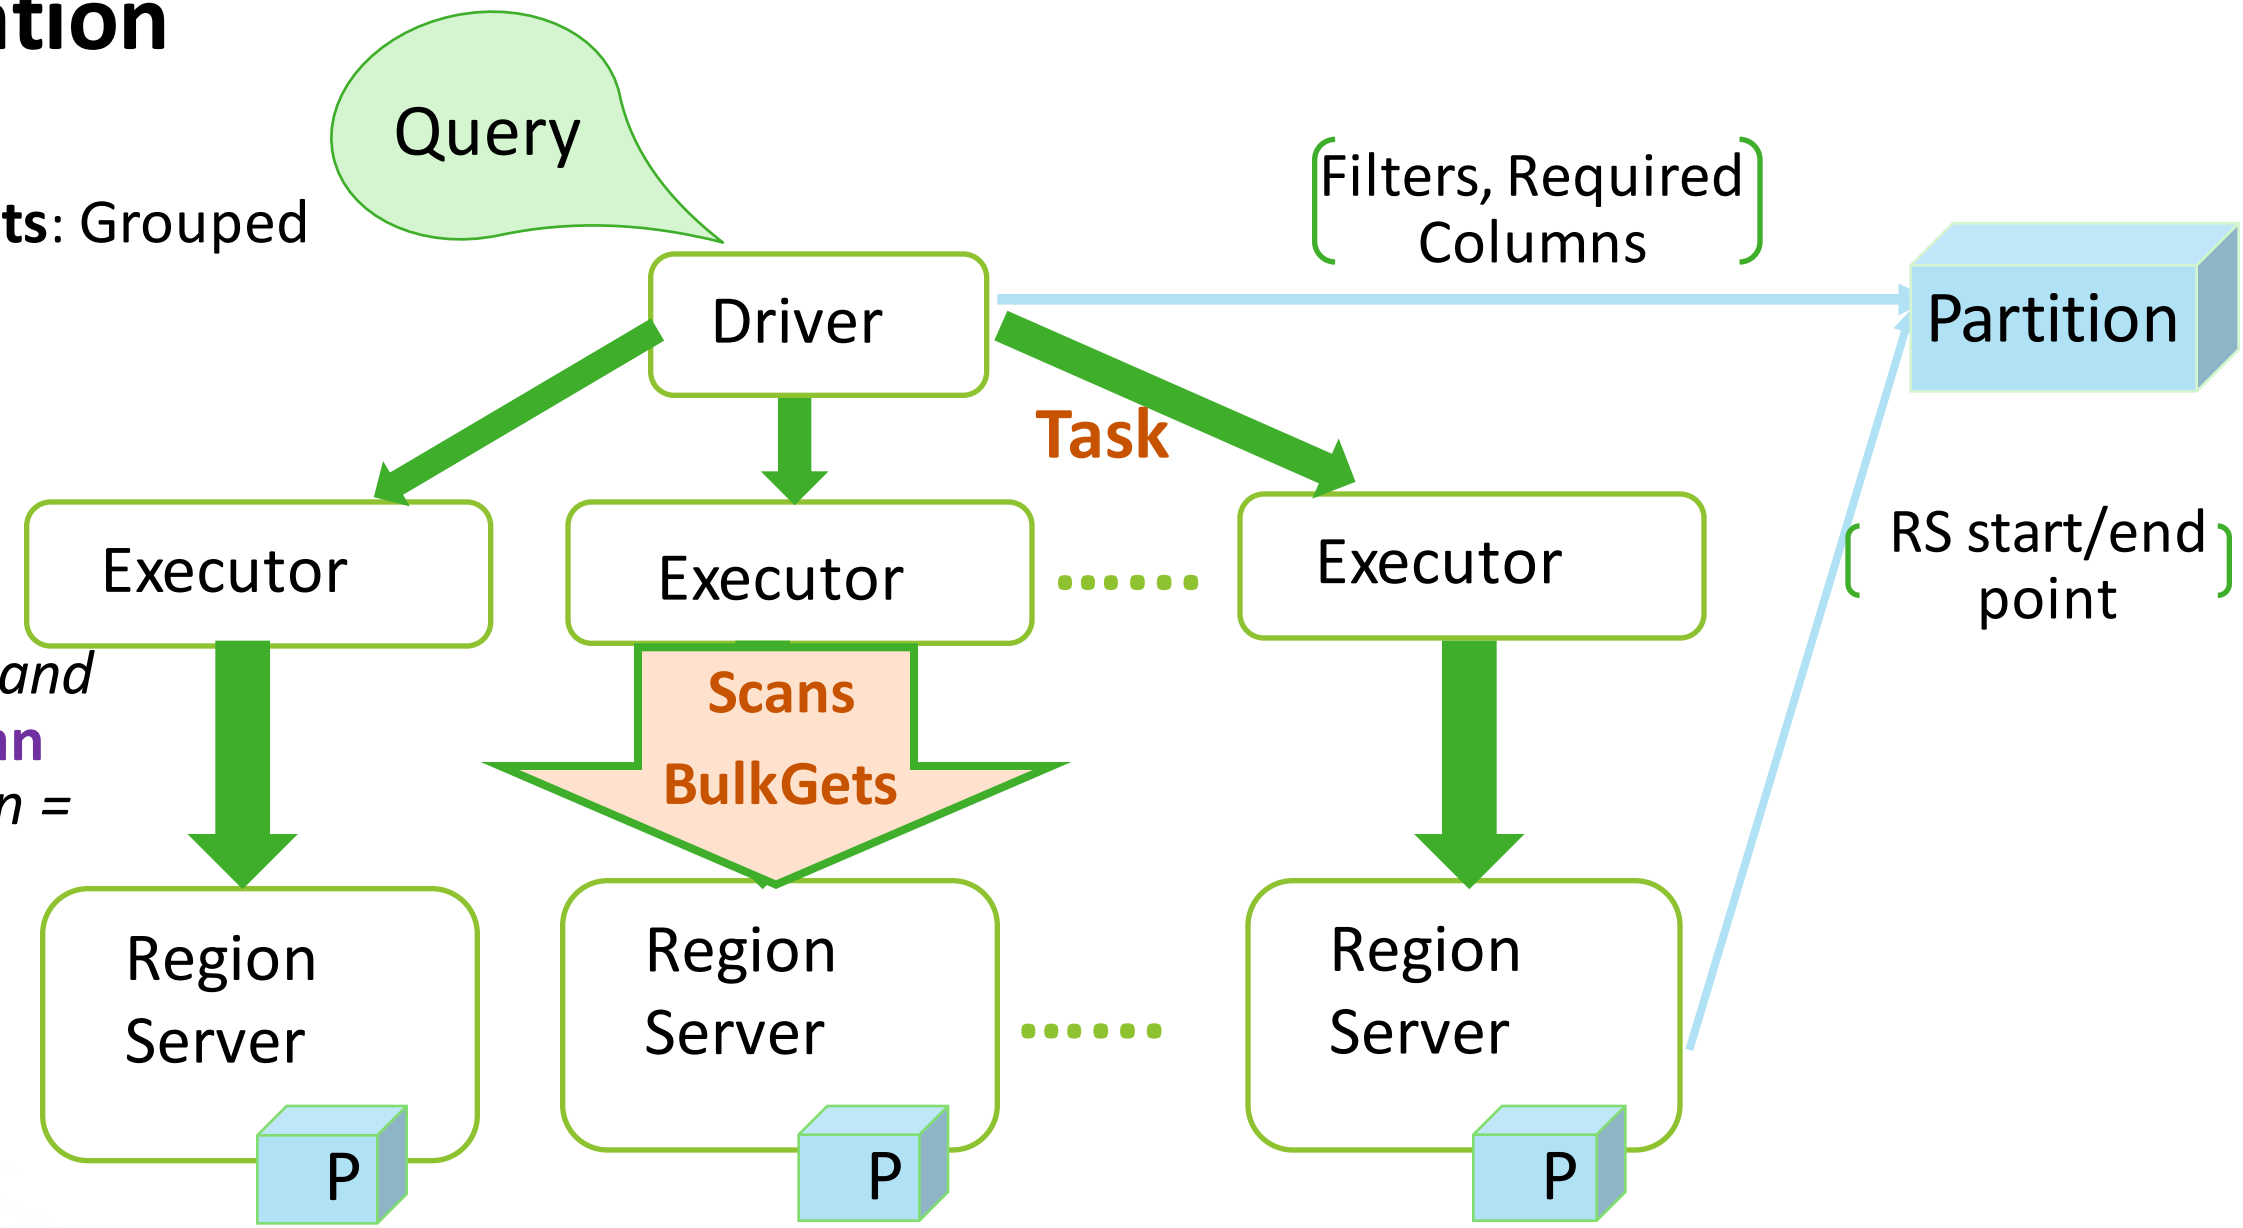


Picture 1. SHC architecture

Implementation

Scan and BulkGets: Grouped by region server.

*WHERE column > x and column < y for **scan** and WHERE column = x for **get**.*



Picture 1. SHC architecture

Usage

How to Use SHC?

- ◆ Github
 - <https://github.com/hortonworks-spark/shc>
- ◆ SHC Examples
 - <https://github.com/hortonworks-spark/shc/tree/master/examples>
- ◆ Running in Secure Cluster
 - Using SHCCredentialsManager for Accessing Multiple Secure HBase Clusters
- ◆ Demo

Usage

Define the catalog for the schema mapping:

```
def catalog = s"""{
  |"table":{"namespace":"default", "name":"phoenixTable",
  |  "tableCoder":"Phoenix", "version":"2.0"},
  |"rowkey":"key",
  |"columns":{
  |"col0":{"cf": "rowkey", "col":"key", "type":"string"},
  |"col1":{"cf":"cf1", "col":"col1", "type":"boolean"},
  |"col2":{"cf":"cf2", "col":"col2", "type":"double"},
  |"col3":{"cf":"cf3", "col":"col3", "type":"float"},
  |"col4":{"cf":"cf4", "col":"col4", "type":"int"},
  |"col5":{"cf":"cf5", "col":"col5", "type":"bigint"},
  |"col6":{"cf":"cf6", "col":"col6", "type":"smallint"},
  |"col7":{"cf":"cf7", "col":"col7", "type":"string"},
  |"col8":{"cf":"cf8", "col":"col8", "type":"tinyint"}
  |}
  |}""".stripMargin
```

Usage

◆ Prepare the data and populate the HBase table

```
val data = (0 to 255).map { i => HBaseRecord(i, "extra") }
```

```
sc.parallelize(data).toDF.write.options(  
  Map(HBaseTableCatalog.tableCatalog -> catalog, HBaseTableCatalog.newTable -> "5"))  
  .format("org.apache.spark.sql.execution.datasources.hbase")  
  .save()
```

Usage

◆ Load the DataFrame

```
def withCatalog(cat: String): DataFrame = {  
  sqlContext  
    .read  
    .options(Map(HBaseTableCatalog.tableCatalog->cat))  
    .format("org.apache.spark.sql.execution.datasources.hbase")  
    .load()  
}
```

```
val df = withCatalog(catalog)
```

Usage

◆ Query

Language integrated query:

```
val s = df.filter((( $"col0" <= "row050" && $"col0" > "row040" ) ||  
    $"col0" === "row005" && ( $"col4" === 1 || $"col4" === 42 ) )  
    .select("col0", "col1", "col4")
```

SQL:

```
val s = df.filter((( $"col0" <= "row050" && $"col0" > "row040" )  
df.registerTempTable("table")  
sqlContext.sql("select count(col1) from table").show
```


Usage

◆ Work with different data sources

// Part 1: write data into Hive table and read data from it
`val df1 = sql("SELECT * FROM shcHiveTable")`

// Part 2: read data from Hbase table
`val df2 = withCatalog(cat)`

// Part 3: join the two dataframes
`val s1 = df1.filter($"key" <= "40").select("key", "col1")`
`val s2 = df2.filter($"key" <= "20" && $"key" >= "1").select("key", "col2")`
`val result = s1.join(s2, Seq("key"))`
`result.show()`

Acknowledgement

- ◆ HBase Community & Spark Community
- ◆ All Spark-HBase Contributors, Zhan Zhang

Reference

◆ Hortonworks Public Repo

- <http://repo.hortonworks.com/content/repositories/releases/com/hortonworks/>

◆ Apache Spark

- <http://spark.apache.org/>

◆ Apache HBase

- <https://hbase.apache.org/>

Thanks

Q & A

Emails:
wyang@hortonworks.com

BACKUP

Kerberos Cluster

- ◆ Kerberos Ticket
 - `kinit -kt foo.keytab foouser` or Principle/Keytab
- ◆ Long Running Service
 - `--principal, --keytab`
- ◆ Multiple Secure HBase Clusters
 - Spark only Supports Single Secure HBase Cluster
 - Use SHC Credential Manager
 - Refer [*LRJobAccessing2Clusters Example in github*](#)