

Cursors in Apache Phoenix

HBaseCon West 2017
June 12, 2017

Anirudha Jadhav
ajadhav2@bloomberg.net

Biju Nair
bnair10@bloomberg.net

techatbloomberg.com

© 2017 Bloomberg Finance L.P. All rights reserved

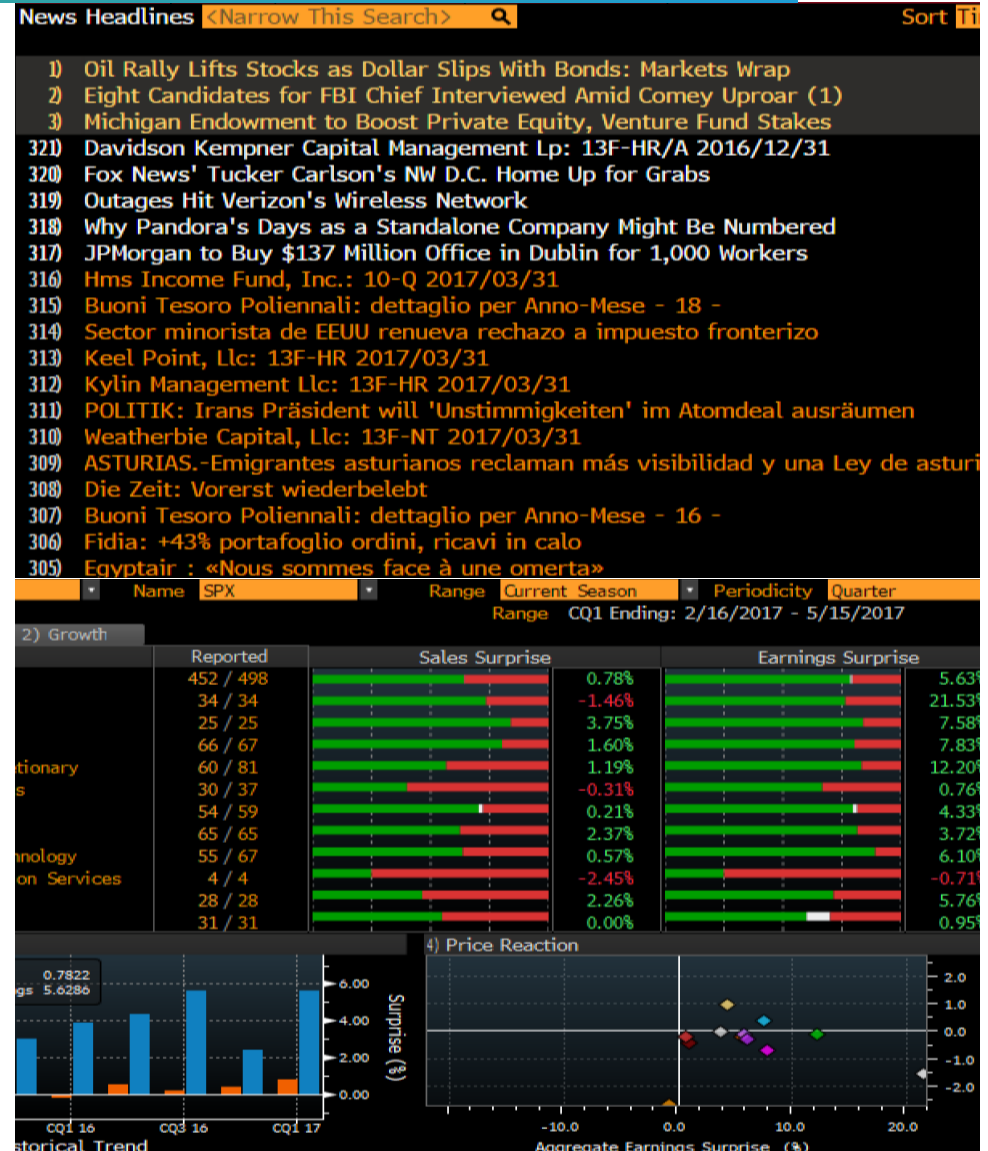
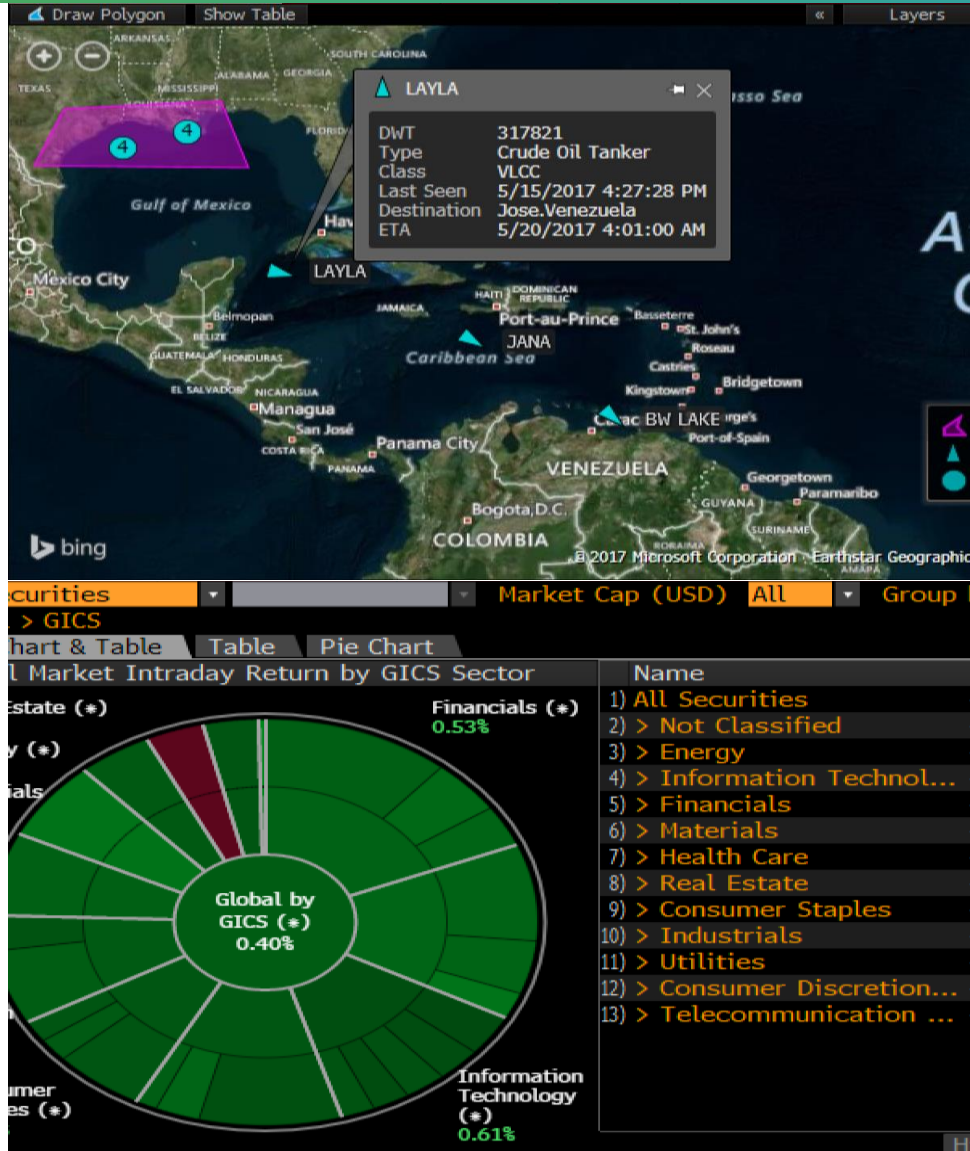
Bloomberg

Bloomberg

Leading data and analytics provider for the financial industry



Bloomberg is a data company



Reality of working with data

- The data model changes over time
- Users querying the data model don't necessarily change
- Alternate query patterns for the same dataset
- Data infrastructure usage needs to be simple

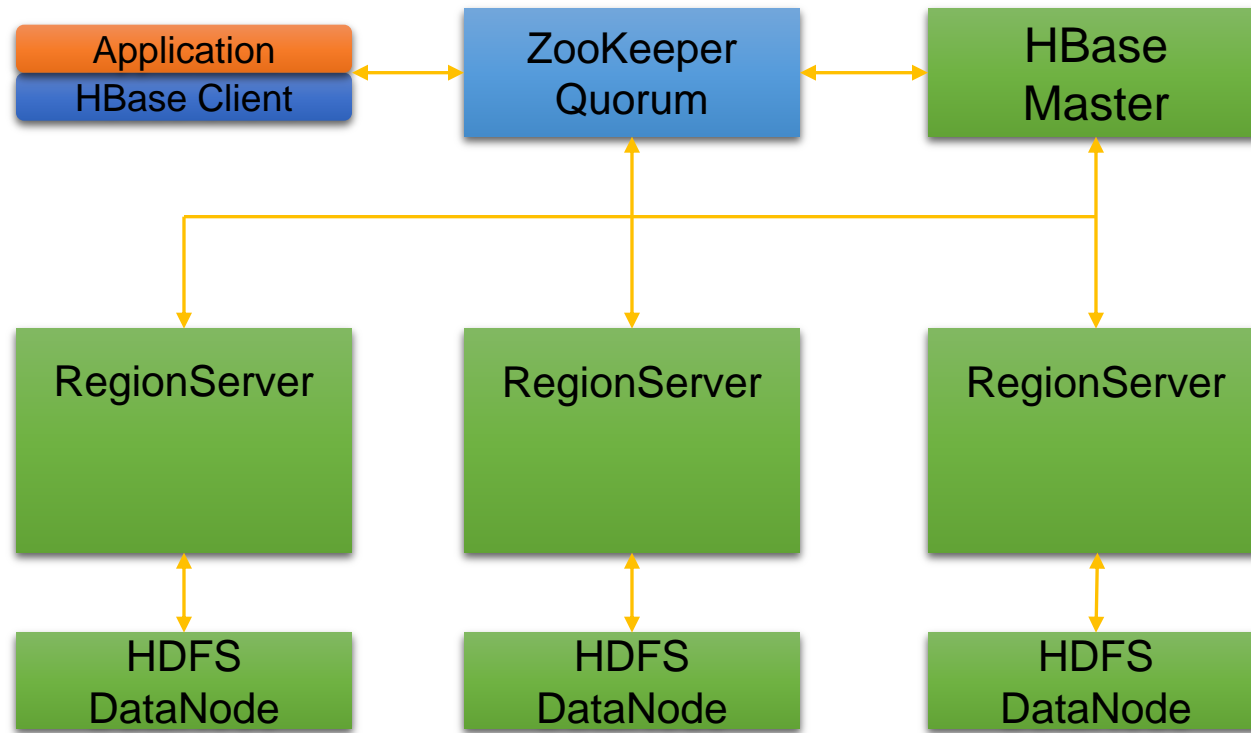
Apache Phoenix

- Recipes of best practices for using HBase over a familiar SQL'ish grammar
- It is so much more than SQL
 - User defined functions for push-down
 - Secondary indices
 - Statistics collections, optimizations based on heuristics
 - ORM libraries
 - JDBC, ODBC support with Query servers
 - Integrations: Spark, Kafka, MR and others

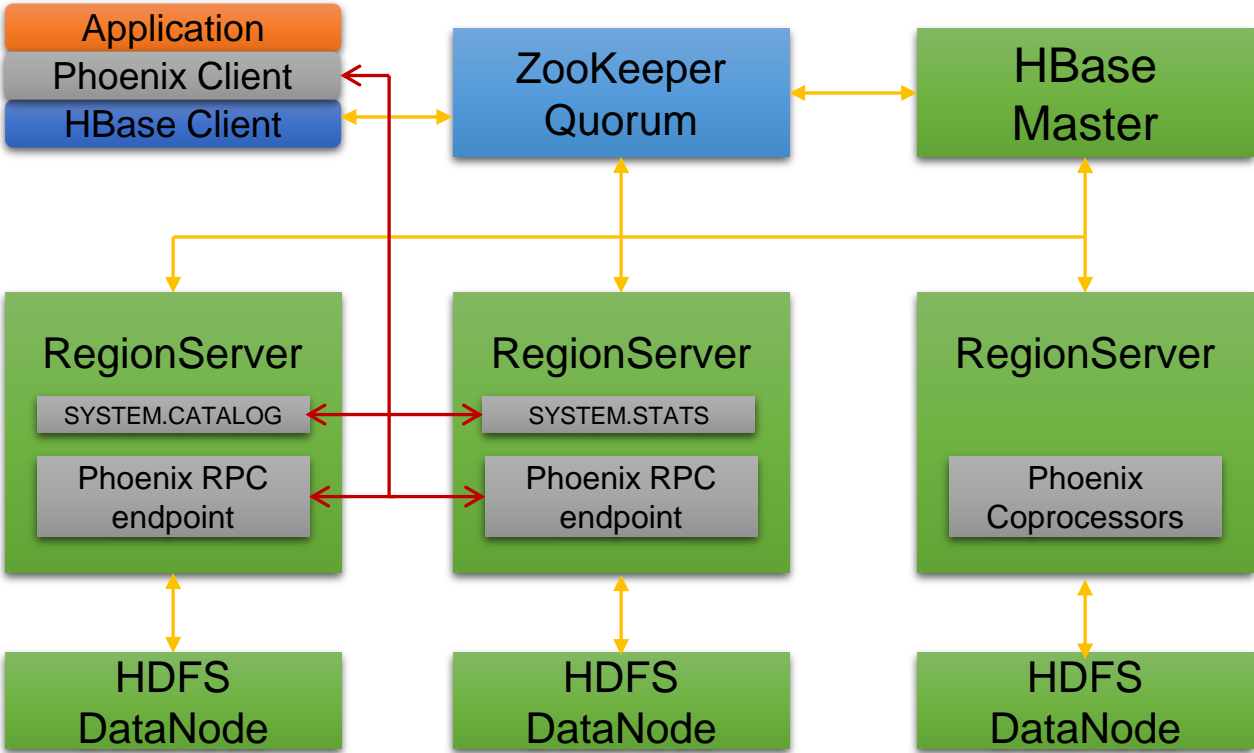
Extending Apache Phoenix

- A very active and helpful community
- Our ongoing work
 - Apache Calcite
 - Distributed tests and nightly performance build
 - Multi-DC replication
 - Deep paging with cursor implementation

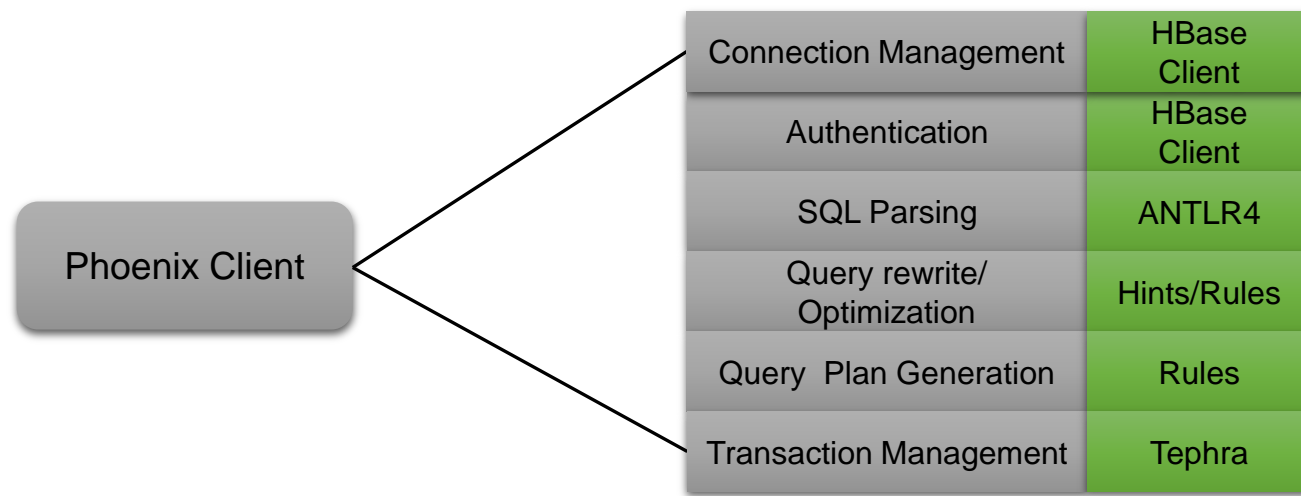
HBase



Phoenix



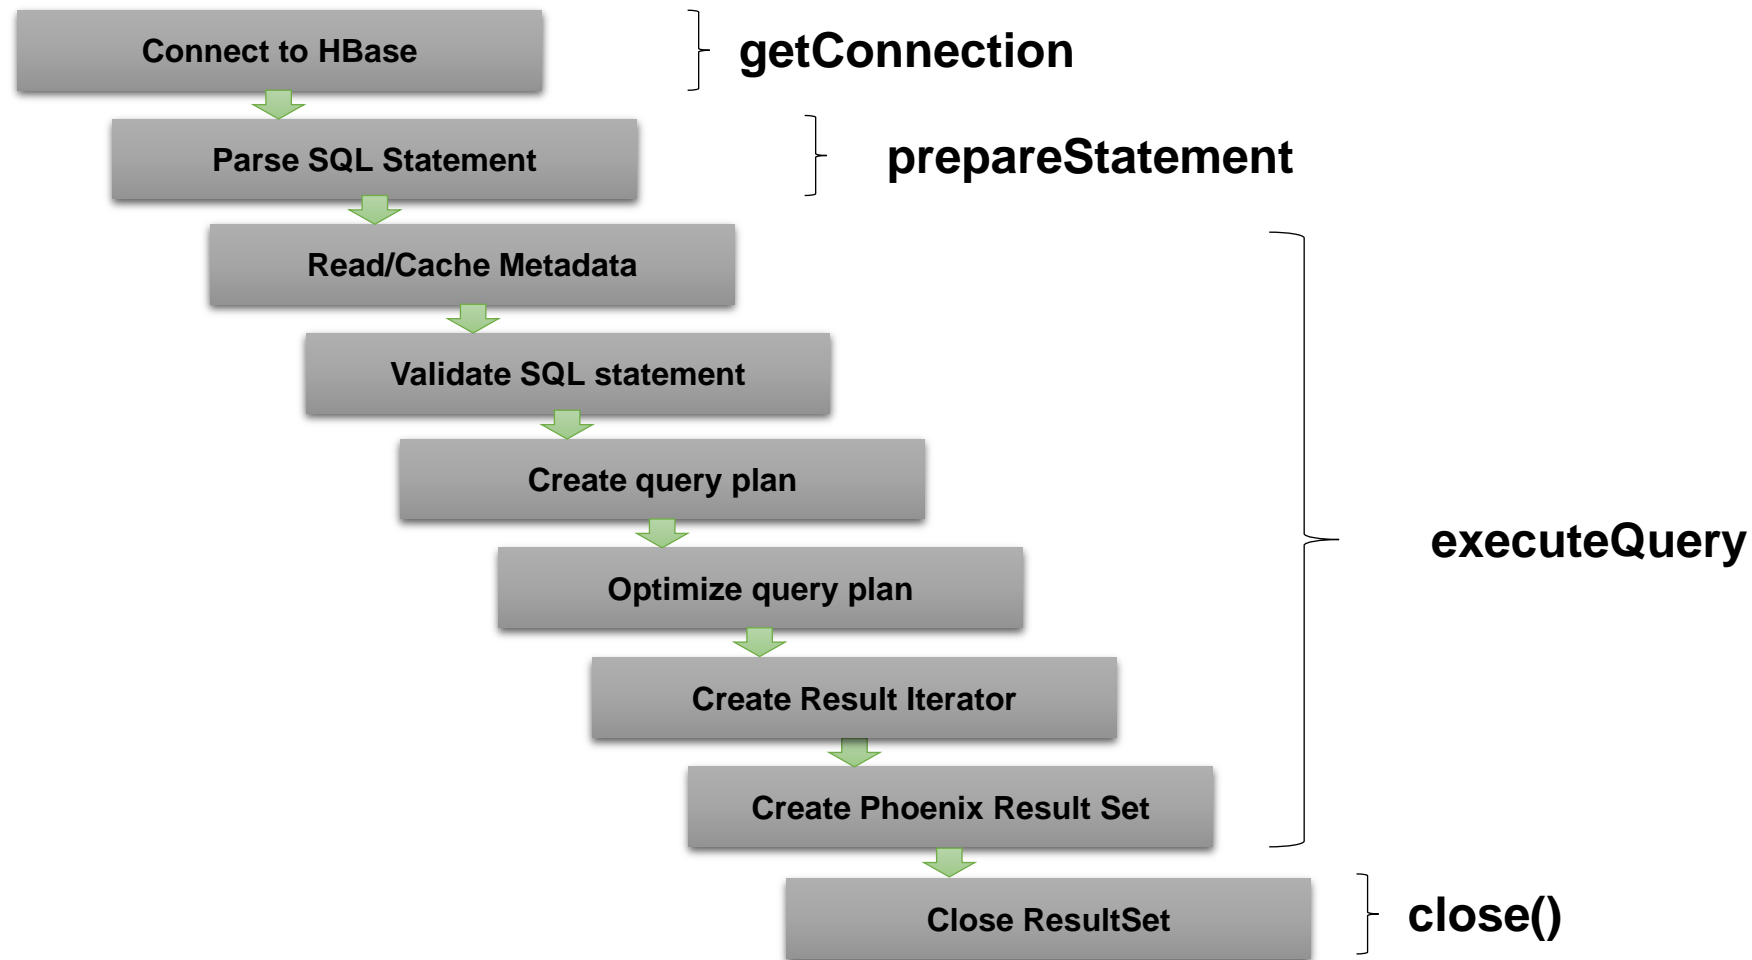
Phoenix Client



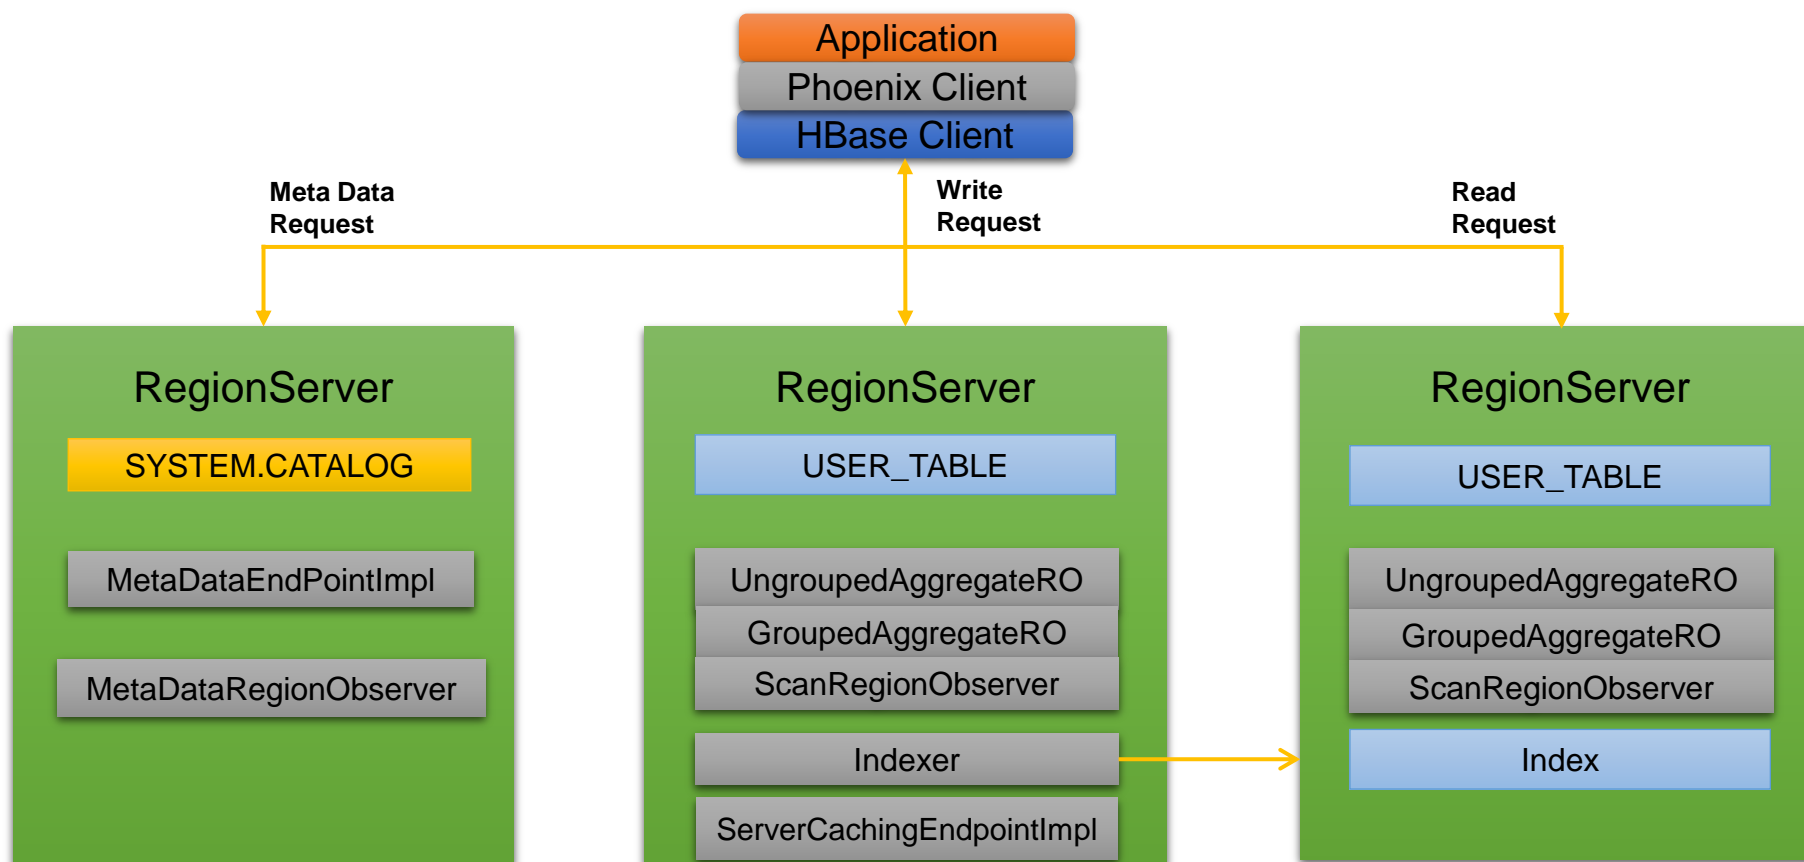
Phoenix query execution

```
Connection con =  
    DriverManager.getConnection("jdbc:phoenix:zkquorum:2181:/hbase:principal:keytabfile")  
    ;  
  
...  
PreparedStatement statement = con.prepareStatement("select * from TBL");  
  
...  
ResultSet rset = statement.executeQuery();  
  
...  
while (rset.next() != null)  
  
...  
rset.close()  
  
...
```

Phoenix query execution



Phoenix Server



Cursors

- To support row pagination
 - Should support forward and backward traversal
- Support required for select queries only
- Data needs to be consistent during traversal

Cursors

- DECLARE tCursor CURSOR FOR SELECT * FROM TBL
- OPEN tCursor
- FETCH NEXT 10 ROWS FROM tCursor
- FETCH PRIOR 5 ROWS FROM tCursor
- CLOSE tCursor

Implementation options

- [PHOENIX-2606](#)
- Use row value constructors
 - Query rewrite and complex
- Wrapper over available query Resultsets
 - Can leverage Resultsets and so relatively simple

Cursor Lifecycle

```
PreparedStatement statement = con.prepareStatement("DECLARE tCursor CURSOR  
FOR SELECT * FROM TBL");
```

```
statement.execute();
```

```
...
```

```
statement = con.prepareStatement("OPEN tCursor");
```

```
statement = con.prepareStatement("FETCH NEXT FROM tCursor");
```

```
ResultSet rset = statement.execute();
```

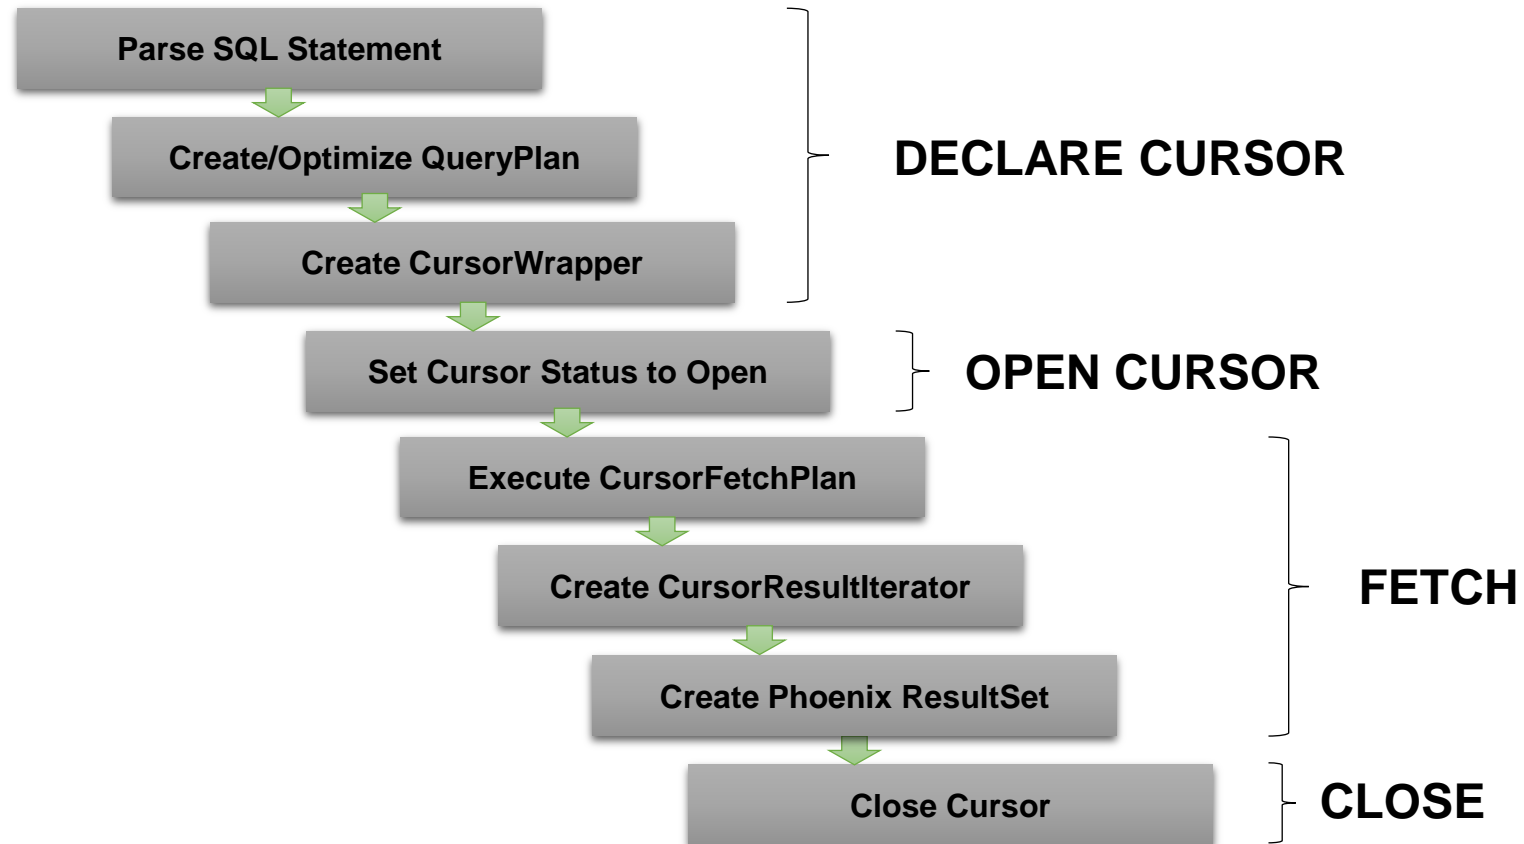
```
while (rset.next != null)
```

```
...
```

```
statement = con.prepareStatement("CLOSE tCursor");
```

```
statement.execute();
```

Cursor lifecycle



Cursor Challenges

- Data Consistency
 - Query start timestamp provides snapshot consistency
- Optimization
 - Use Scan object for non aggregate queries
- Cache sizing
 - Dynamic sizing

Contributors

- Gabriel Jimenez (MIT)
- Anirudha Jadhav (Bloomberg)
- Biju Nair (Bloomberg)
- Ankit Singhal (Hortonworks)

Q&A

Thank You

Hardening Apache Phoenix @ PhoenixCon tomorrow 2 PM PT

techatbloomberg.com

© 2017 Bloomberg Finance L.P. All rights reserved

Bloomberg