


YAHOO!

Accordion: HBase Breathes with In-Memory Compaction

Eshcar Hillel, Anastasia Braginsky, Edward Bortnikov | HBaseCon, Jun 12, 2017



The Team



Edward Bortnikov



Eshcar Hillel
(committer)



Anastasia Braginsky
(committer)




Michael Stack



Anoop Sam John

Ramkrishna Vasudevan

Quest: The User's Holy Grail



What is Accordion?

Novel Write-Path Algorithm

Better Performance of Write-Intensive Workloads

Write Throughput ↗, Read Latency ↘

Better Disk Use

Write amplification ↘

GA in HBase 2.0 (becomes default MemStore implementation)

In a Nutshell

Inspired by Log-Structured-Merge (LSM) Tree Design


Transforms random I/O to sequential I/O (efficient!)

Governs the HBase storage organization

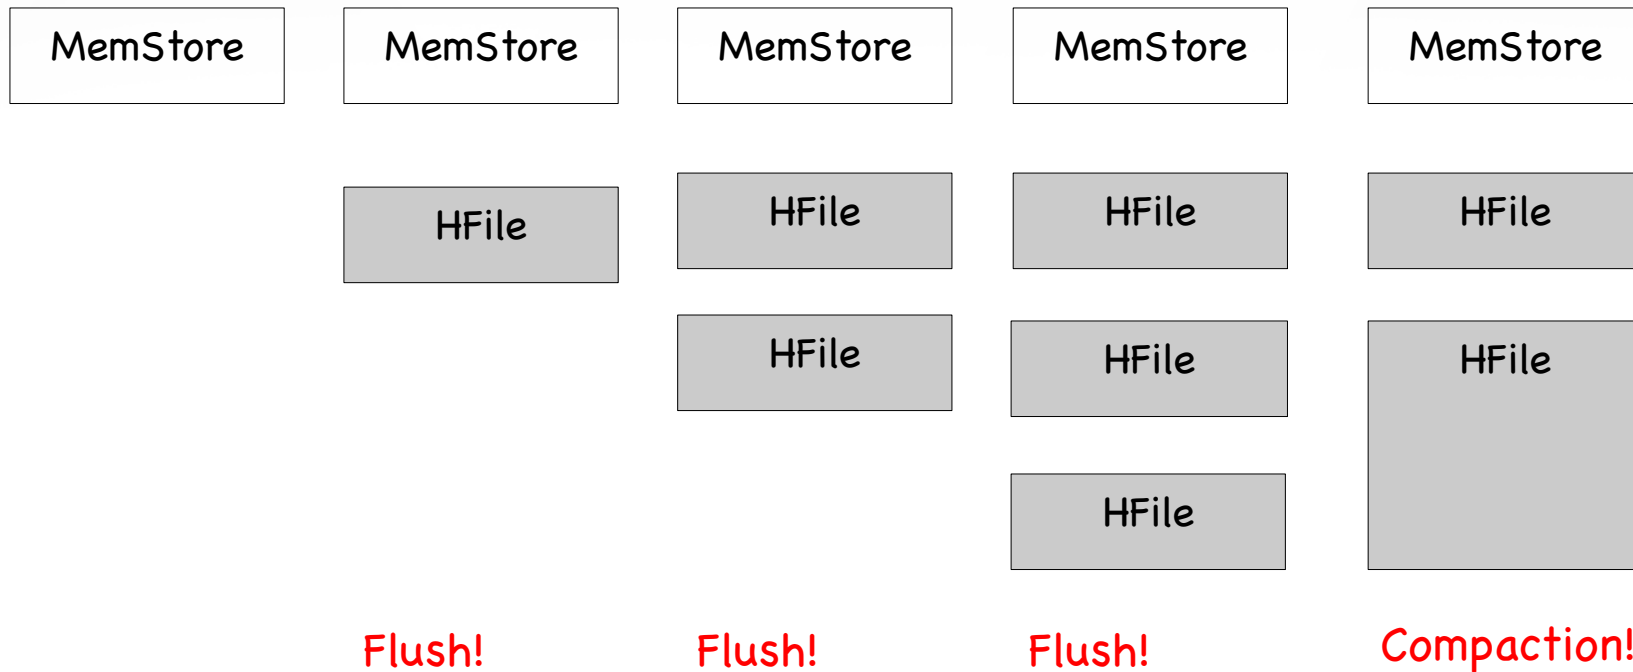
Accordian reapplies the LSM Tree design to RAM data

- Efficient resource use – data lives in memory longer
- Less disk I/O
- Ultimately, higher speed

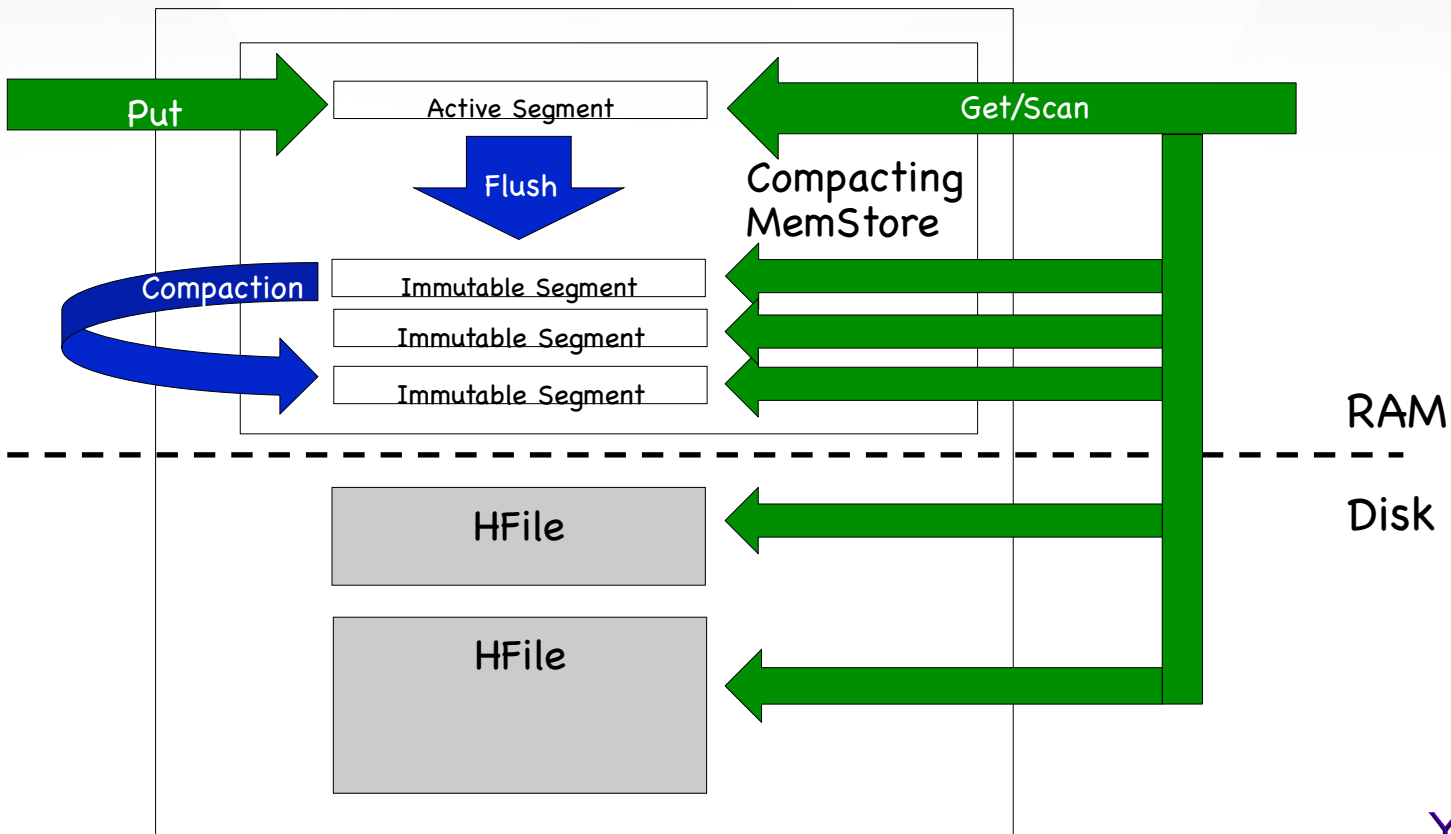
How LSM Trees Work



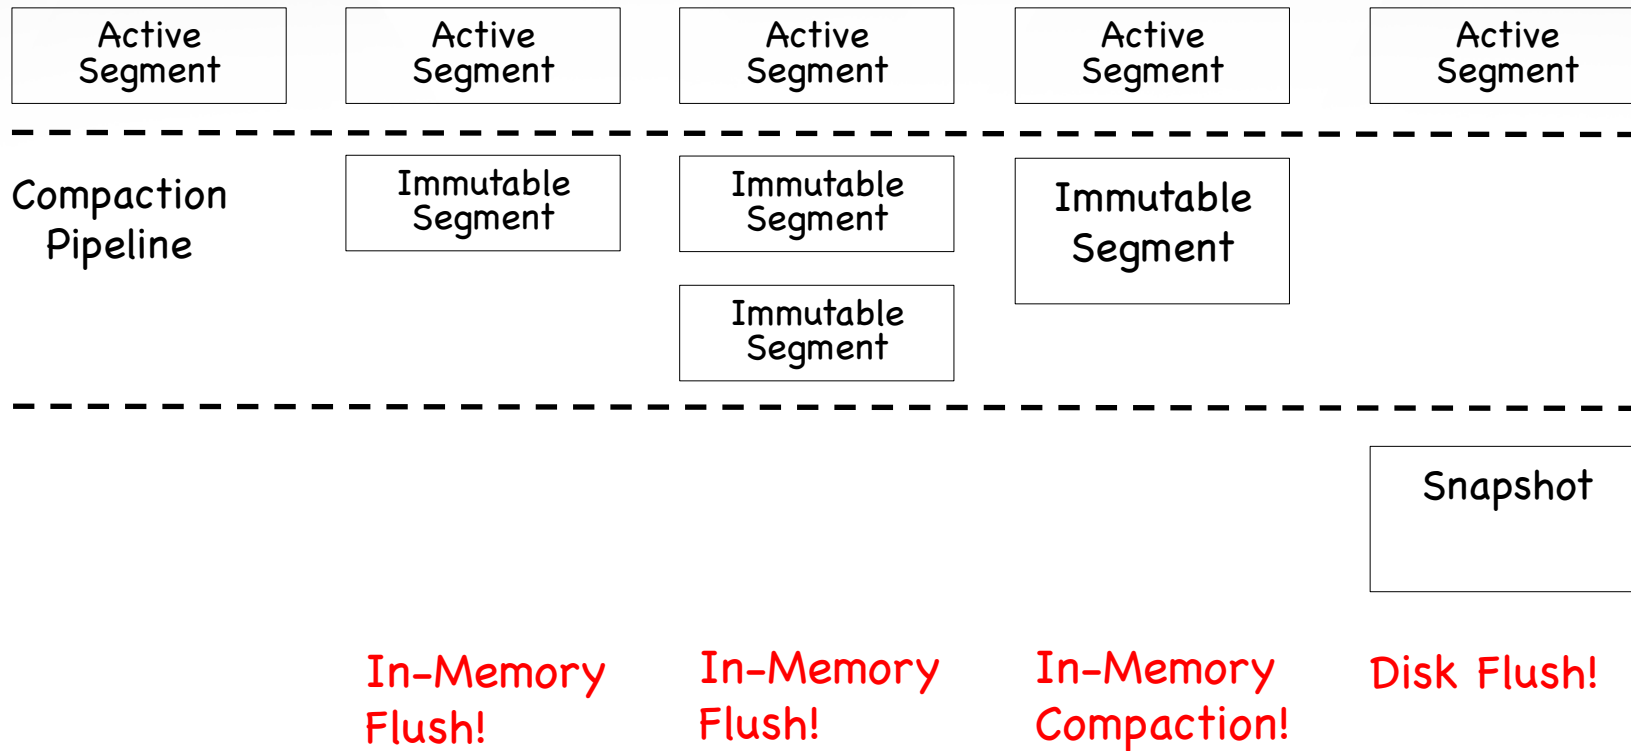
LSM Trees in Action



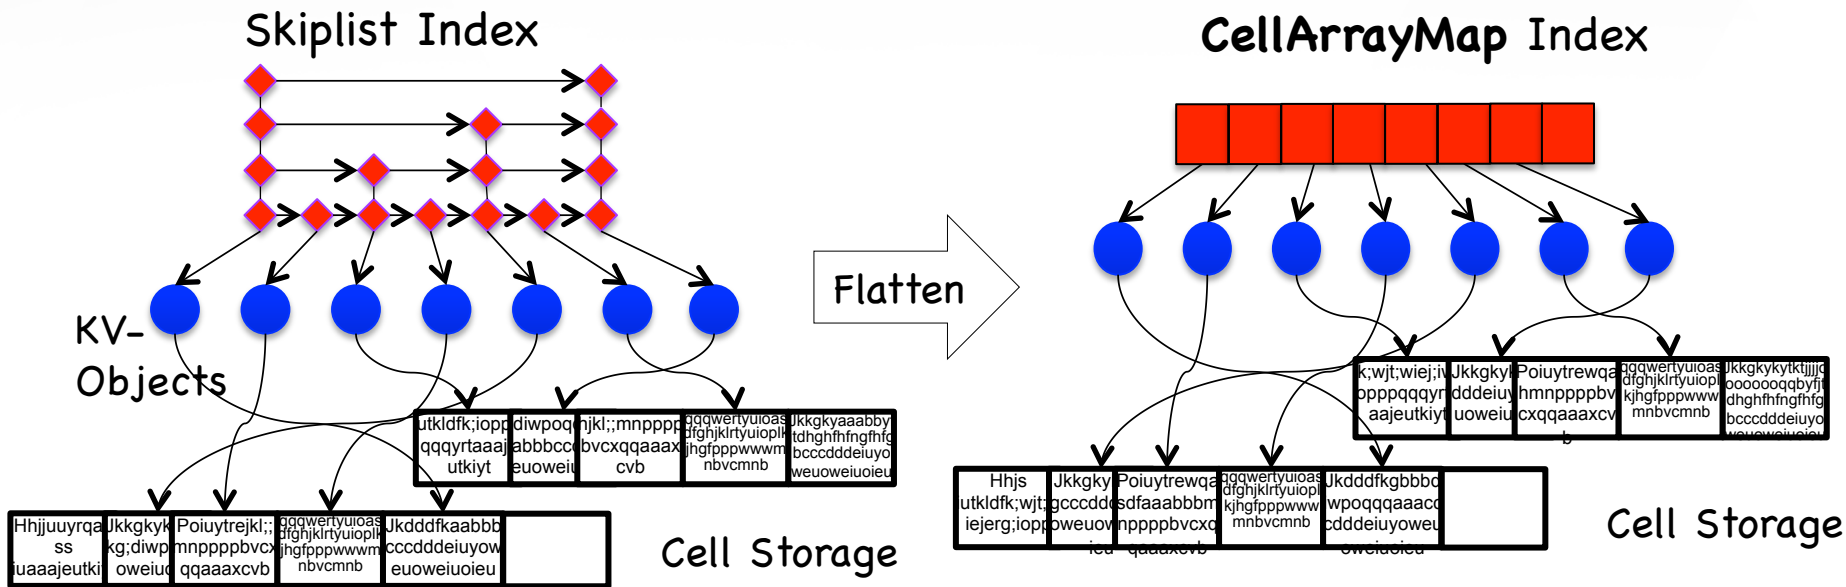
Accordion: In-Memory LSM Tree



Accordion in Action



Flat Immutable Segment Index



Lean footprint – the smaller the cells the better!

Redundancy Elimination

In-Memory Compaction **merges** the pipelined segments

Get access latency under control (less segments to scan)

BASIC compaction

Multiple indexes merged into one, cell data remains in place

EAGER compaction

Redundant data versions eliminated (SQM scan)

BASIC vs EAGER

BASIC: universal optimization, avoids physical data copy

EAGER: high value for highly redundant workloads

- SQM scan is expensive

- Data relocation cost may be high (think MSLAB!)

Configuration

- BASIC is default, EAGER may be configured

- Future implementation may figure out the right mode automatically

Compaction Pipeline: Correctness & Performance

Shared Data Structure

Read access: Get, Scan, in-memory compaction

Write access: in-memory flush, in-memory compaction, disk flush

Design Choice: Non-Blocking Reads

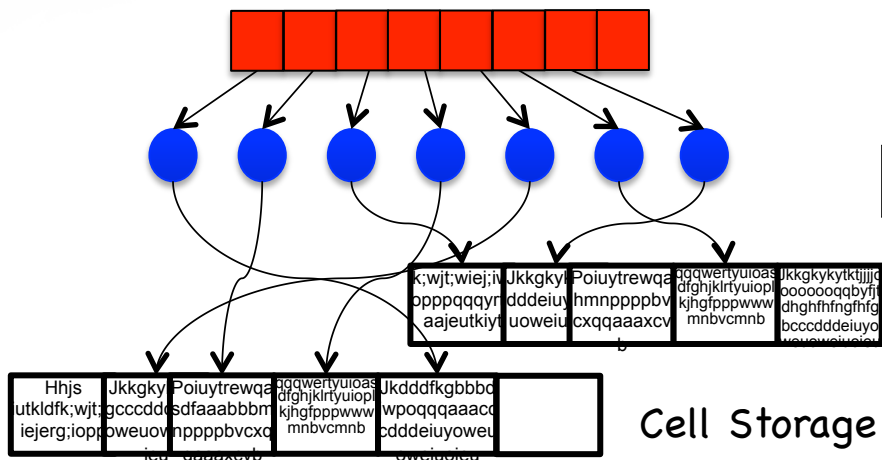
Read-Only pipeline clone – no synchronization upon read access

Copy-on-Write upon modification

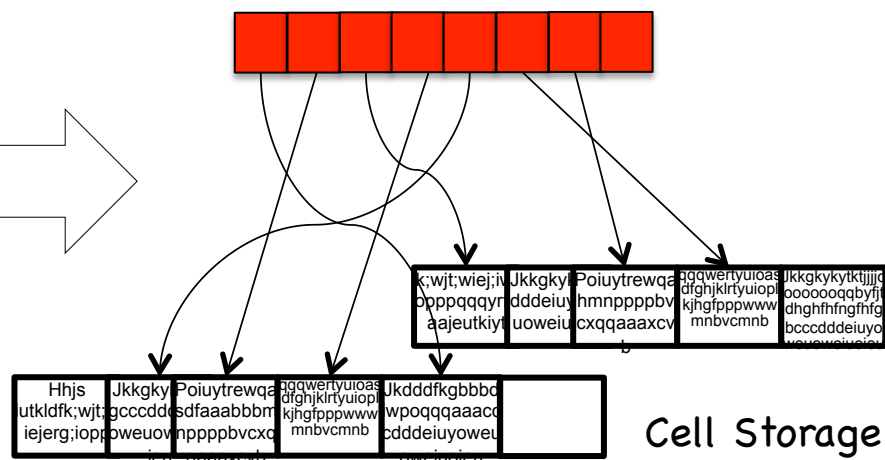
Versioning prevents compaction concurrent to other updates

More Memory Efficiency - KV Object Elimination

CellArrayMap Index



CellChunkMap Index



Lean Footprint (no KV-Objects). Friendly to Off-Heap Implementation.

The Software Side: What's New?

CompactingMemStore: BASIC and EAGER configurations

DefaultMemStore: NONE configuration

Segment Class Hierarchy: Mutable, Immutable, Composite

NavigableMap Implementations: CellArrayMap, CellChunkMap

MemStoreCompactor: compaction algorithms implementation

CellChunkMap Support (Experimental)

Cell objects embedded directly into CellChunkMap (CCM)

New cell type - reference data by unique **ChunkID**

ChunkCreator: Chunk allocation + ChunkID management

Stores mapping of ChunkID's to Chunk references

Strong references to chunks managed by CCM's, **weak** to the rest

The CCM's themselves are allocated via the same mechanism

Some exotic use cases

E.g., jumbo cells allocated in one-time chunks outside chunk pools

Evaluation Setup

System

2-node HBase on top of 3-node HDFS, 1Gbps interconnect

Intel Xeon E5620 (12-core), 2.8TB SSD storage, 48GB RAM

RS config: 16GB RAM (40% Cache/40% MemStore), on-heap, no MSLAB

Data

1 table (100 regions, 50 columns), 30GB-100GB

Workload Driver

YCSB (1 node, 12 threads)

Batched (async) writes (10KB buffer)

Experiments

Metrics

Write throughput, read latency (distribution), disk footprint/amplification

Workloads (varied at client side)

Write-Only (100% Put) vs Mixed (50% Put/50% Get)

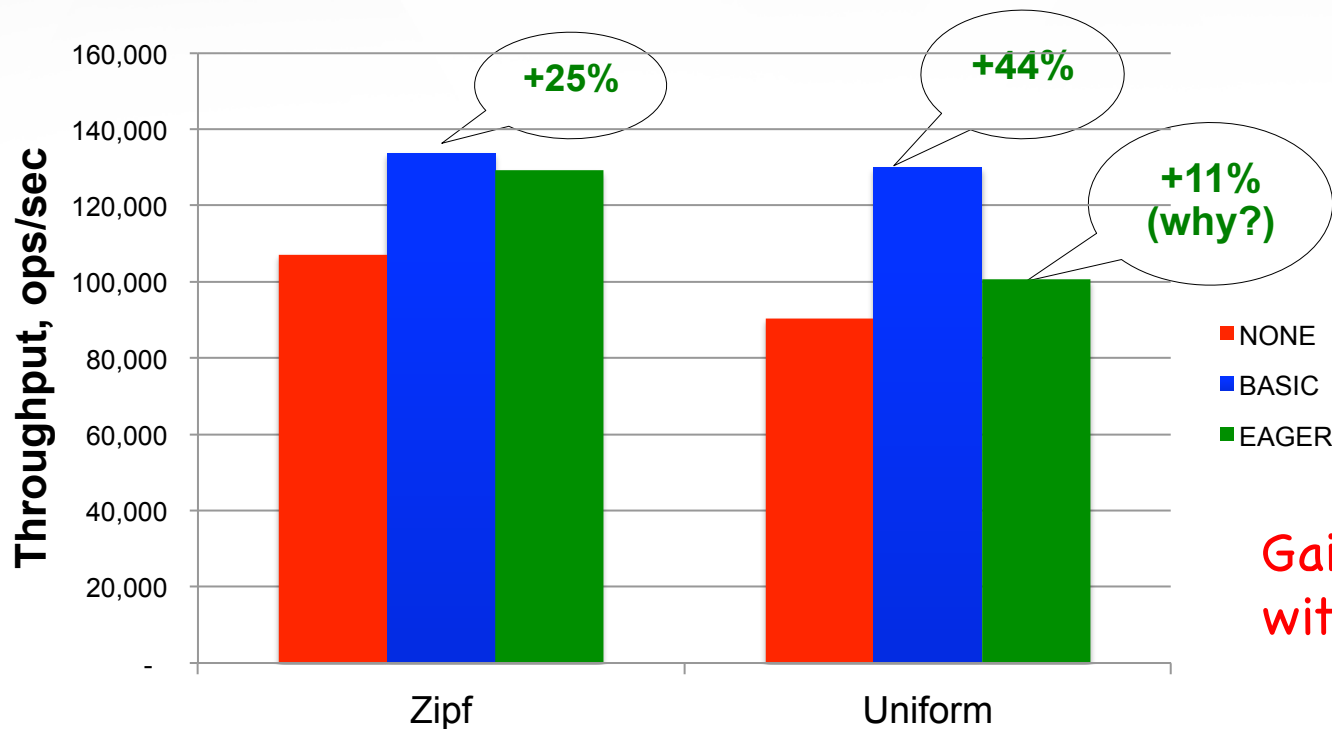
Uniform vs Zipfian Key Distributions

Small Values (100B) vs Big Values (1K)

Configurations (varied at server side)

Most experiments exercise Async WAL

Write Throughput




100GB Dataset
100% Writes
100B Values

Every write updates
a single column


Gains less pronounced
with big values (1KB)

Single-Key Write Latency



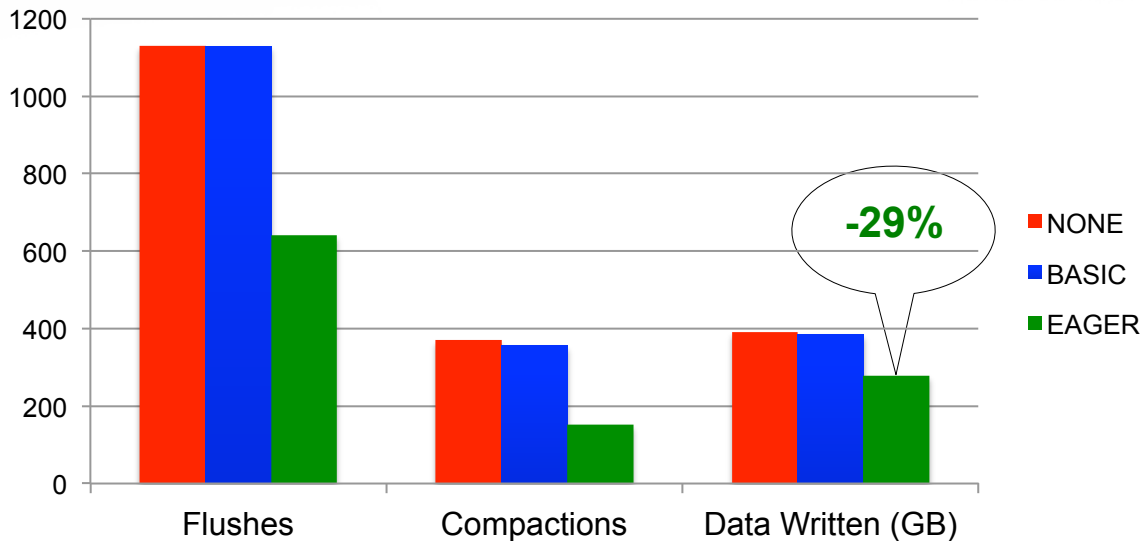
100GB Dataset
Zipf distribution
100% Writes
100B Values

Single-Key Read Latency



30GB Dataset
Zipf Distribution
50% Writes/50% Reads
100B Values

Disk Footprint/Write Amplification



100GB Dataset
Zipf Distribution
100% Writes
100B Values

Status

In-Memory Compaction GA in HBase 2.0

Master JIRA [HBASE-14918](#) complete (~20 subtasks)

Major refactoring/extension of the MemStore code

Many details in [Apache HBase blog](#) posts

CellChunkMap Index, Off-Heap support in progress

Master JIRA [HBASE-16421](#)

Summary

Accordion = a leaner and faster write path

Space-Efficient Index + Redundancy Elimination → **less I/O**

Less Frequent Flushes → **increased write throughput**

Less On-Disk Compaction → **reduced write amplification**

Data stays longer in RAM → **reduced tail read latency**

Edging Closer to In-Memory Database Performance

Thanks to Our Partners for Being Awesome