



# Transactions in HBase

Andreas Neumann  
Gokul Gunasekaran  
HbaseCon June 2017

gokul at cask.co  
anew at apache.org  
@caskoid

# Goals of this Talk

- Why transactions?
- Optimistic Concurrency Control
- Three Apache projects: Omid, Tephra, Trafodion
- How are they different?



# Transactions in noSQL?

## History


- SQL: RDBMS, EDW, ...
- noSQL: MapReduce, HDFS, HBase, ...
- n(ot)o(nly)SQL: Hive, Phoenix, ...

## Motivation:

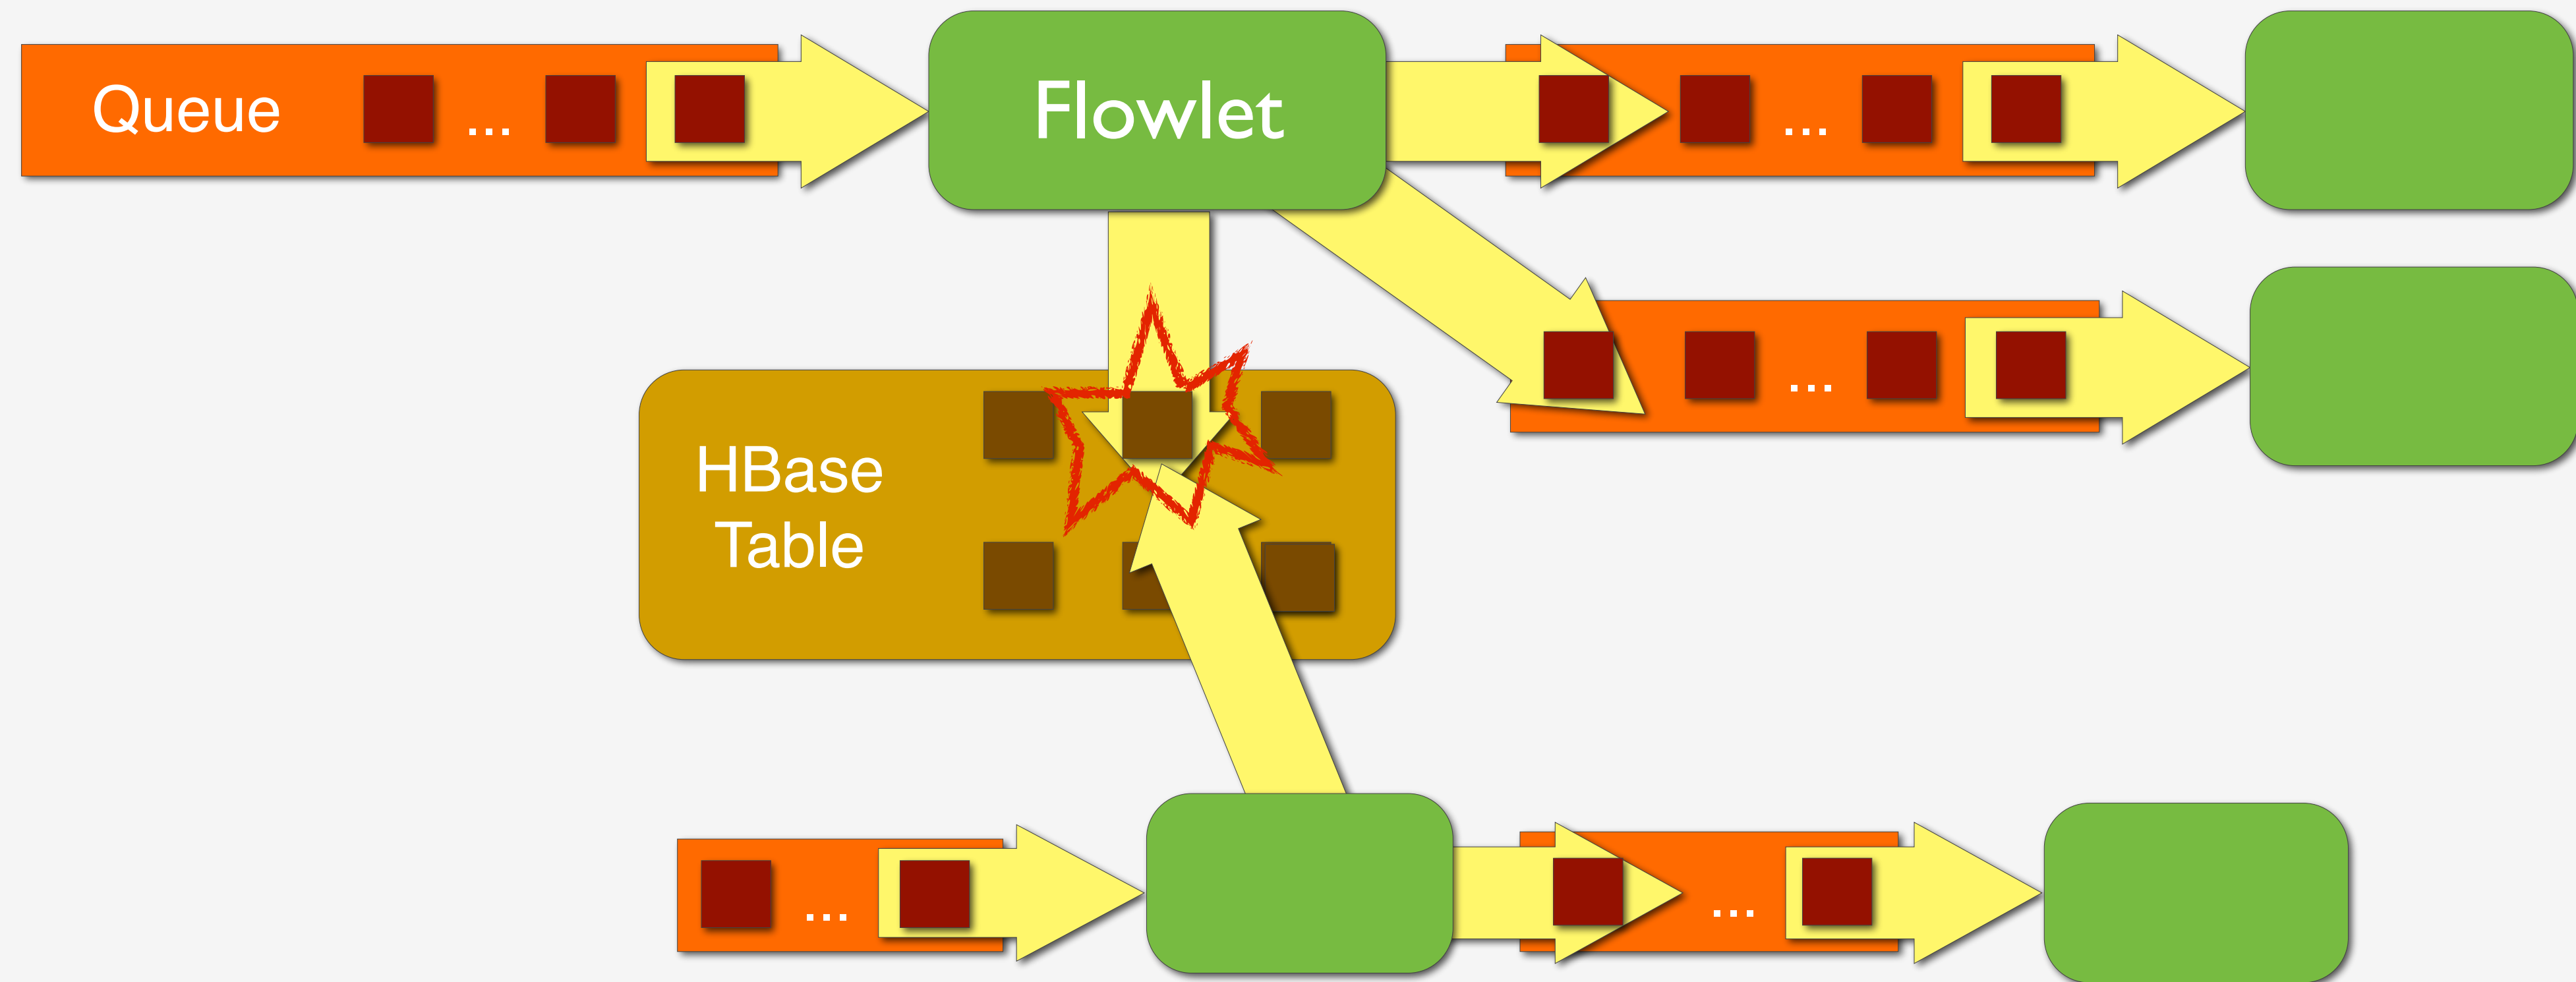
- Data consistency under highly concurrent loads
- Partial outputs after failure
- Consistent view of data for long-running jobs
- (Near) real-time processing



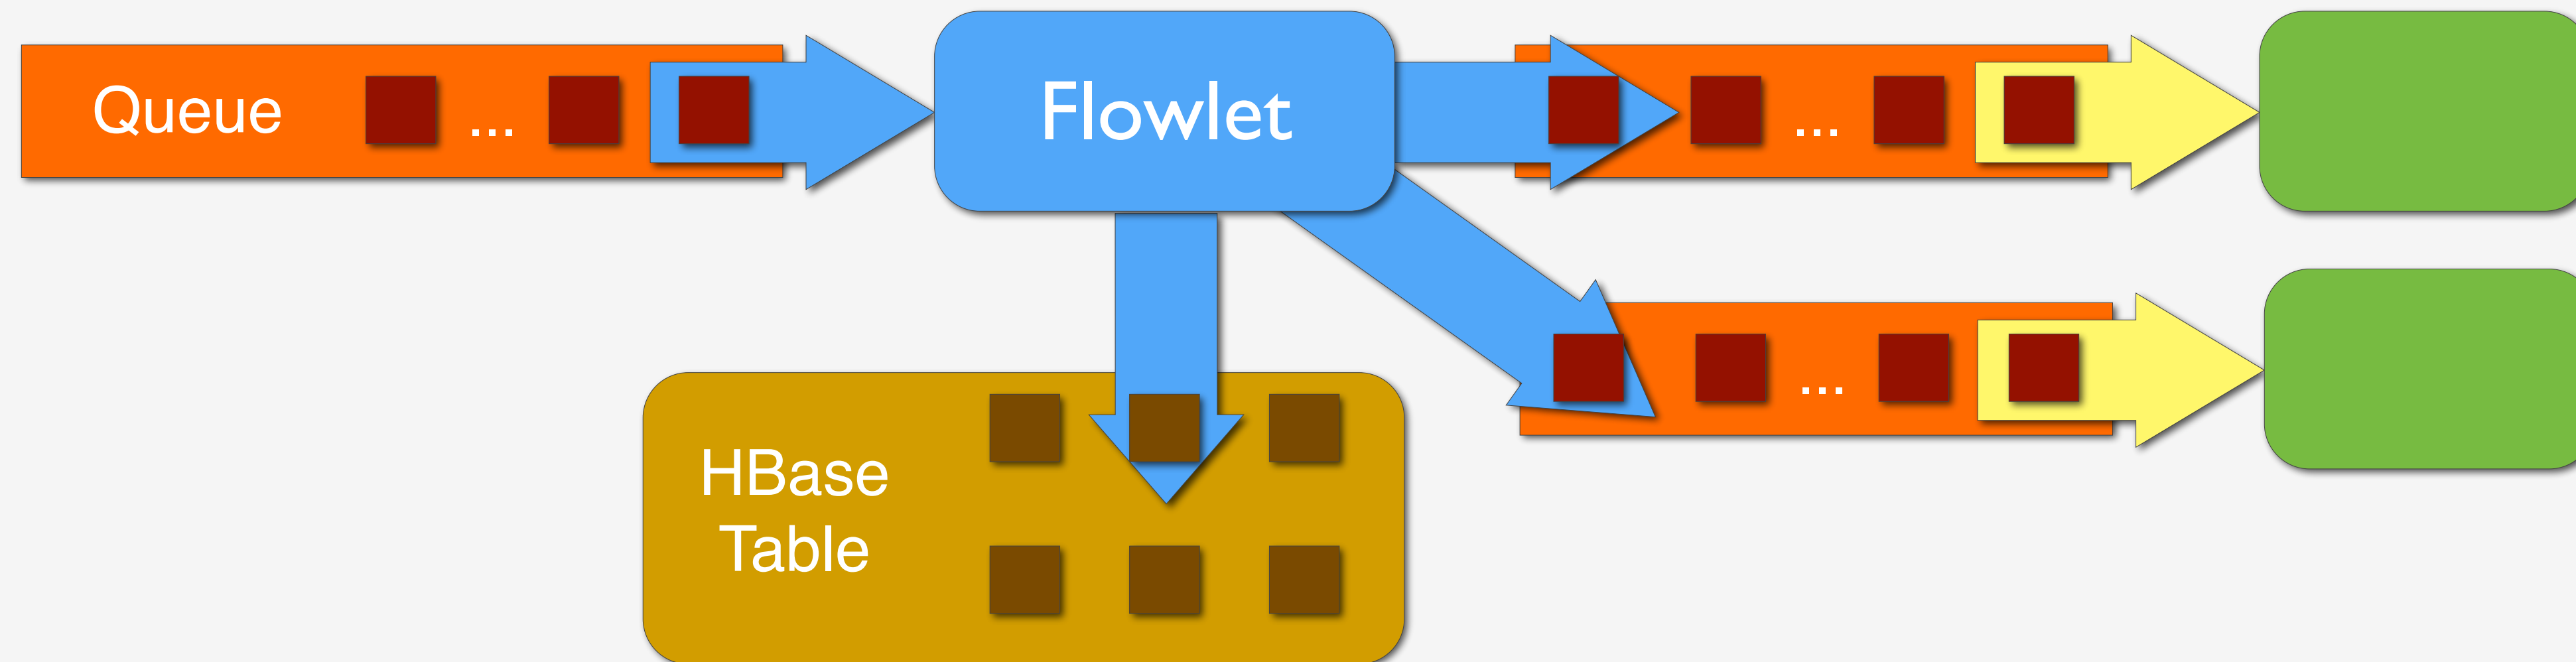
# Stream Processing



# Write Conflict!



# Transactions to the Rescue



- Atomicity of all writes involved
- Protection from concurrent update



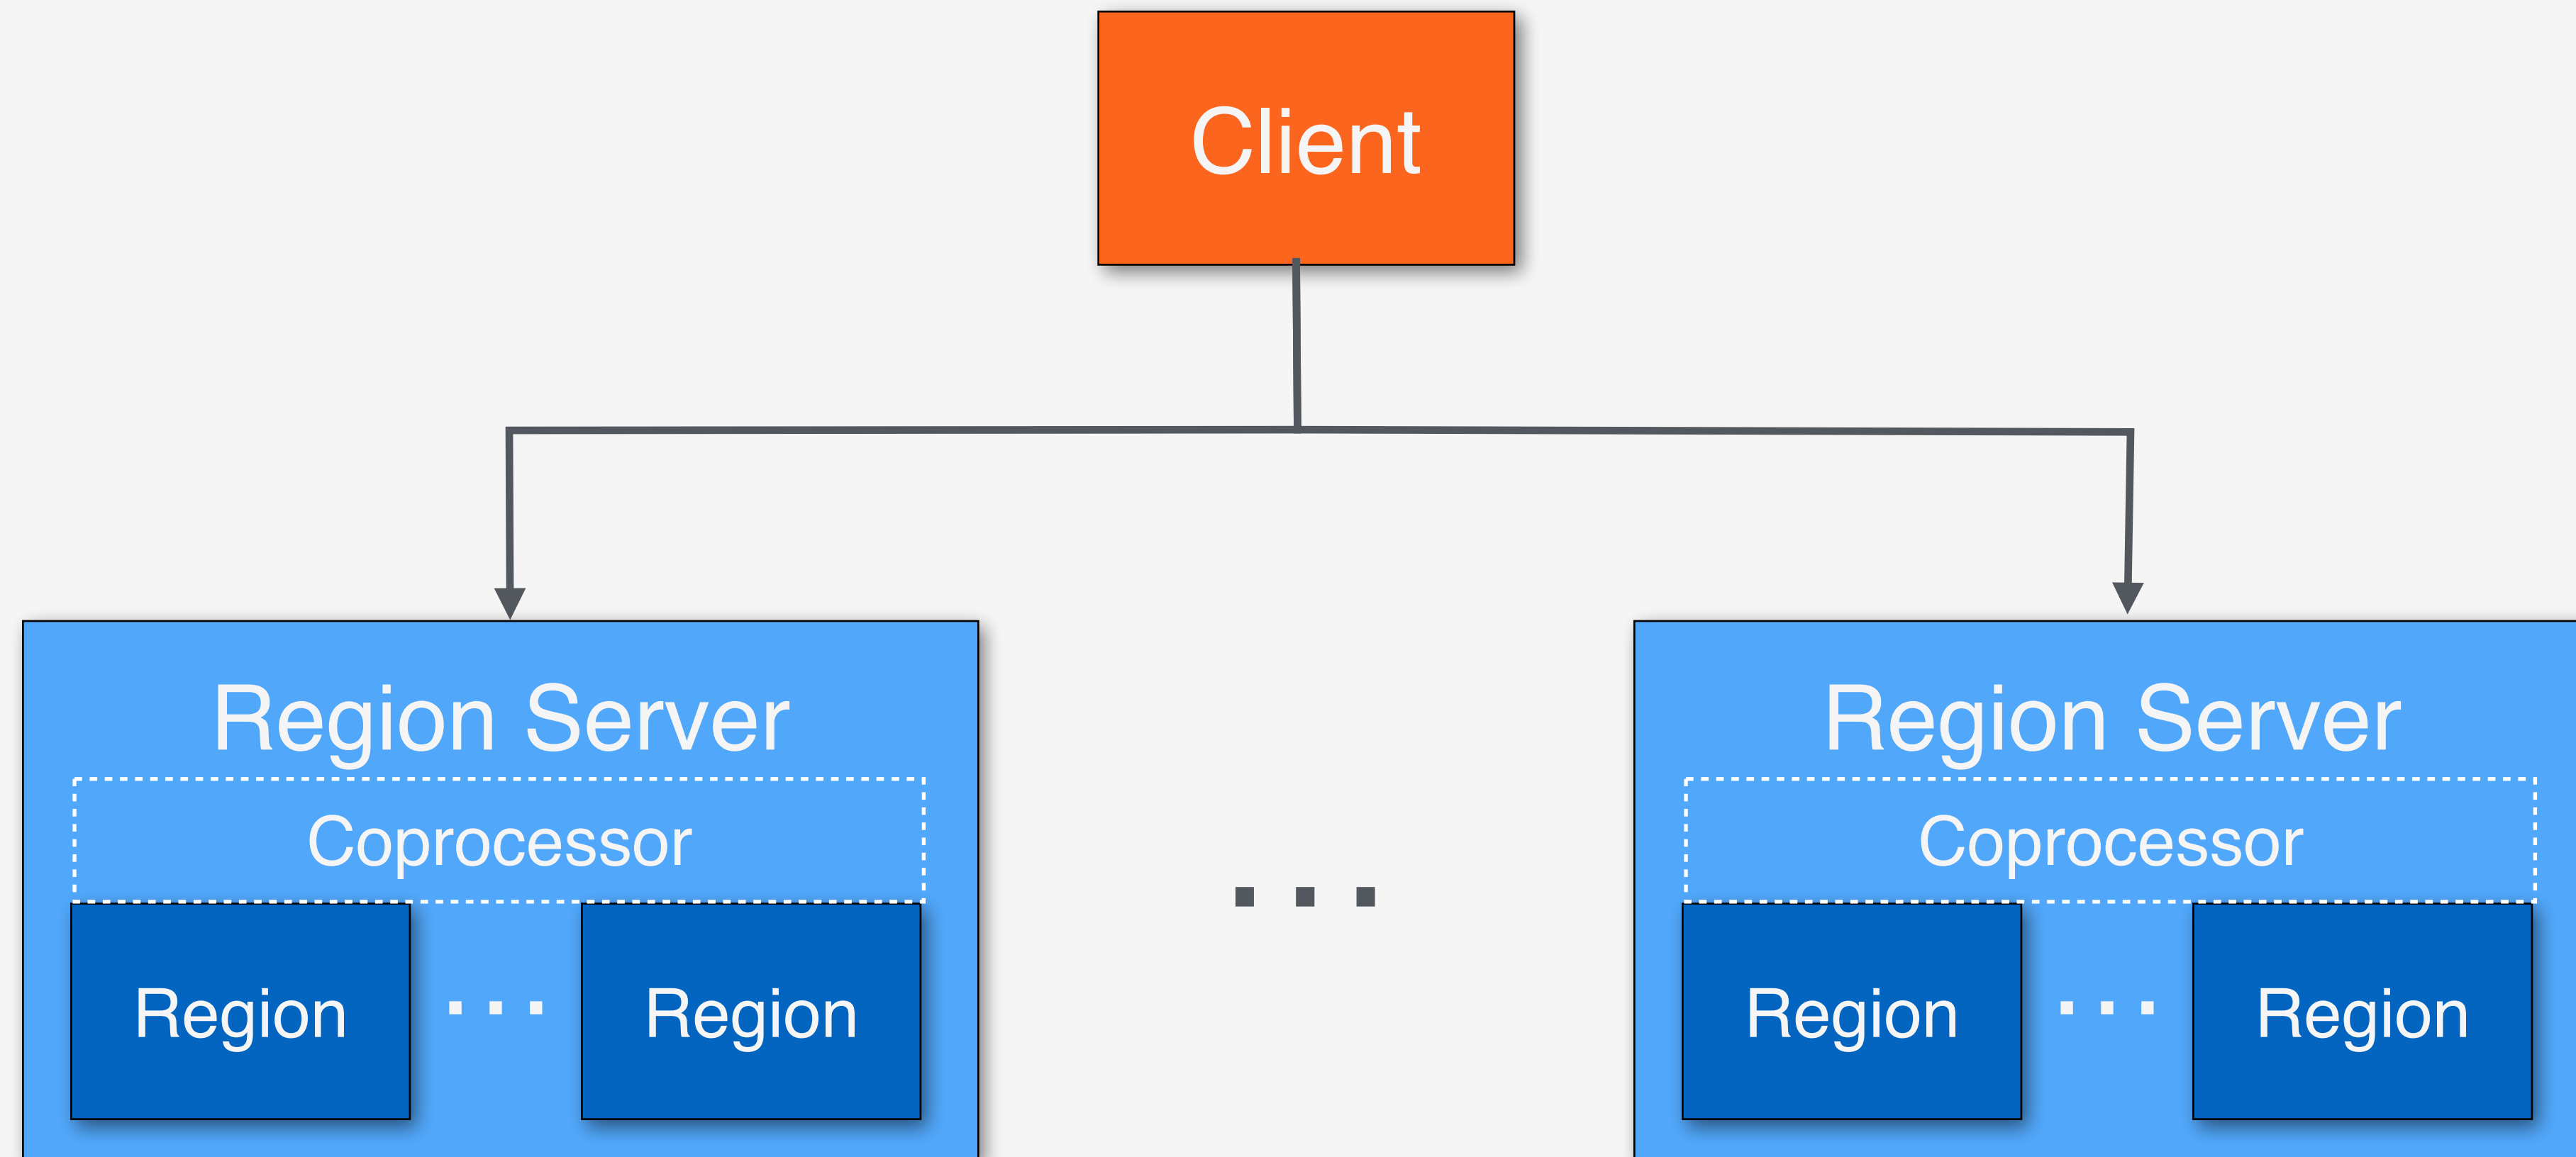
# ACID Properties

From good old SQL:

- **Atomic** - Entire transaction is committed as one
- **Consistent** - No partial state change due to failure
- **Isolated** - No dirty reads, transaction is only visible after commit
- **Durable** - Once committed, data is persisted reliably



# What is HBase?





# What is HBase?

Simplified:

- Distributed Key-Value Store
- Key = <row>.<family>.<column>.<timestamp>
- Partitioned into Regions (= continuous range of rows)
- Each Region Server hosts multiple regions
- Optional: Coprocessor in Region Server
- Durable writes



# ACID Properties in HBase

- **Atomic**
  - At cell, row, and region level
  - Not across regions, tables or multiple calls
- **Consistent** - No built-in rollback mechanism
- **Isolated** - Timestamp filters provide some level of isolation
- **Durable** - Once committed, data is persisted reliably

How to implement full ACID?

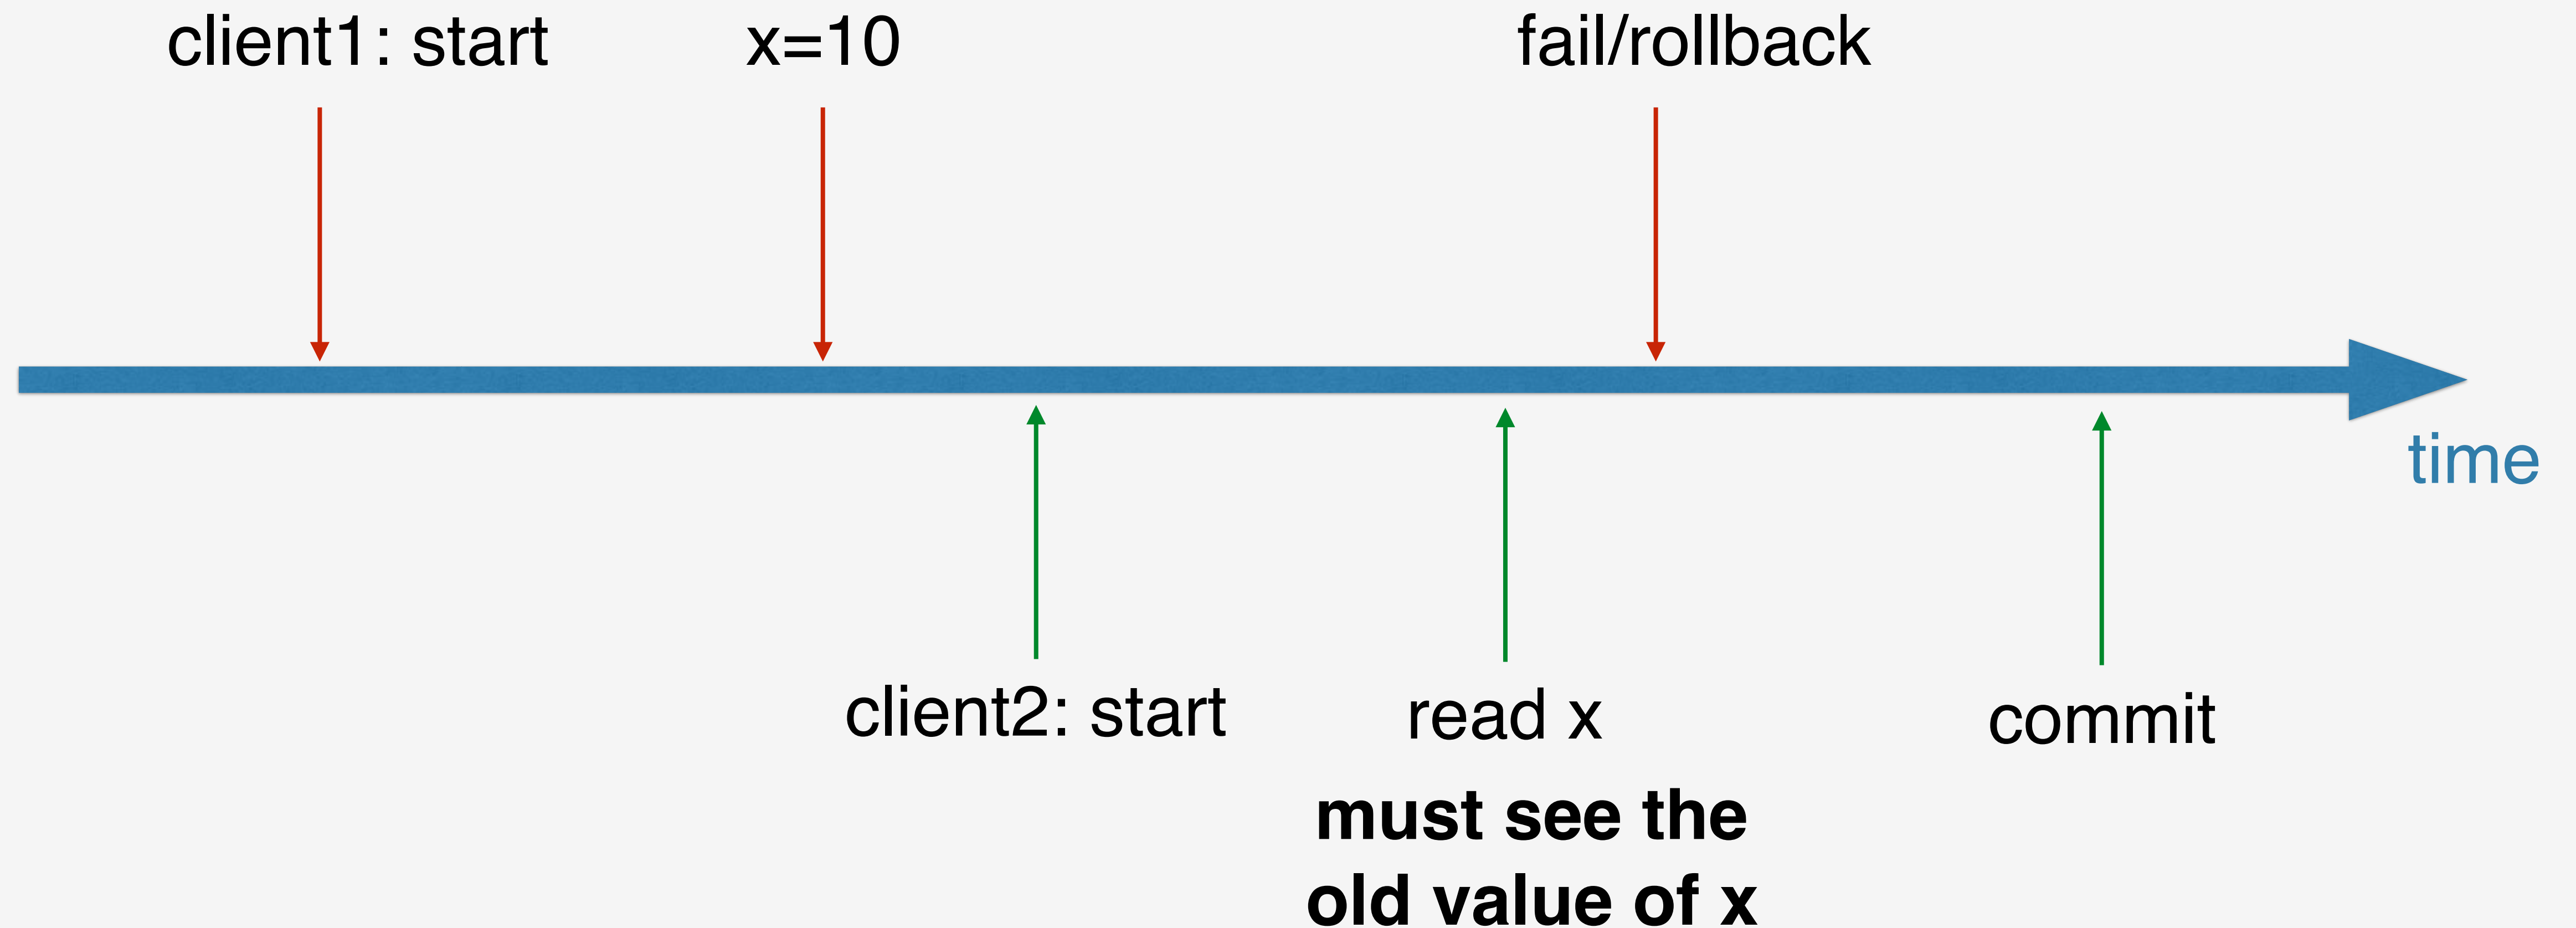


# Implementing Transactions


- Traditional approach (RDBMS): locking
  - May produce deadlocks
  - Causes idle wait
  - complex and expensive in a distributed env
- Optimistic Concurrency Control
  - lockless: allow concurrent writes to go forward
  - on commit, detect conflicts with other transactions
  - on conflict, roll back all changes and retry
- Snapshot Isolation
  - Similar to repeatable read
  - Take snapshot of all data at transaction start
  - Read isolation




# Optimistic Concurrency Control



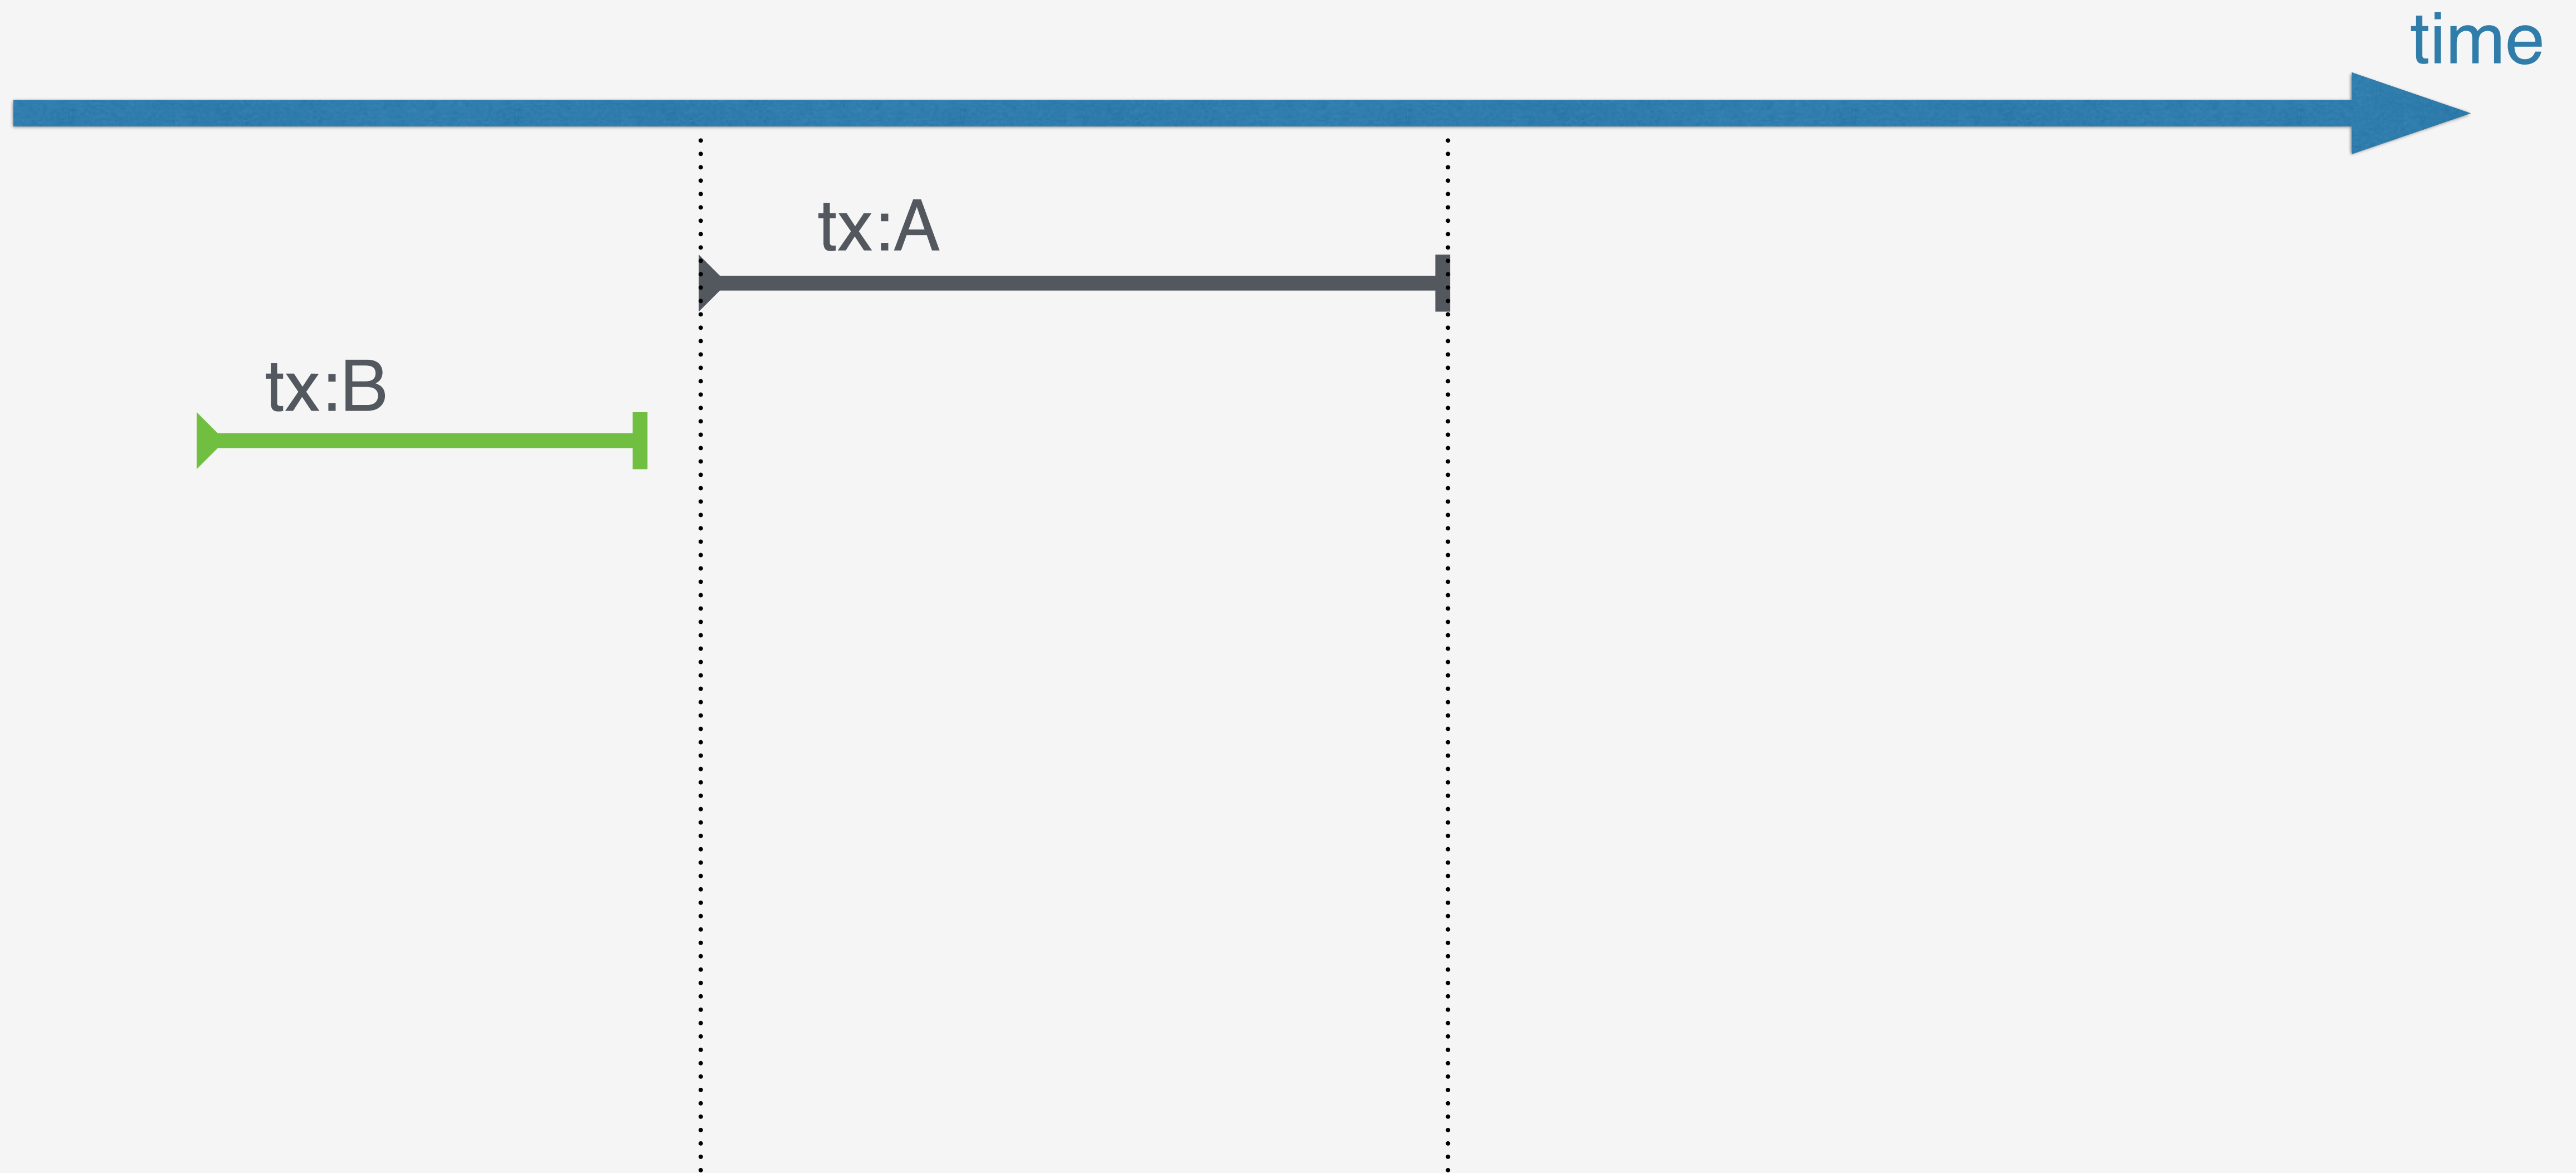
# Optimistic Concurrency Control



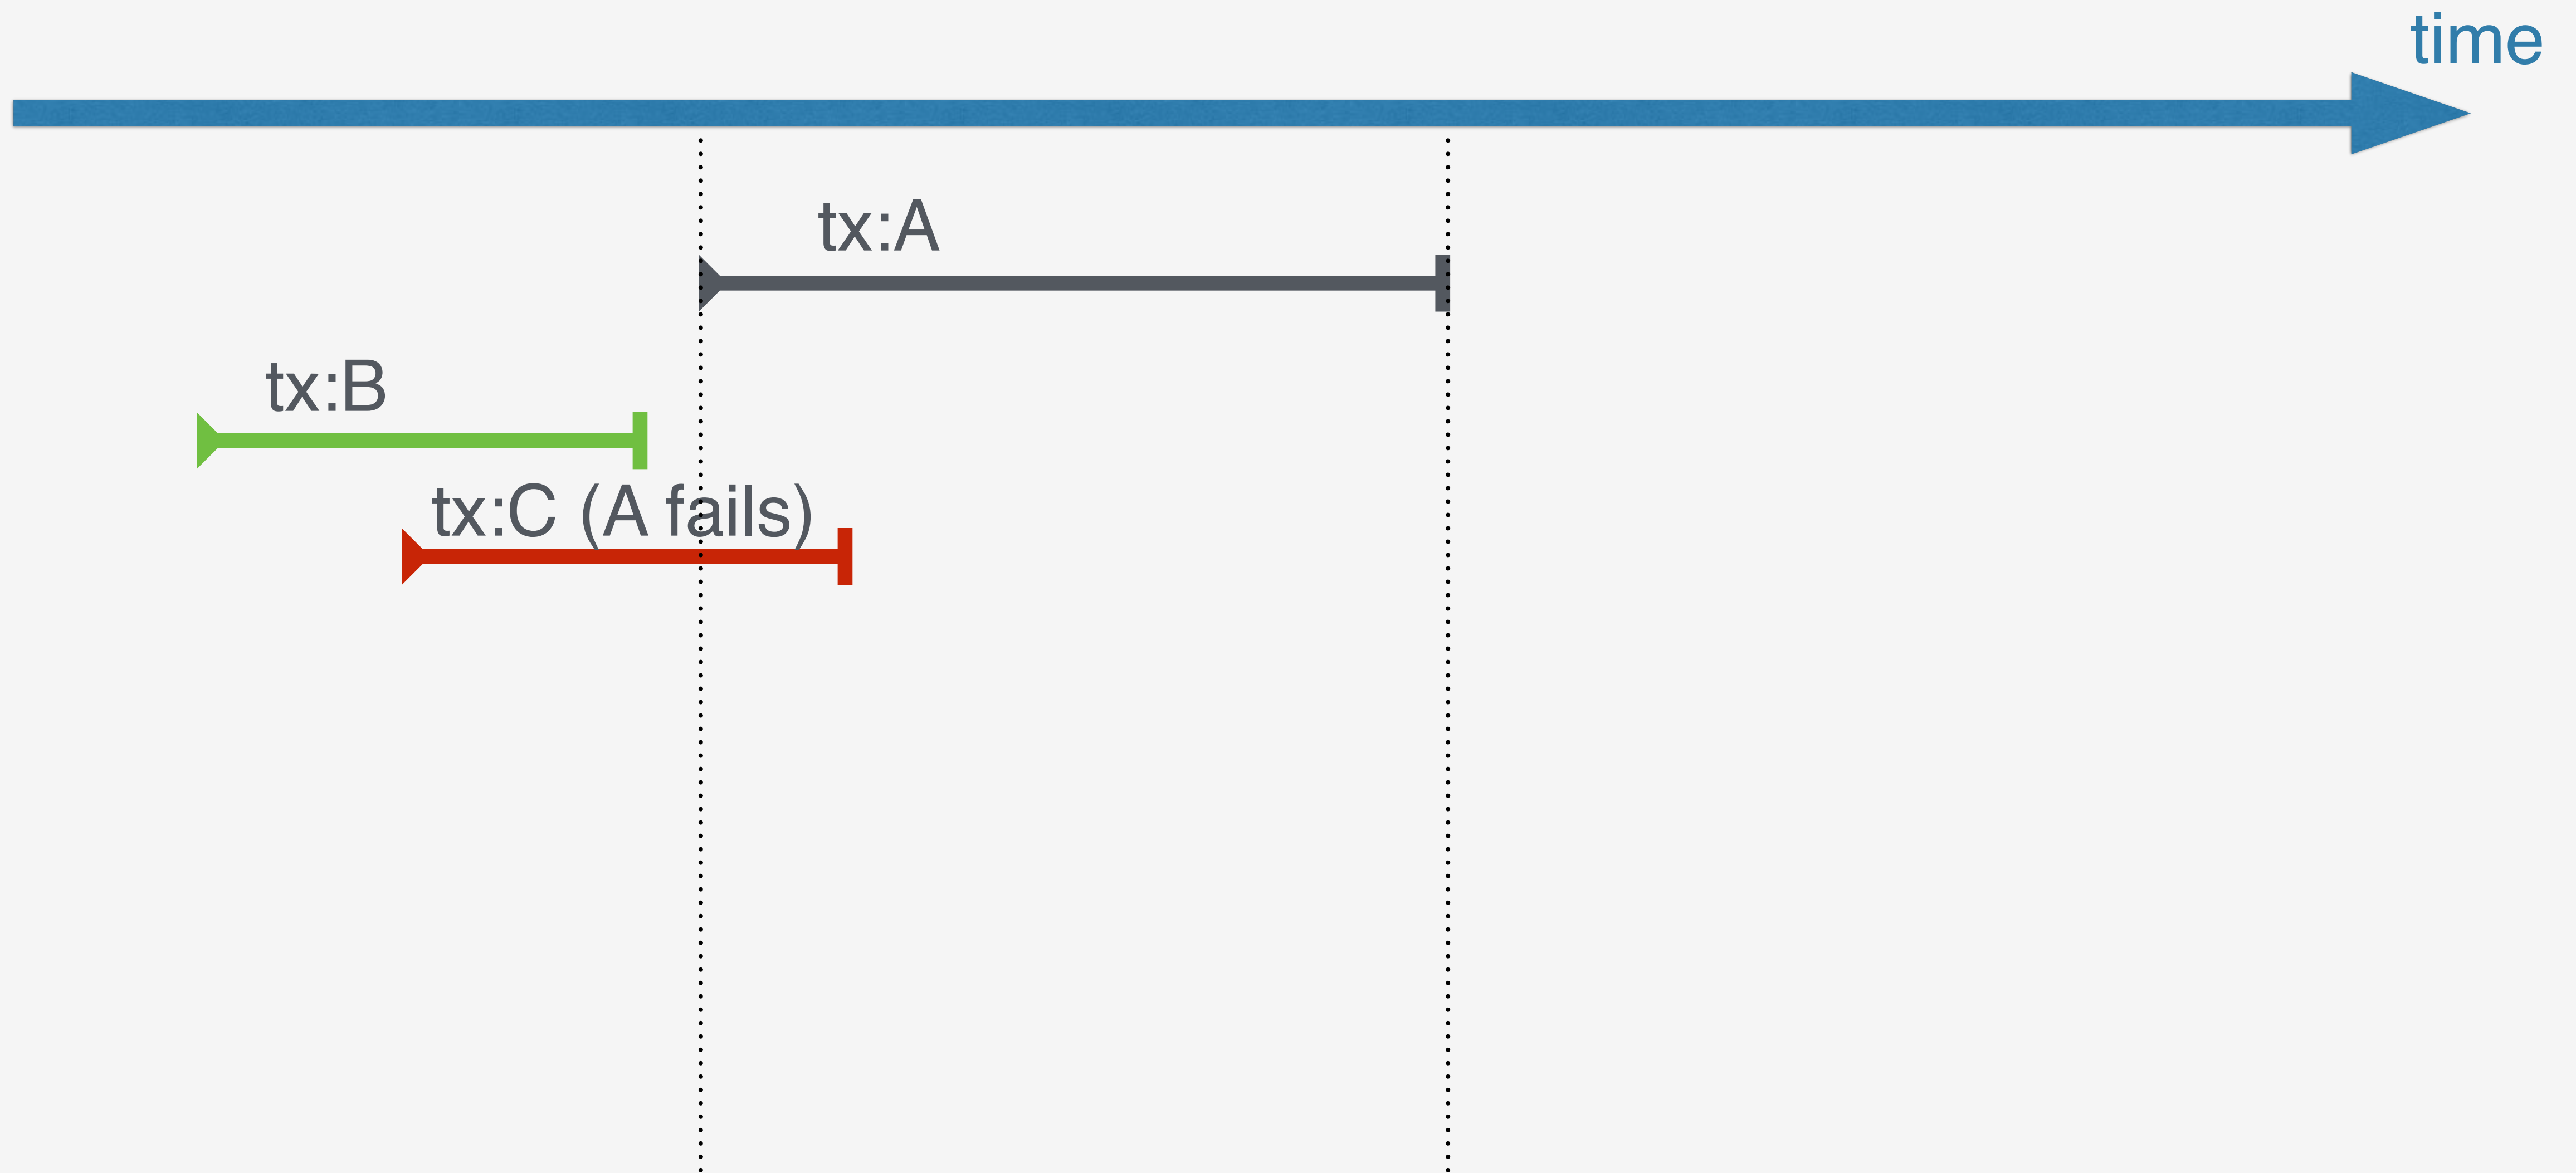
# Conflicting Transactions



# Conflicting Transactions

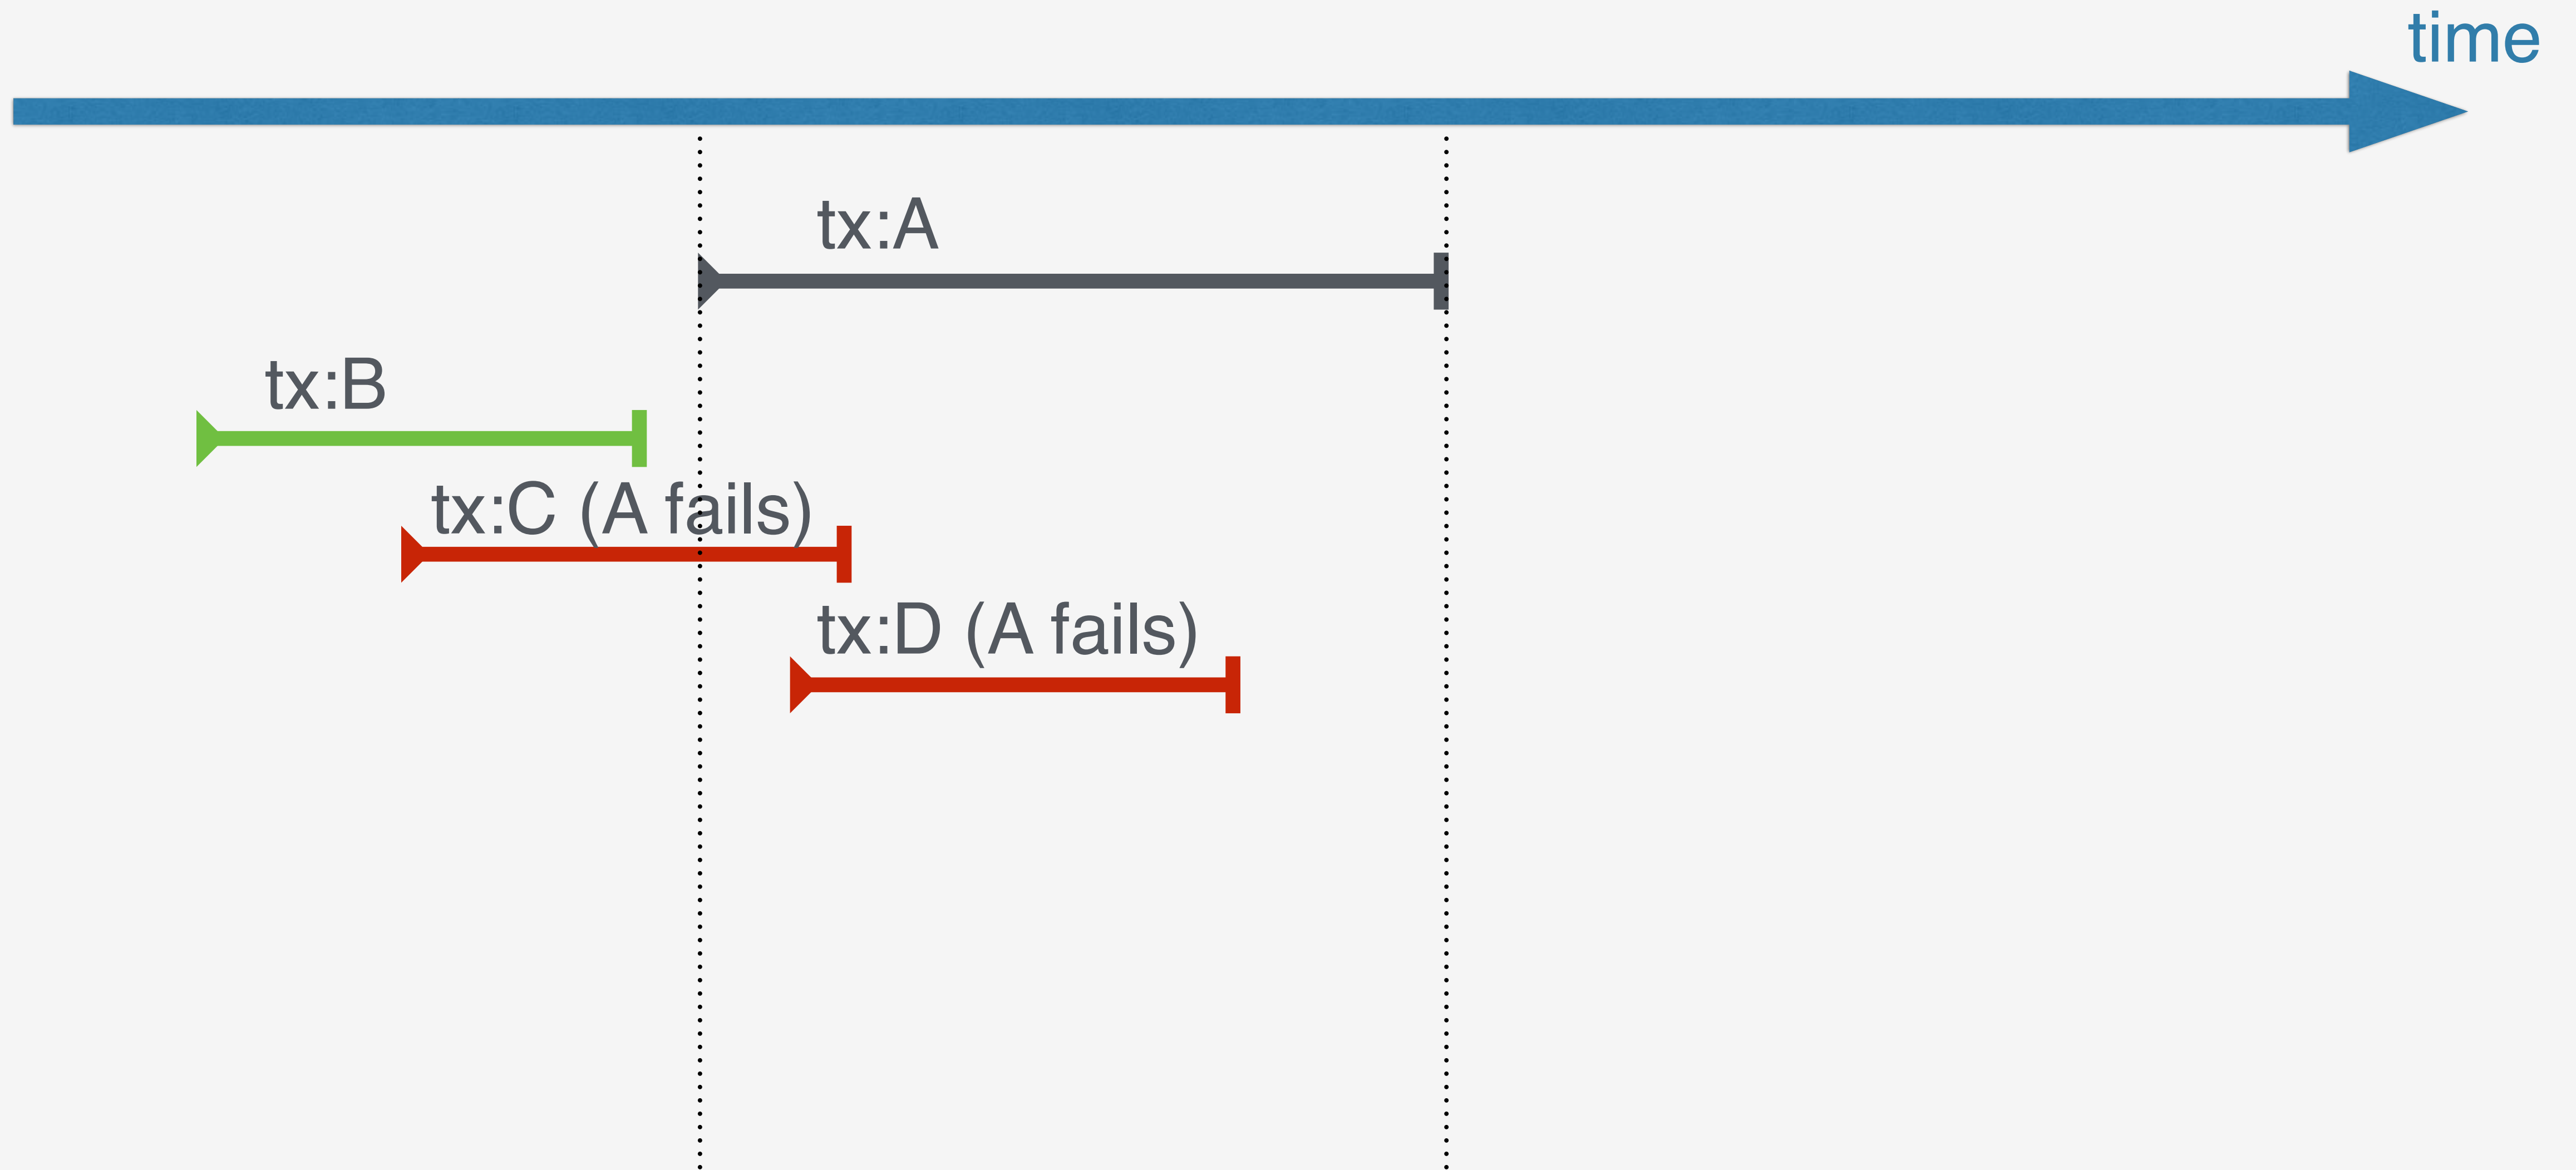


# Conflicting Transactions






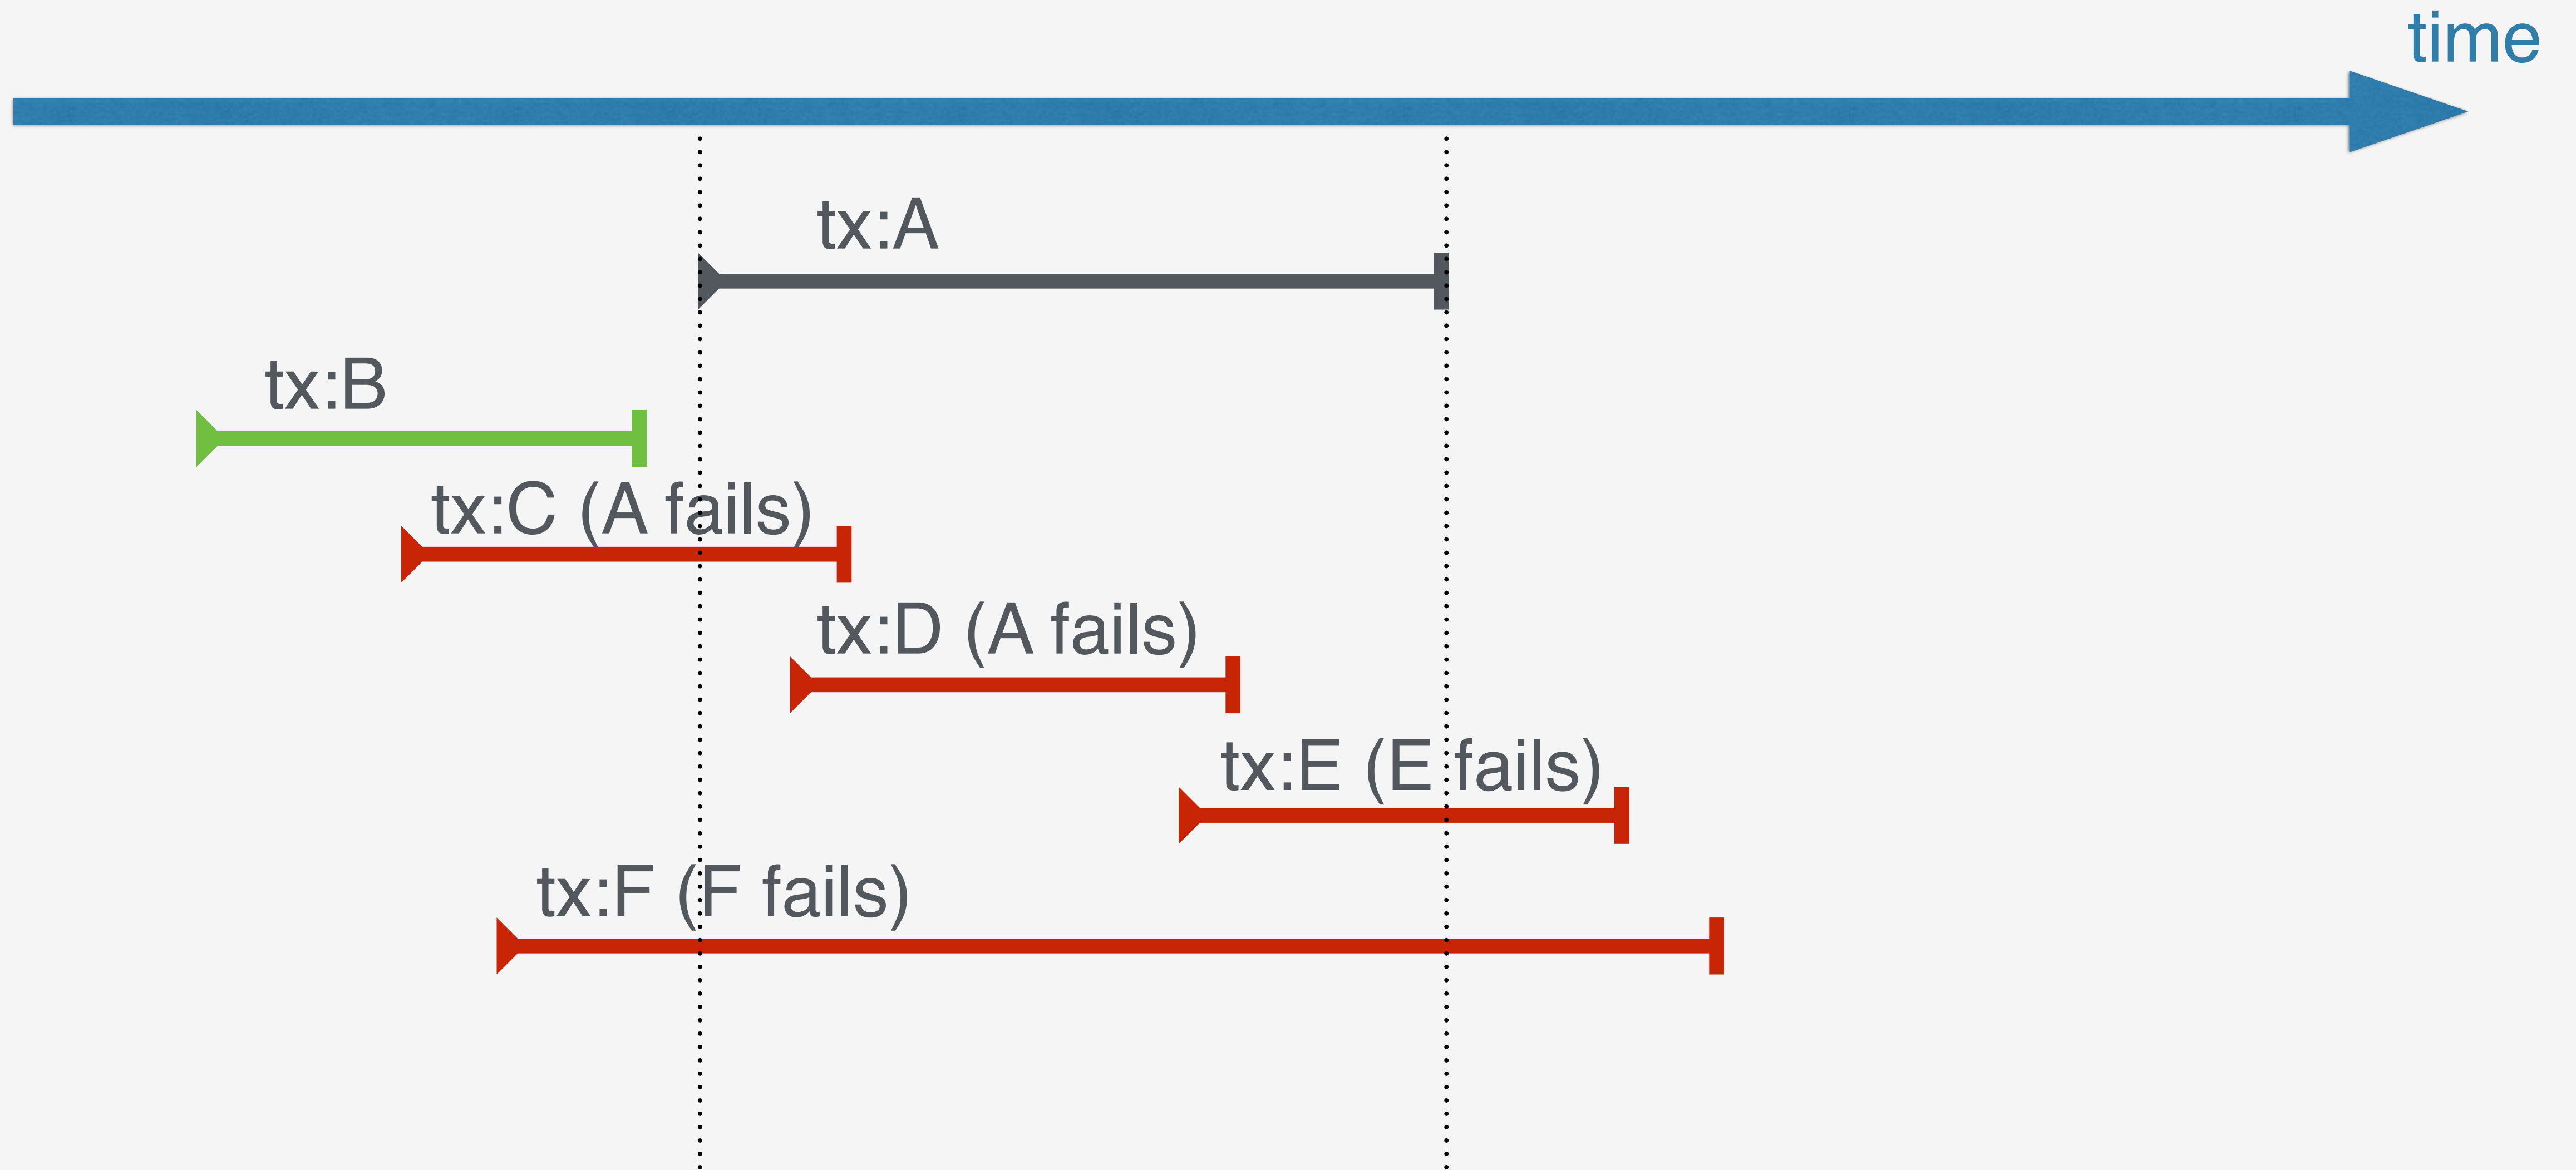
# Conflicting Transactions




# Conflicting Transactions



# Conflicting Transactions



# Conflicting Transactions



# Conflicting Transactions

- Two transactions have a conflict if
  - they write to the same cell
  - they overlap in time
- If two transactions conflict, the one that commits later rolls back
- Active change set = set of transactions  $t$  such that:
  - $t$  is committed, and
  - there is at least one in-flight tx  $t'$  that started before  $t$ 's commit time
- This change set is needed in order to perform conflict detection.



# HBase Transactions in Apache



(incubating)



Apache Omid (incubating)



(incubating)



# In Common

- Optimistic Concurrency Control must:
  - maintain Transaction State:
    - what tx are in flight and committed?
    - what is the change set of each tx? (for conflict detection, rollback)
    - what transactions are invalid (failed to roll back due to crash etc.)
  - generate unique transaction IDs
  - coordinate the life cycle of a transaction
    - start, detect conflicts, commit, rollback
- All of { Omid, Tephra, Trafodion } implement this
  - but vary in how they do it



# Apache Tephra

- Based on the original Omid paper:

Daniel Gómez Ferro, Flavio Junqueira, Ivan Kelly, Benjamin Reed, Maysam Yabandeh:  
*Omid: Lock-free transactional support for distributed data stores*. ICDE 2014.

- Transaction Manager:
  - Issues unique, monotonic transaction IDs
  - Maintains the set of excluded (in-flight and invalid) transactions
  - Maintains change sets for active transactions
  - Performs conflict detection
- Client:
  - Uses transaction ID as timestamp for writes
  - Filters excluded transactions for isolation
  - Performs rollback





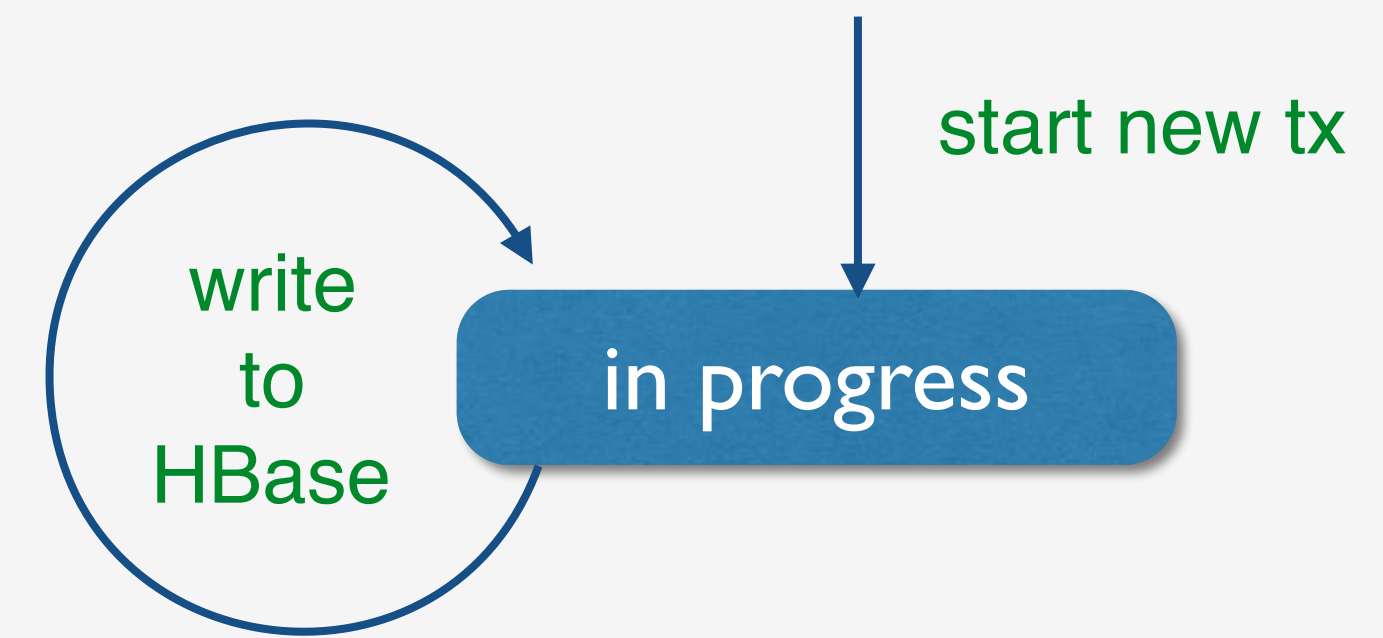
# Transaction Lifecycle




# Transaction Lifecycle



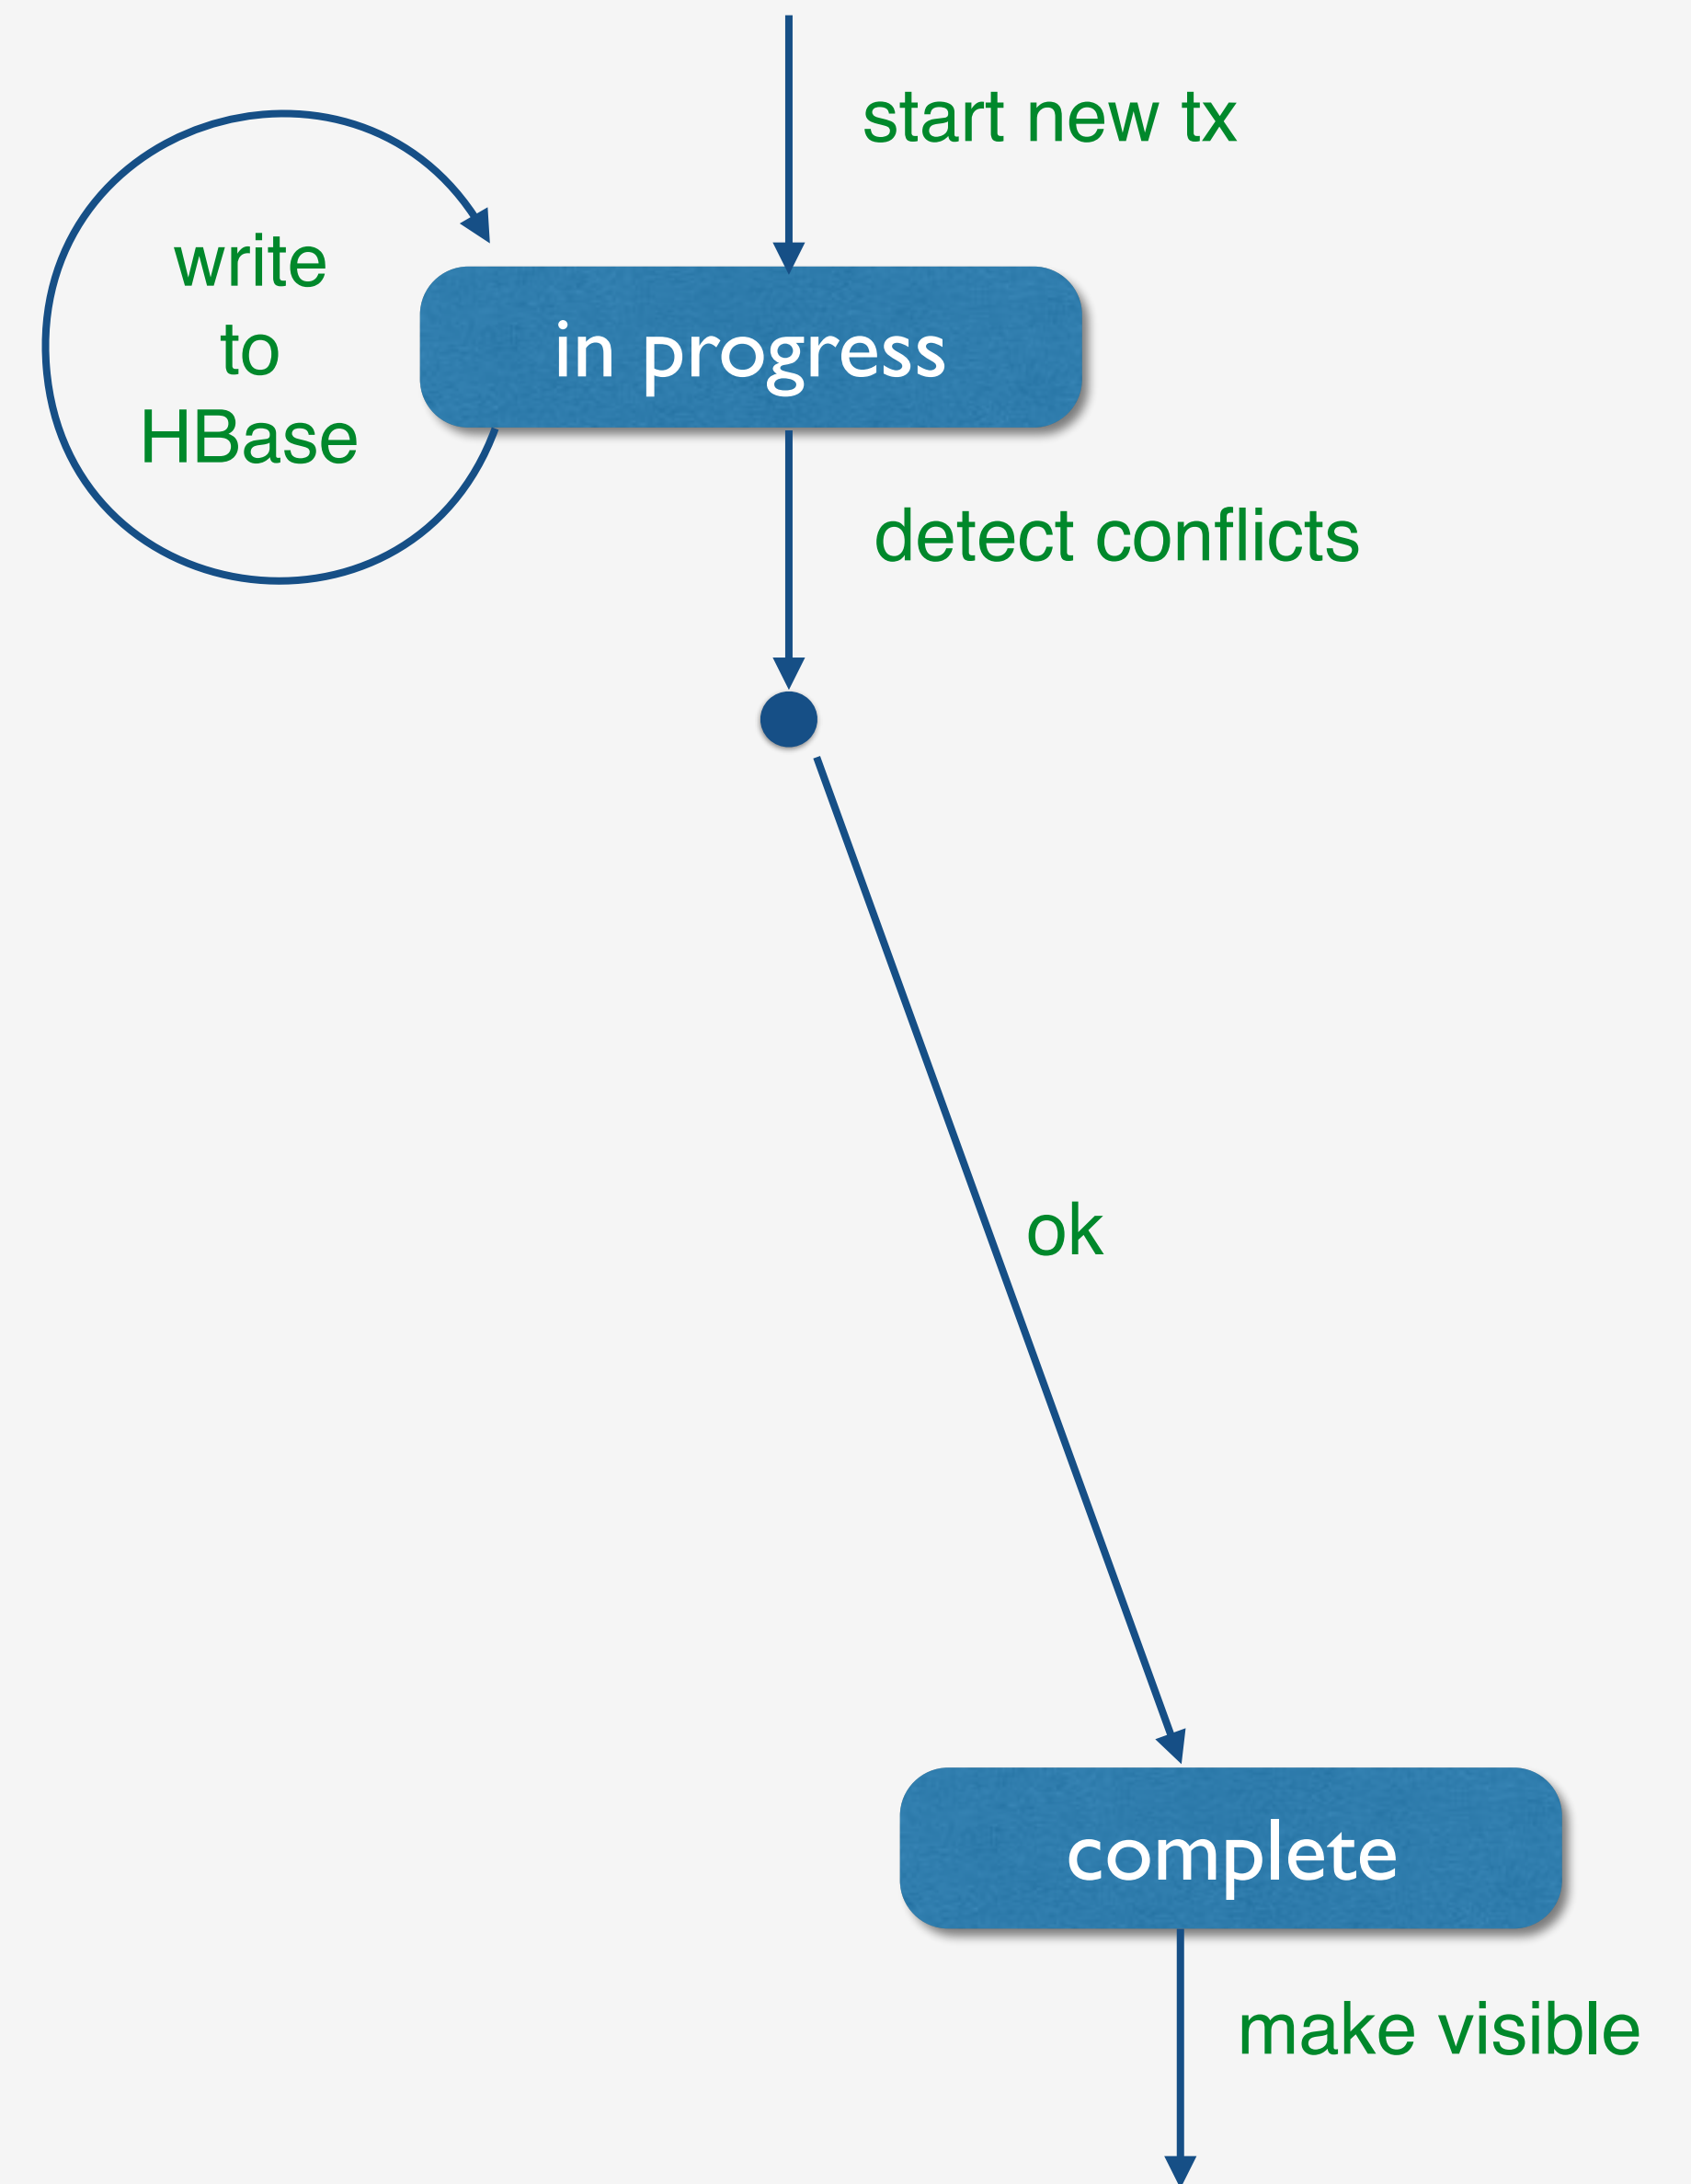
# Transaction Lifecycle



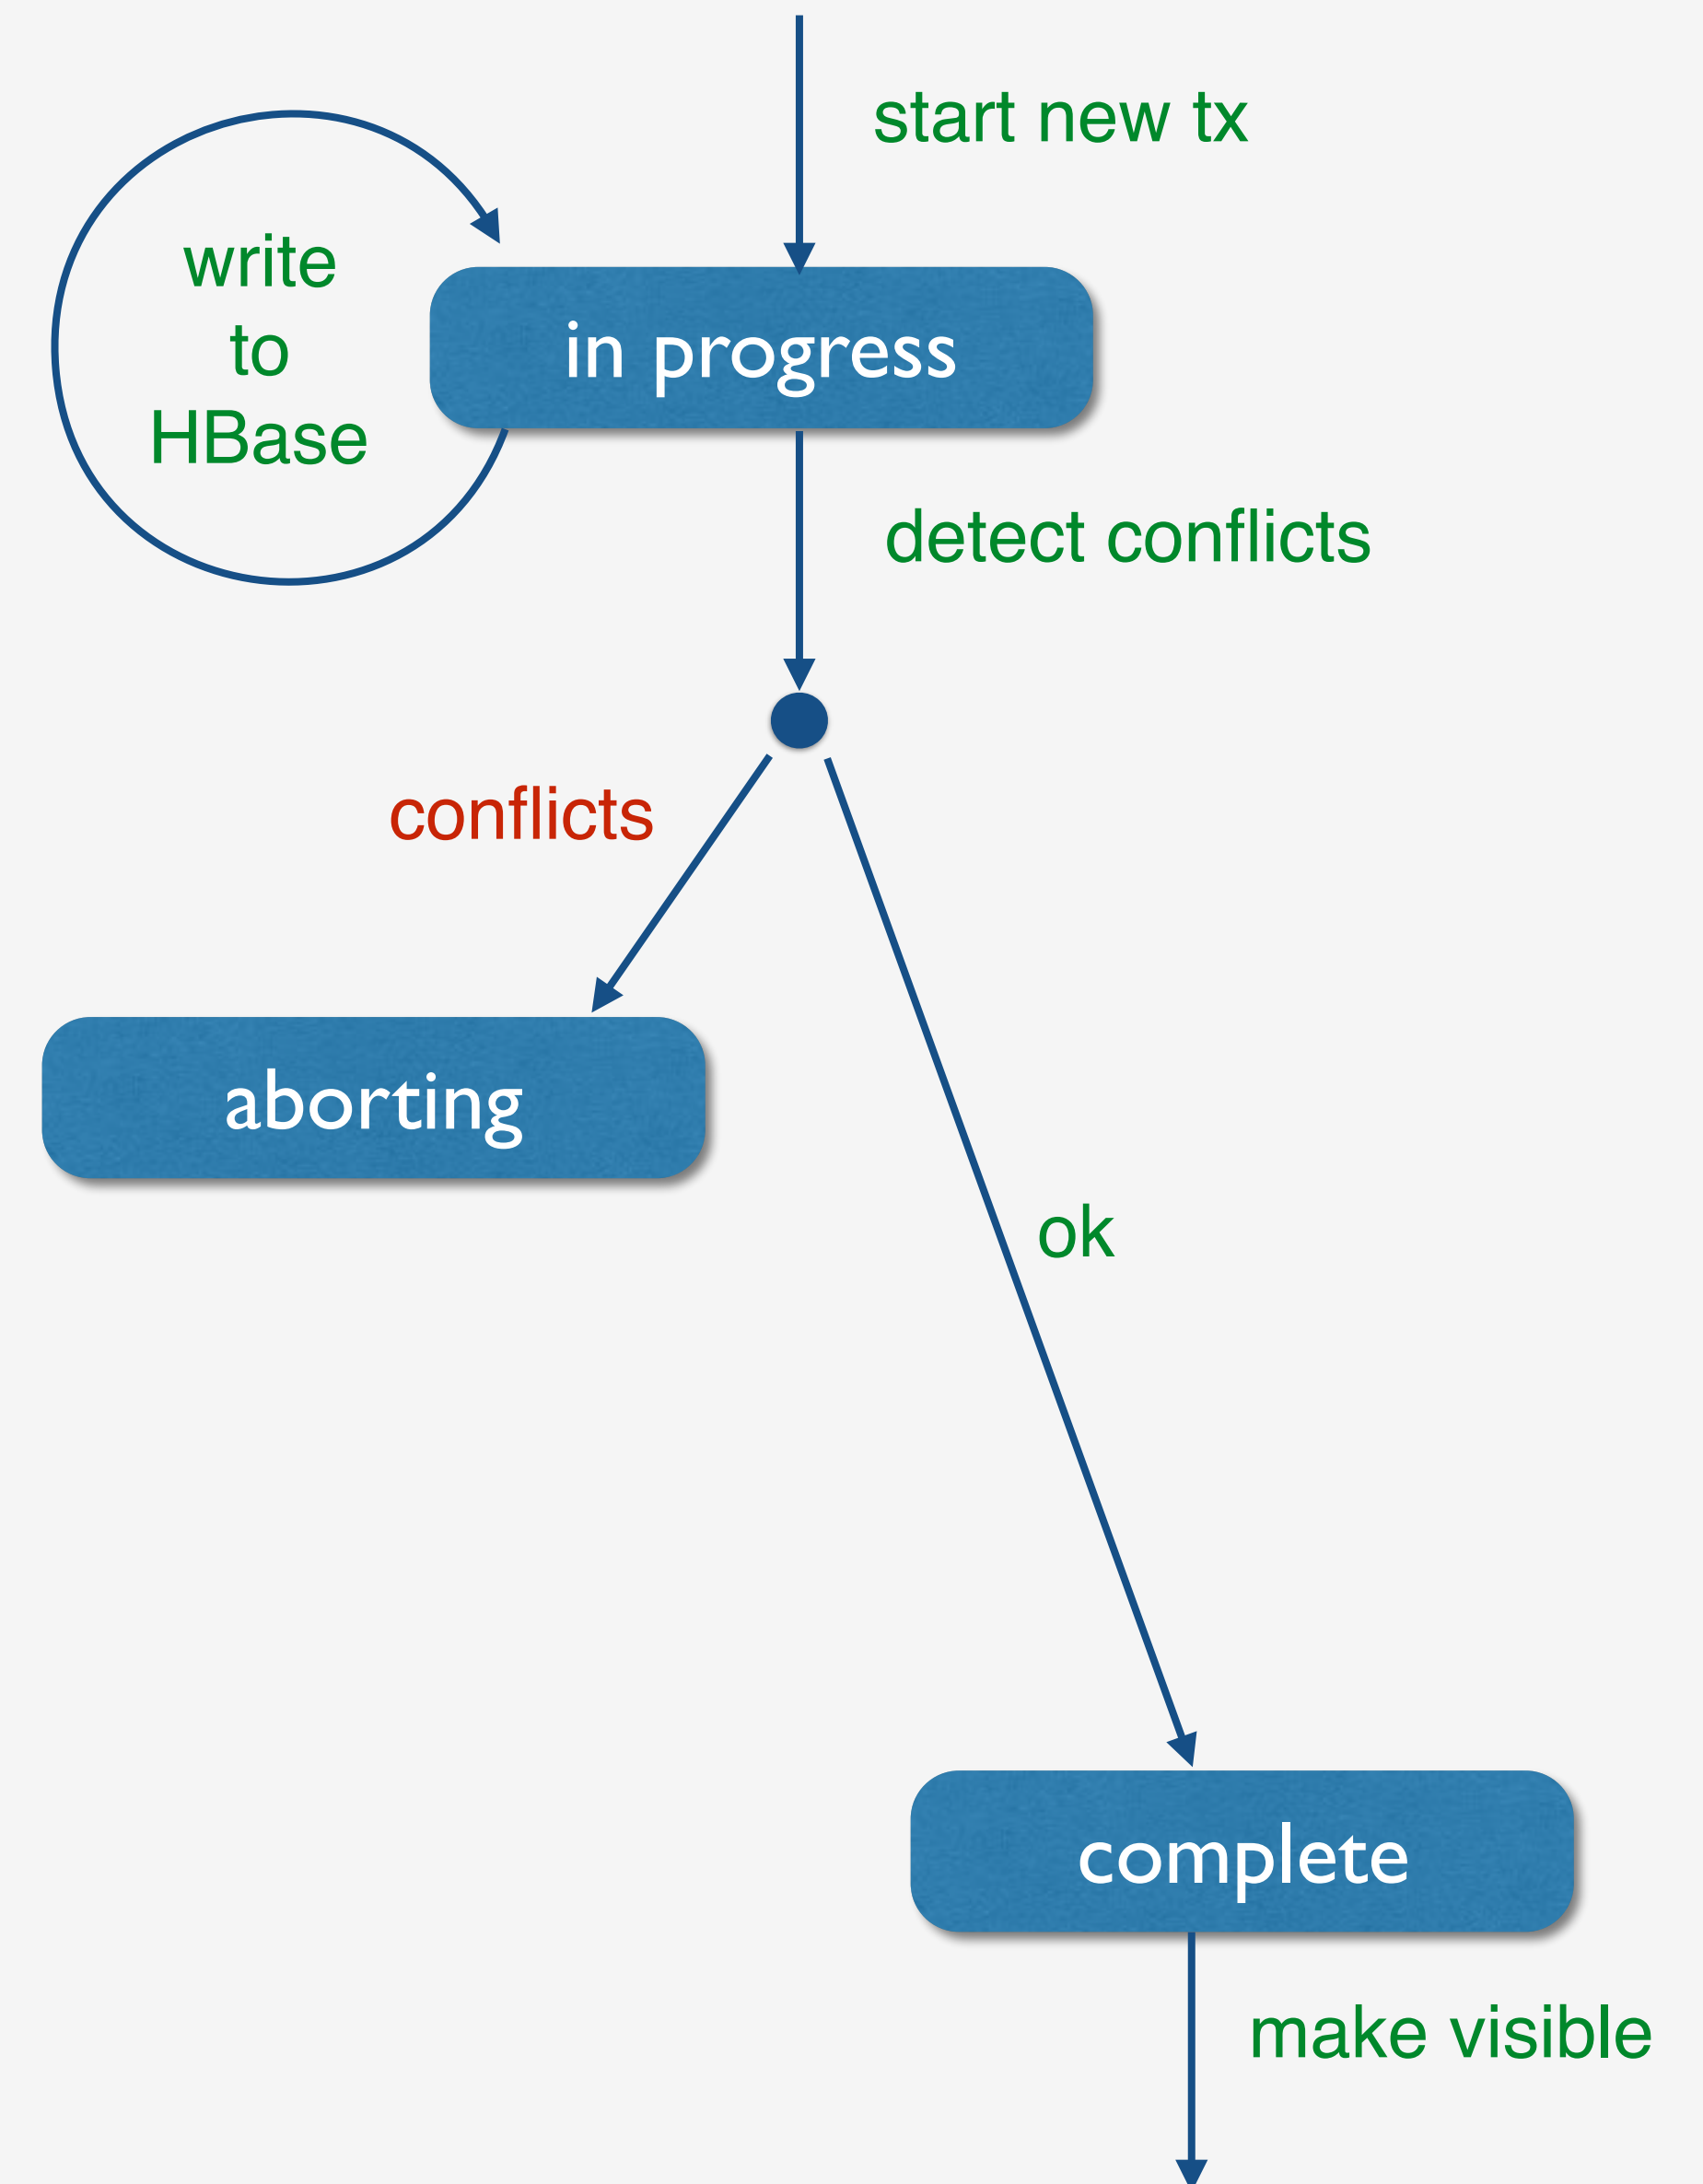
# Transaction Lifecycle



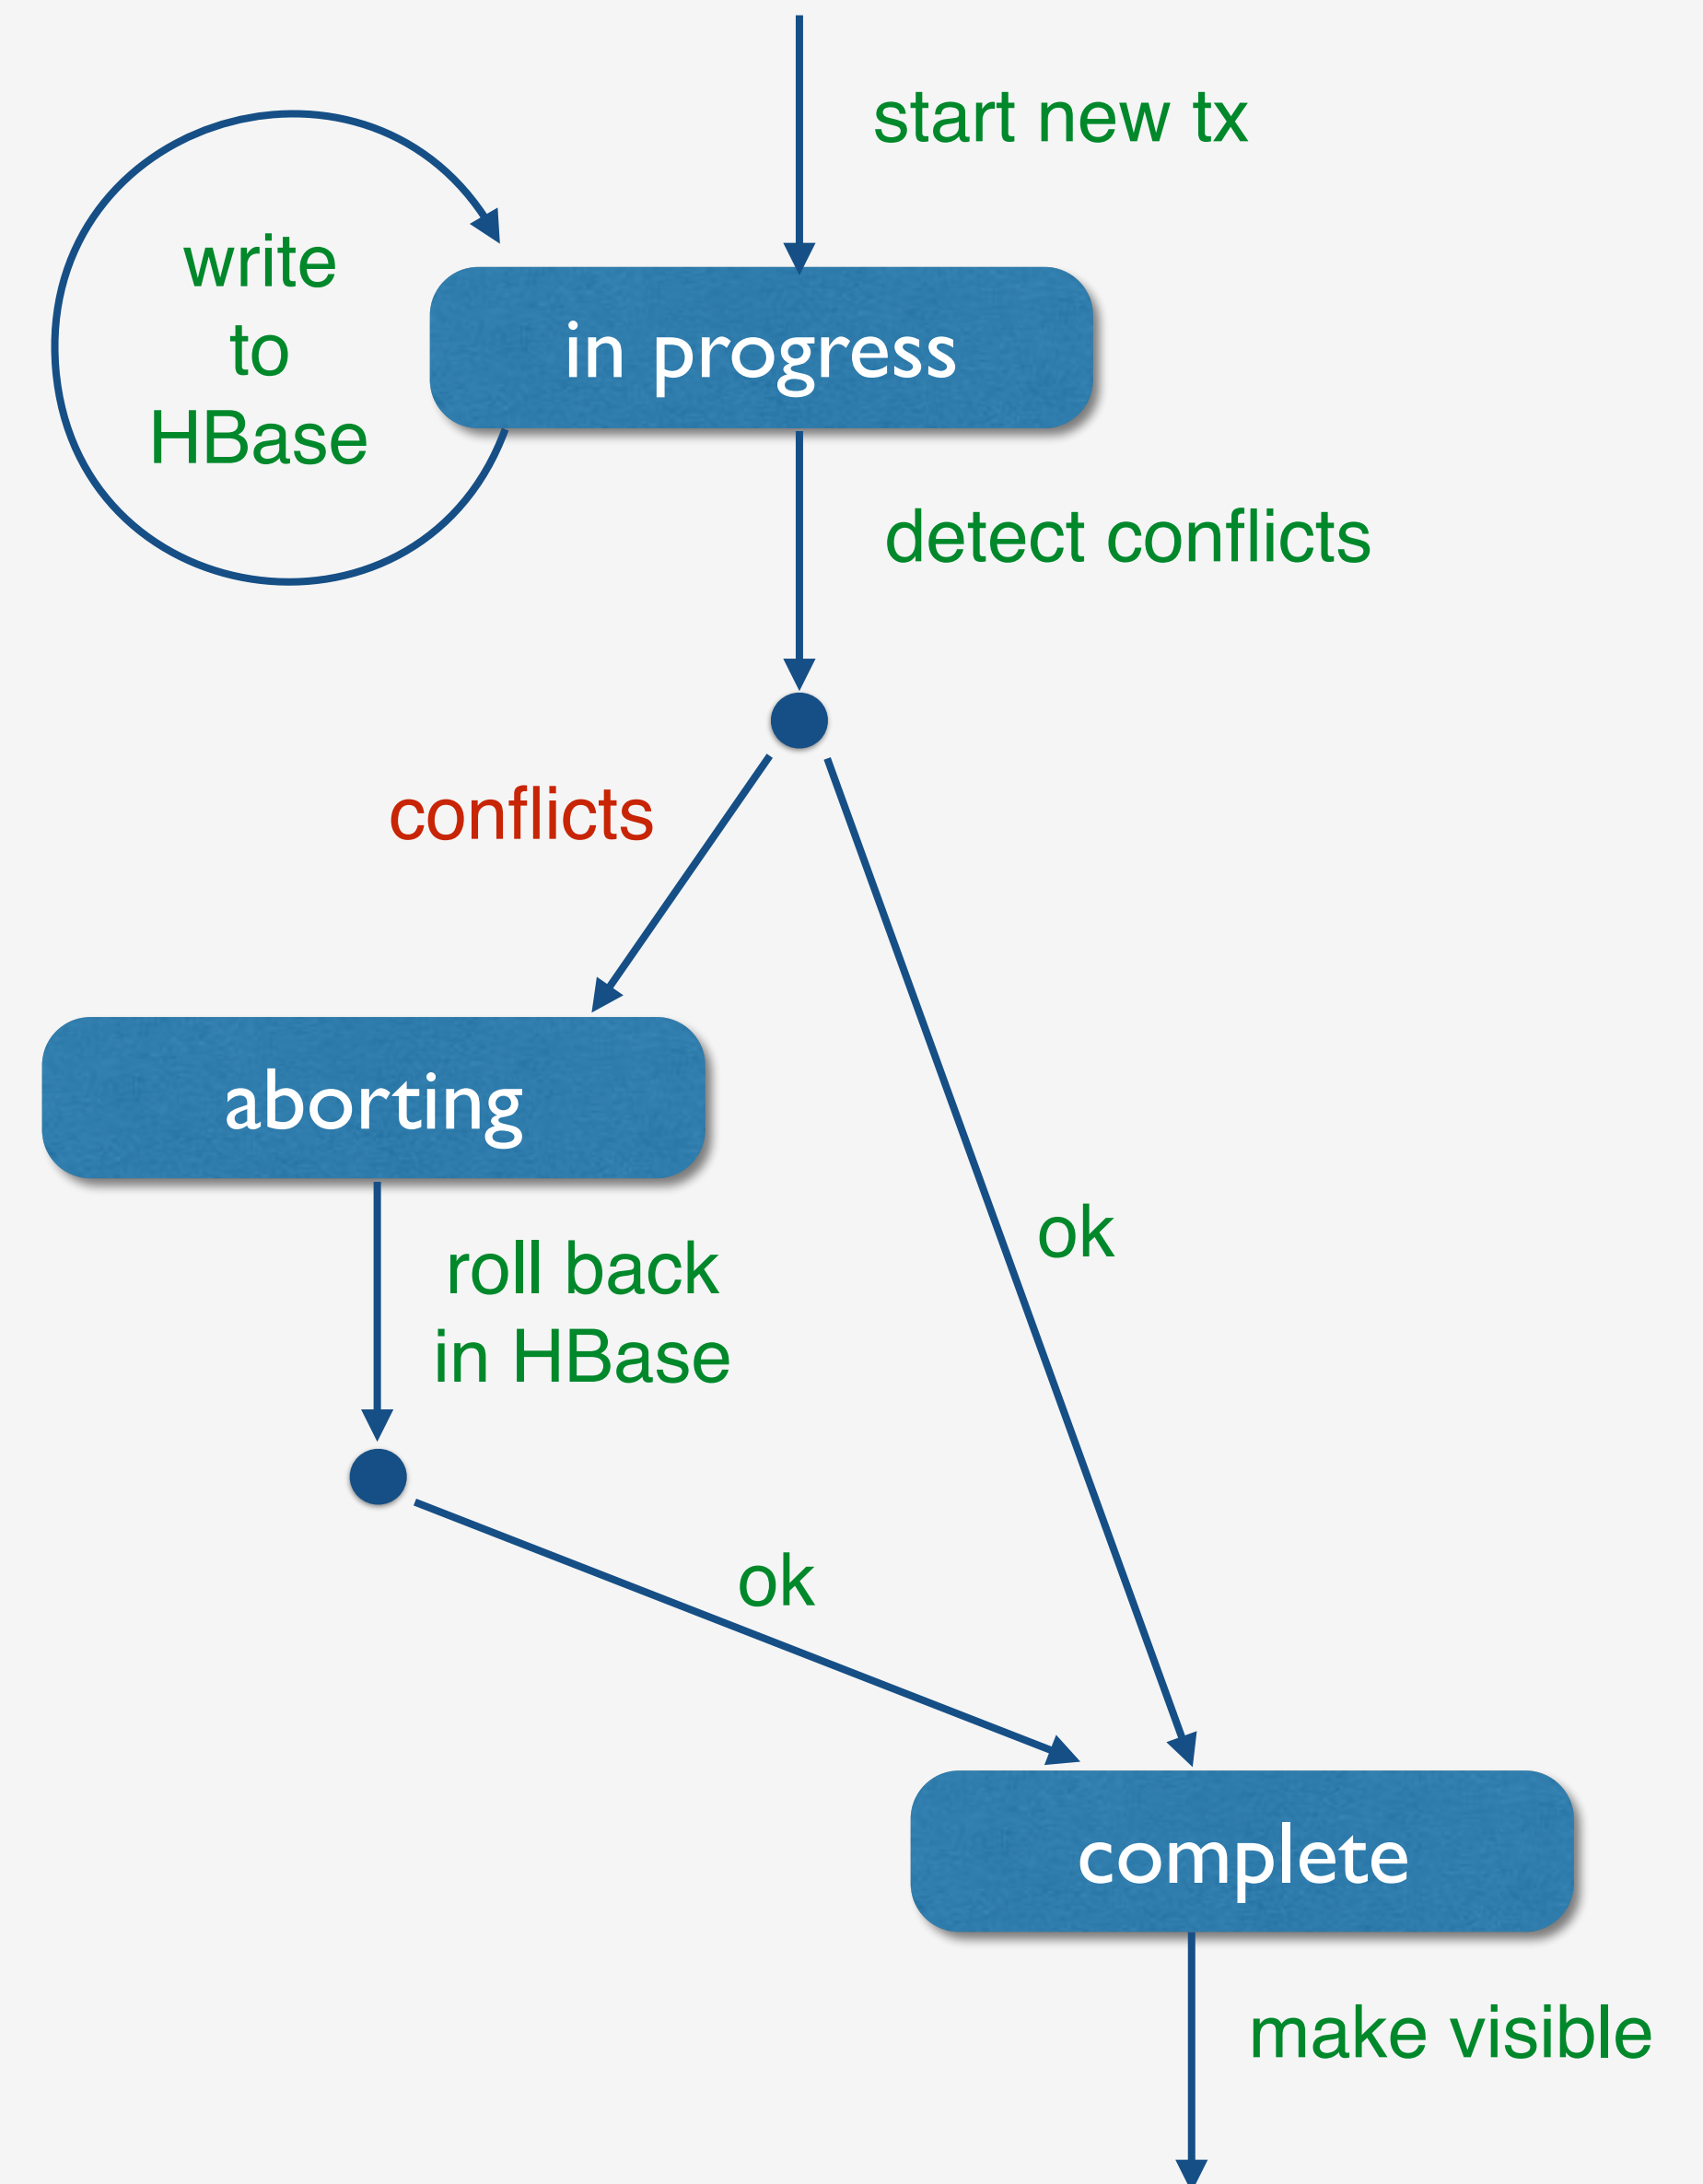
# Transaction Lifecycle



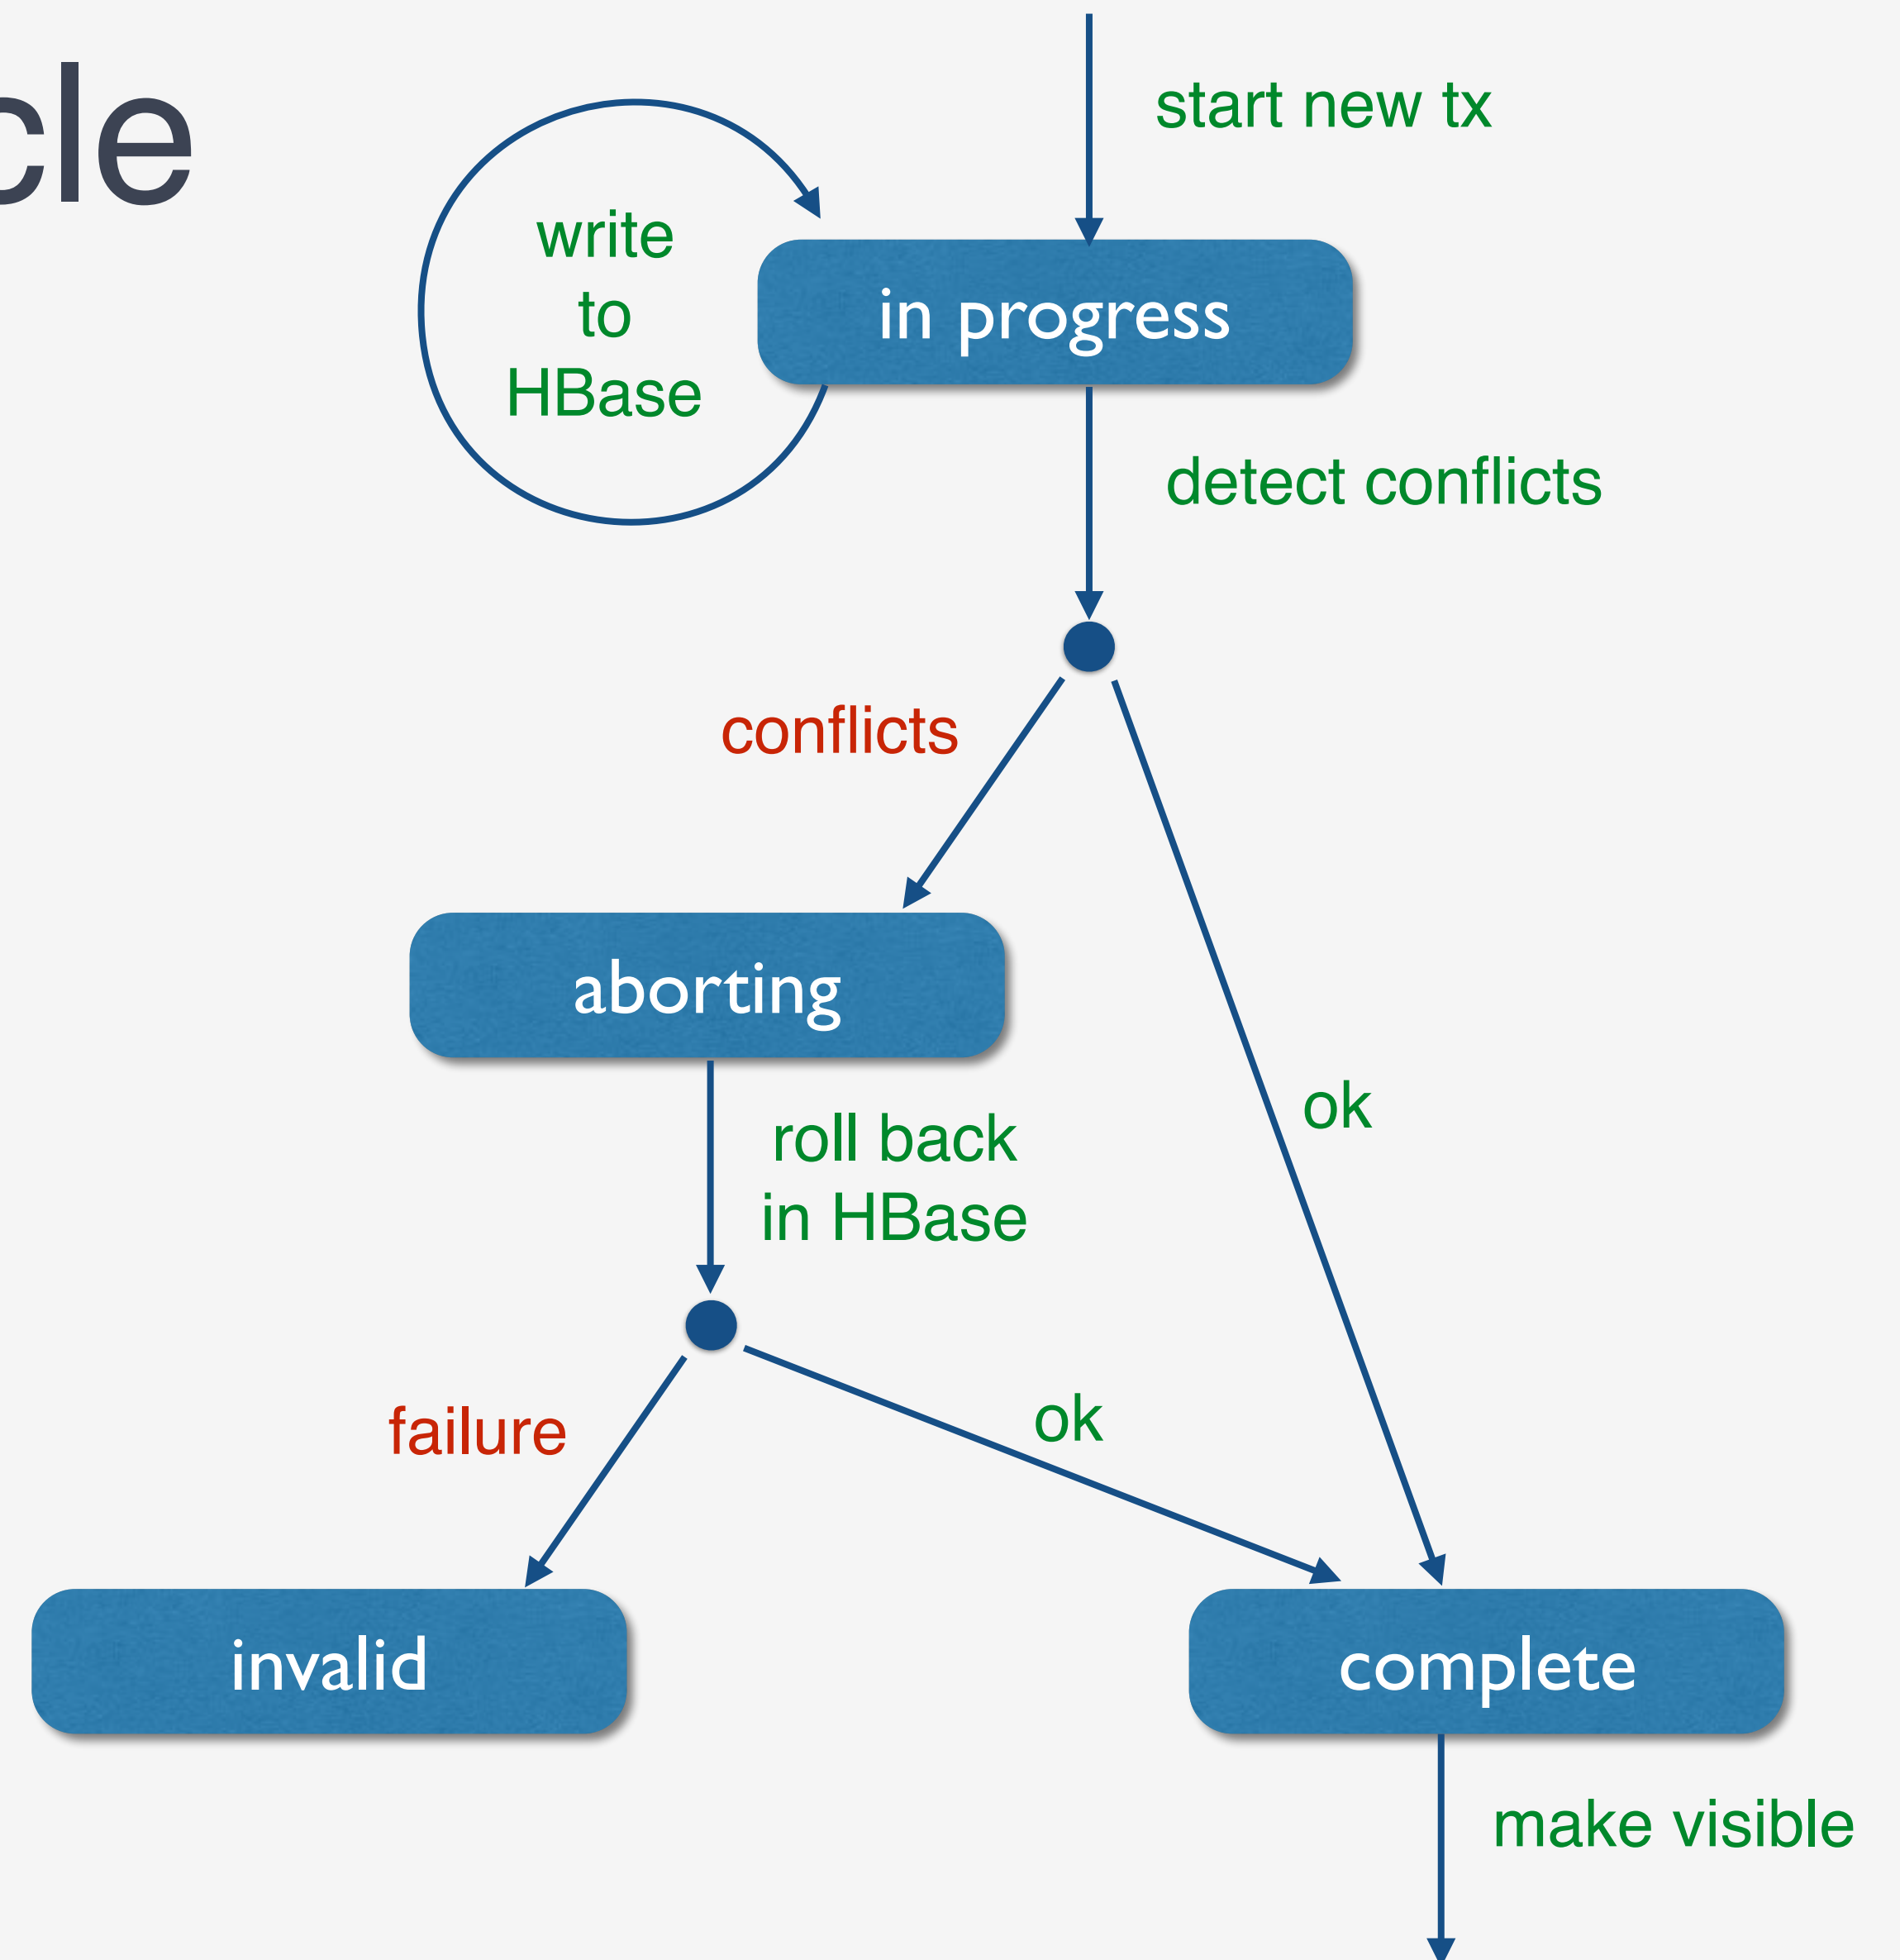
# Transaction Lifecycle



# Transaction Lifecycle

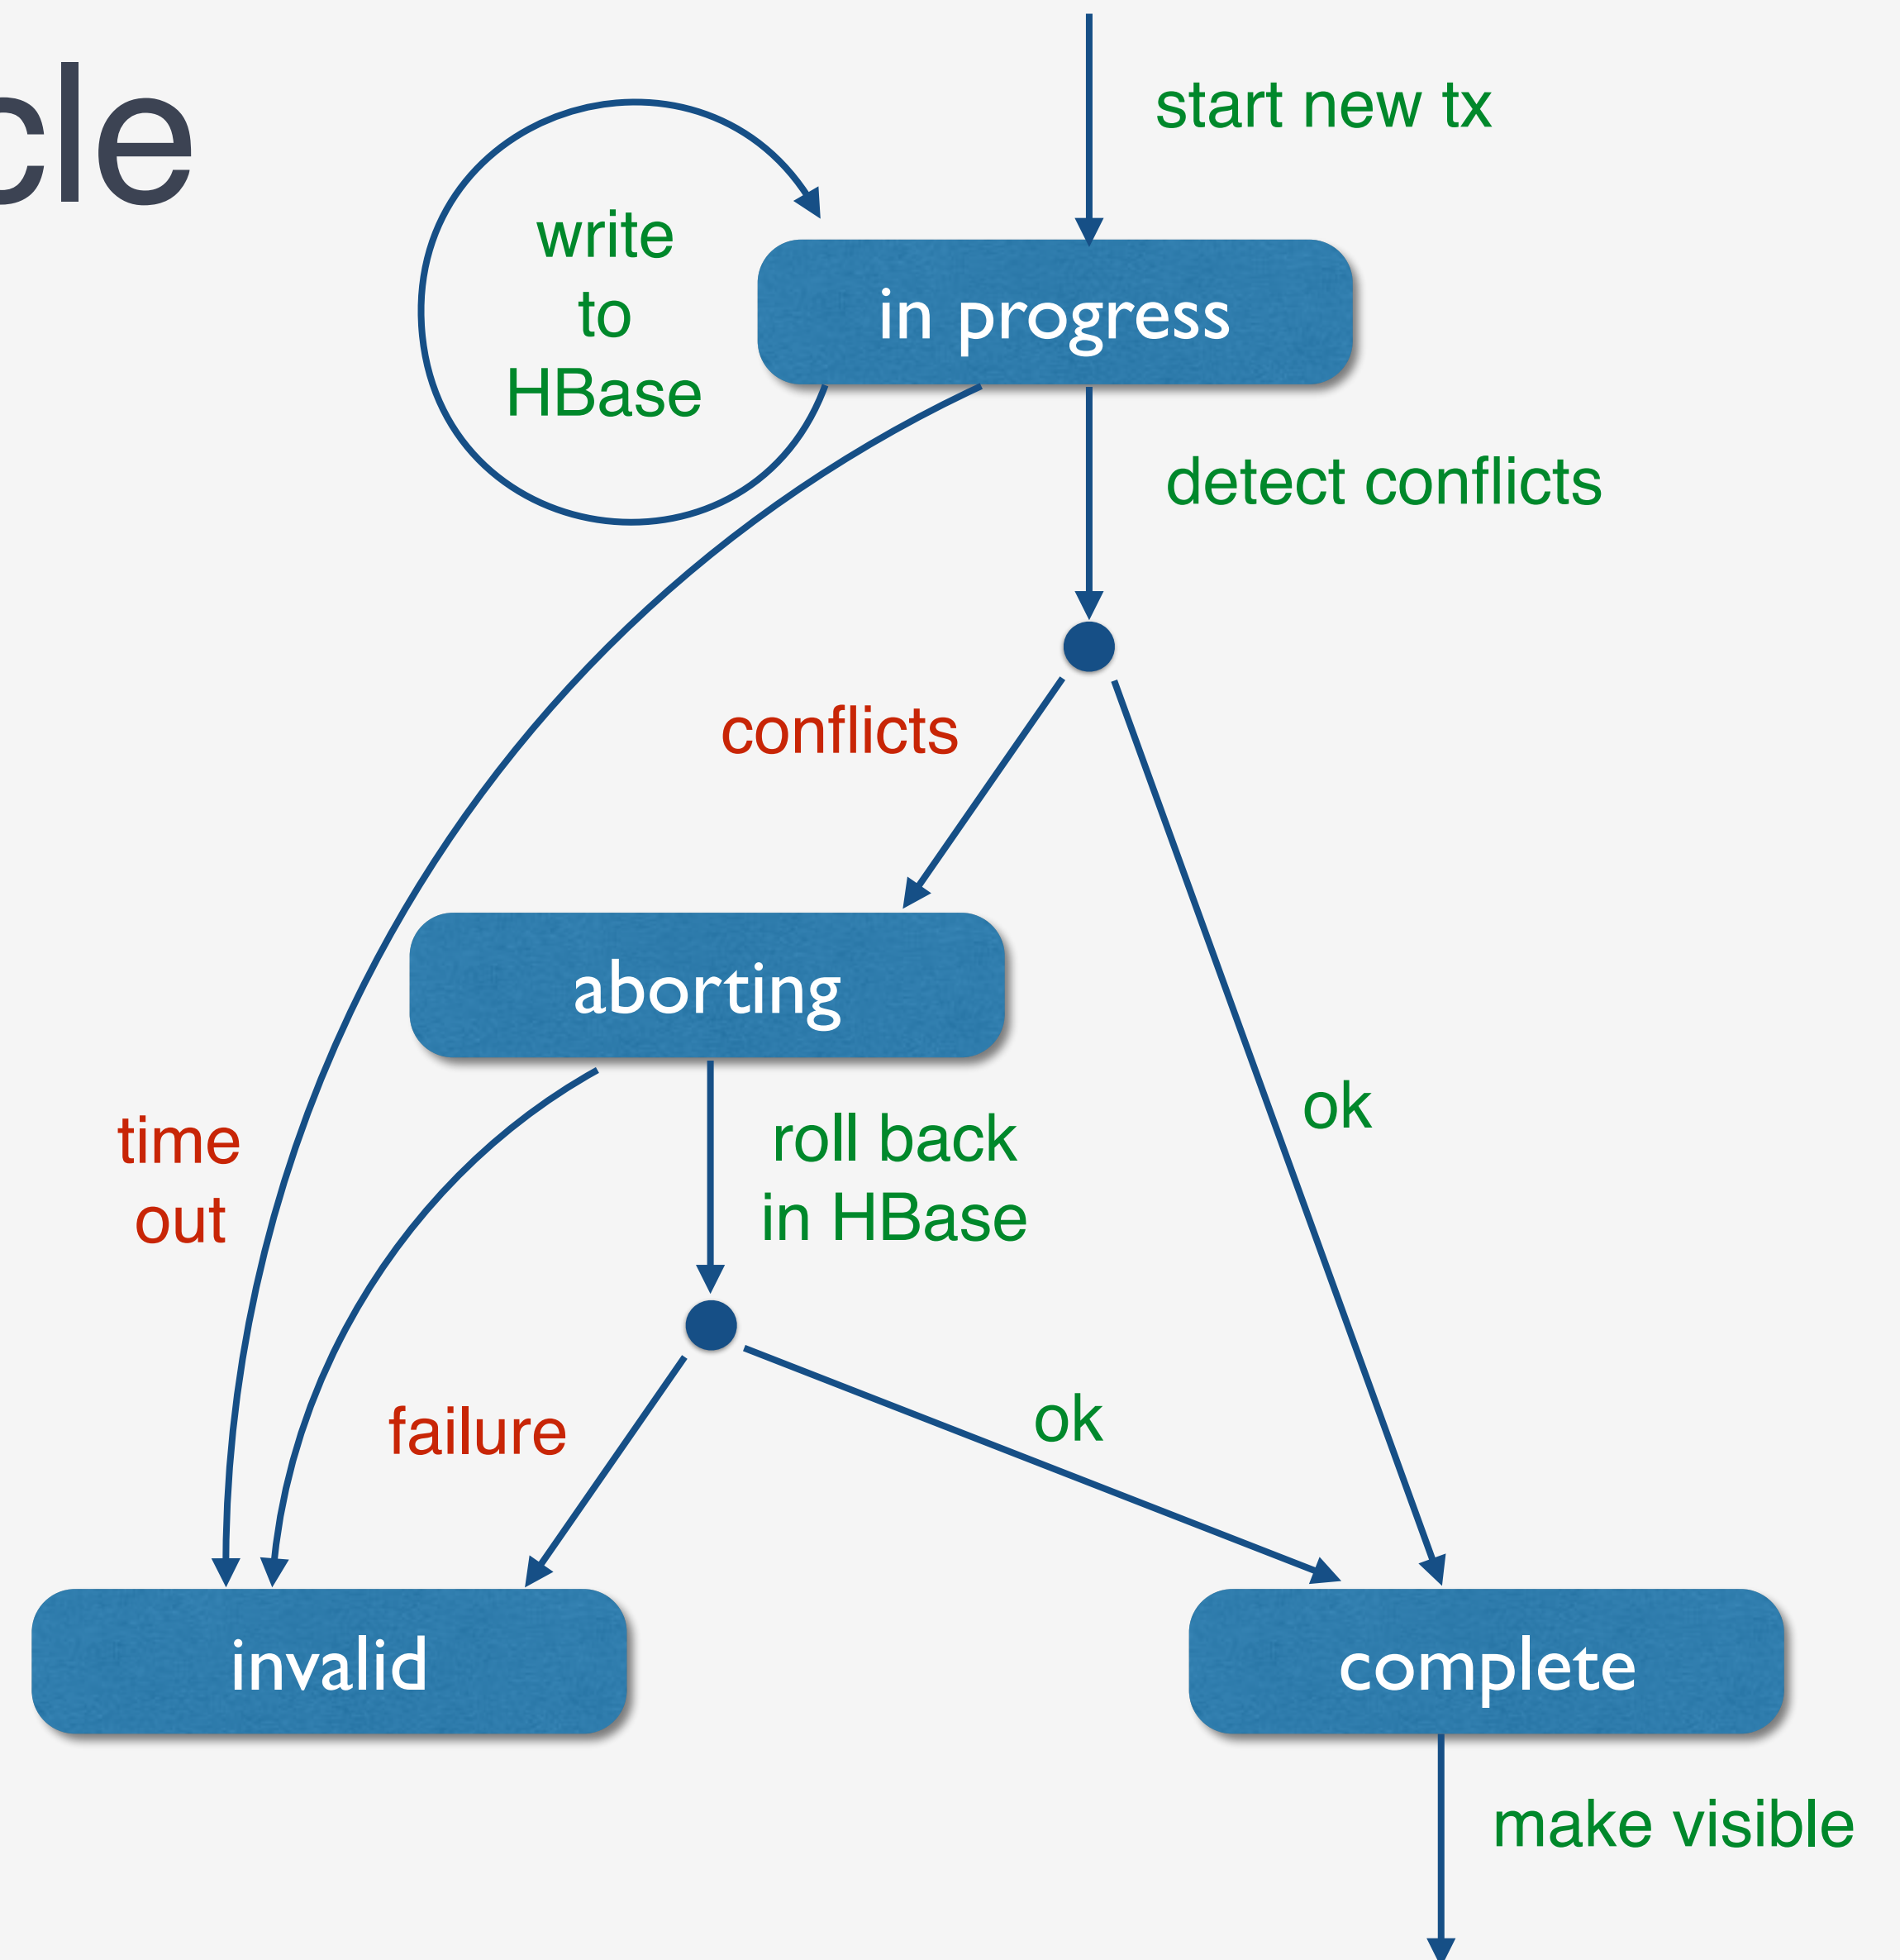


# Transaction Lifecycle



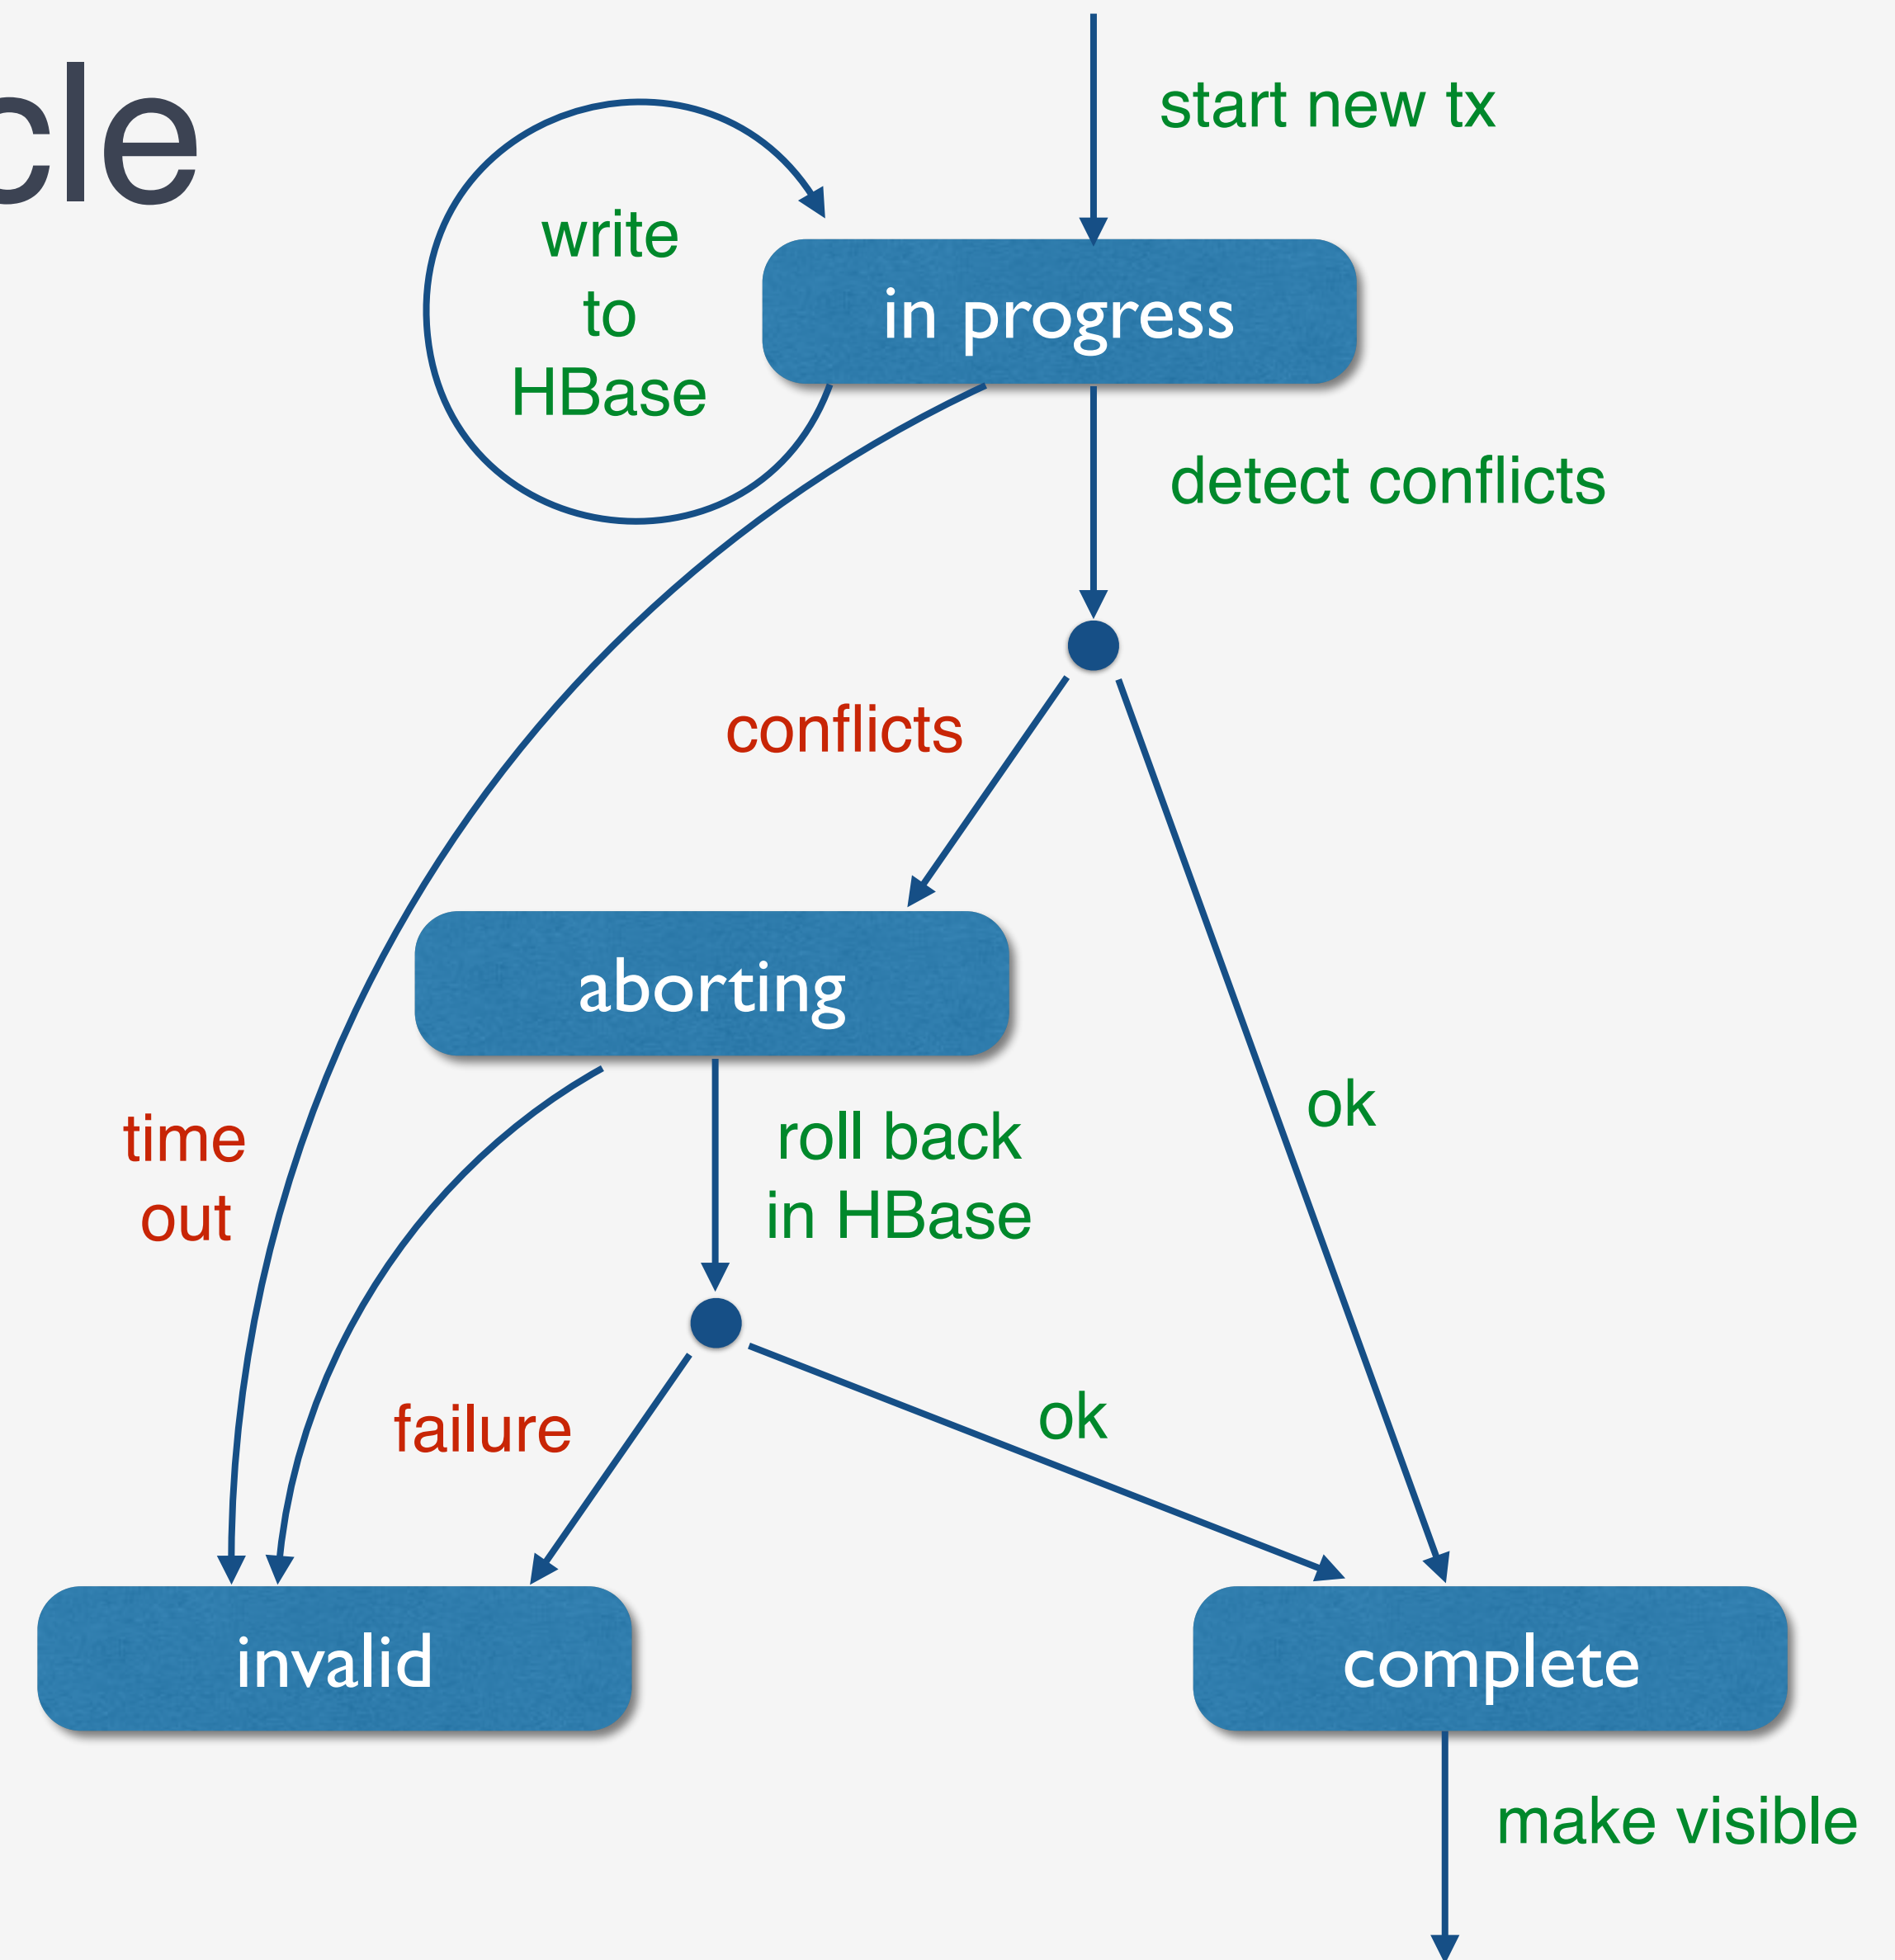


# Transaction Lifecycle



# Transaction Lifecycle

- Transaction consists of:
  - transaction ID (unique timestamp)
  - exclude list (in-flight and invalid tx)
- Transactions that do complete
  - must still participate in conflict detection
  - disappear from transaction state when they do not overlap with in-flight tx
- Transactions that do not complete
  - time out (by transaction manager)
  - added to invalid list



# Apache Tephra

Client A

Tx  
Manager

in-flight:  
...

## HBase

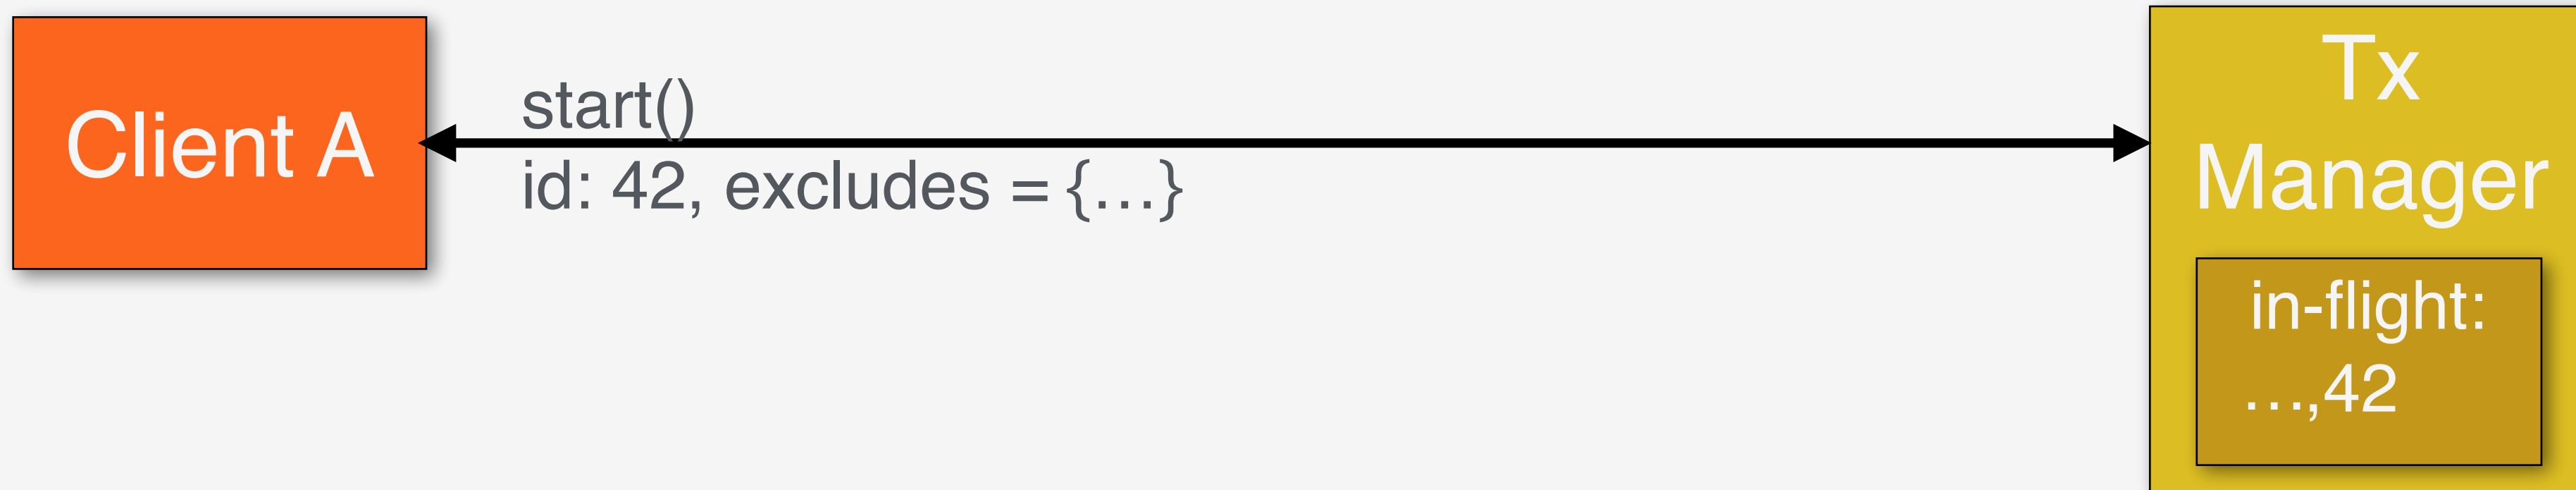
Region Server

x:10 37


Region Server



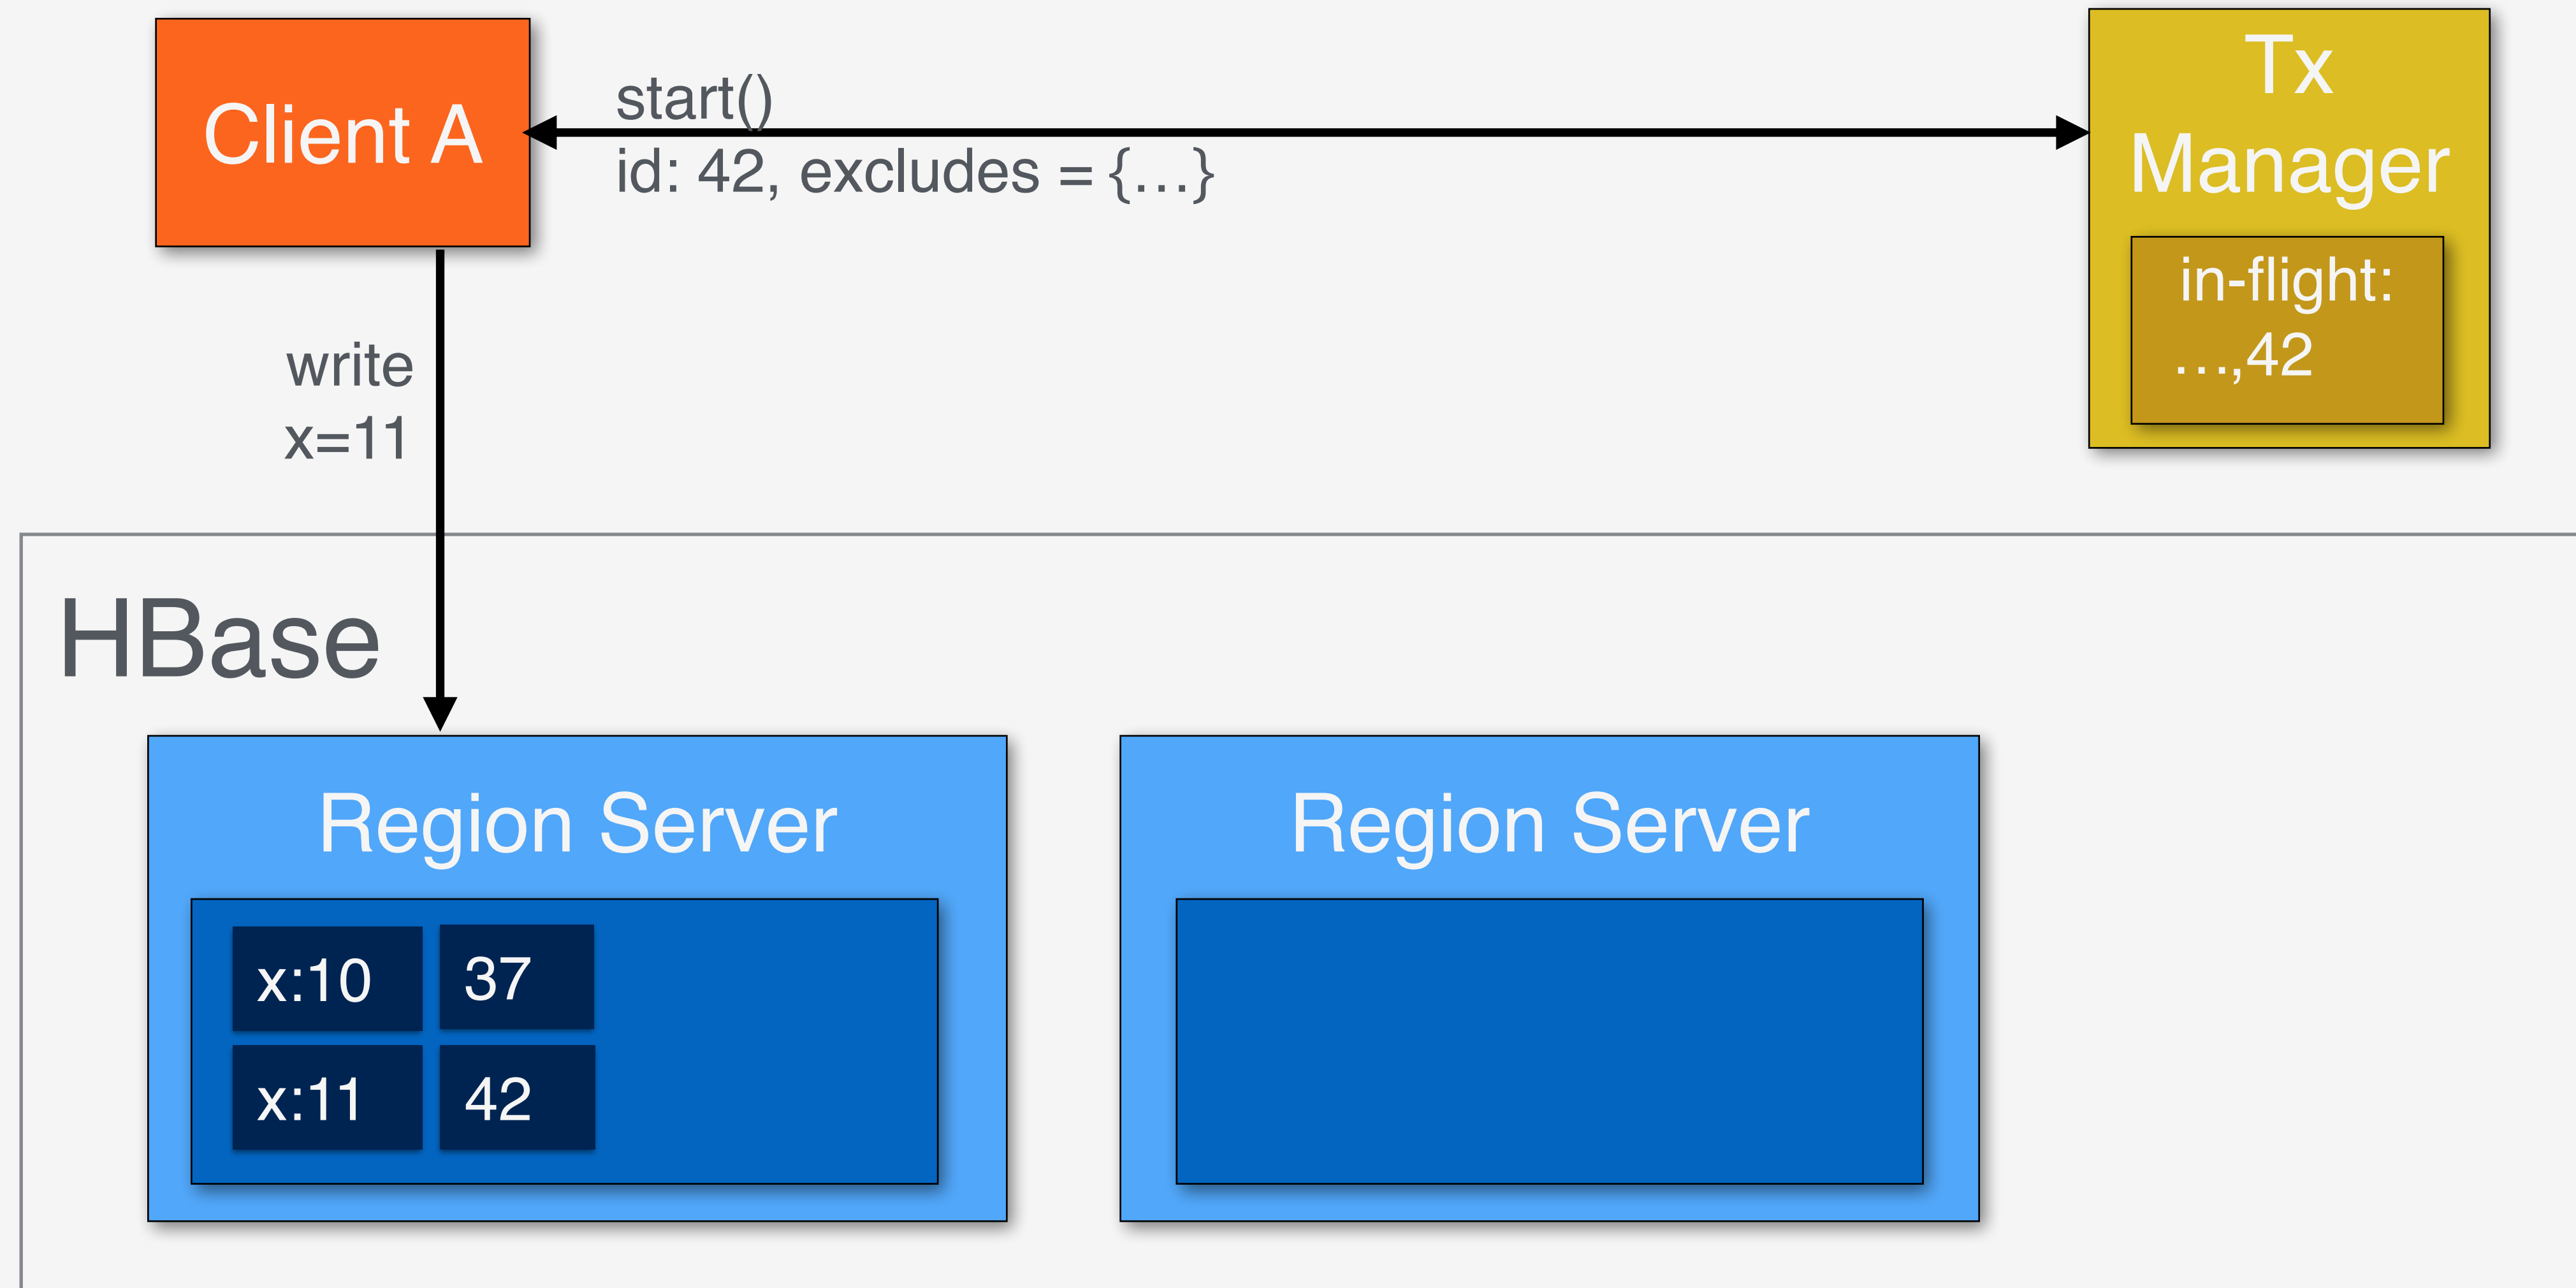
# Apache Tephra




## HBase



# Apache Tephra



# Apache Tephra



# Apache Tephra

Client B

Tx  
Manager

in-flight:  
...,42

## HBase

Region Server

x:10 37

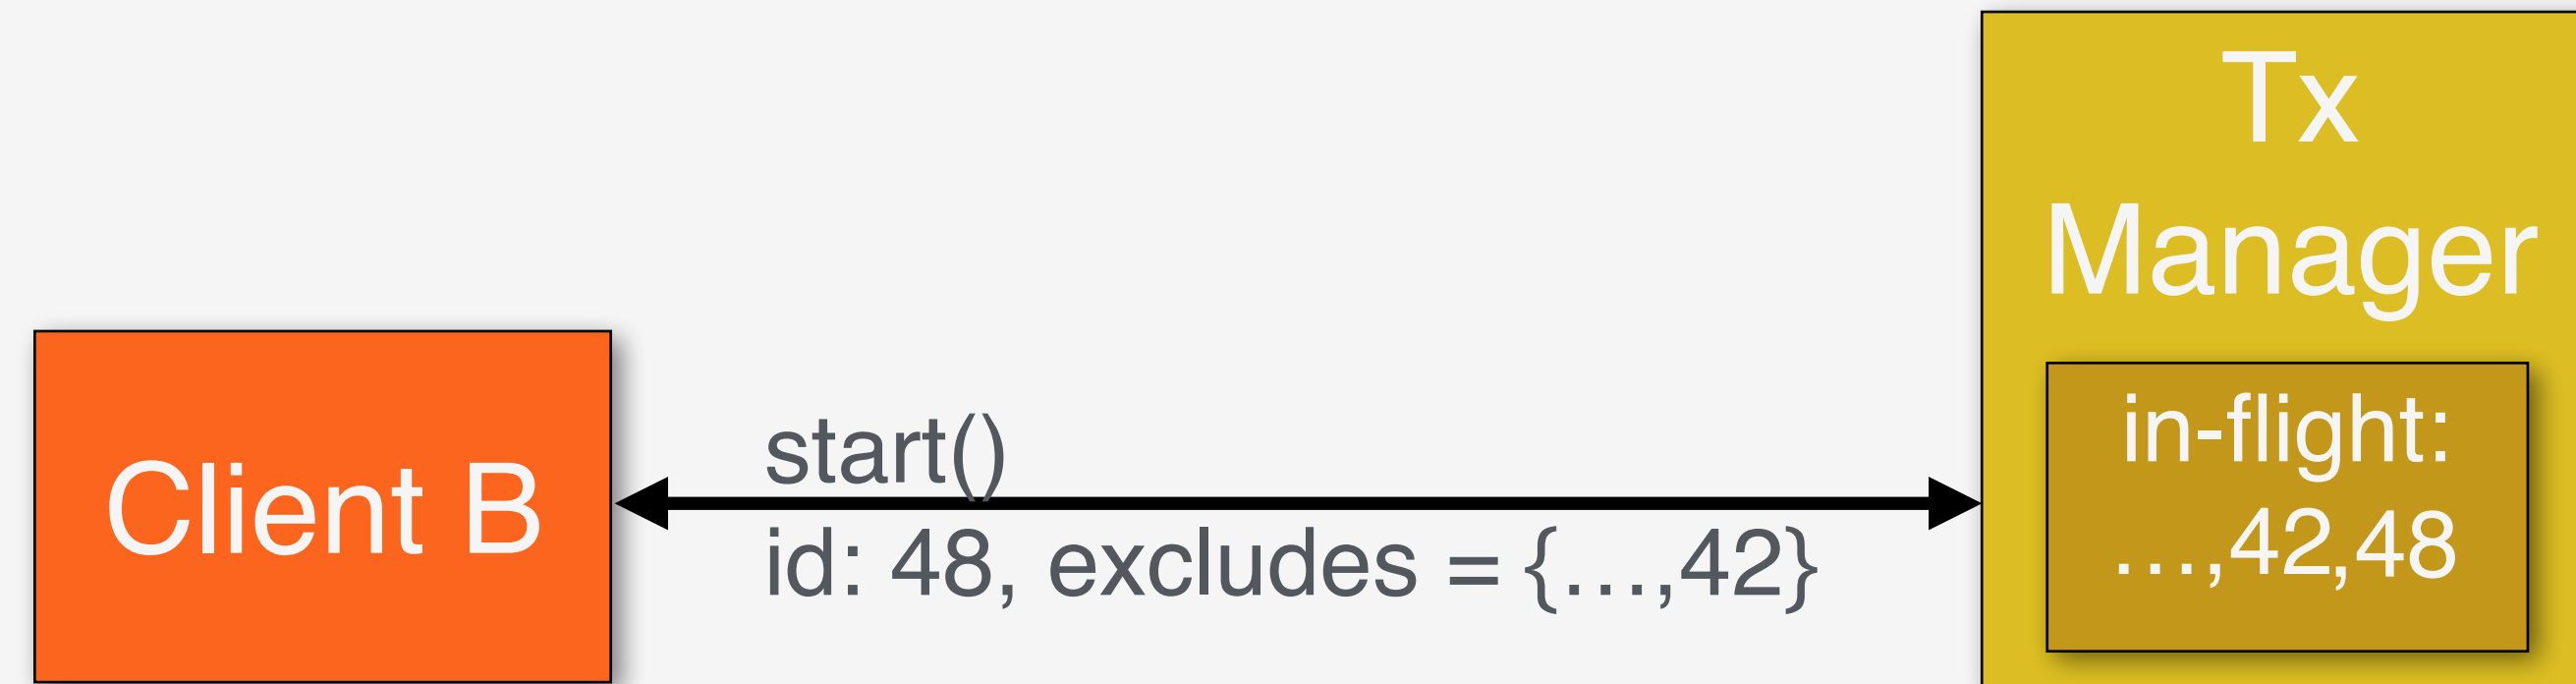
x:11 42

Region Server

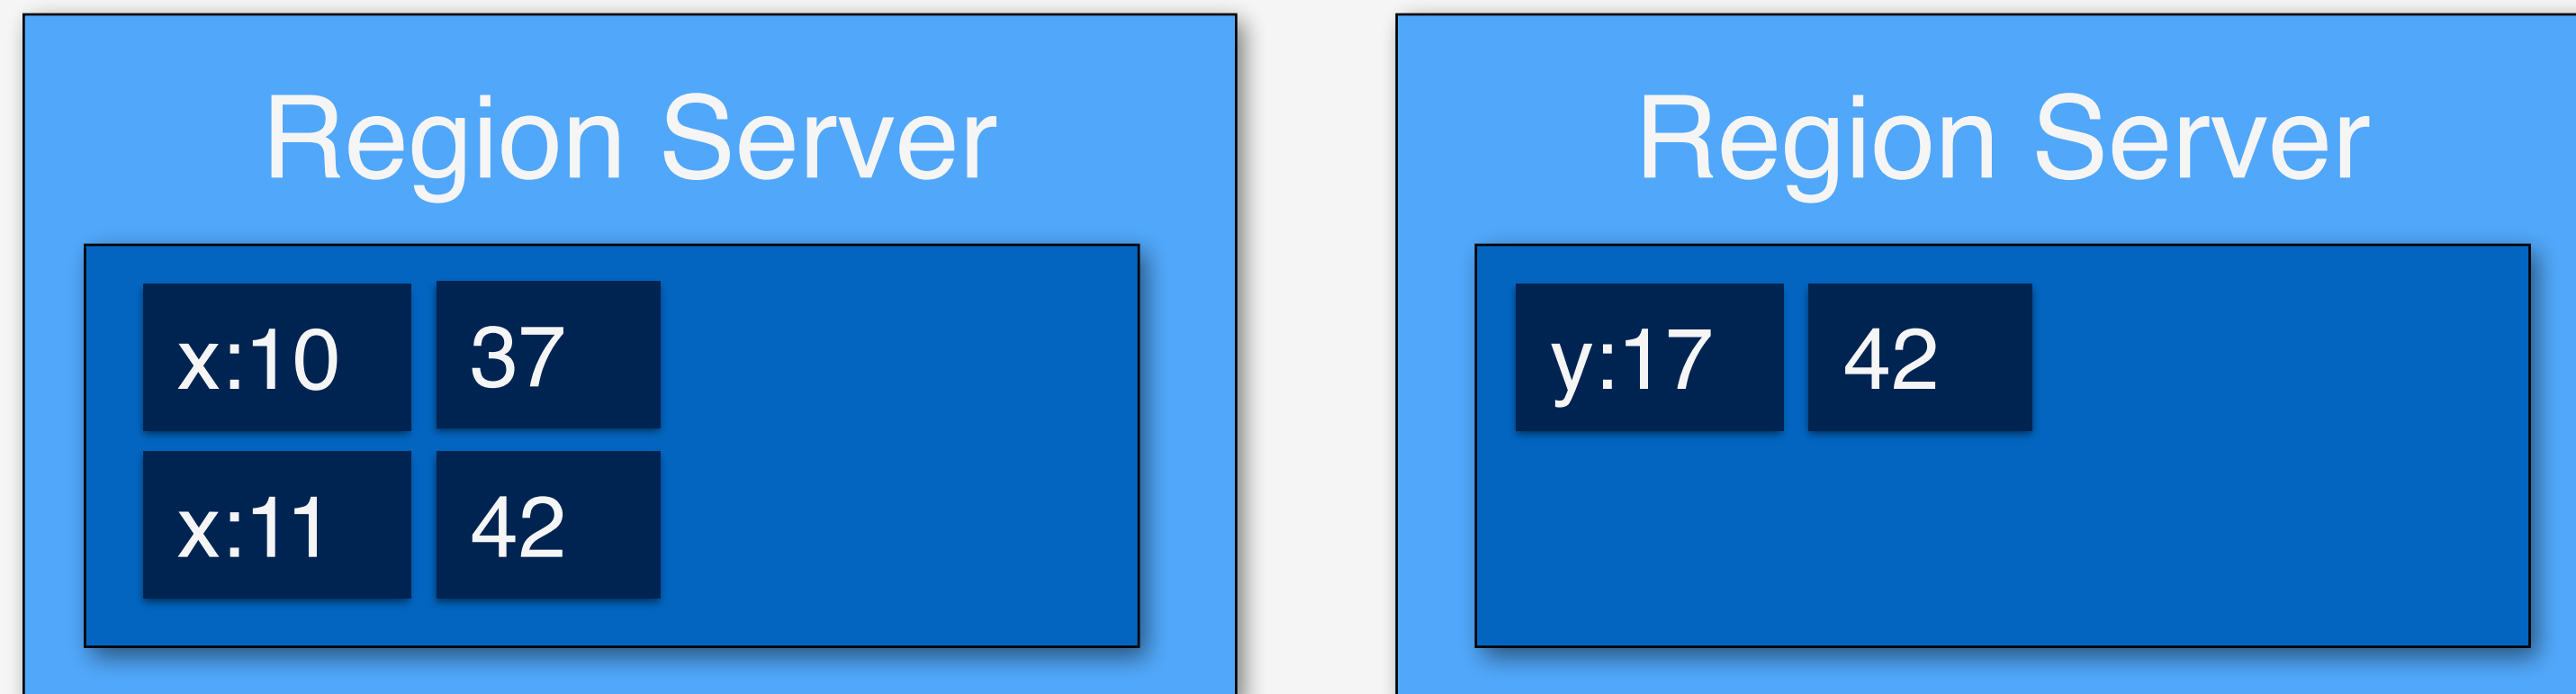
y:17 42



# Apache Tephra

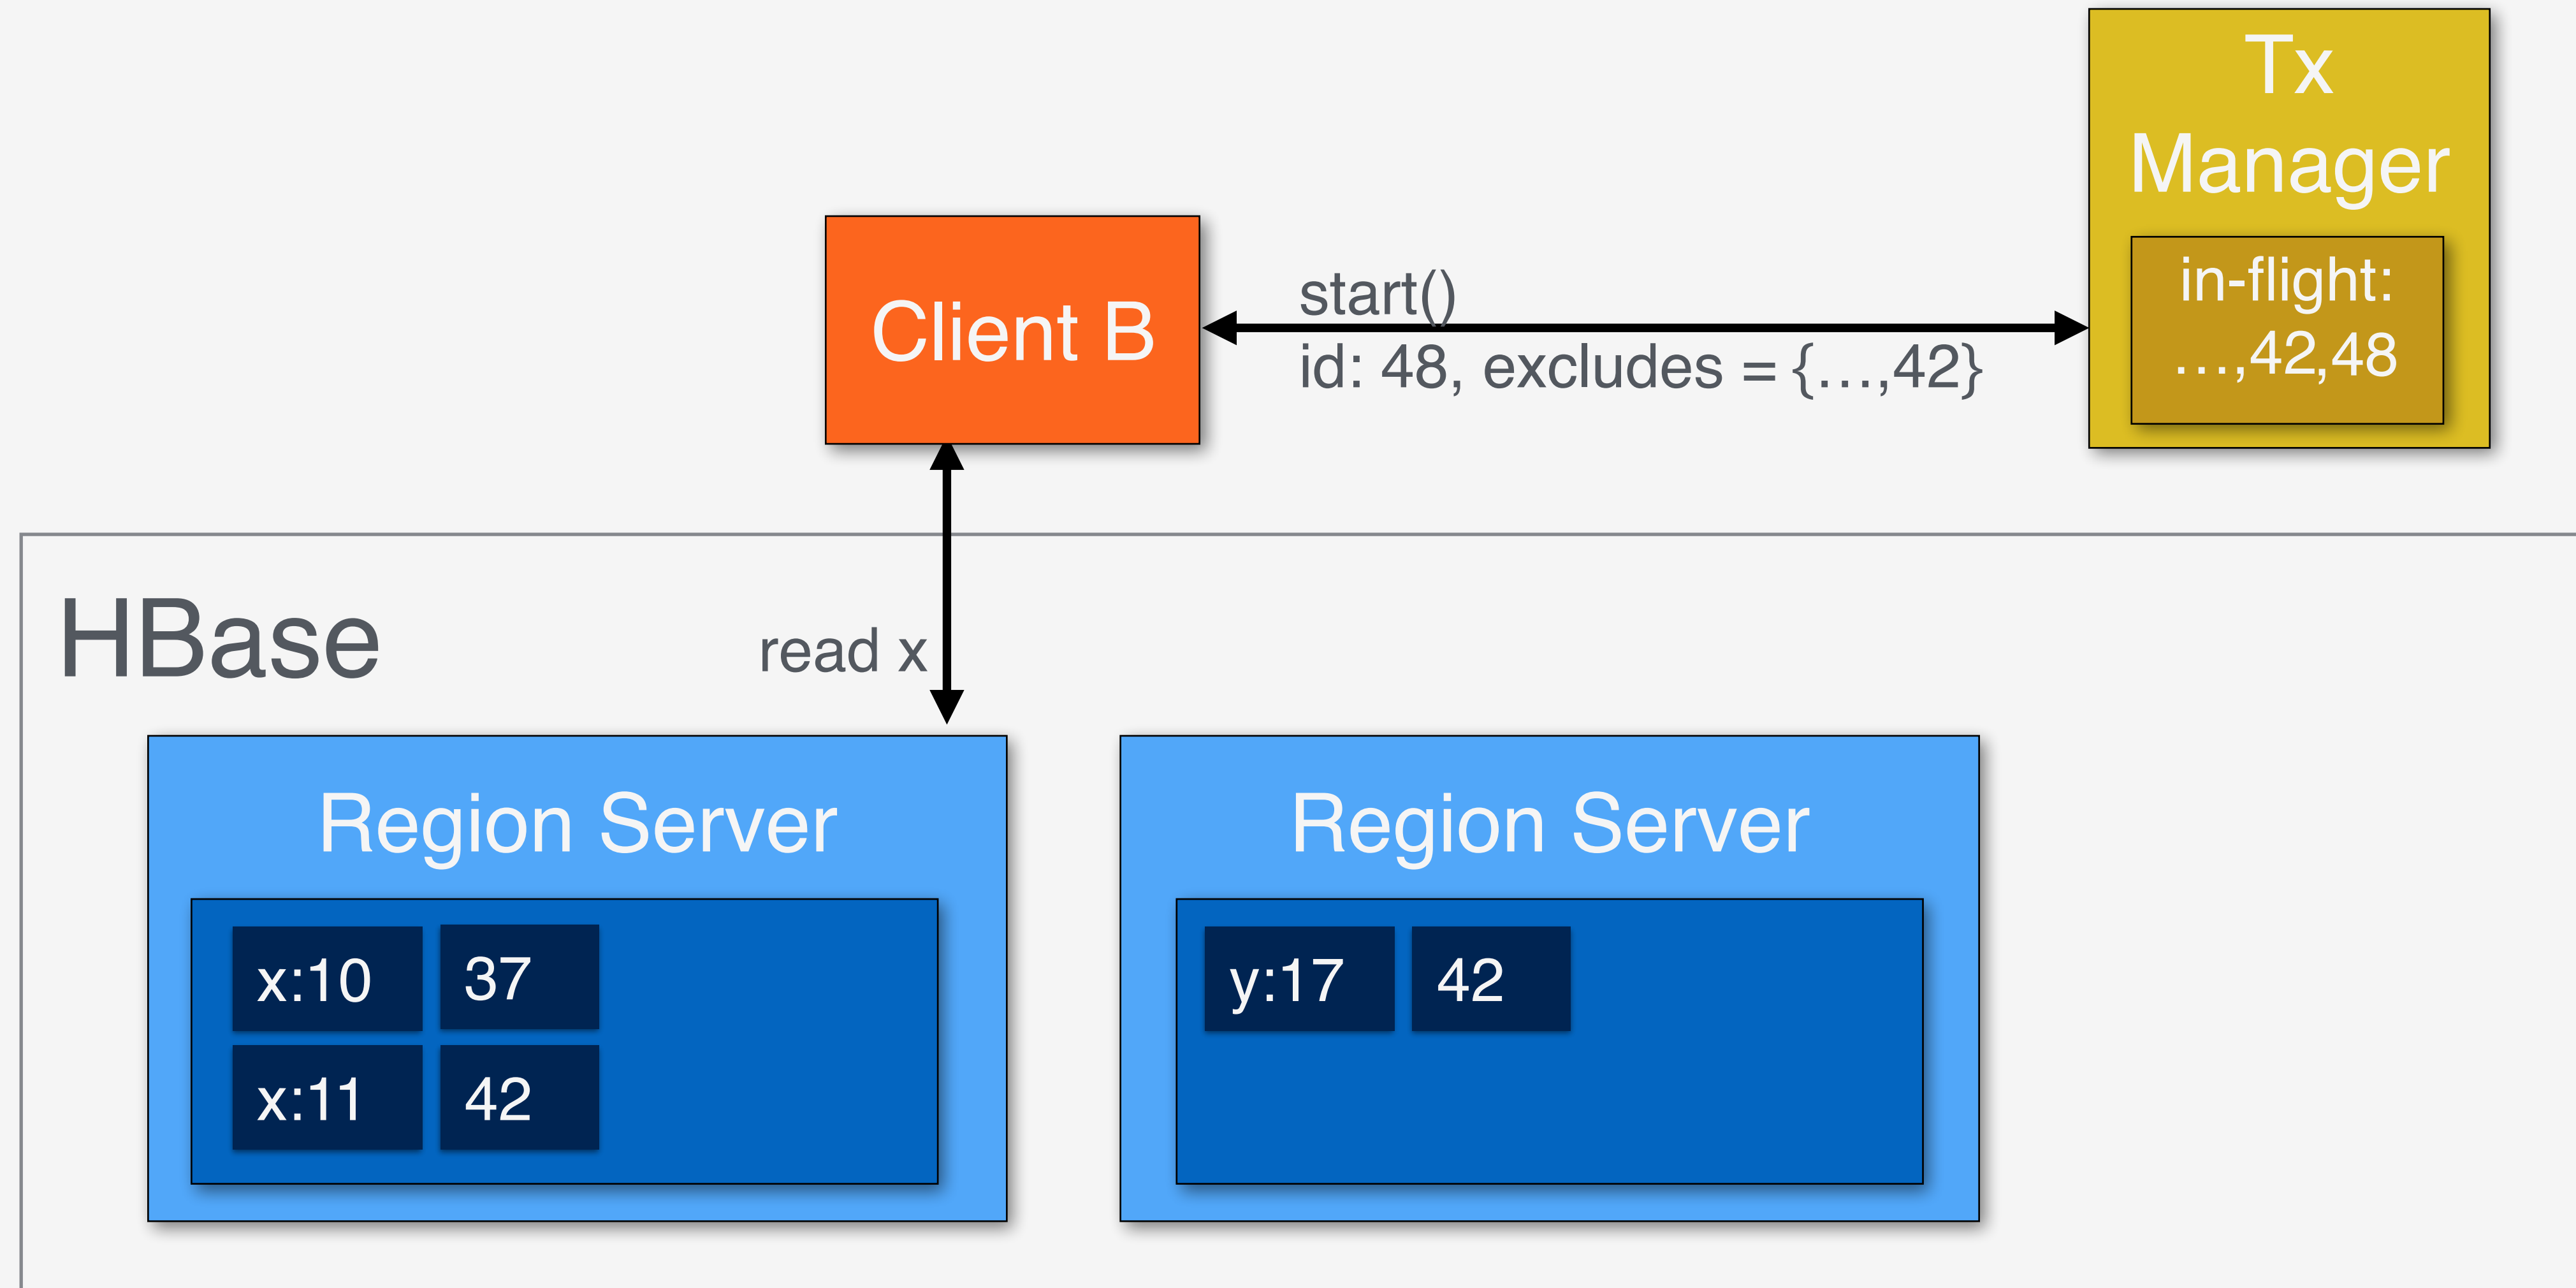


## HBase






# Apache Tephra



# Apache Tephra



# Apache Tephra

Client A

Tx  
Manager

in-flight:  
...,42

## HBase

Region Server

x:10 37

x:11 42

Region Server

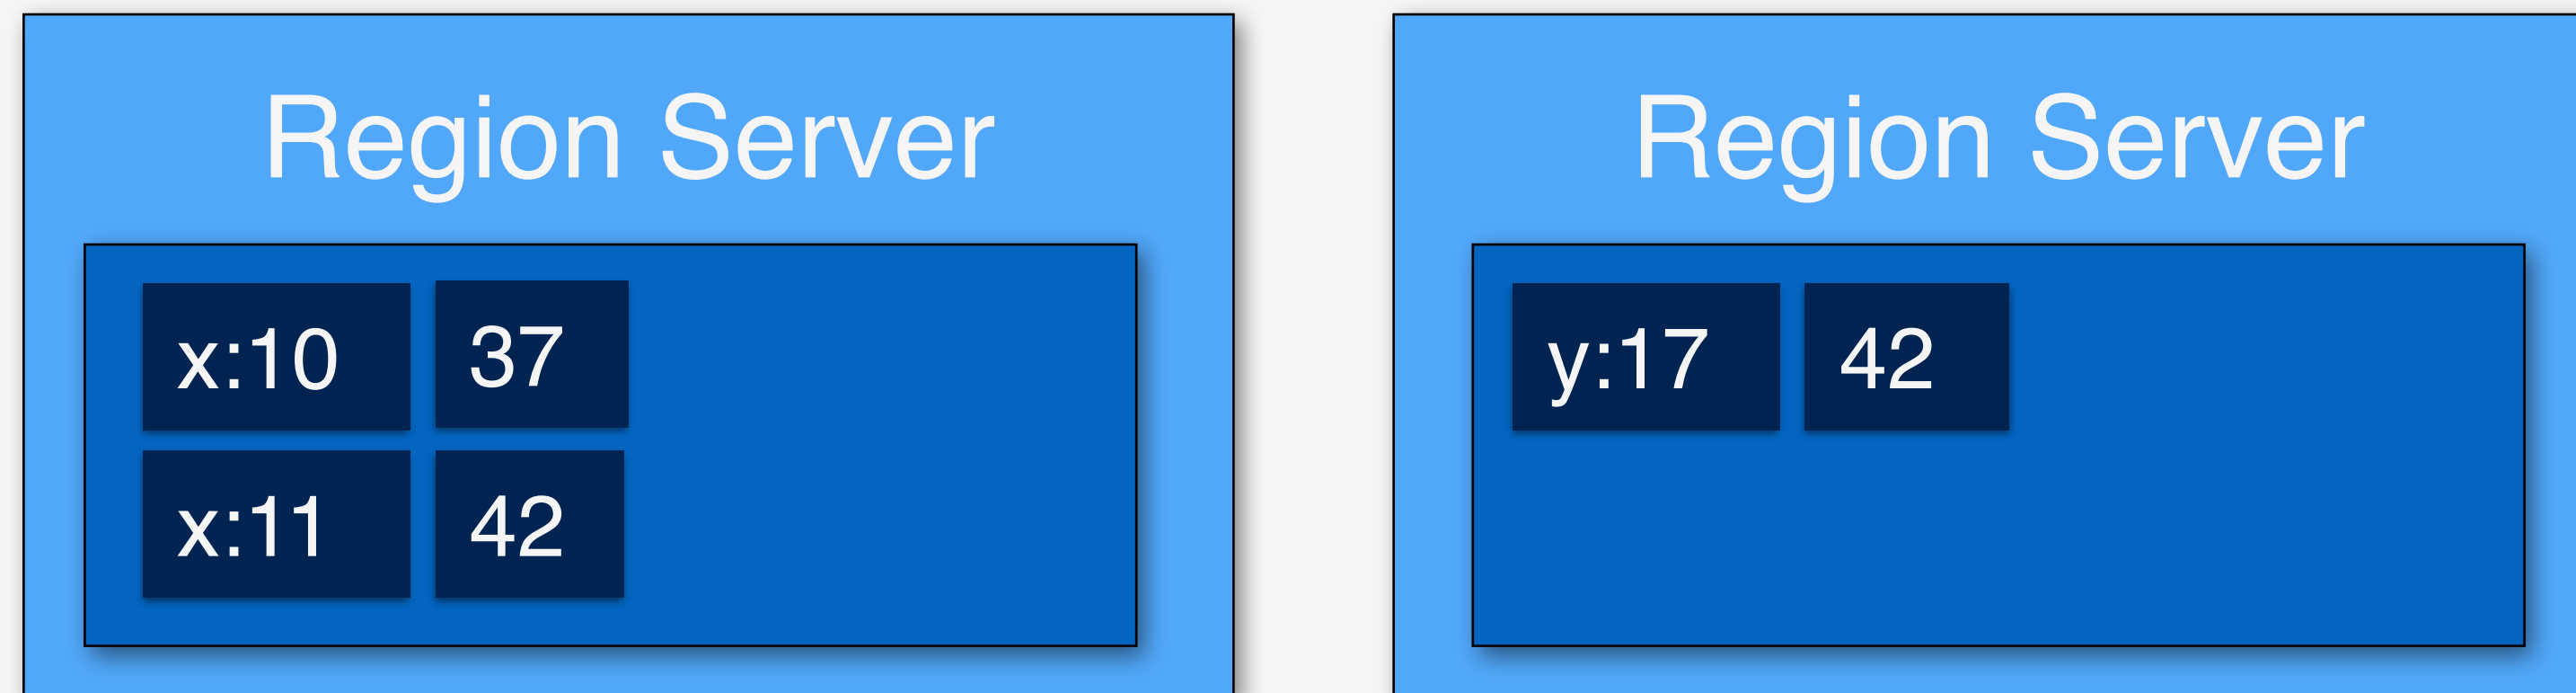
y:17 42



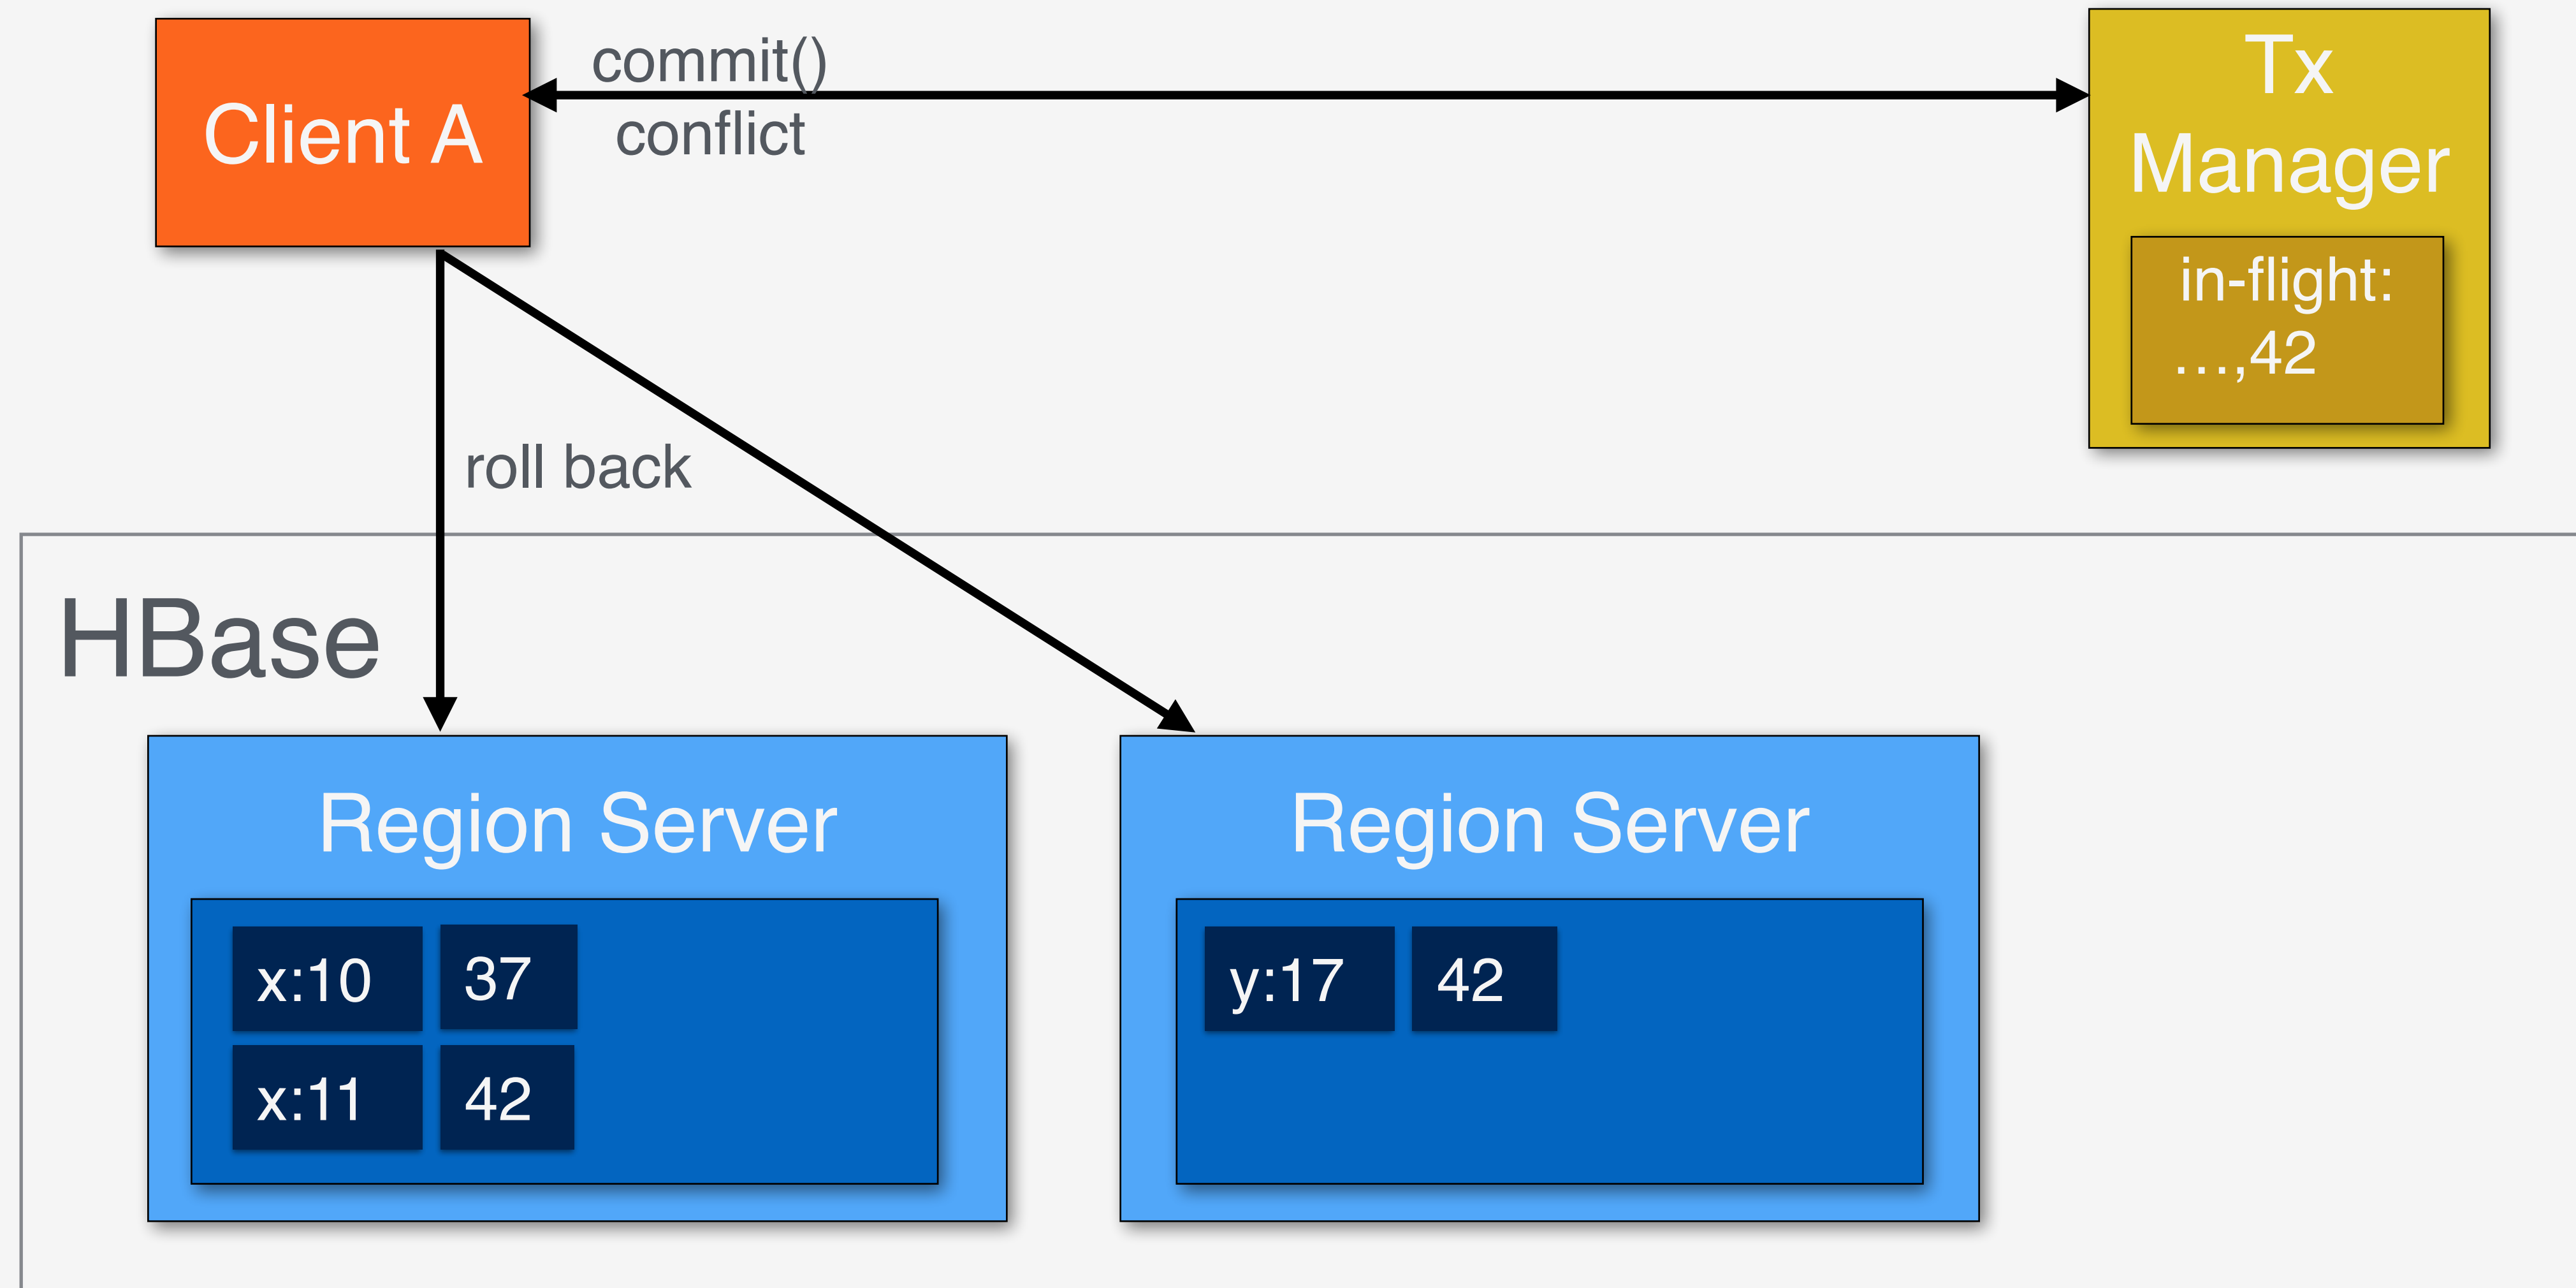
# Apache Tephra



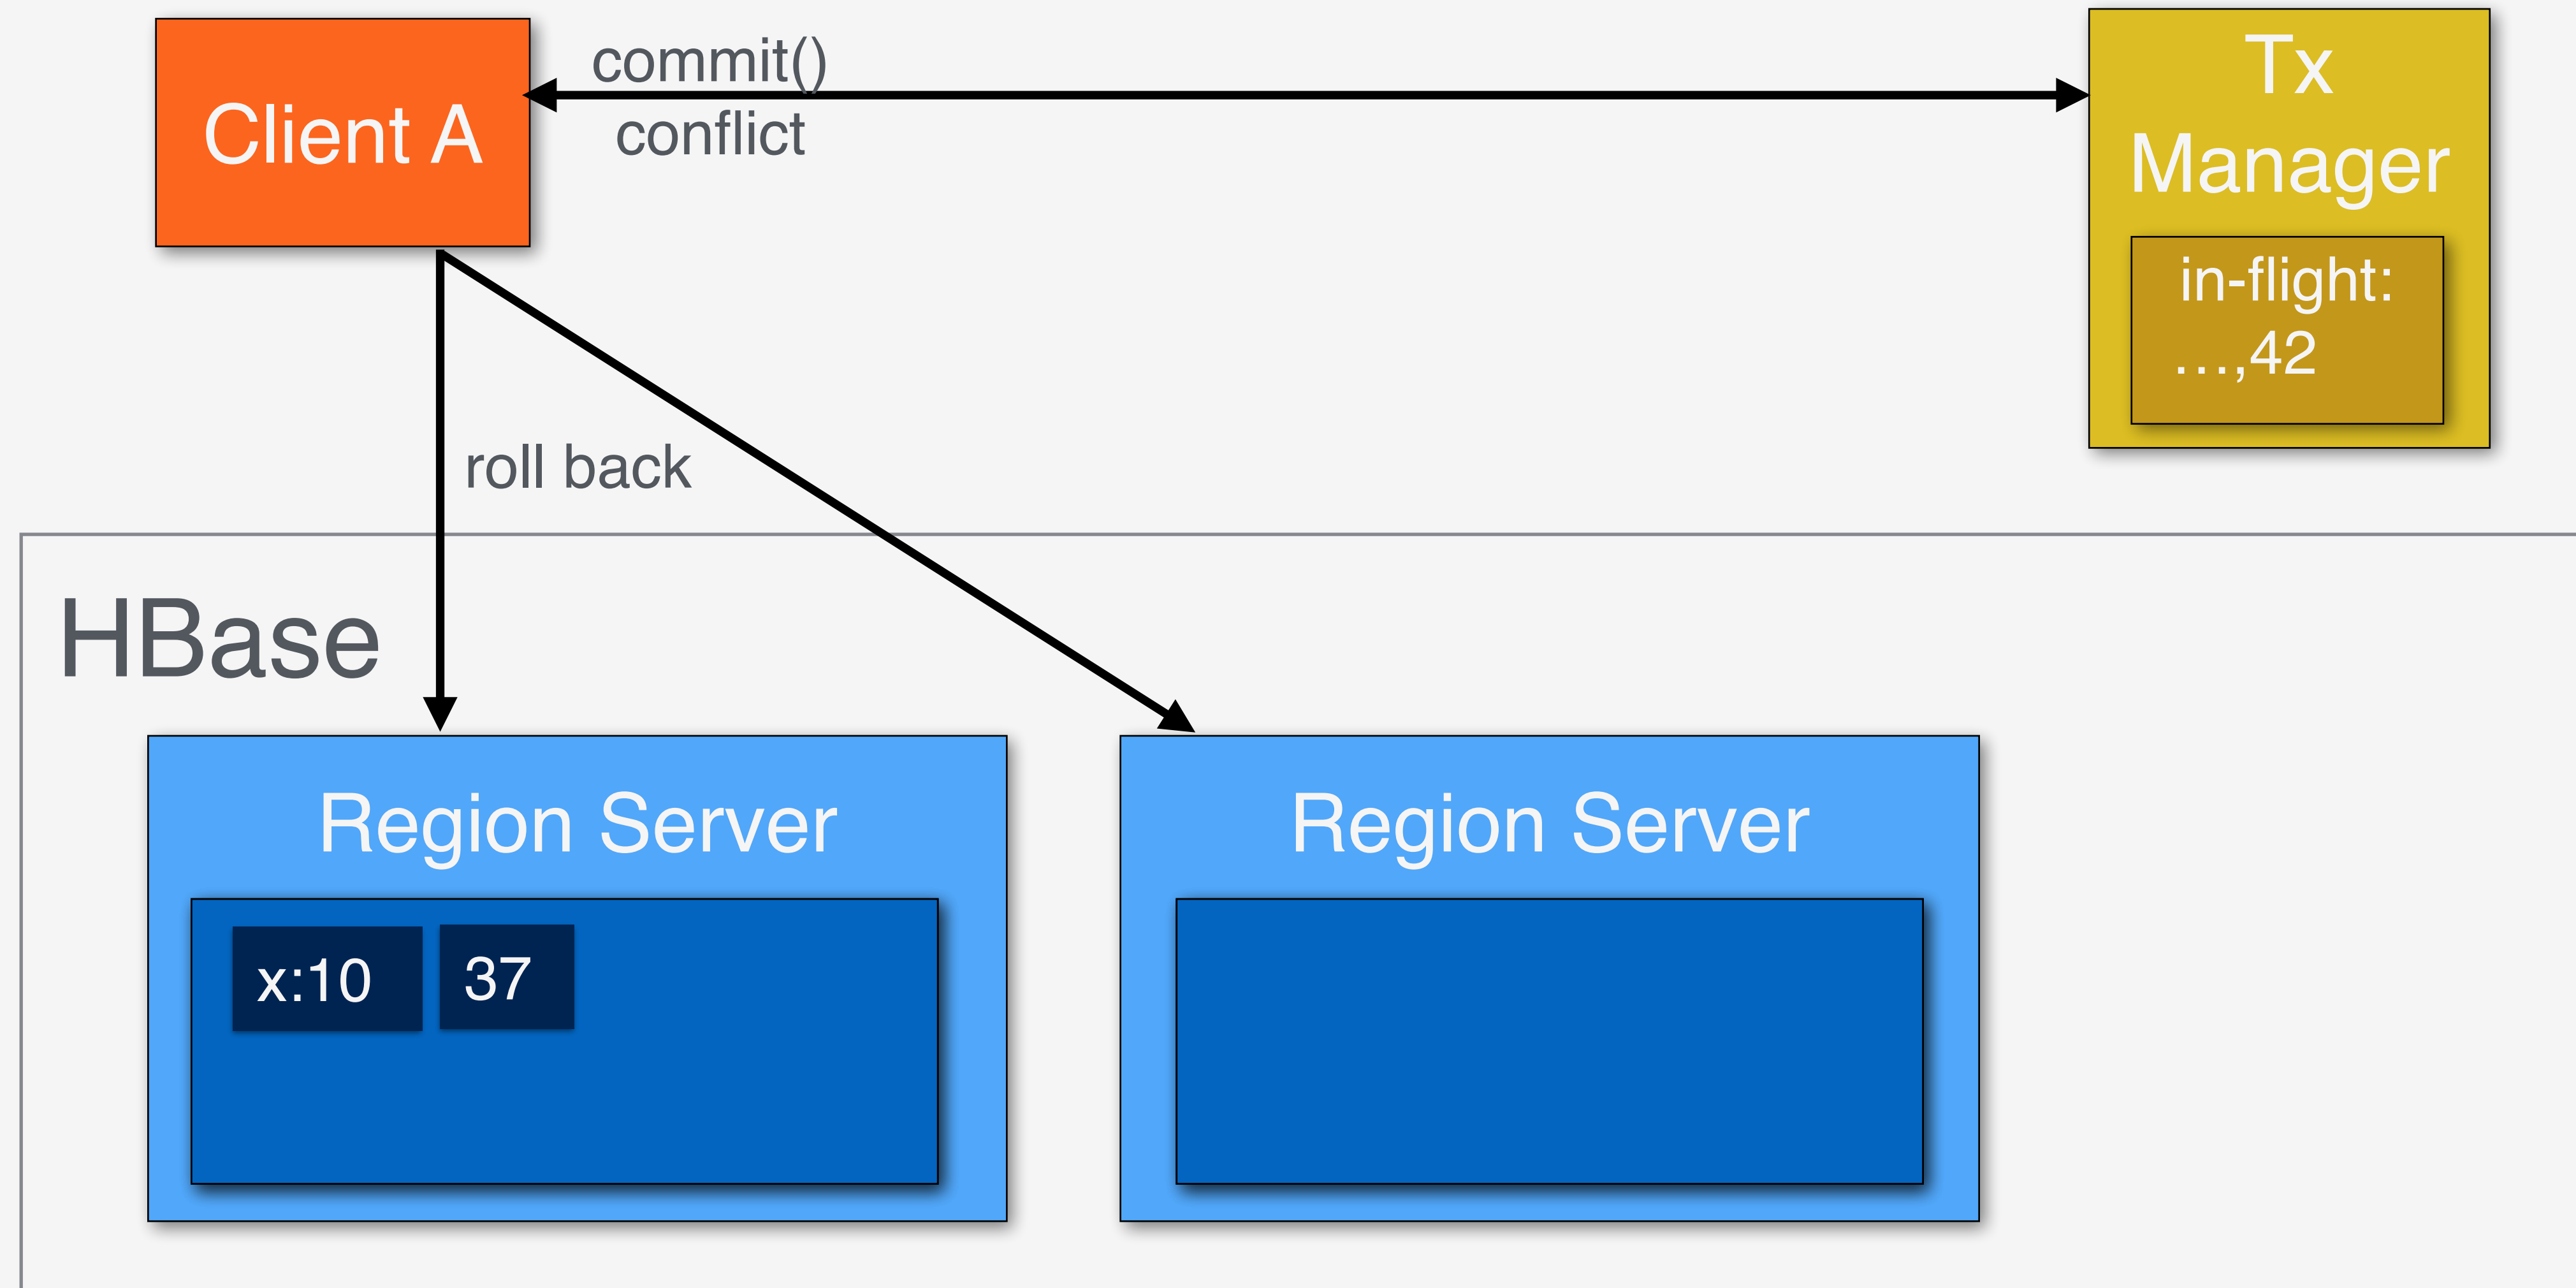
## HBase




# Apache Tephra



# Apache Tephra



# Apache Tephra



# Apache Tephra

Client A

Tx  
Manager

in-flight:  
...,42

## HBase

Region Server

x:10 37

x:11 42

Region Server

y:17 42

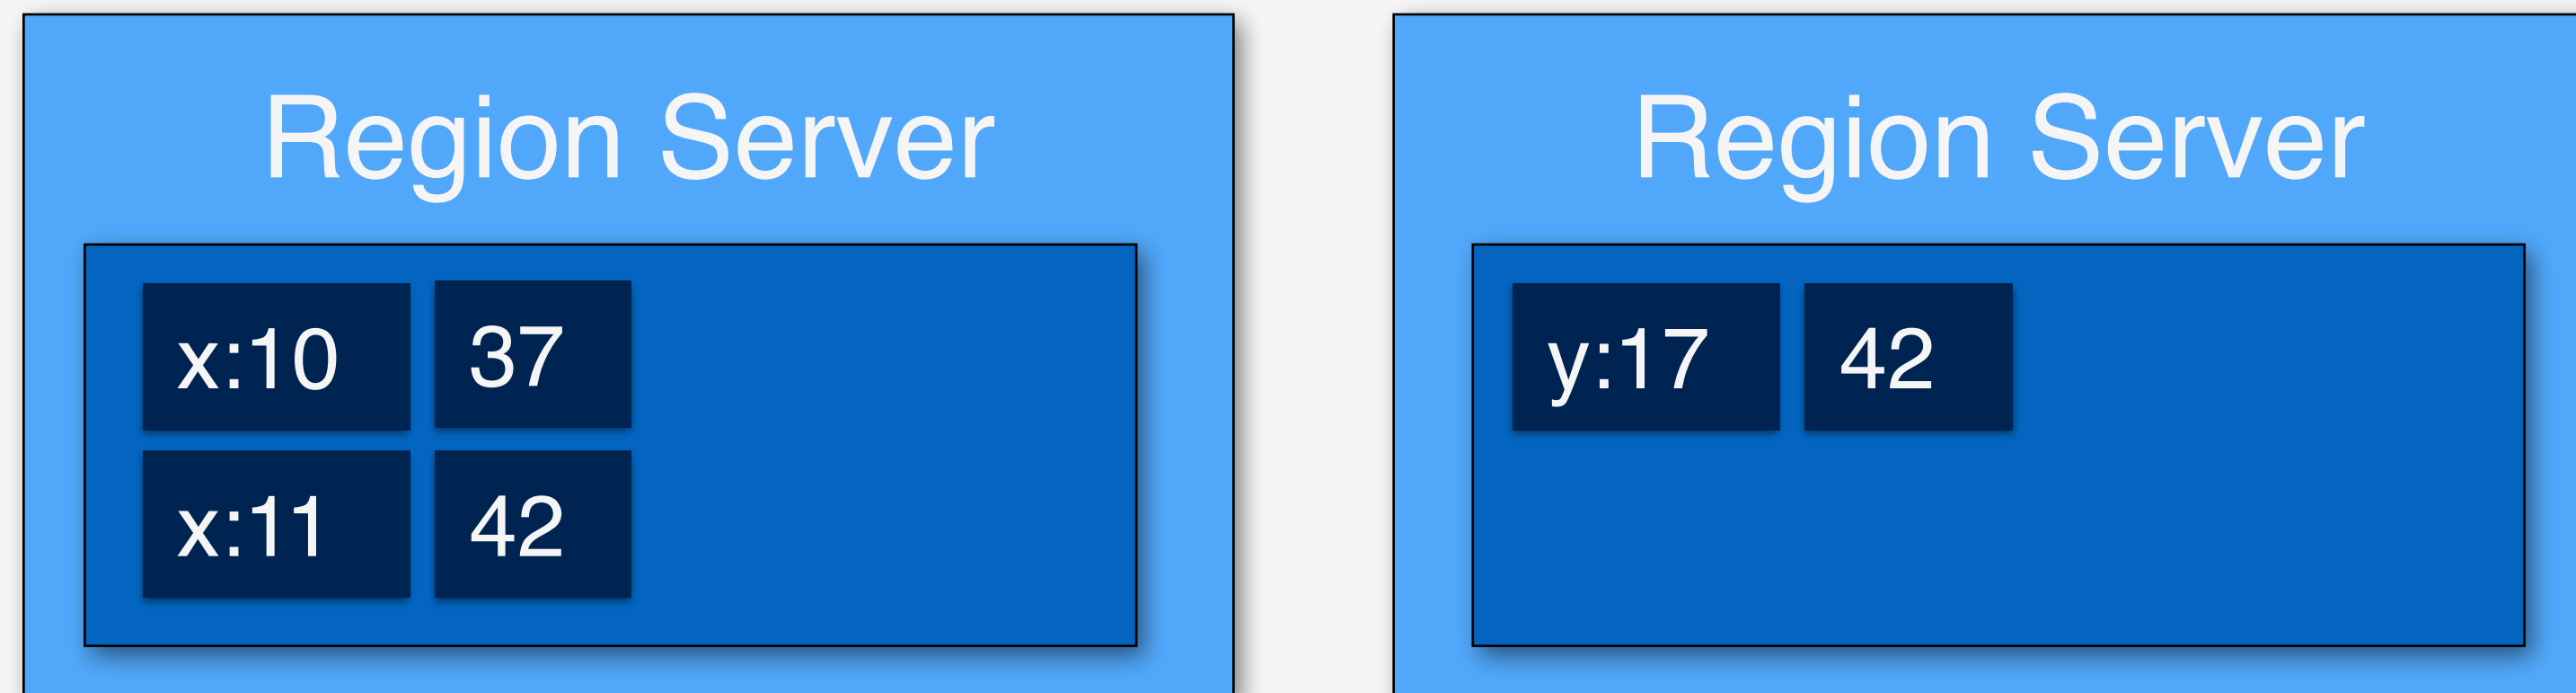




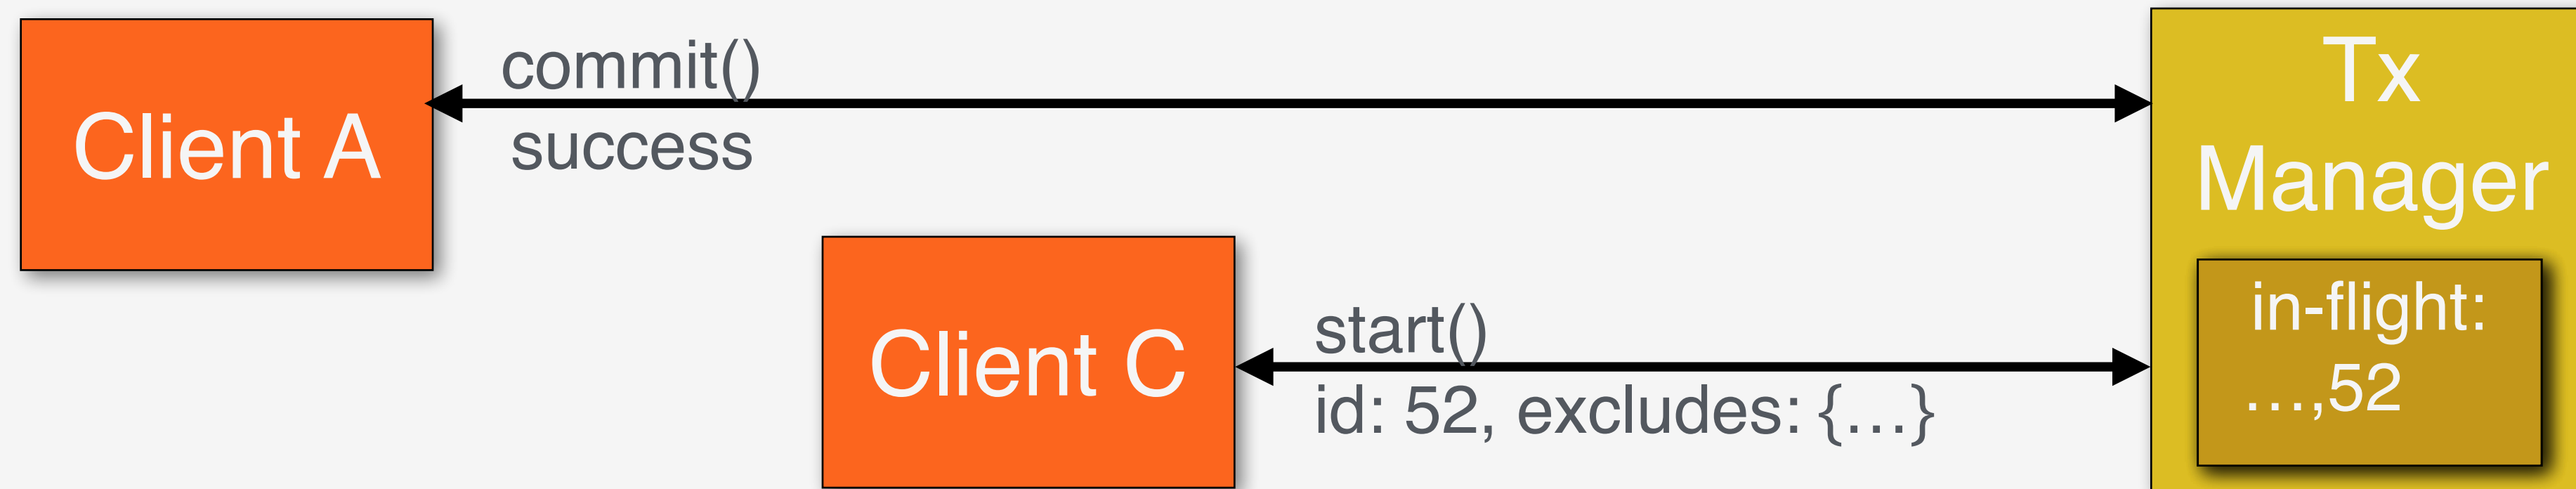
# Apache Tephra




## HBase



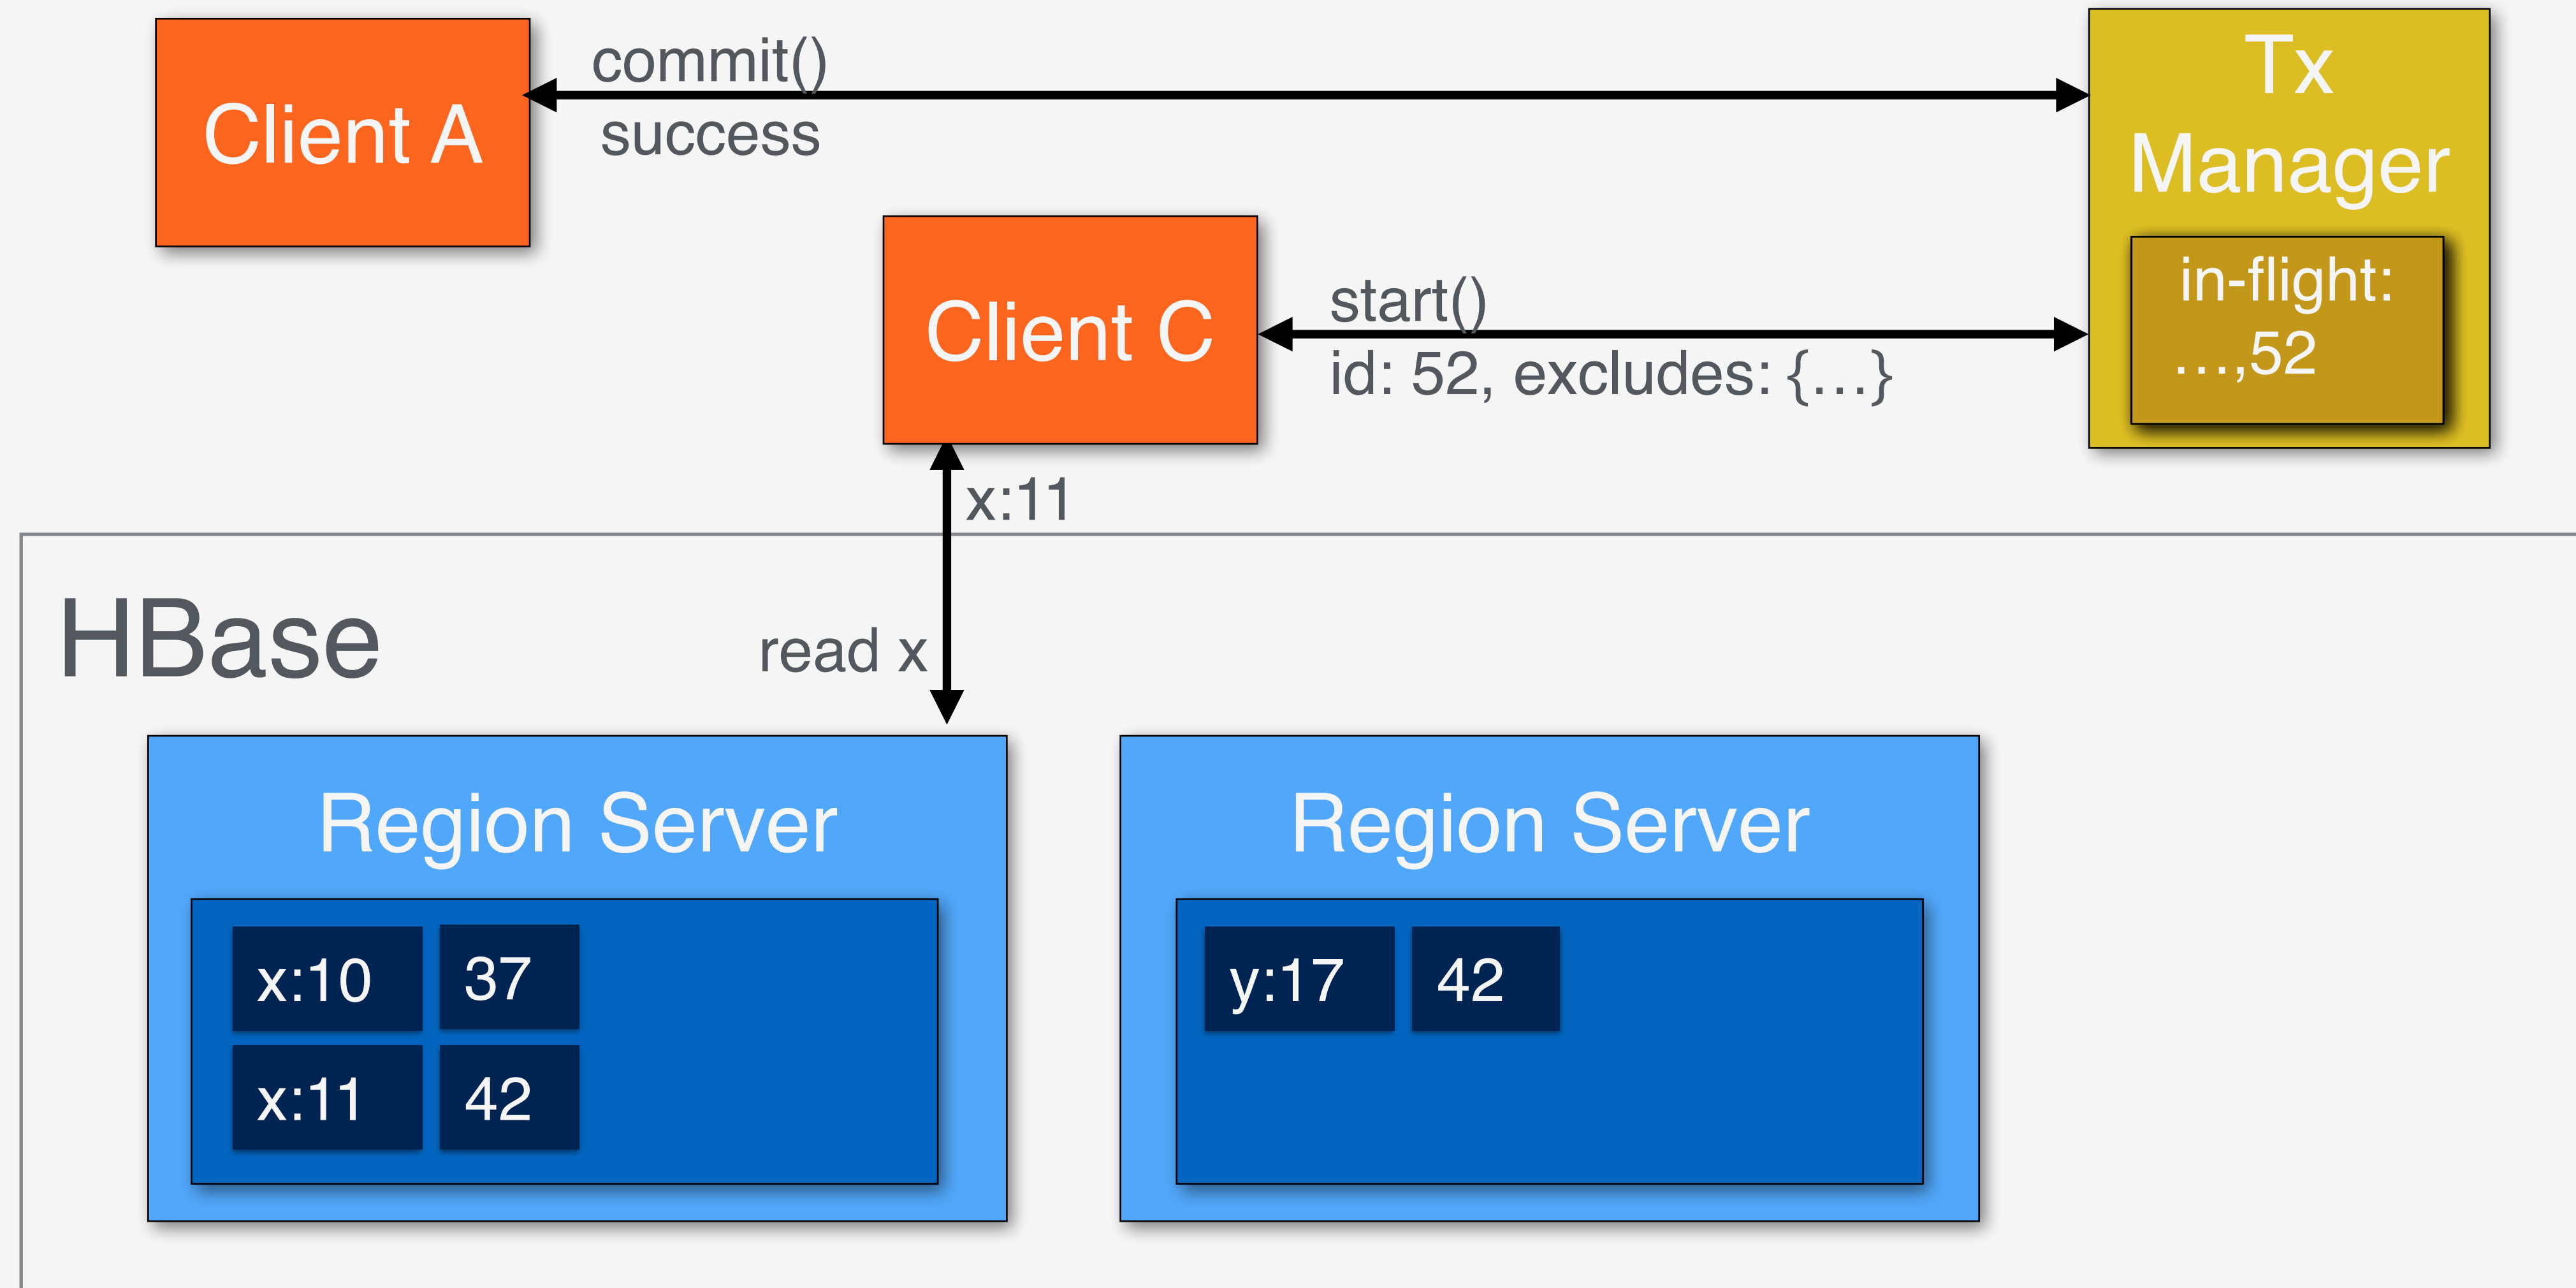
# Apache Tephra




## HBase



# Apache Tephra



# Apache Tephra



# Apache Tephra

- HBase coprocessors
  - For efficient visibility filtering (on region-server side)
  - For eliminating invalid cells on flush and compaction
- Programming Abstraction
  - TransactionalHTable:
    - Implements HTable interface
    - Existing code is easy to port
  - TransactionContext:
    - Implements transaction lifecycle



# Apache Tephra - Example

```
txTable = new TransactionAwareHTable(table);
txContext = new TransactionContext(txClient, txTable);
txContext.start();
try {
    // perform Hbase operations in txTable
    txTable.put(...);
    ...
} catch (Exception e) {
    // throws TransactionFailureException(e)
    txContext.abort(e);
}
// throws TransactionConflictException if so
txContext.finish();
```



# Apache Tephra - Strengths

- Compatible with existing, non-tx data in HBase
- Programming model
  - Same API as HTable, keep existing client code
- Conflict detection granularity
  - Row, Column, Off
  - Special “long-running tx” for MapReduce and similar jobs
- HA and Fault Tolerance
  - Checkpoints and WAL for transaction state, Standby Tx Manager
- Replication compatible
  - Checkpoint to HBase, use HBase replication
- Secure, Multi-tenant



# Apache Tephra - Not-So Strengths

- Exclude list can grow large over time
  - RPC, post-filtering overhead
  - Solution: Invalid tx pruning on compaction - complex!
- Single Transaction Manager
  - performs all lifecycle state transitions, including conflict detection
  - conflict detection requires lock on the transaction state
  - becomes a bottleneck
  - Solution: distributed Transaction Manager with consensus protocol





# Apache Trafodion

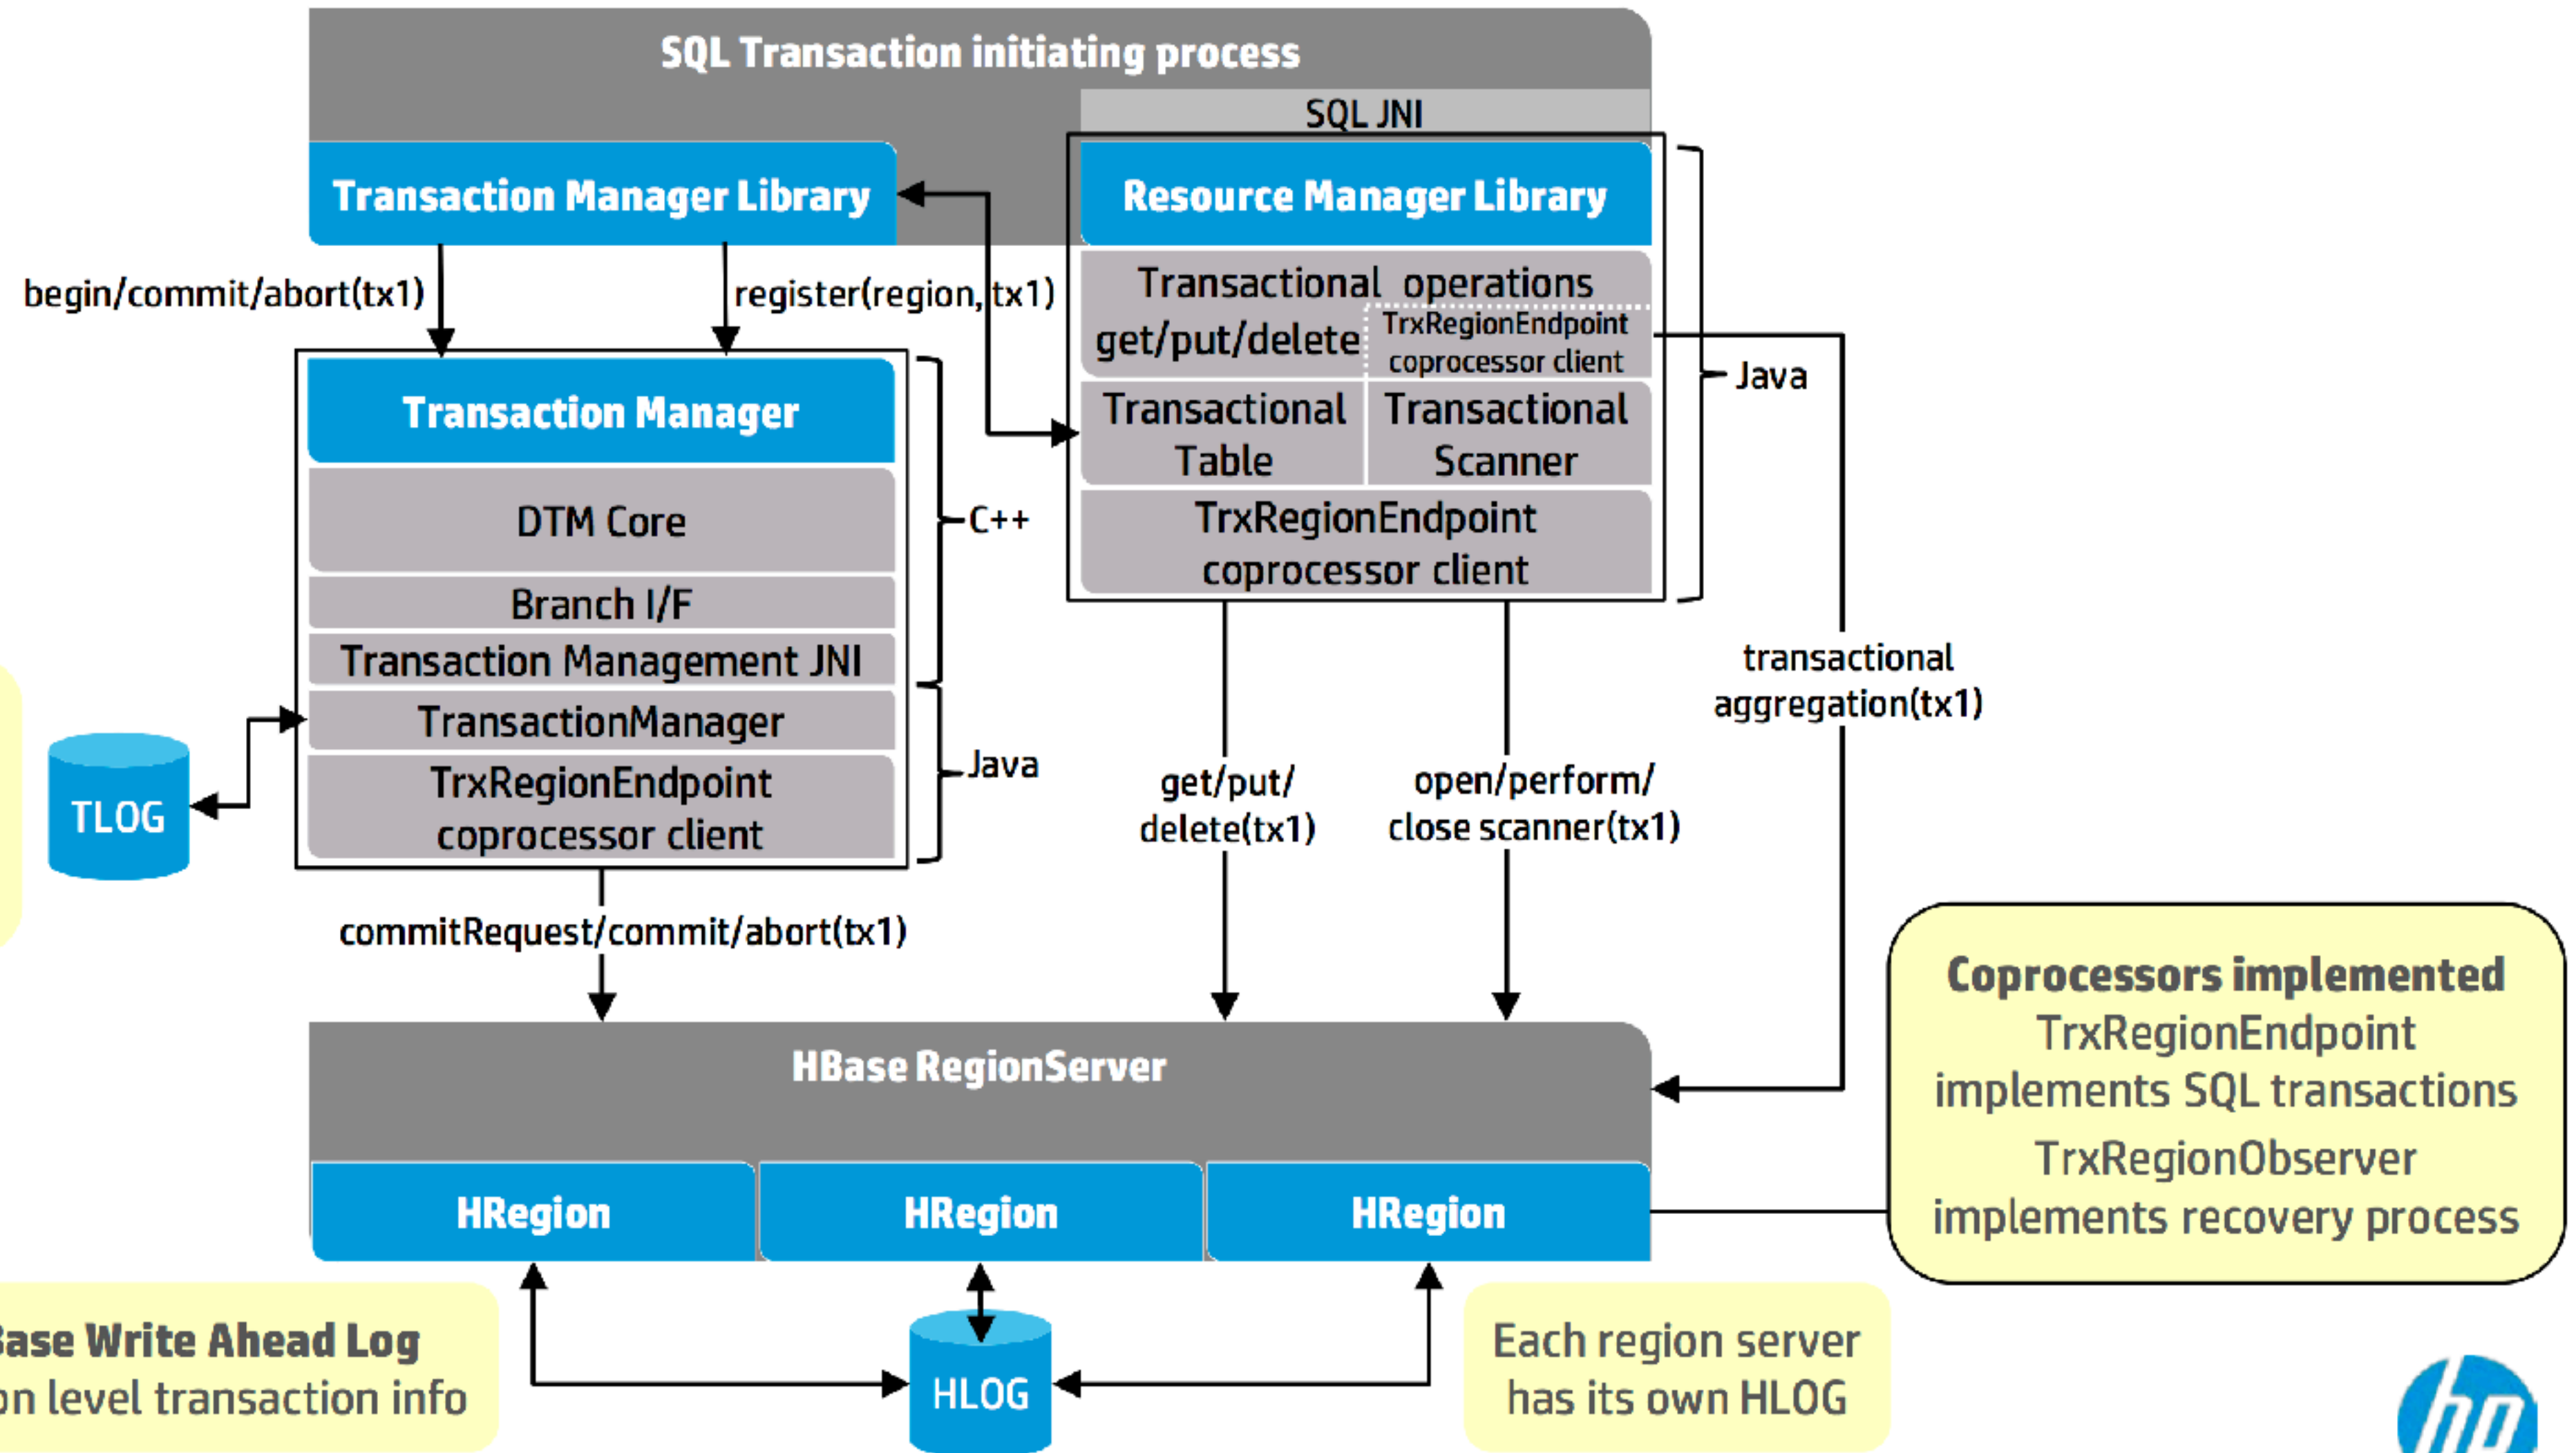
- A complete distributed database (RDBMS)
  - transaction system is not available by itself
  - APIs: jdbc, SQL
- Inspired by original HBase TRX (transactional region server)
  - migrated transaction logic into coprocessors
  - coprocessors cache in-flight data in-memory
  - transaction state (change sets) in coprocessors
  - conflict detection with 2-phase commit
- Transaction Manager
  - orchestrates transaction lifecycle across involved region servers
  - multiple instances, but one per client



(incubating)



# Apache Trafodion



# Apache Trafodion

Client A

Tx  
Manager

in-flight:  
...

## HBase

Region Server

x:10


Region Server



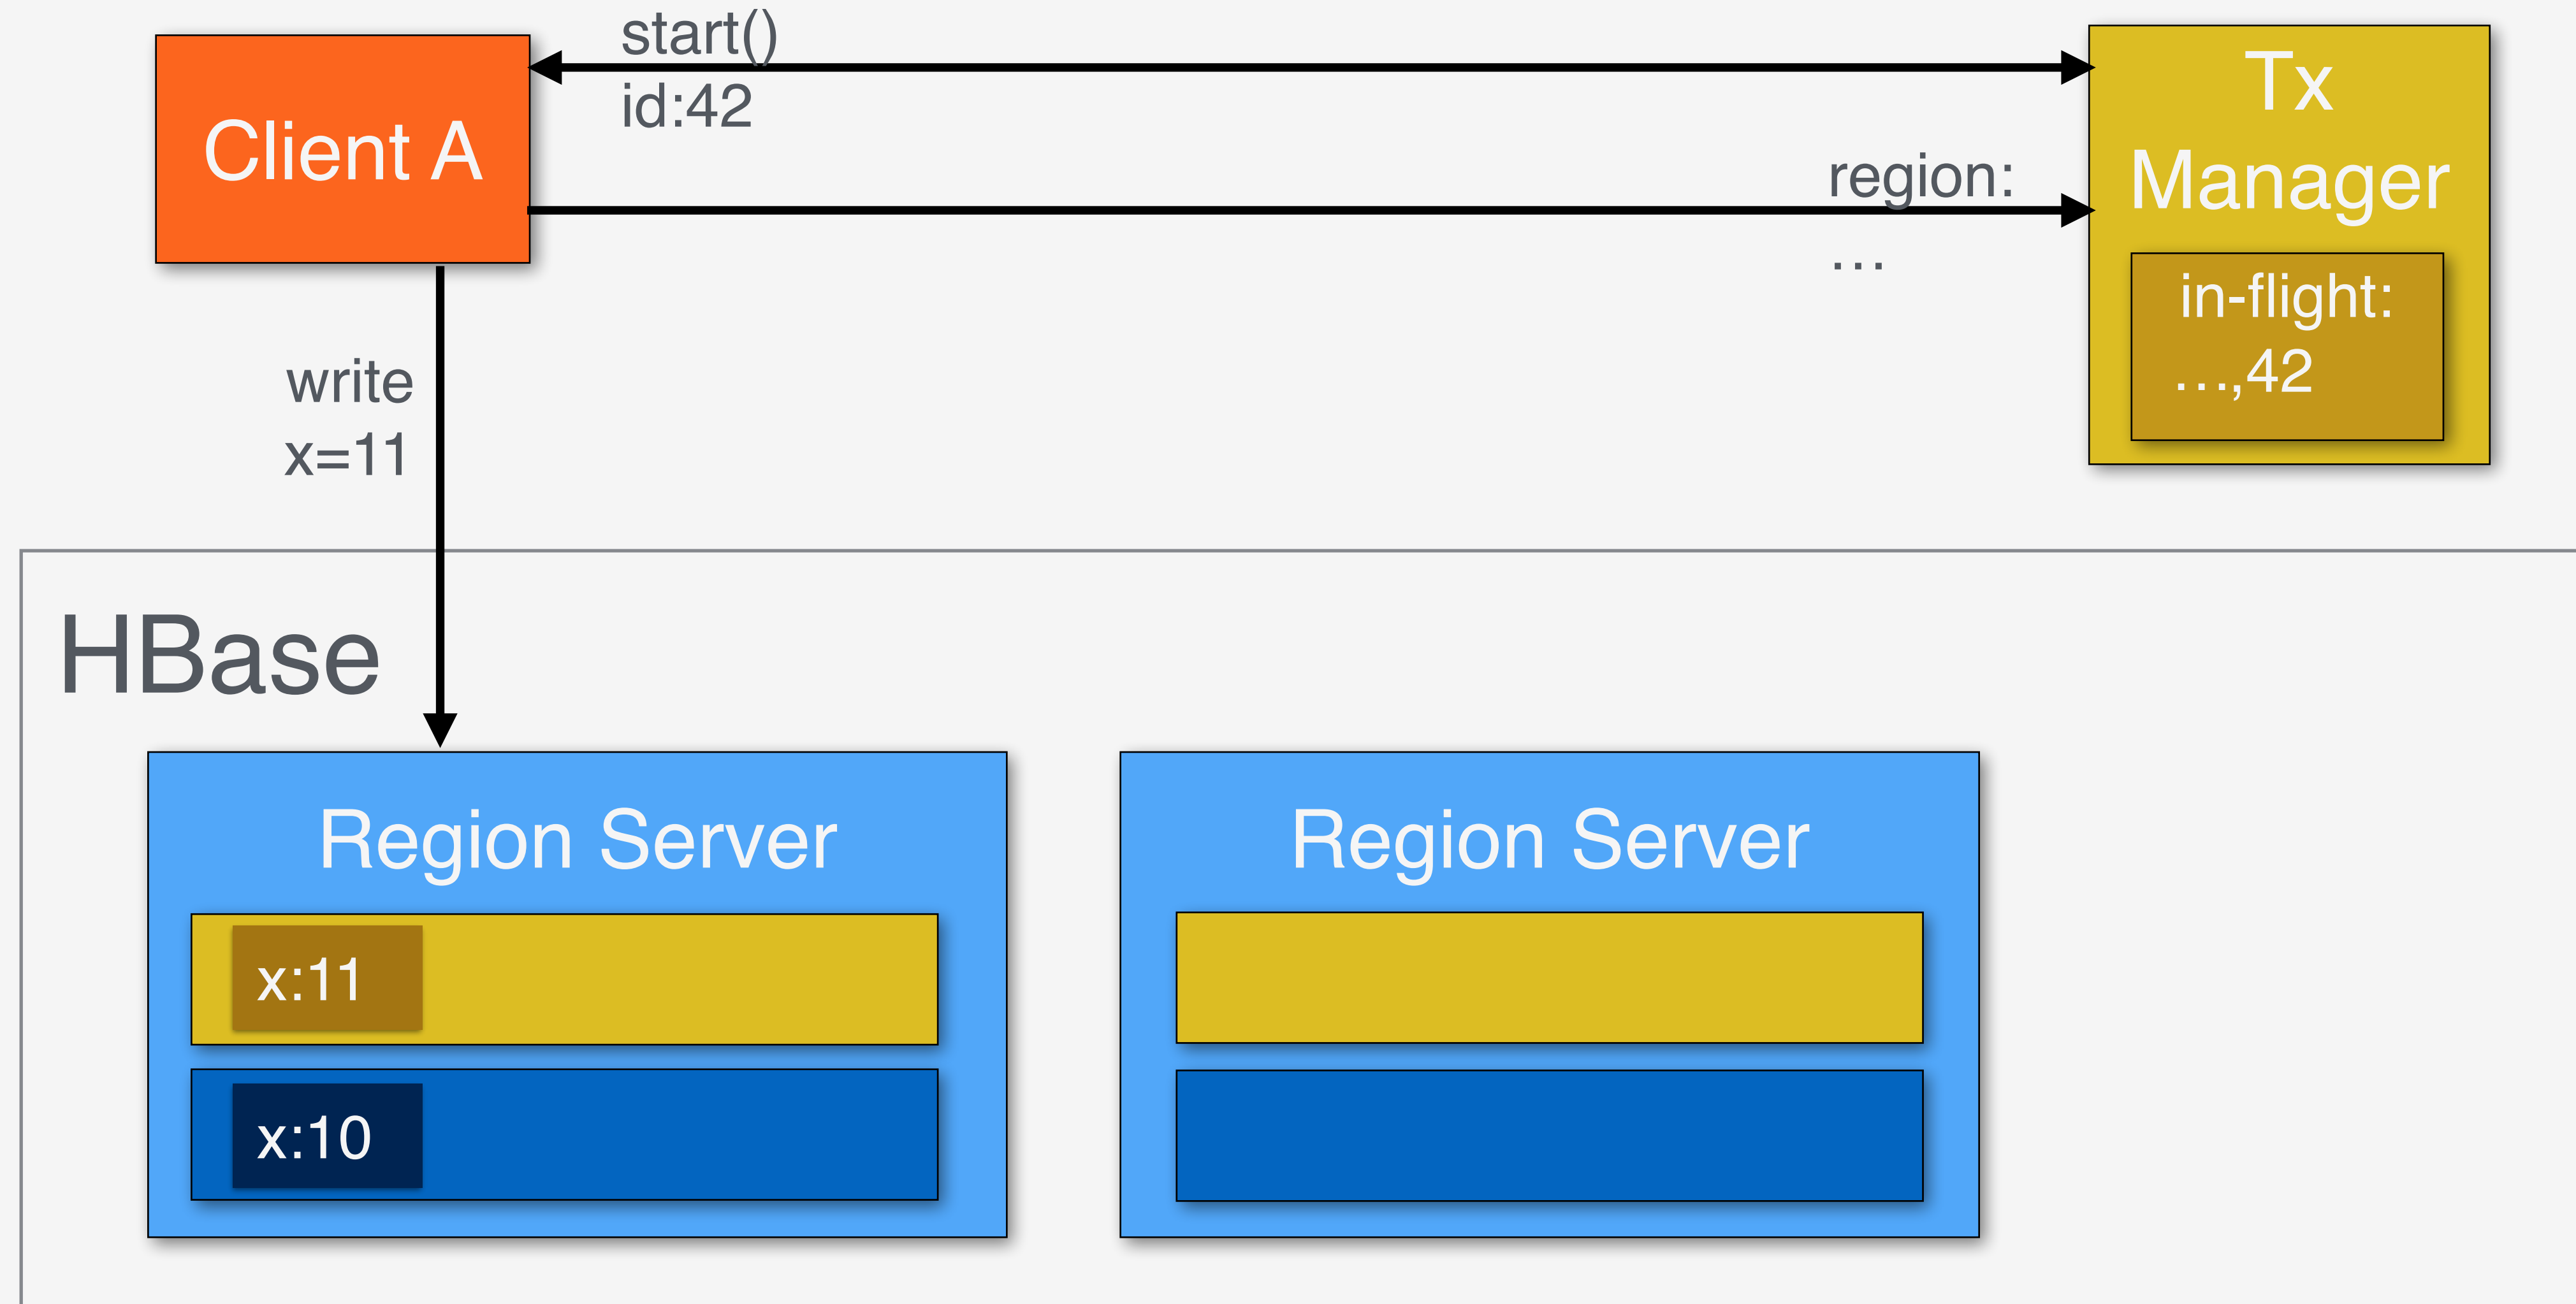
# Apache Trafodion




## HBase



# Apache Trafodion



# Apache Trafodion



# Apache Trafodion

Client B

Tx  
Manager

in-flight:  
...,42

## HBase

Region Server

x:11


x:10

Region Server


y:17



# Apache Trafodion

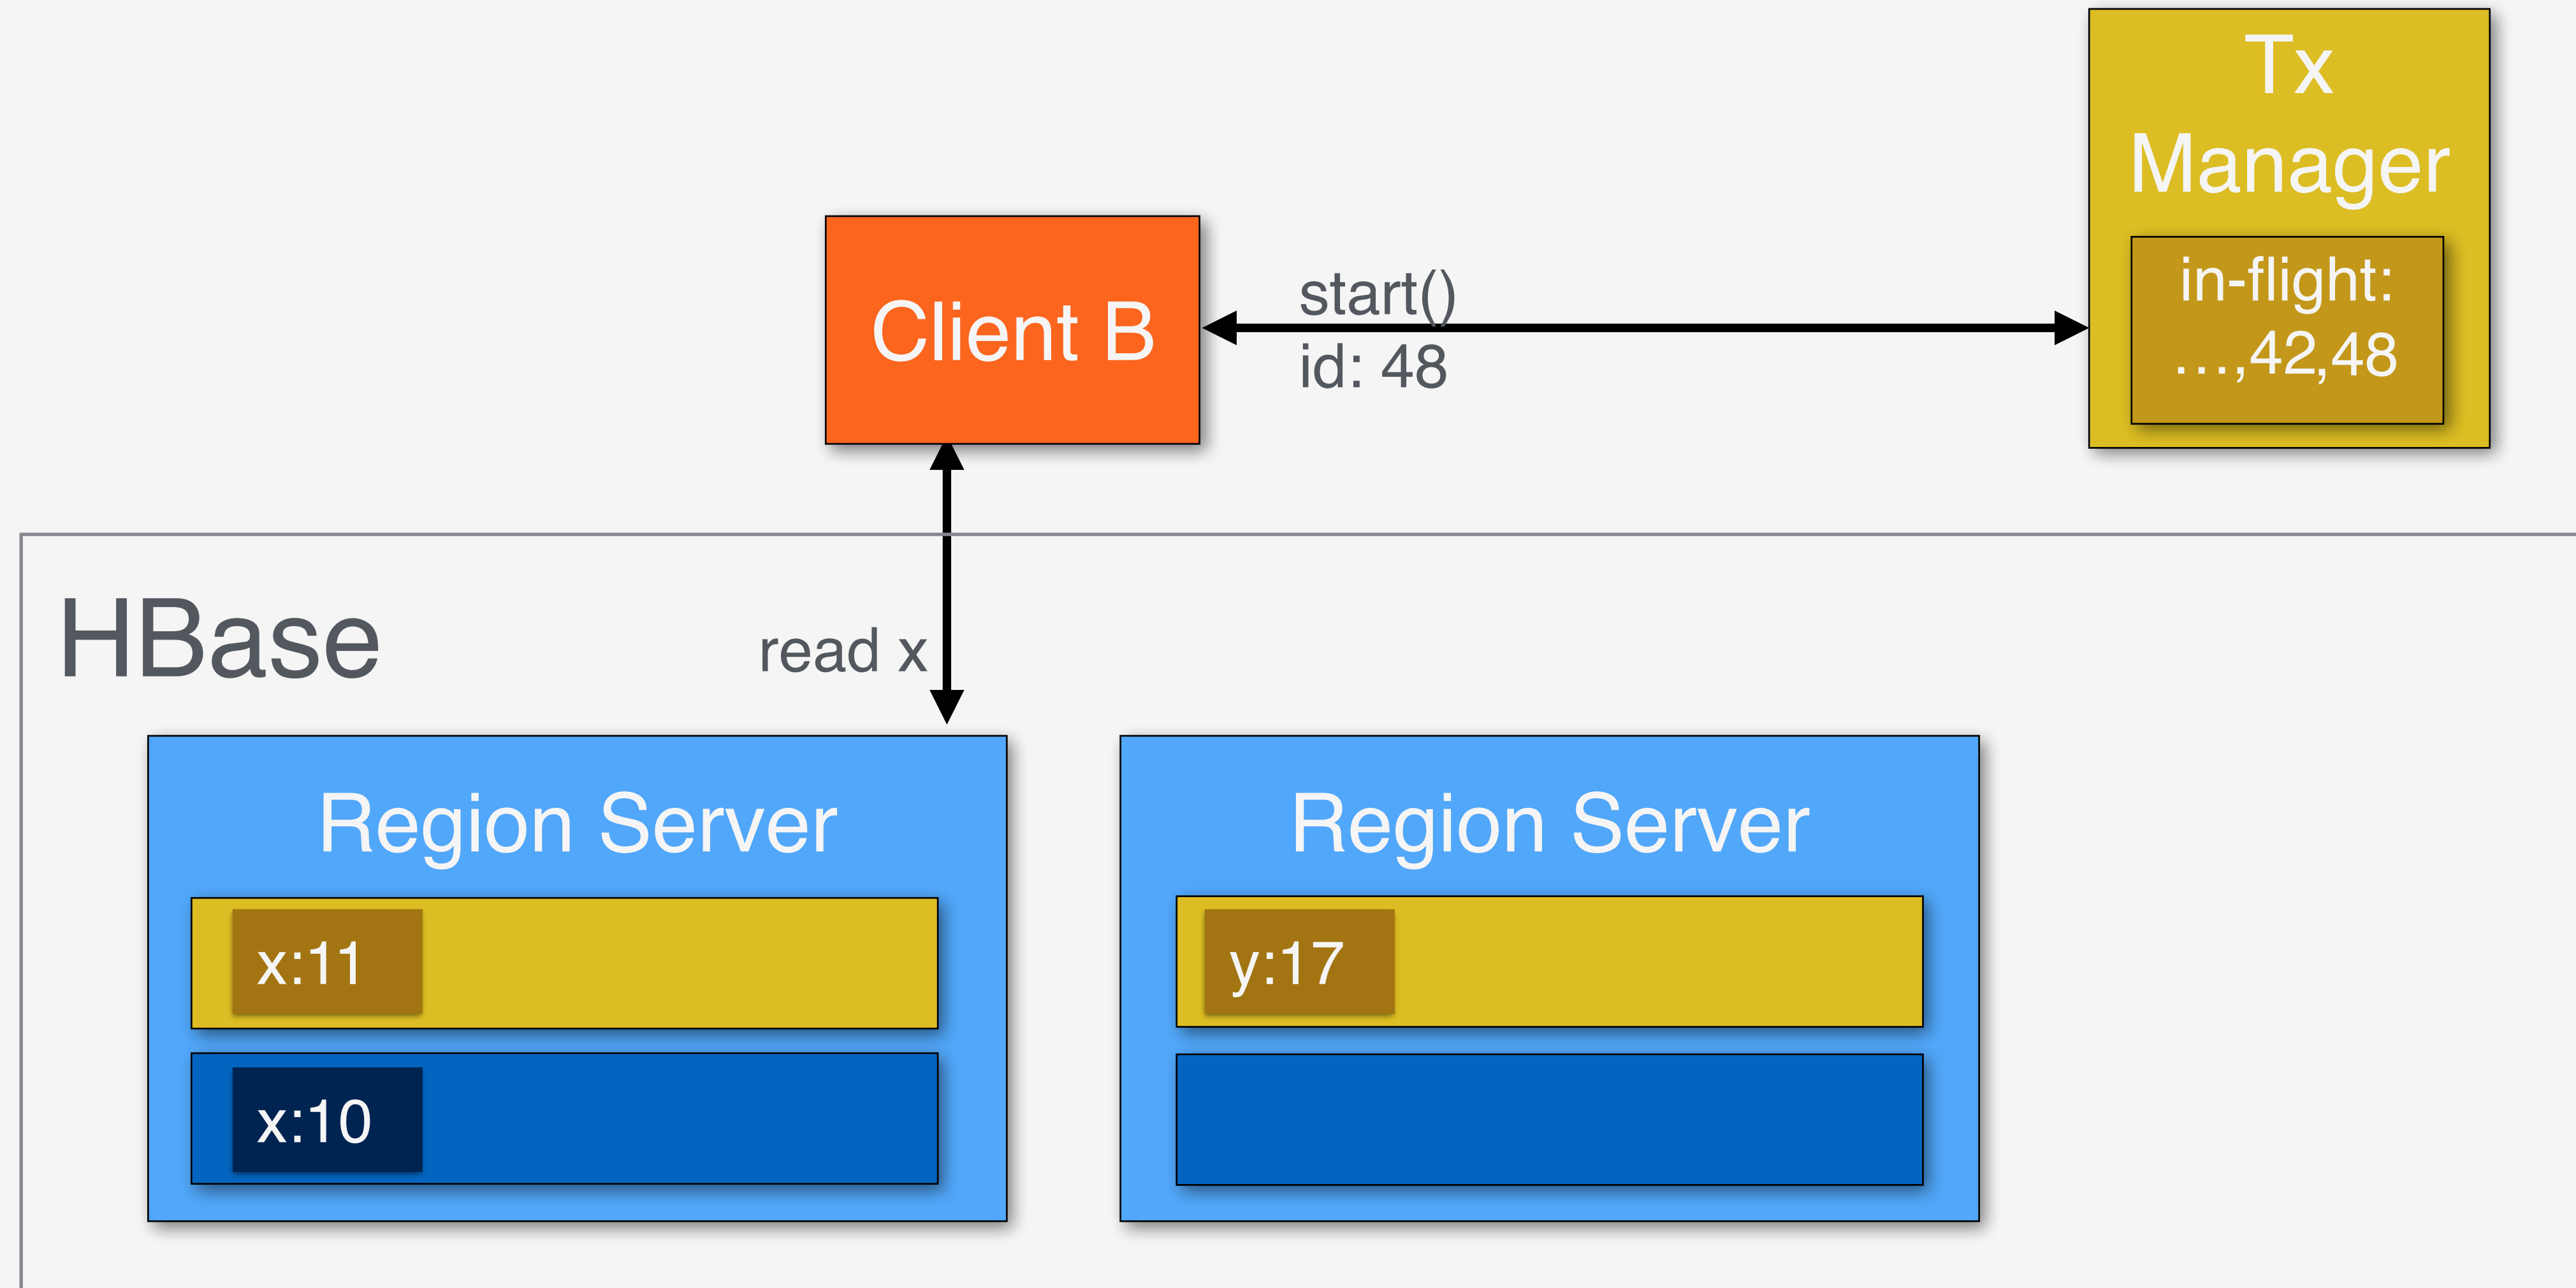


## HBase

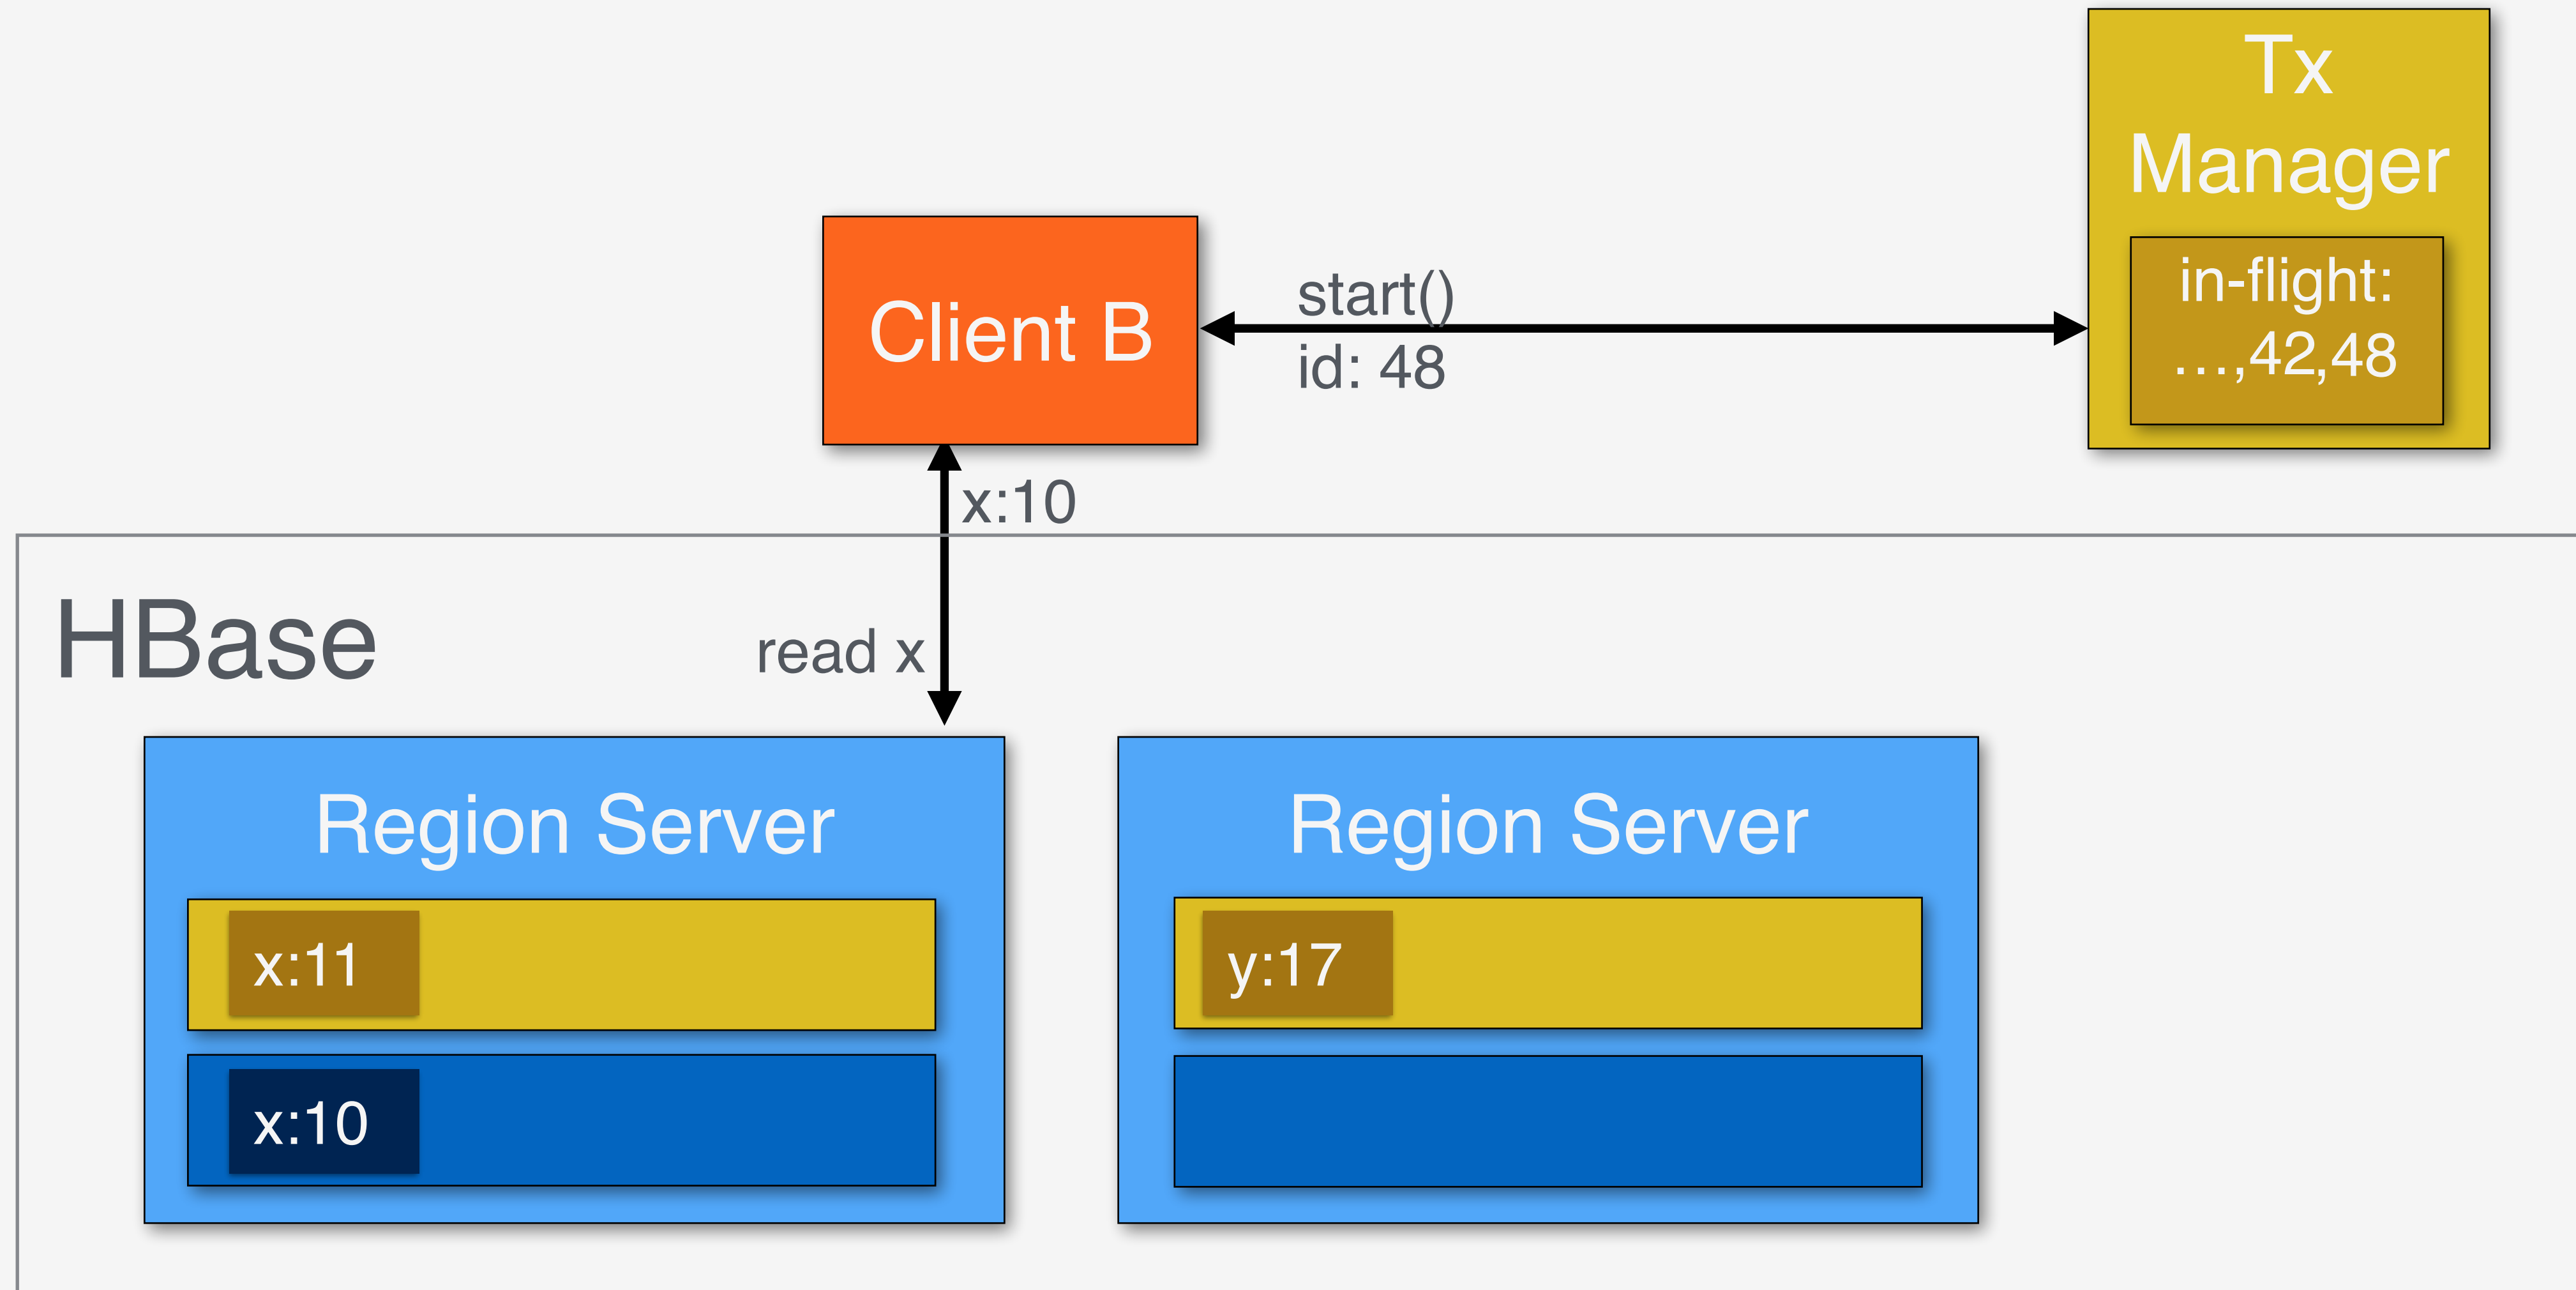




# Apache Trafodion



# Apache Trafodion



# Apache Trafodion

Client A

Tx  
Manager

in-flight:  
...,42

## HBase

Region Server

x:11

x:10

Region Server


y:17




# Apache Trafodion



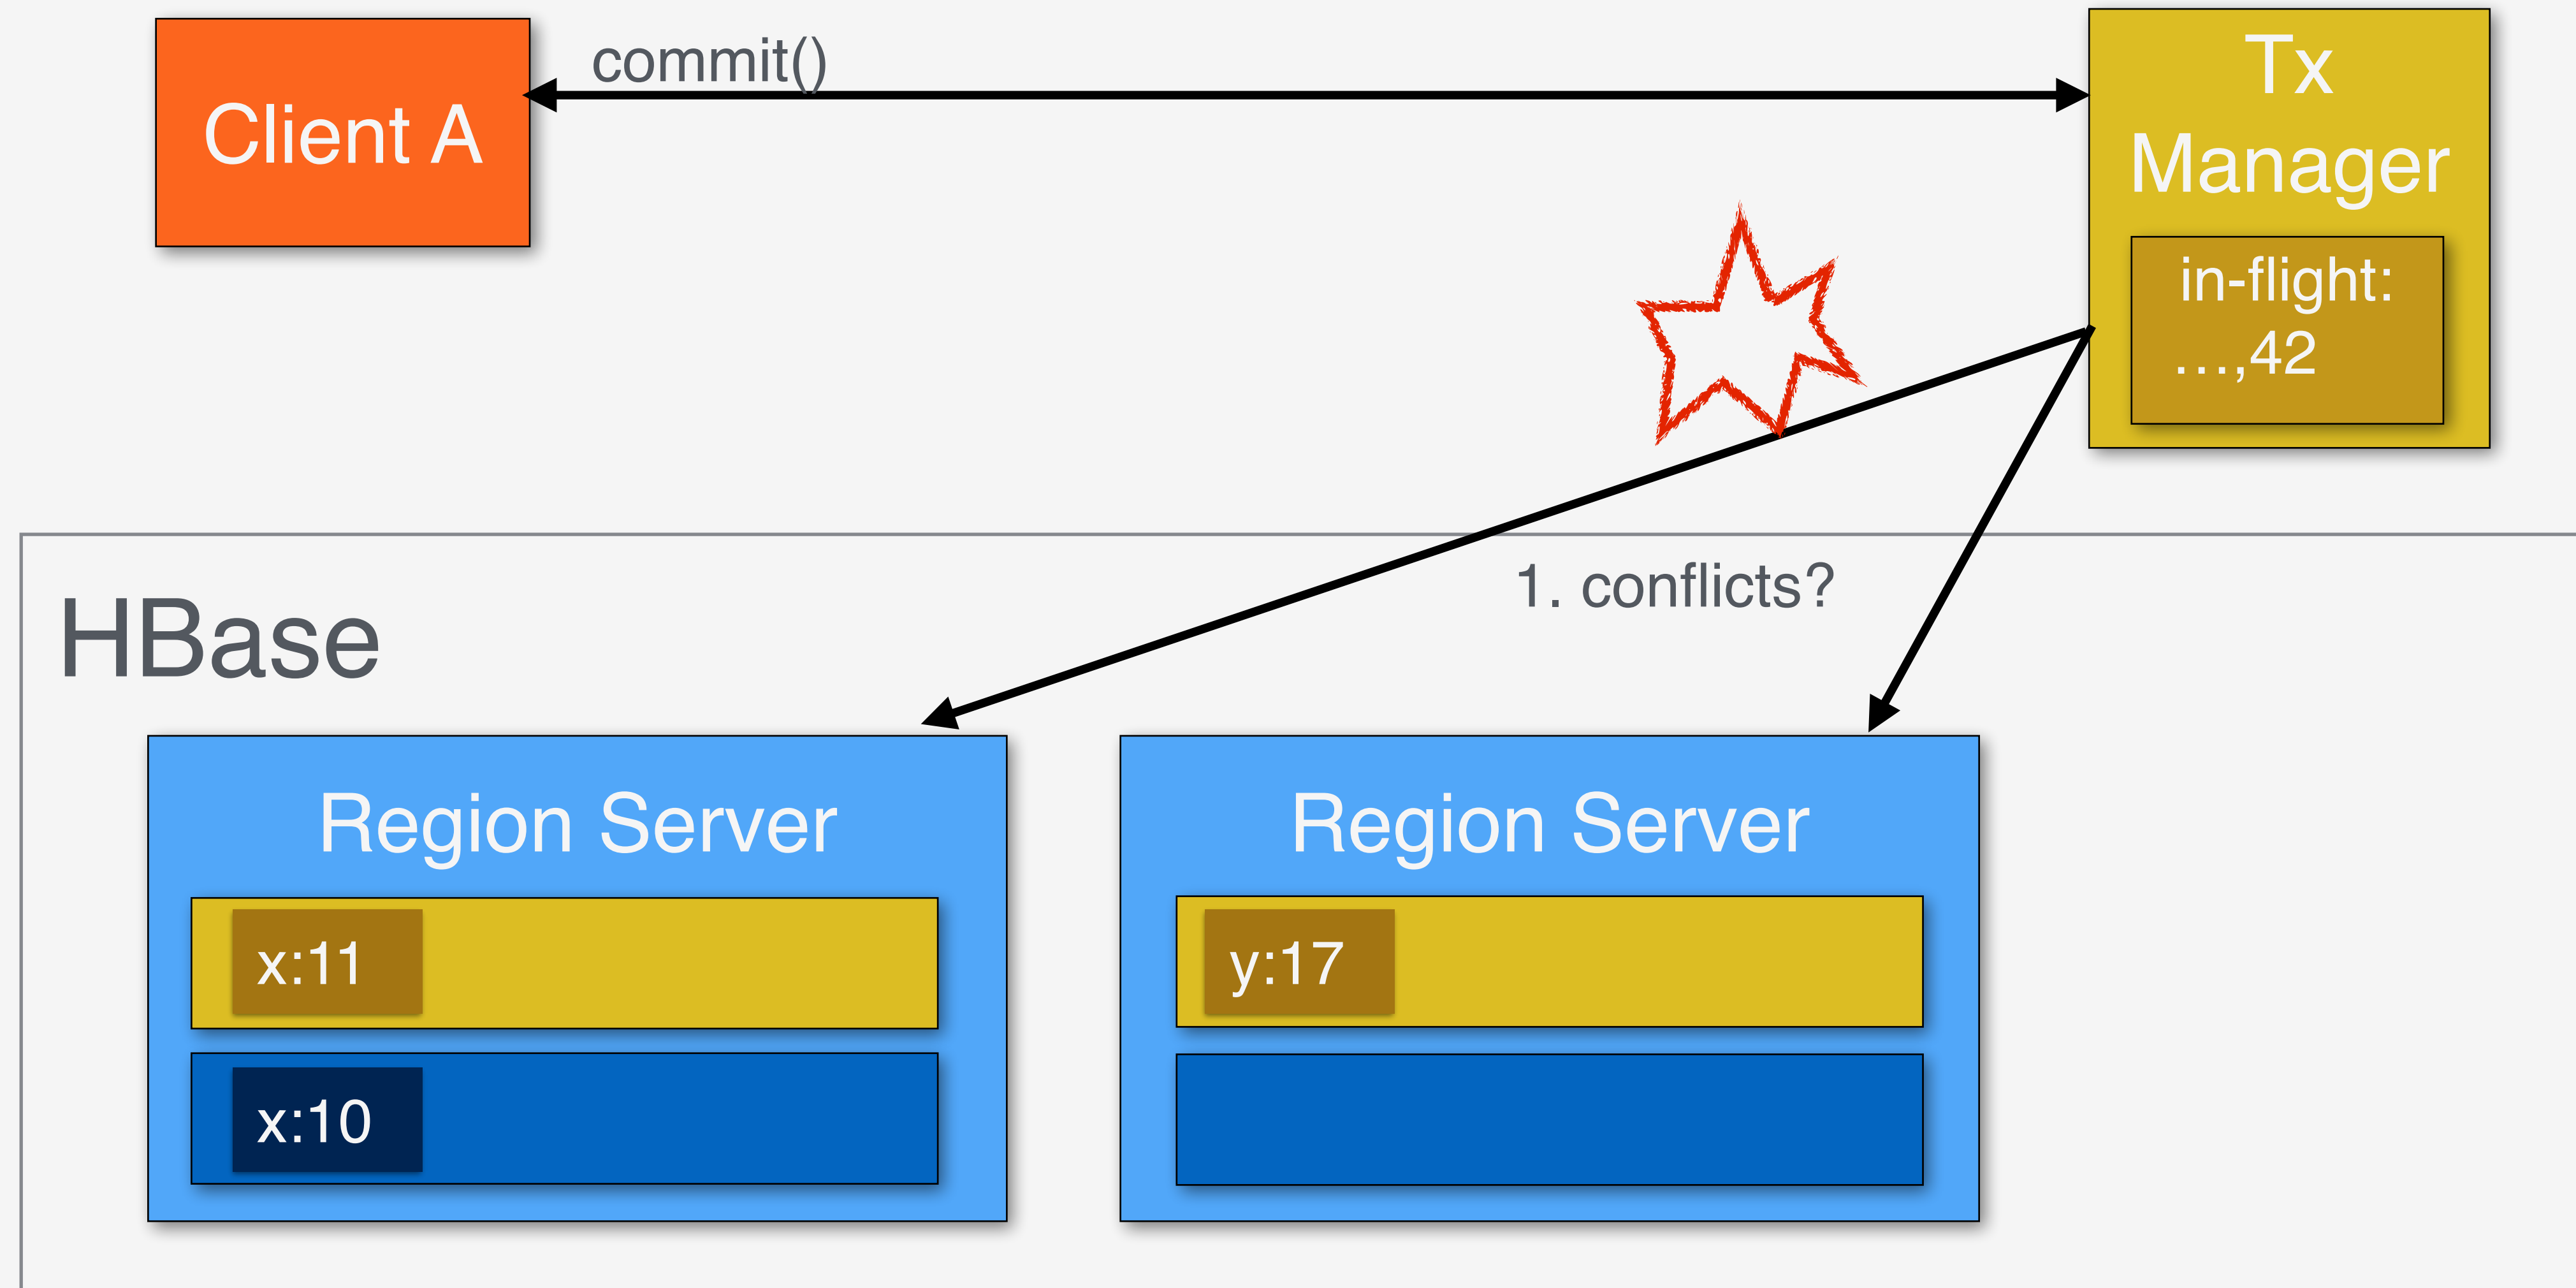
## HBase



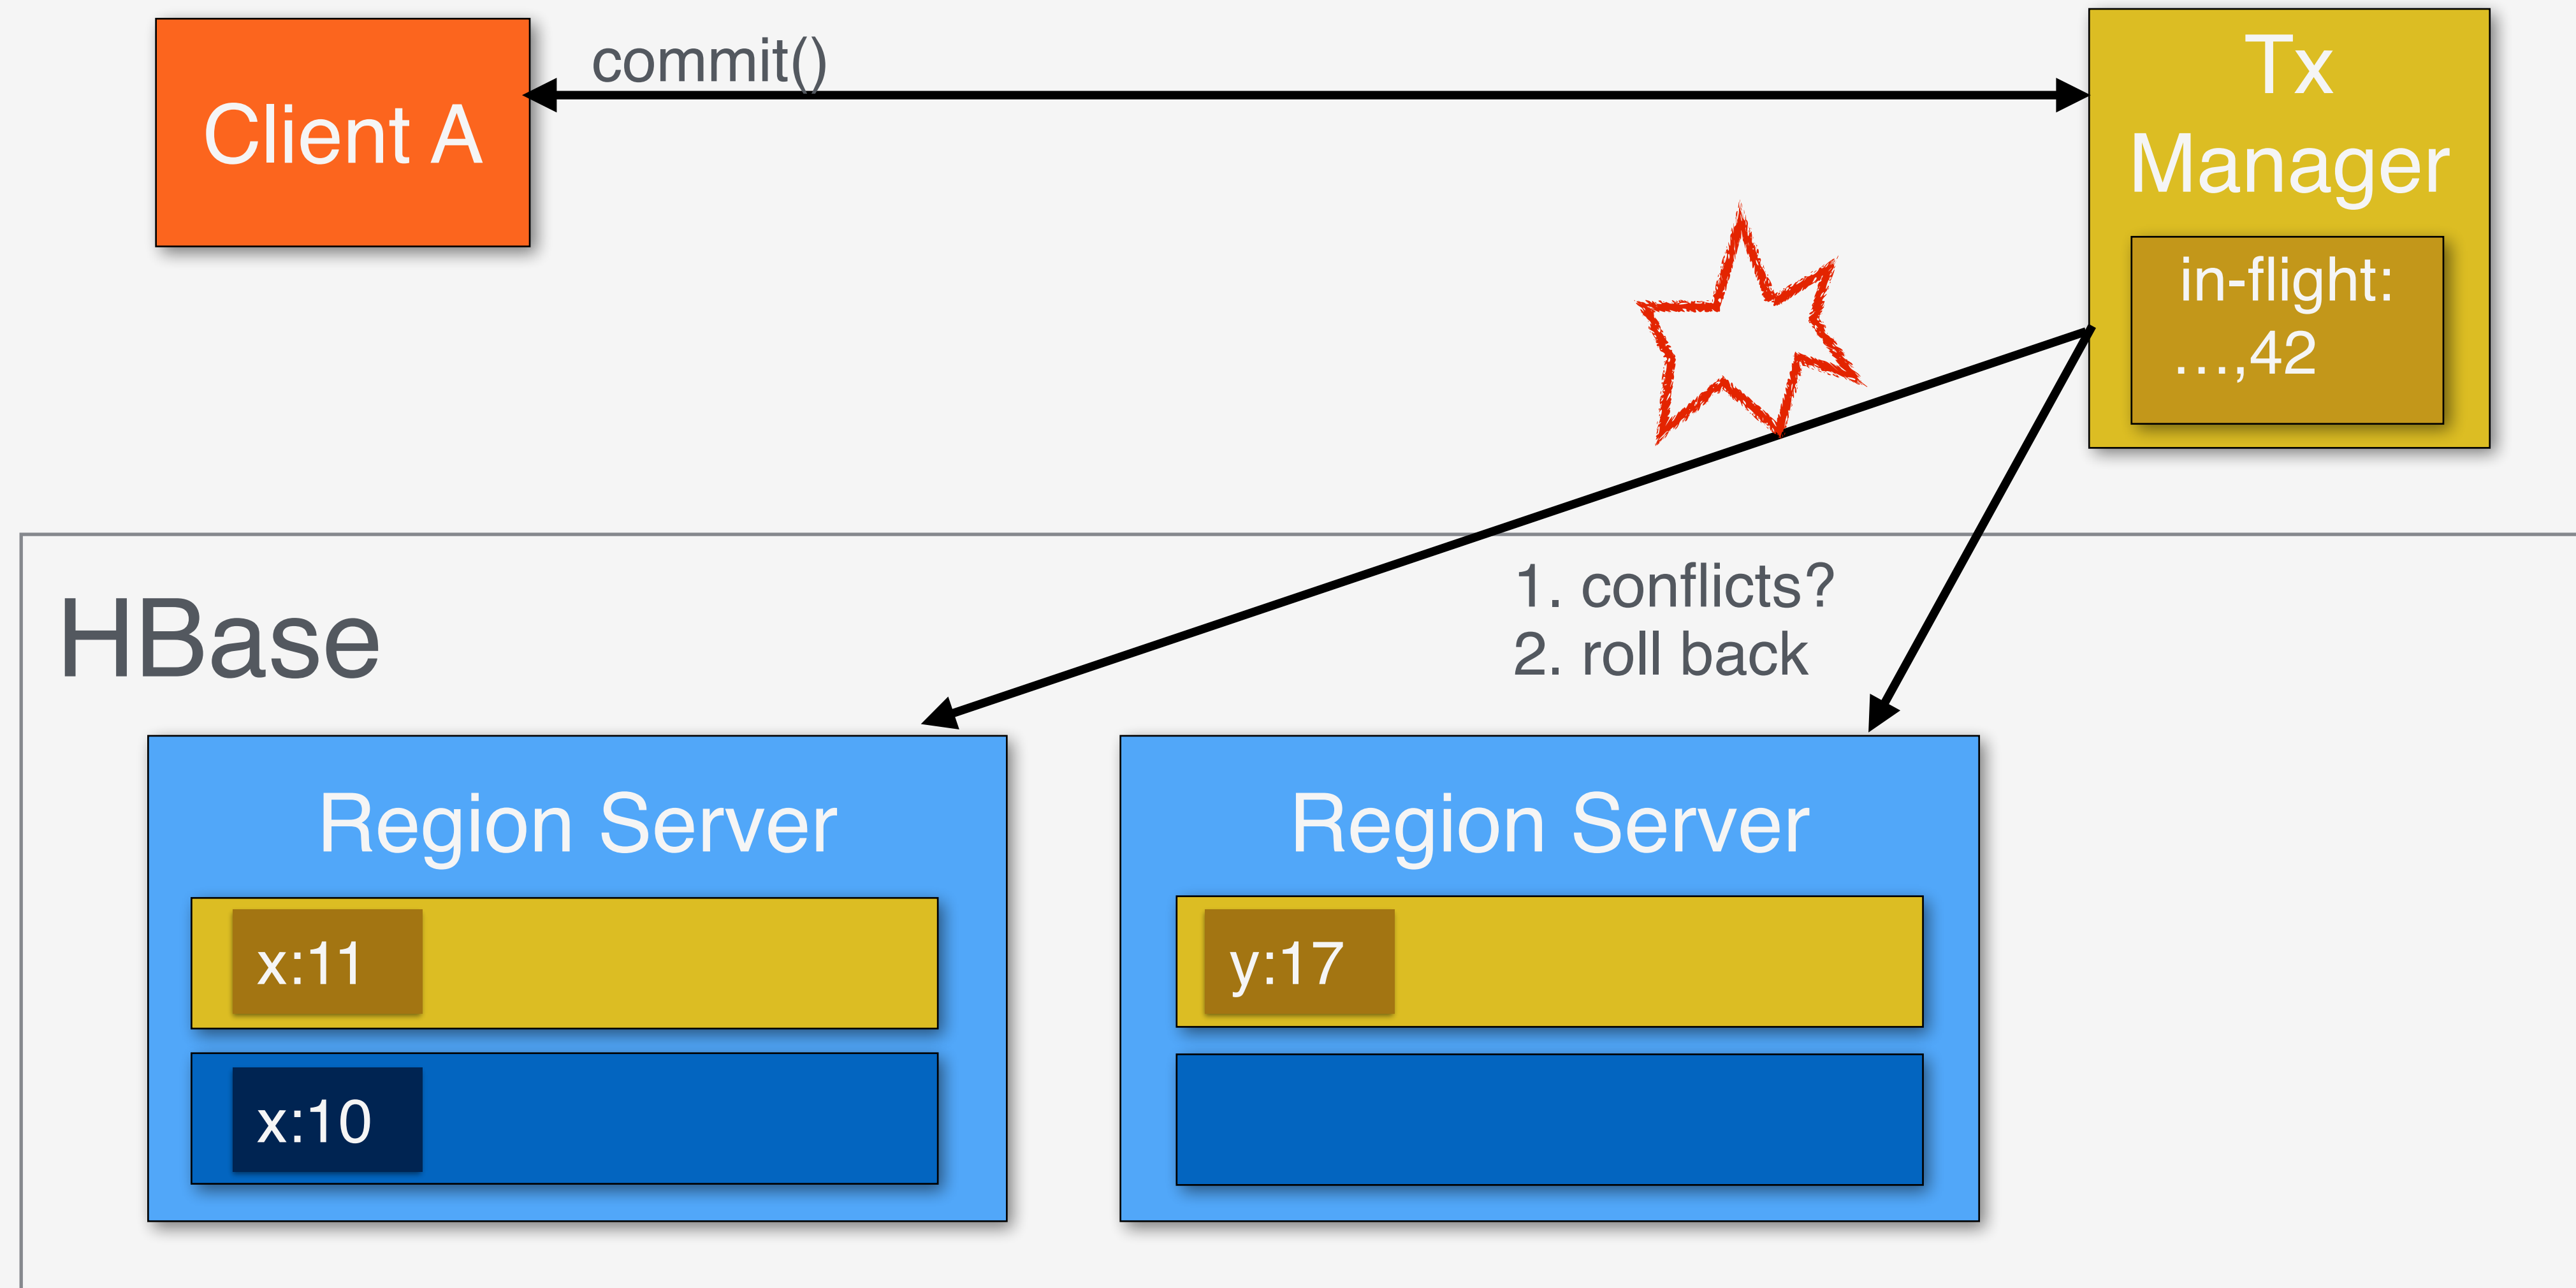
# Apache Trafodion



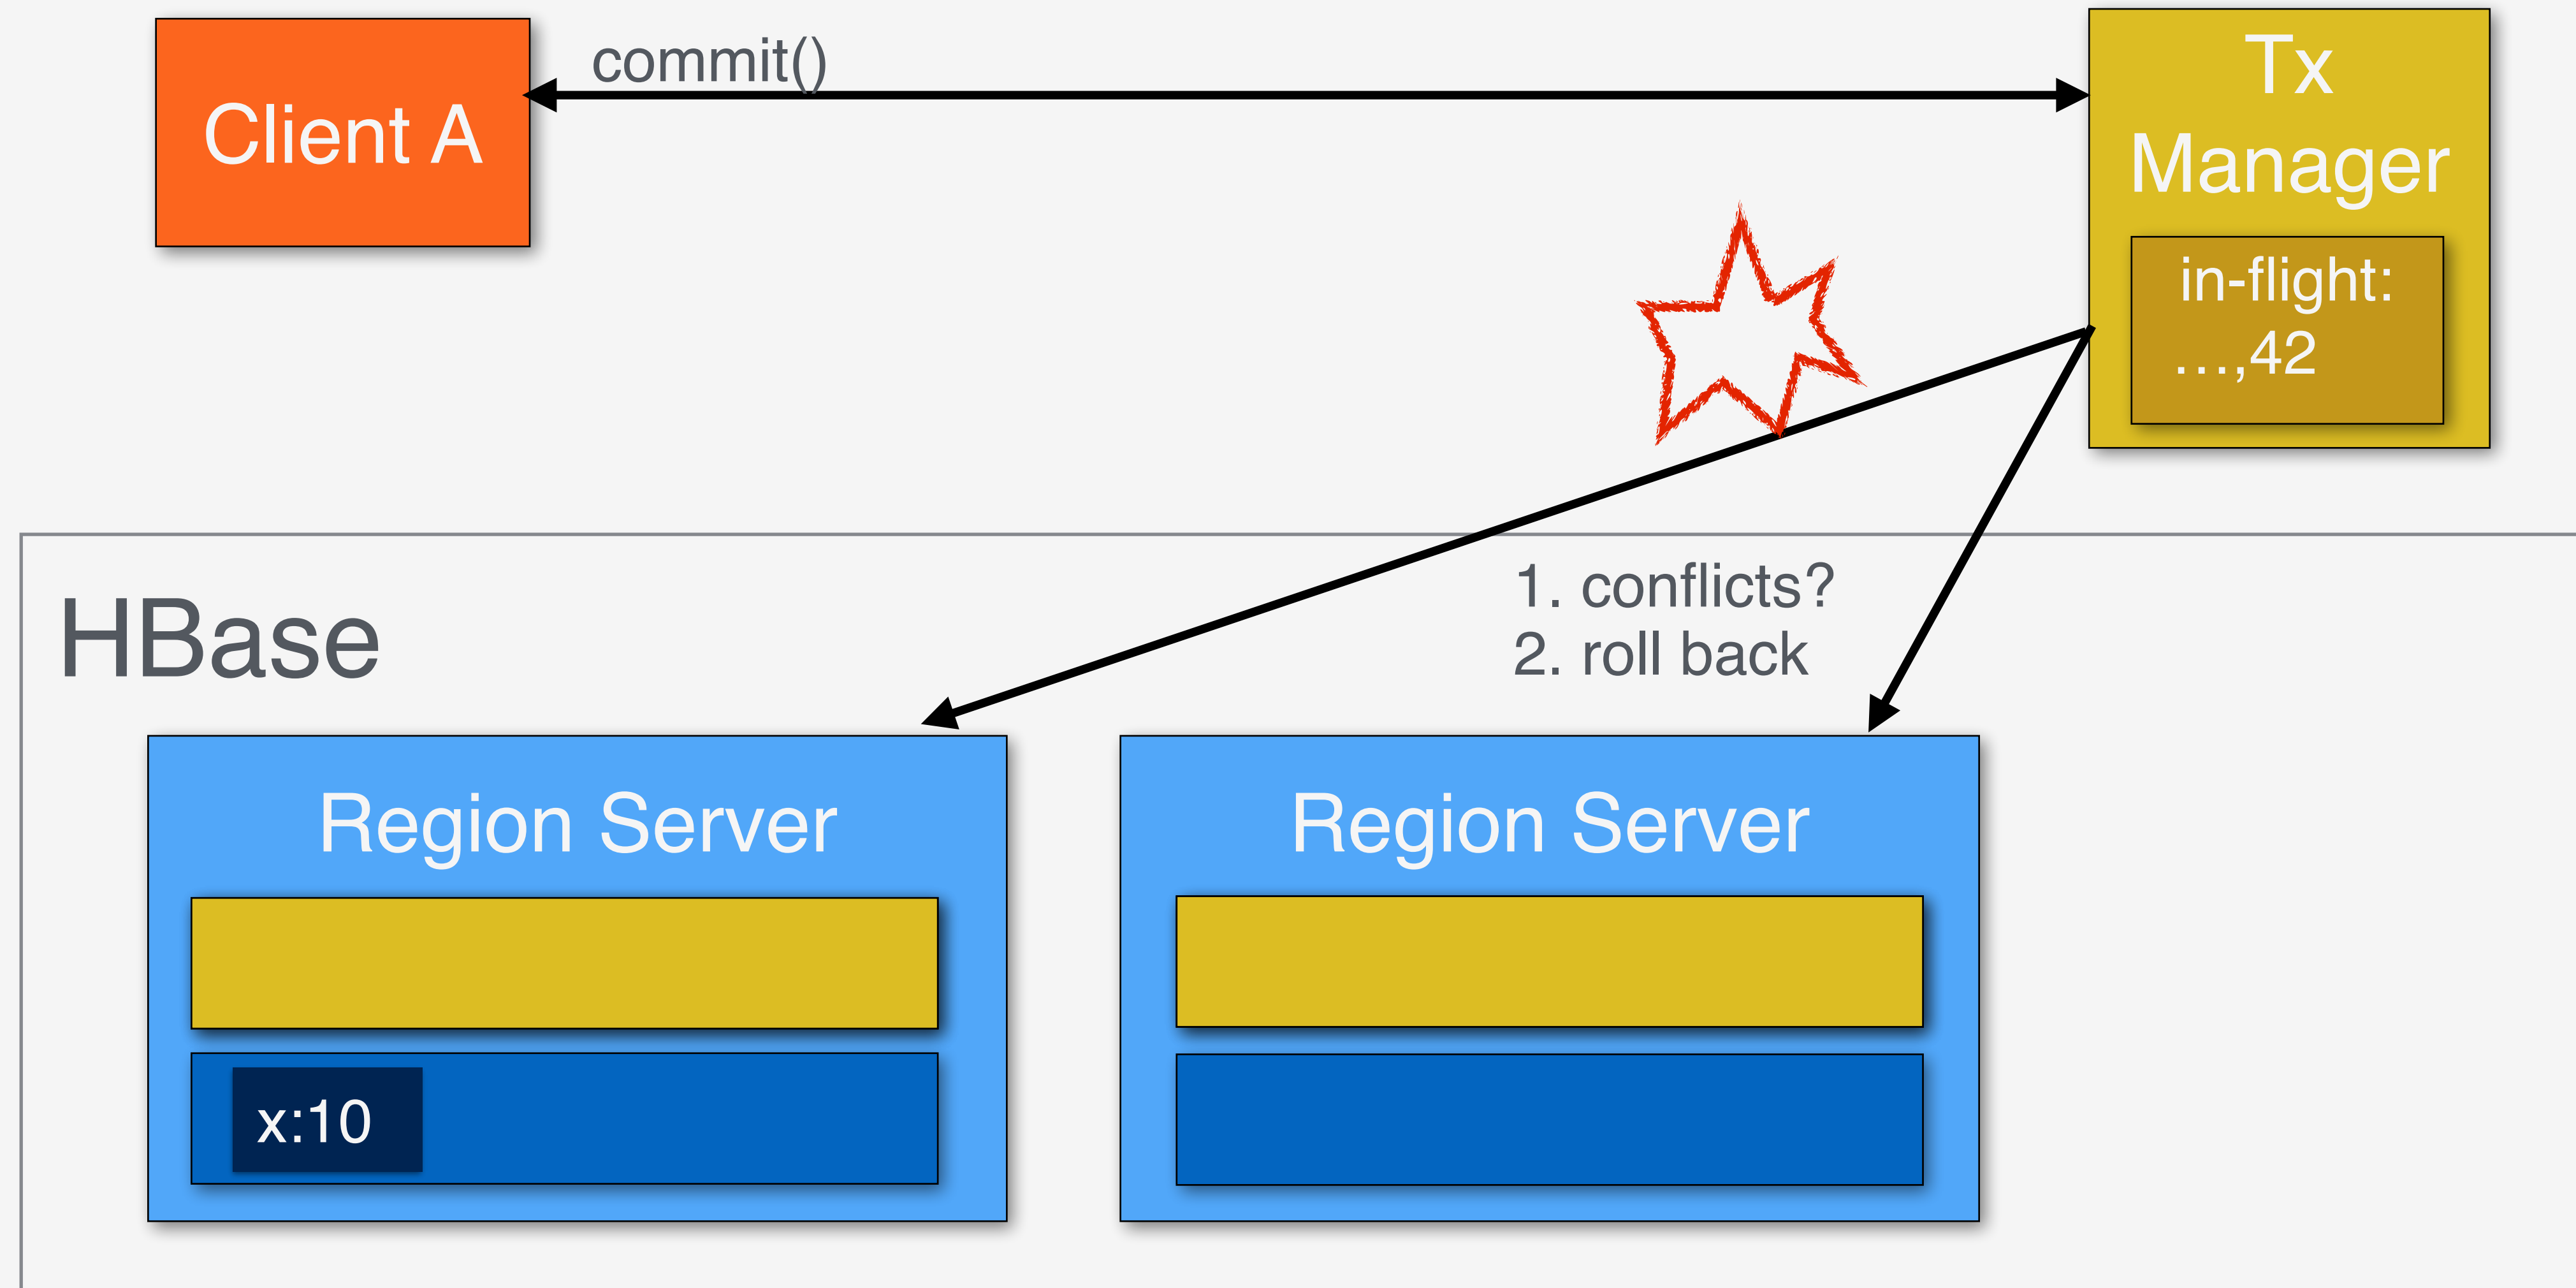
# Apache Trafodion



# Apache Trafodion

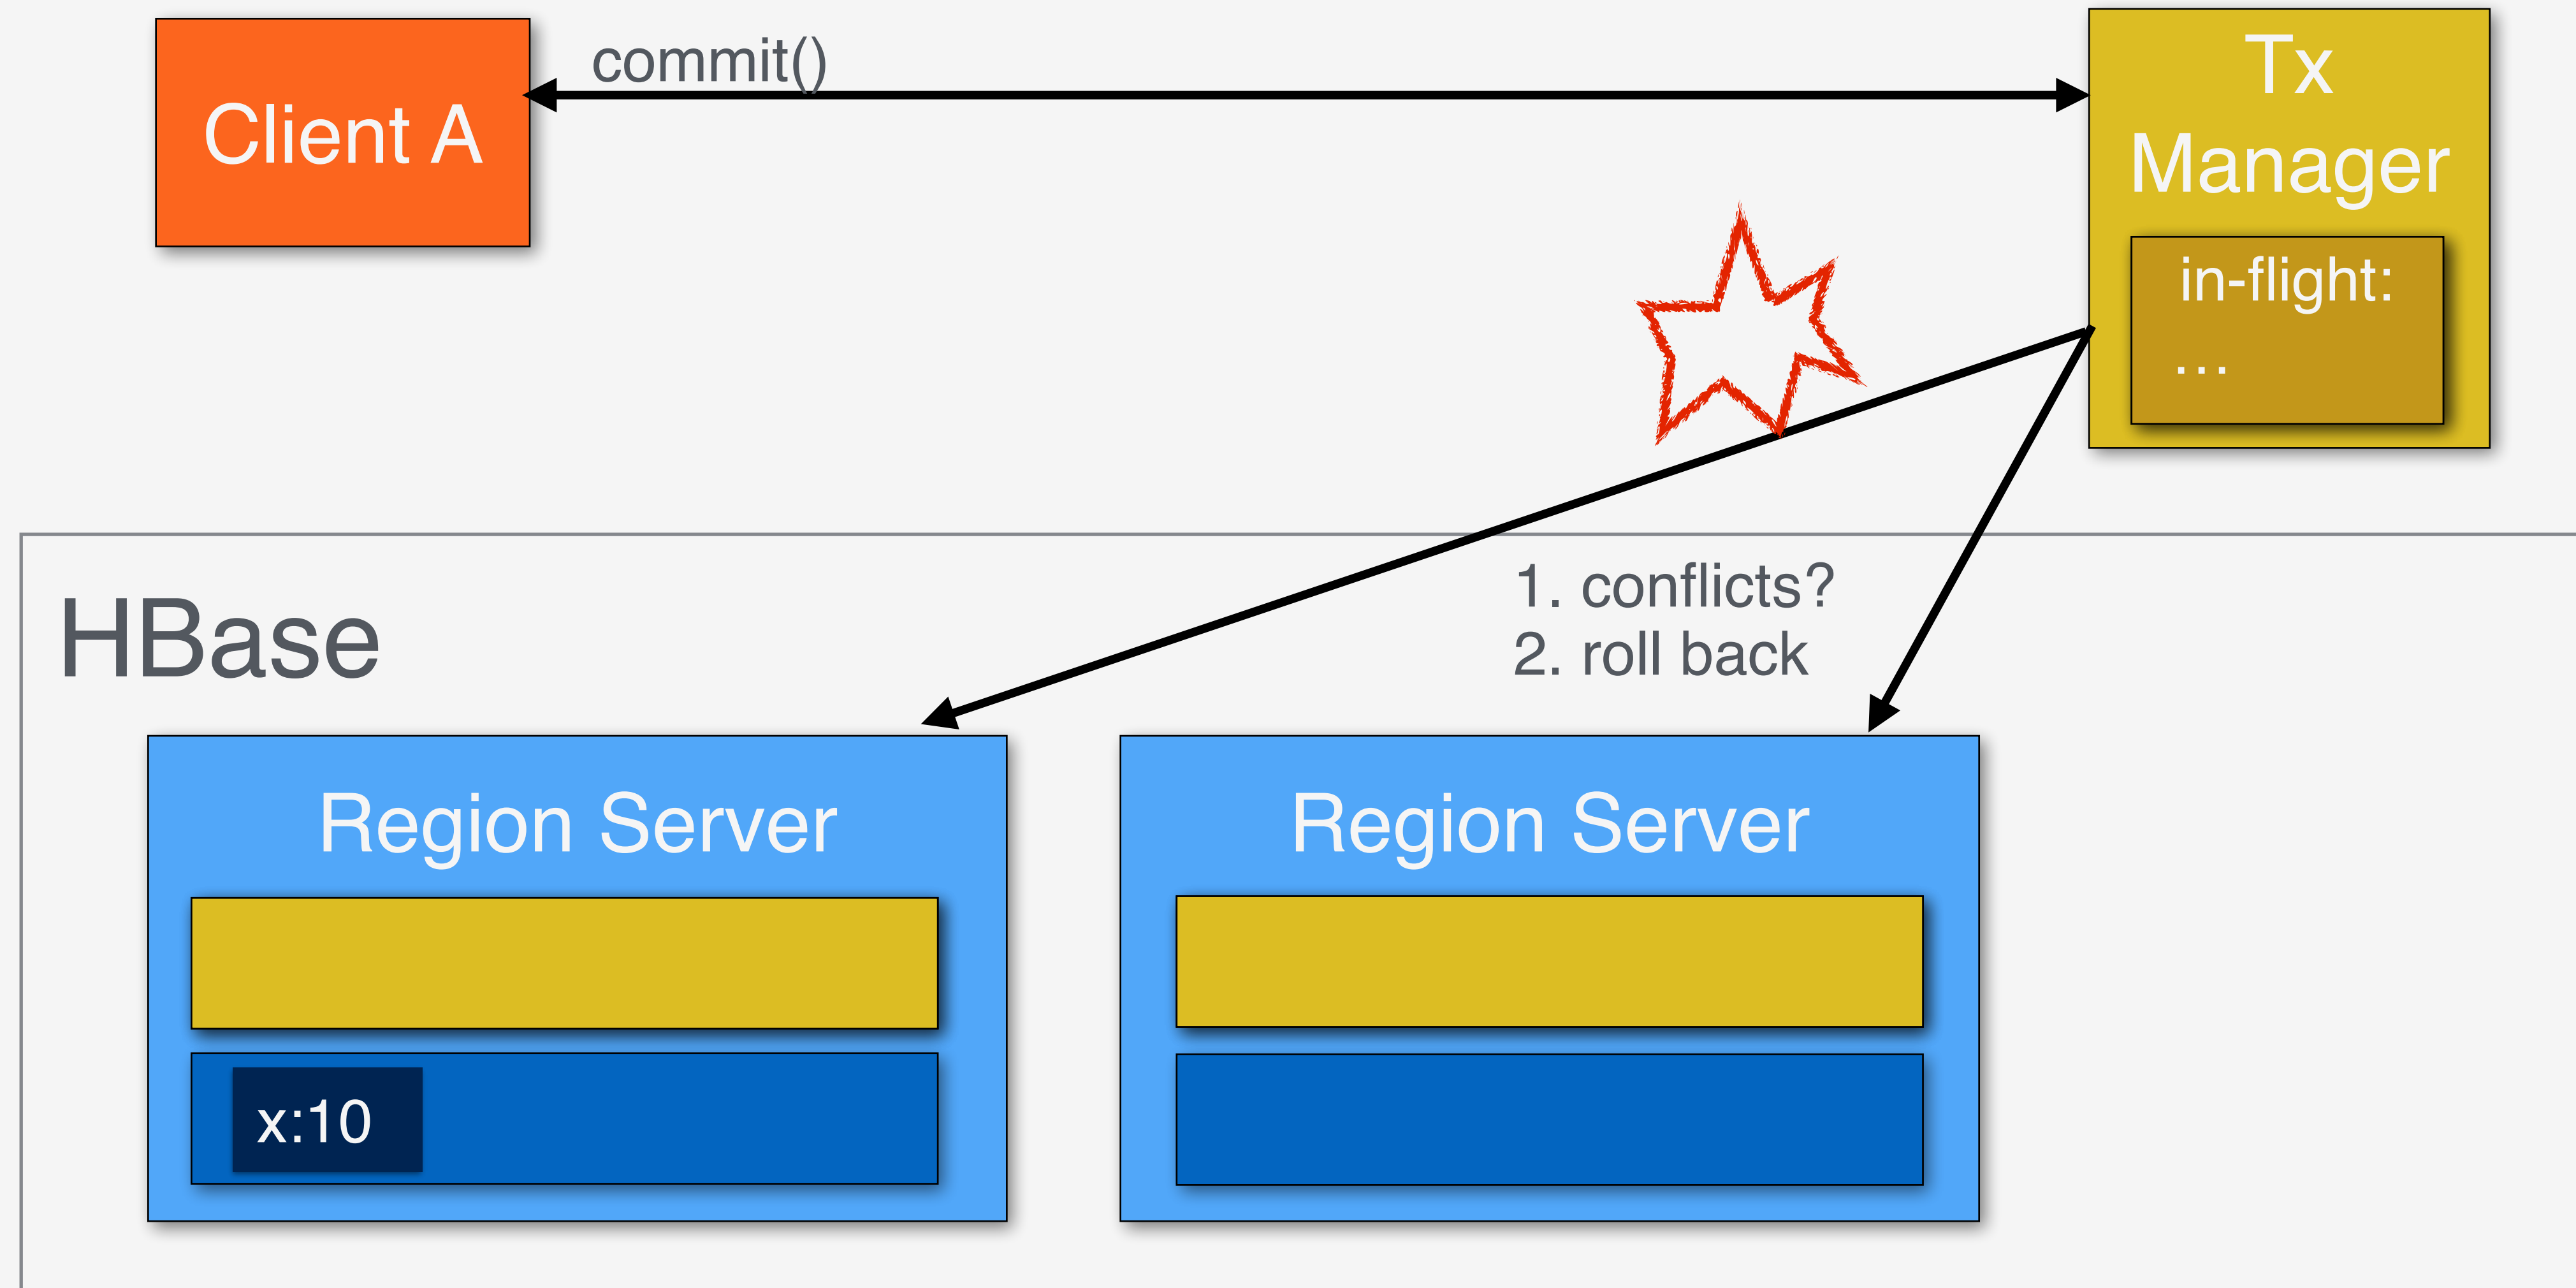


# Apache Trafodion





# Apache Trafodion



# Apache Trafodion

Client A

Tx  
Manager

in-flight:  
...,42

## HBase

Region Server

x:11

x:10

Region Server

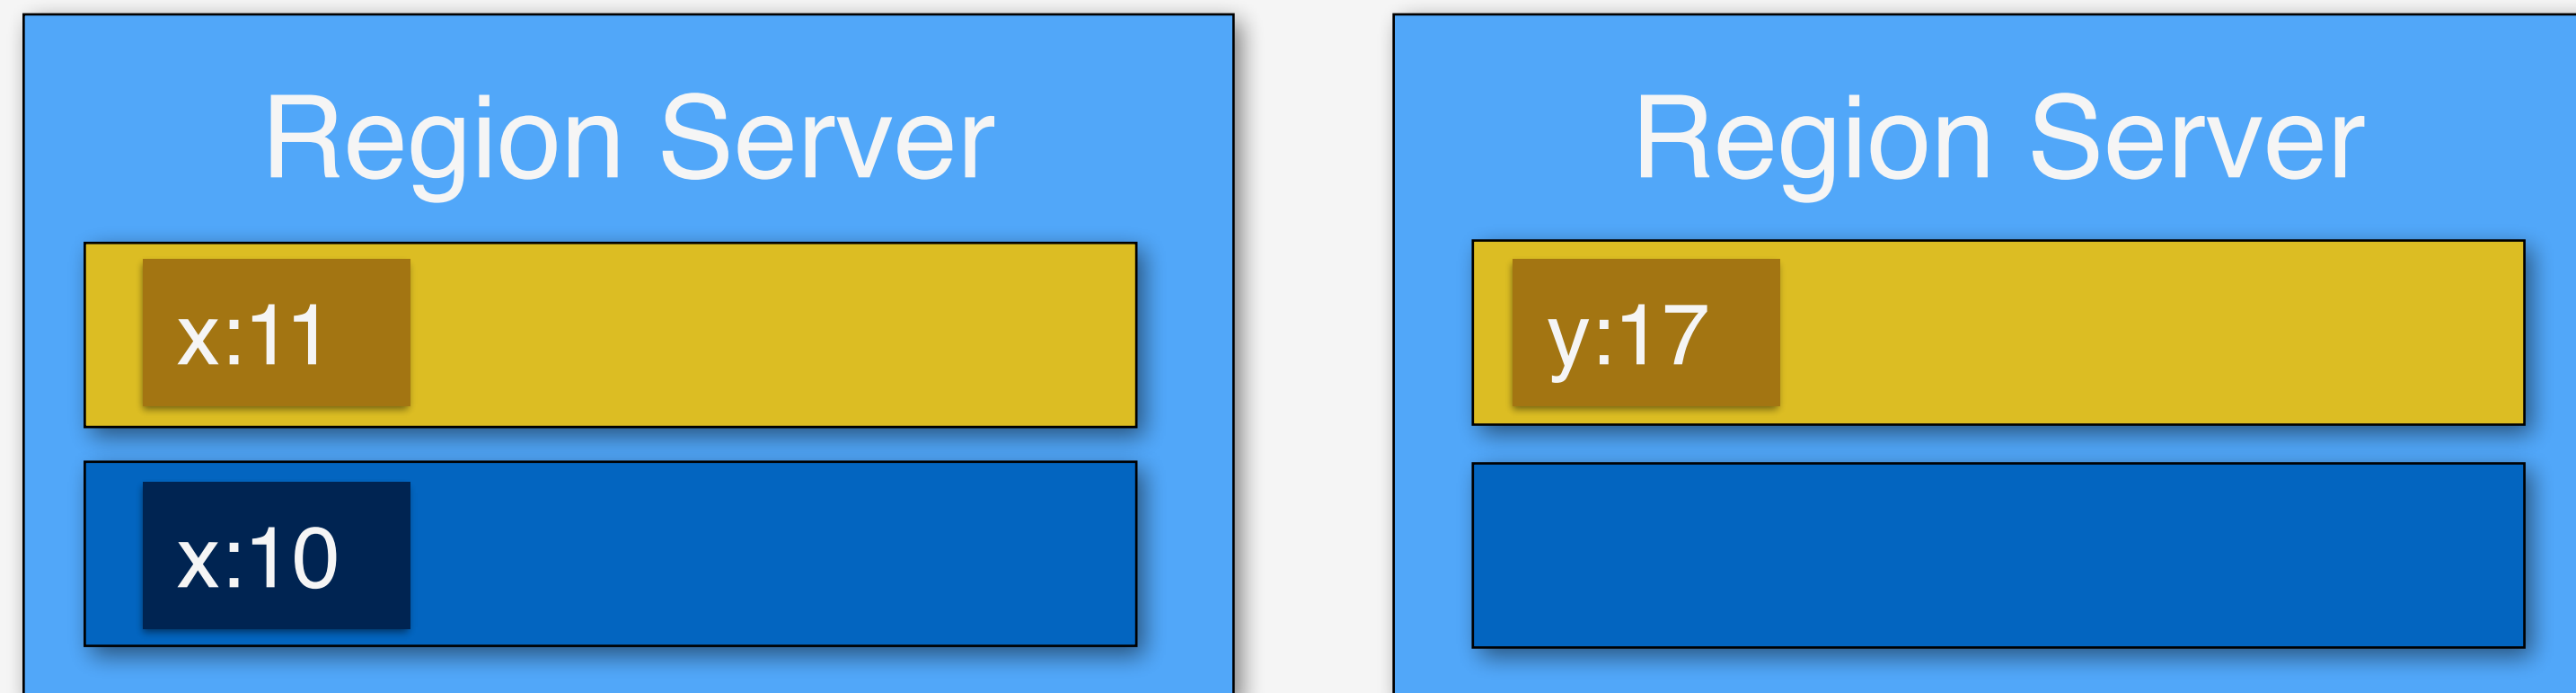
y:17



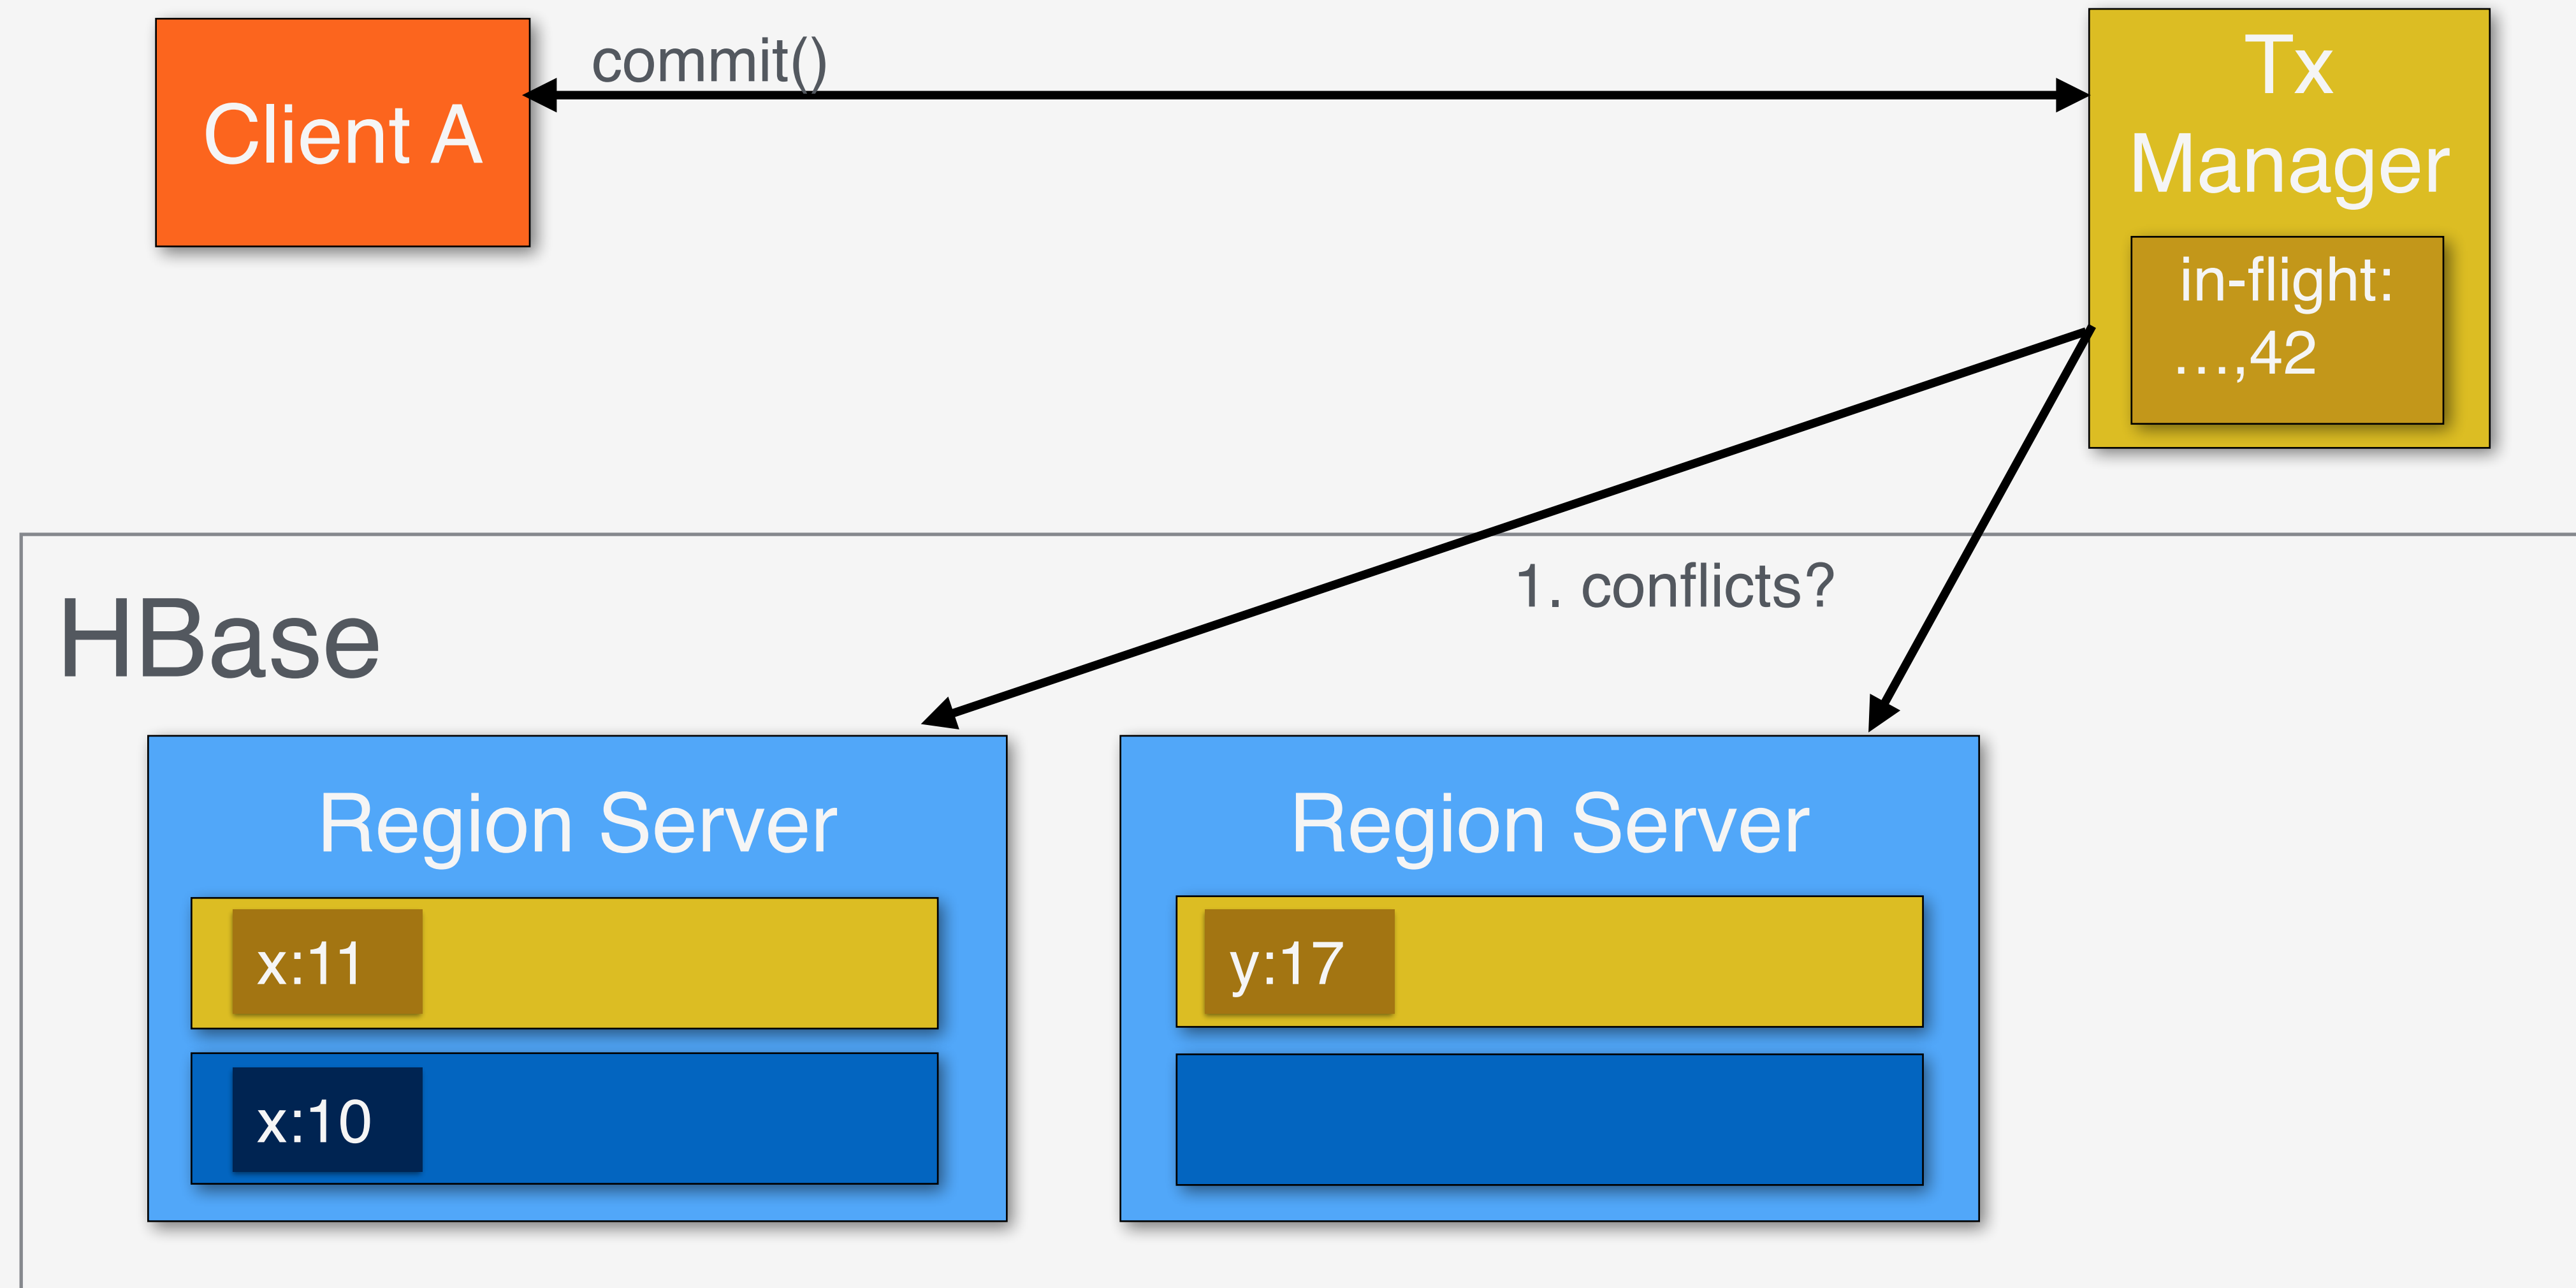
# Apache Trafodion



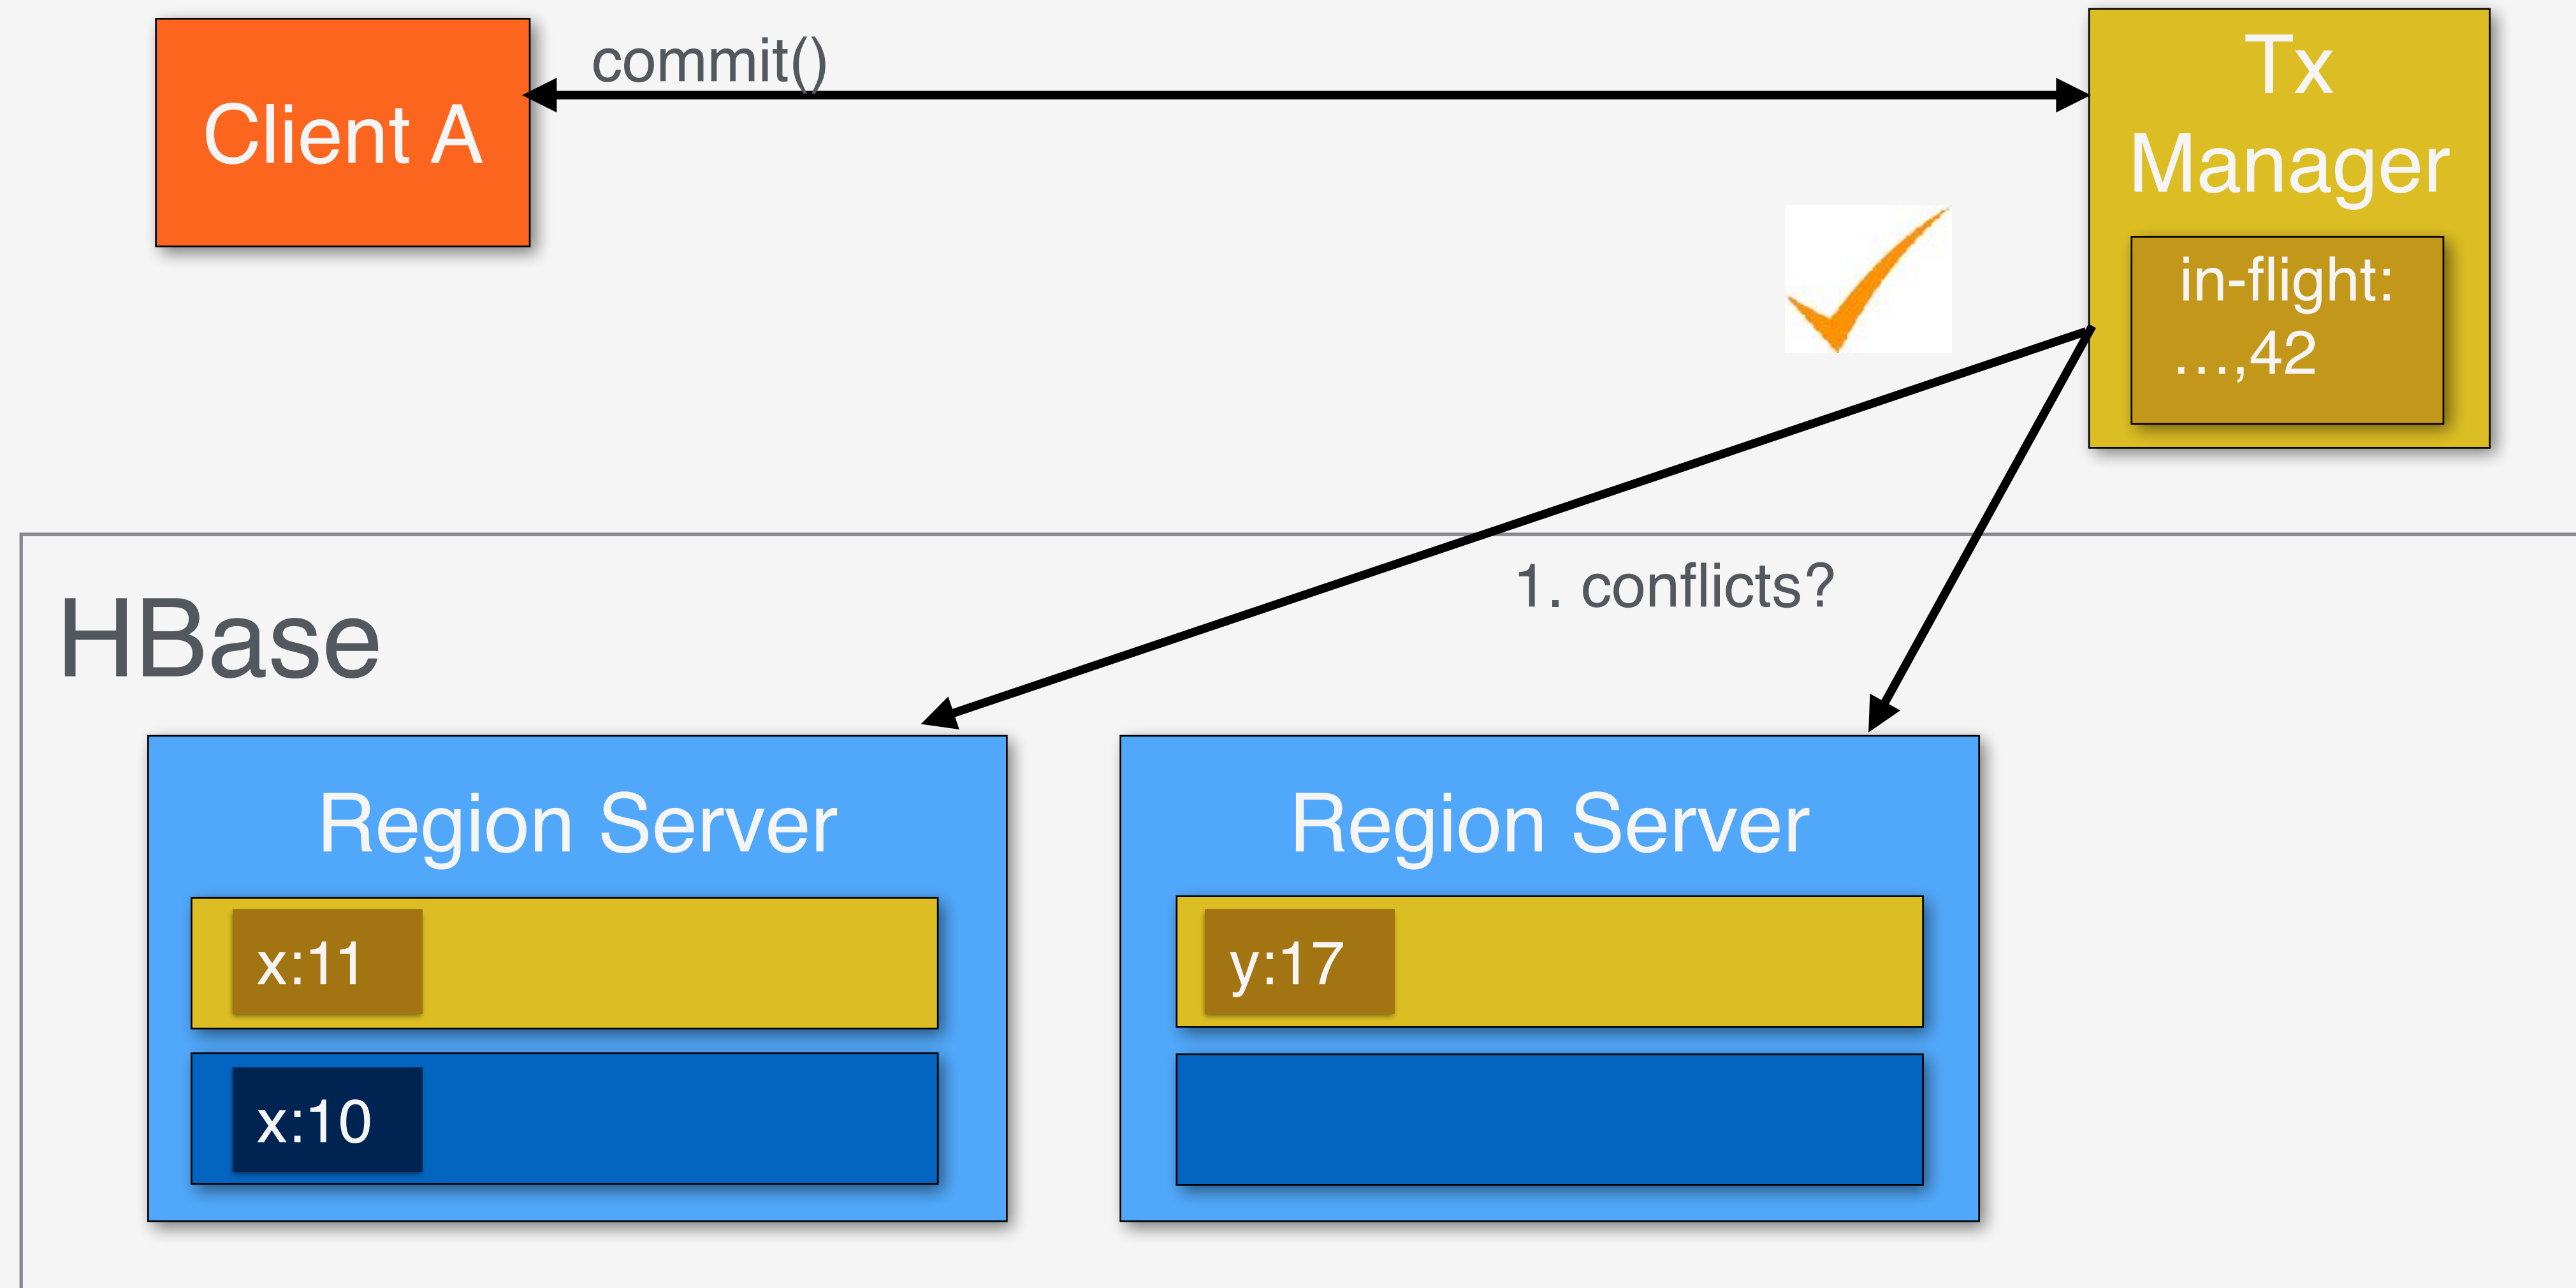
## HBase



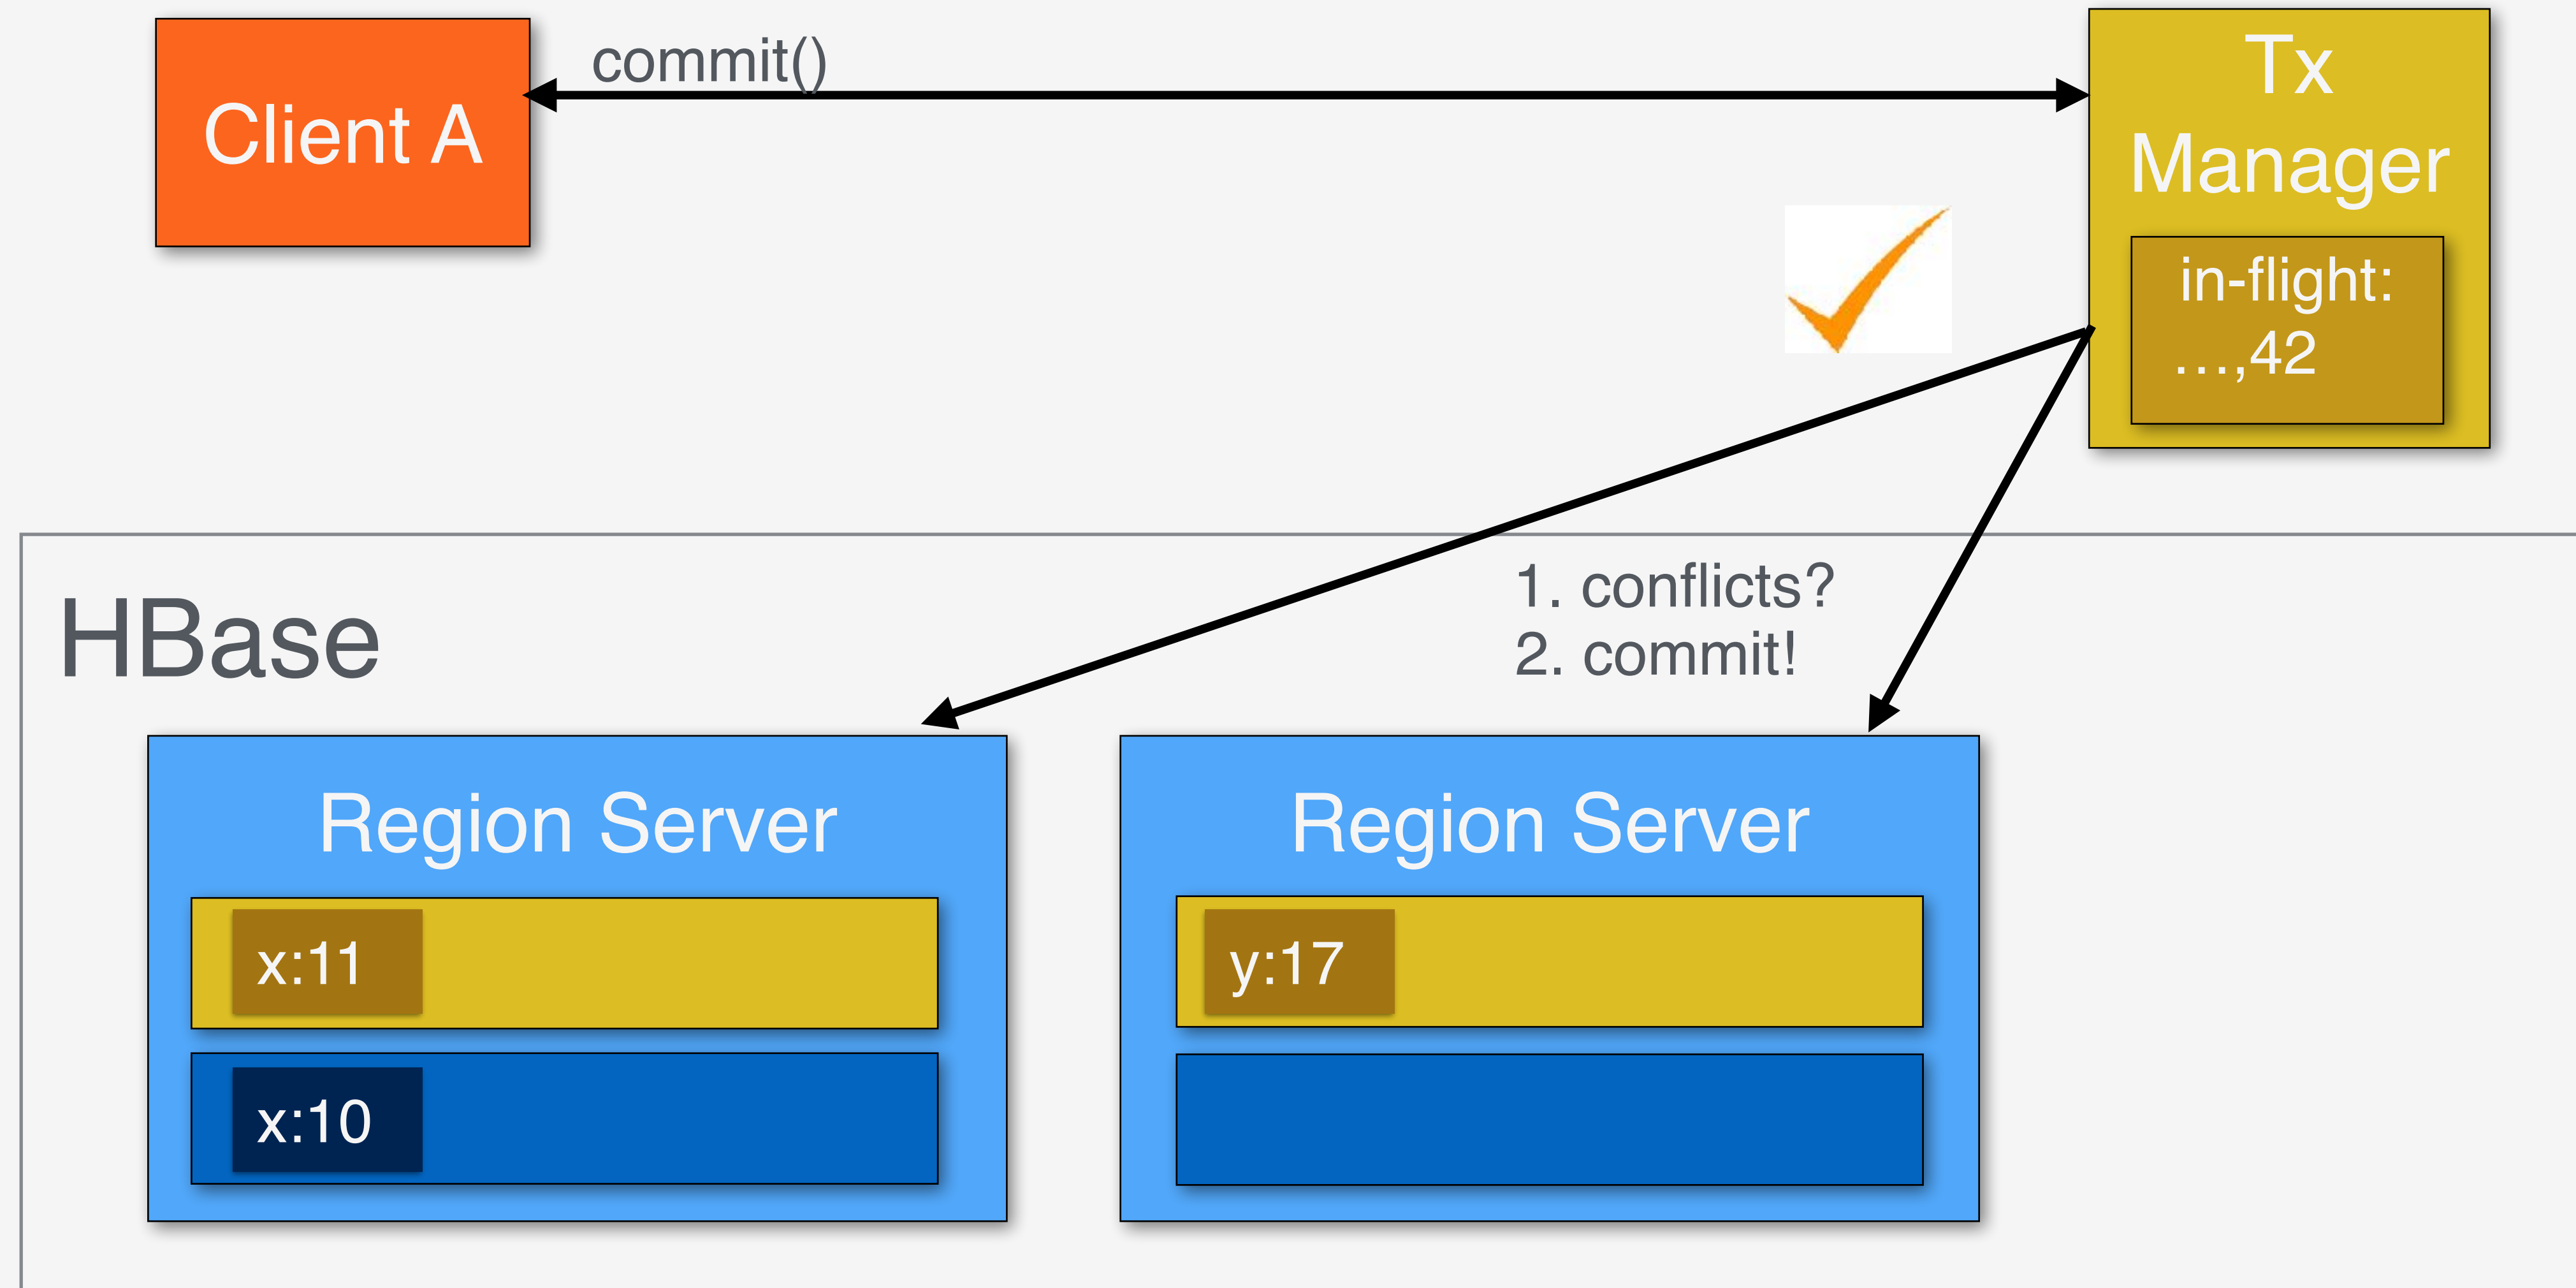
# Apache Trafodion



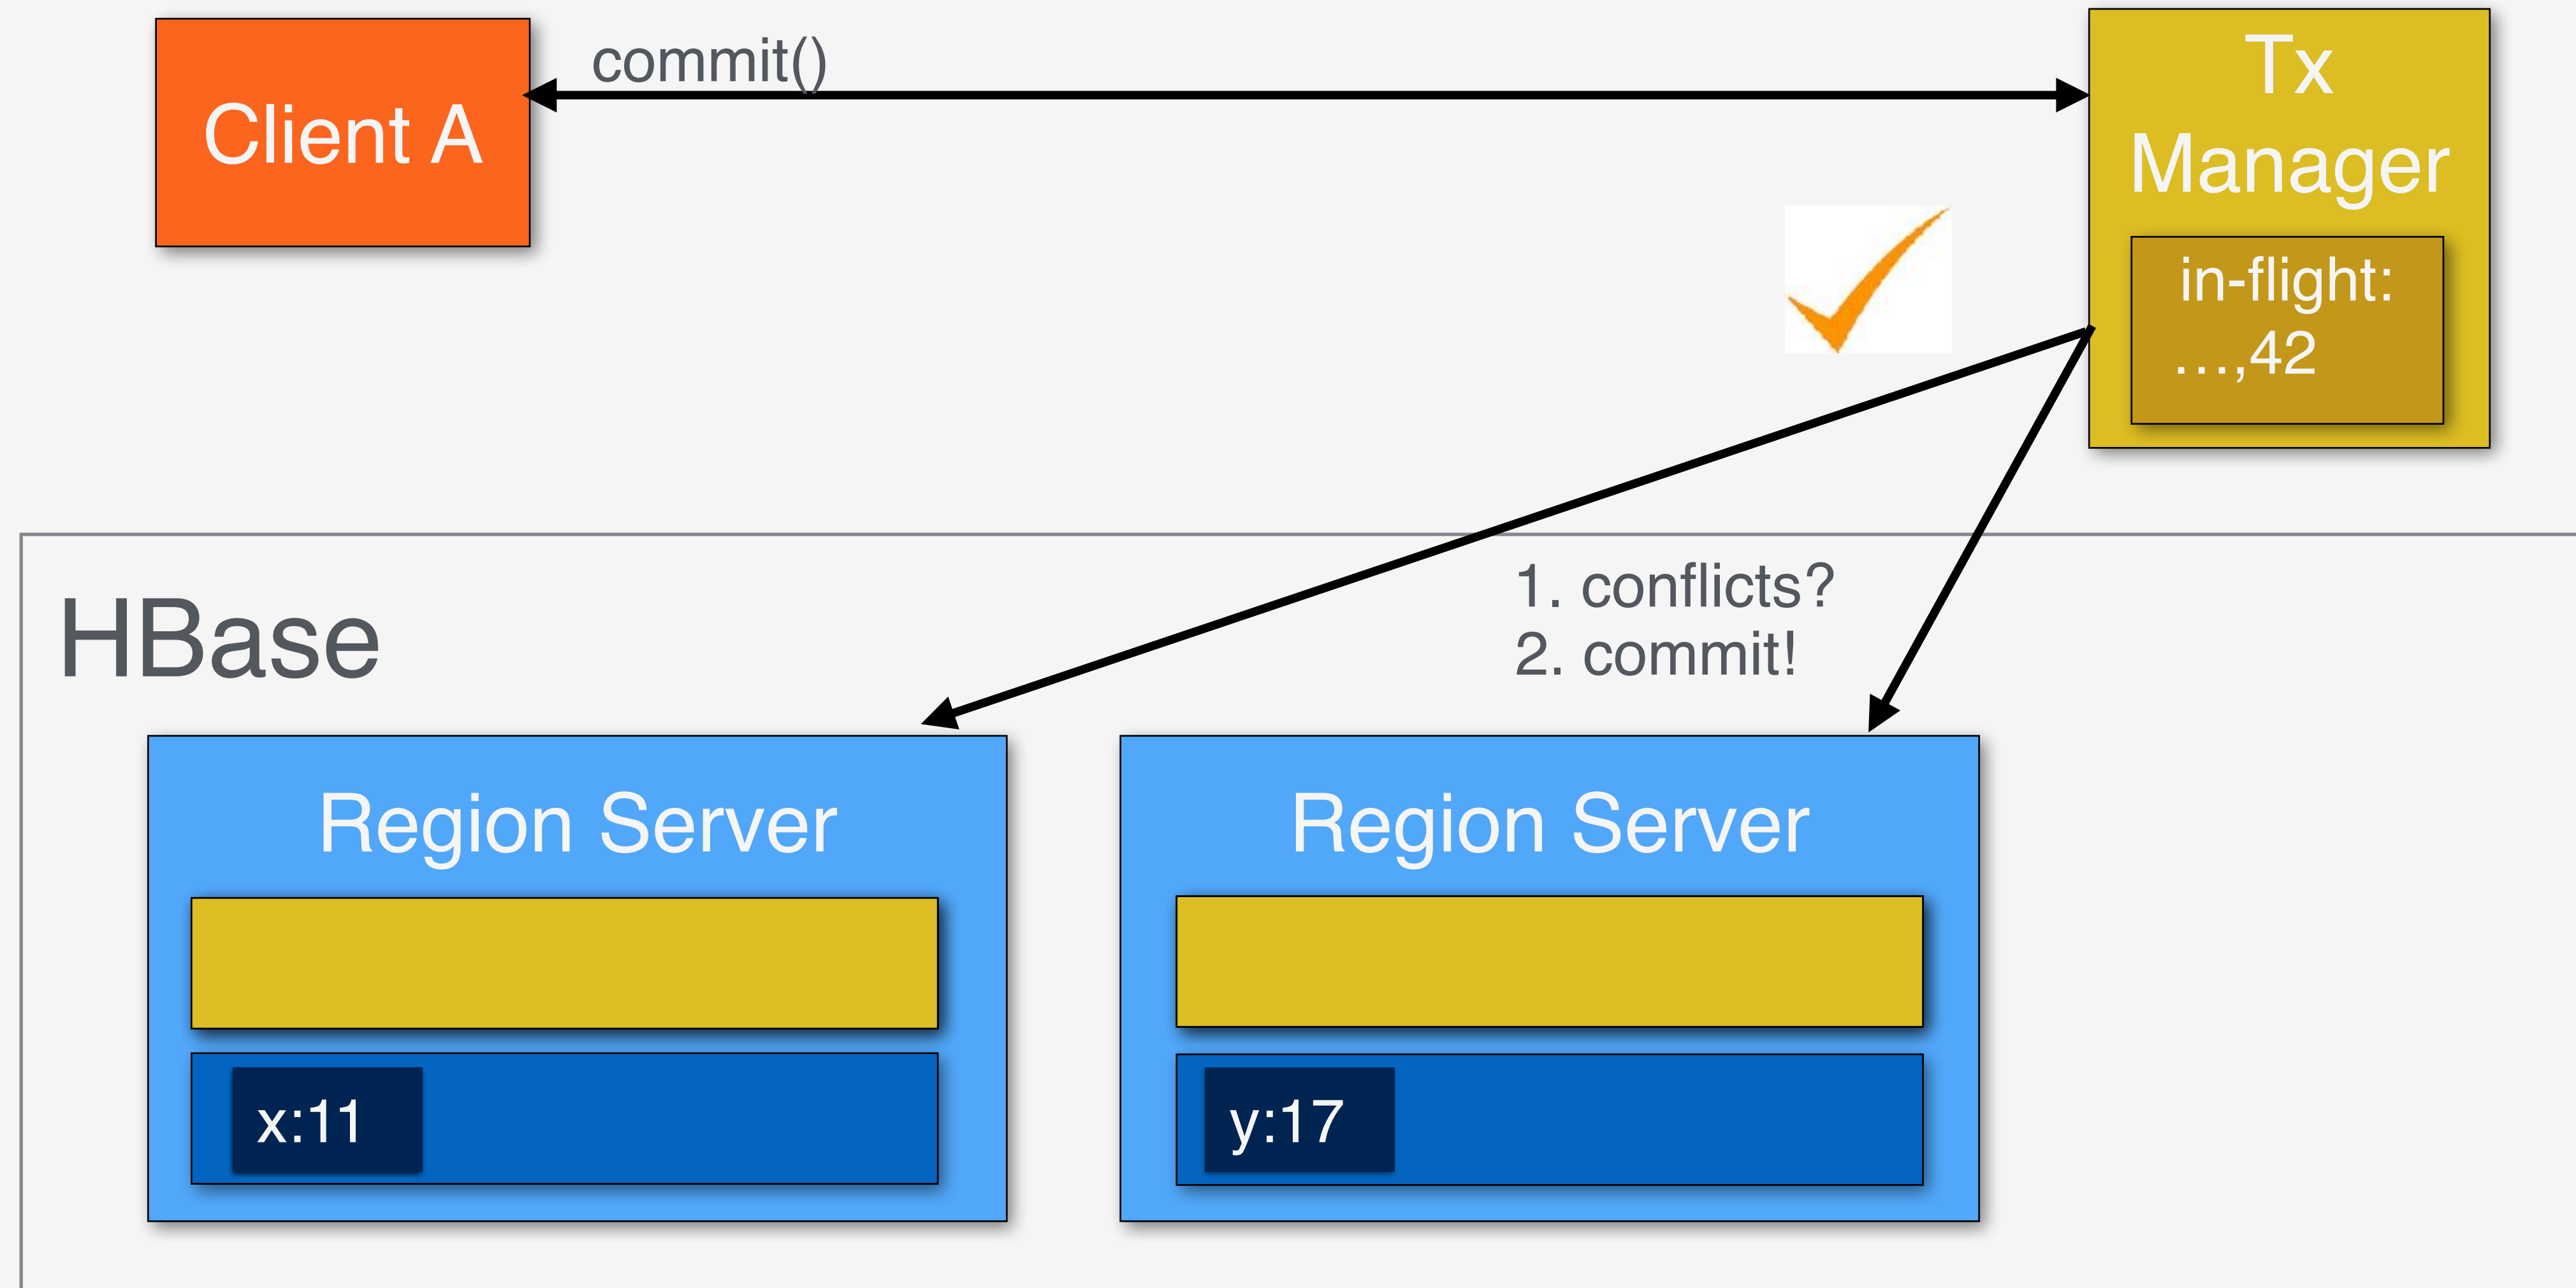
# Apache Trafodion



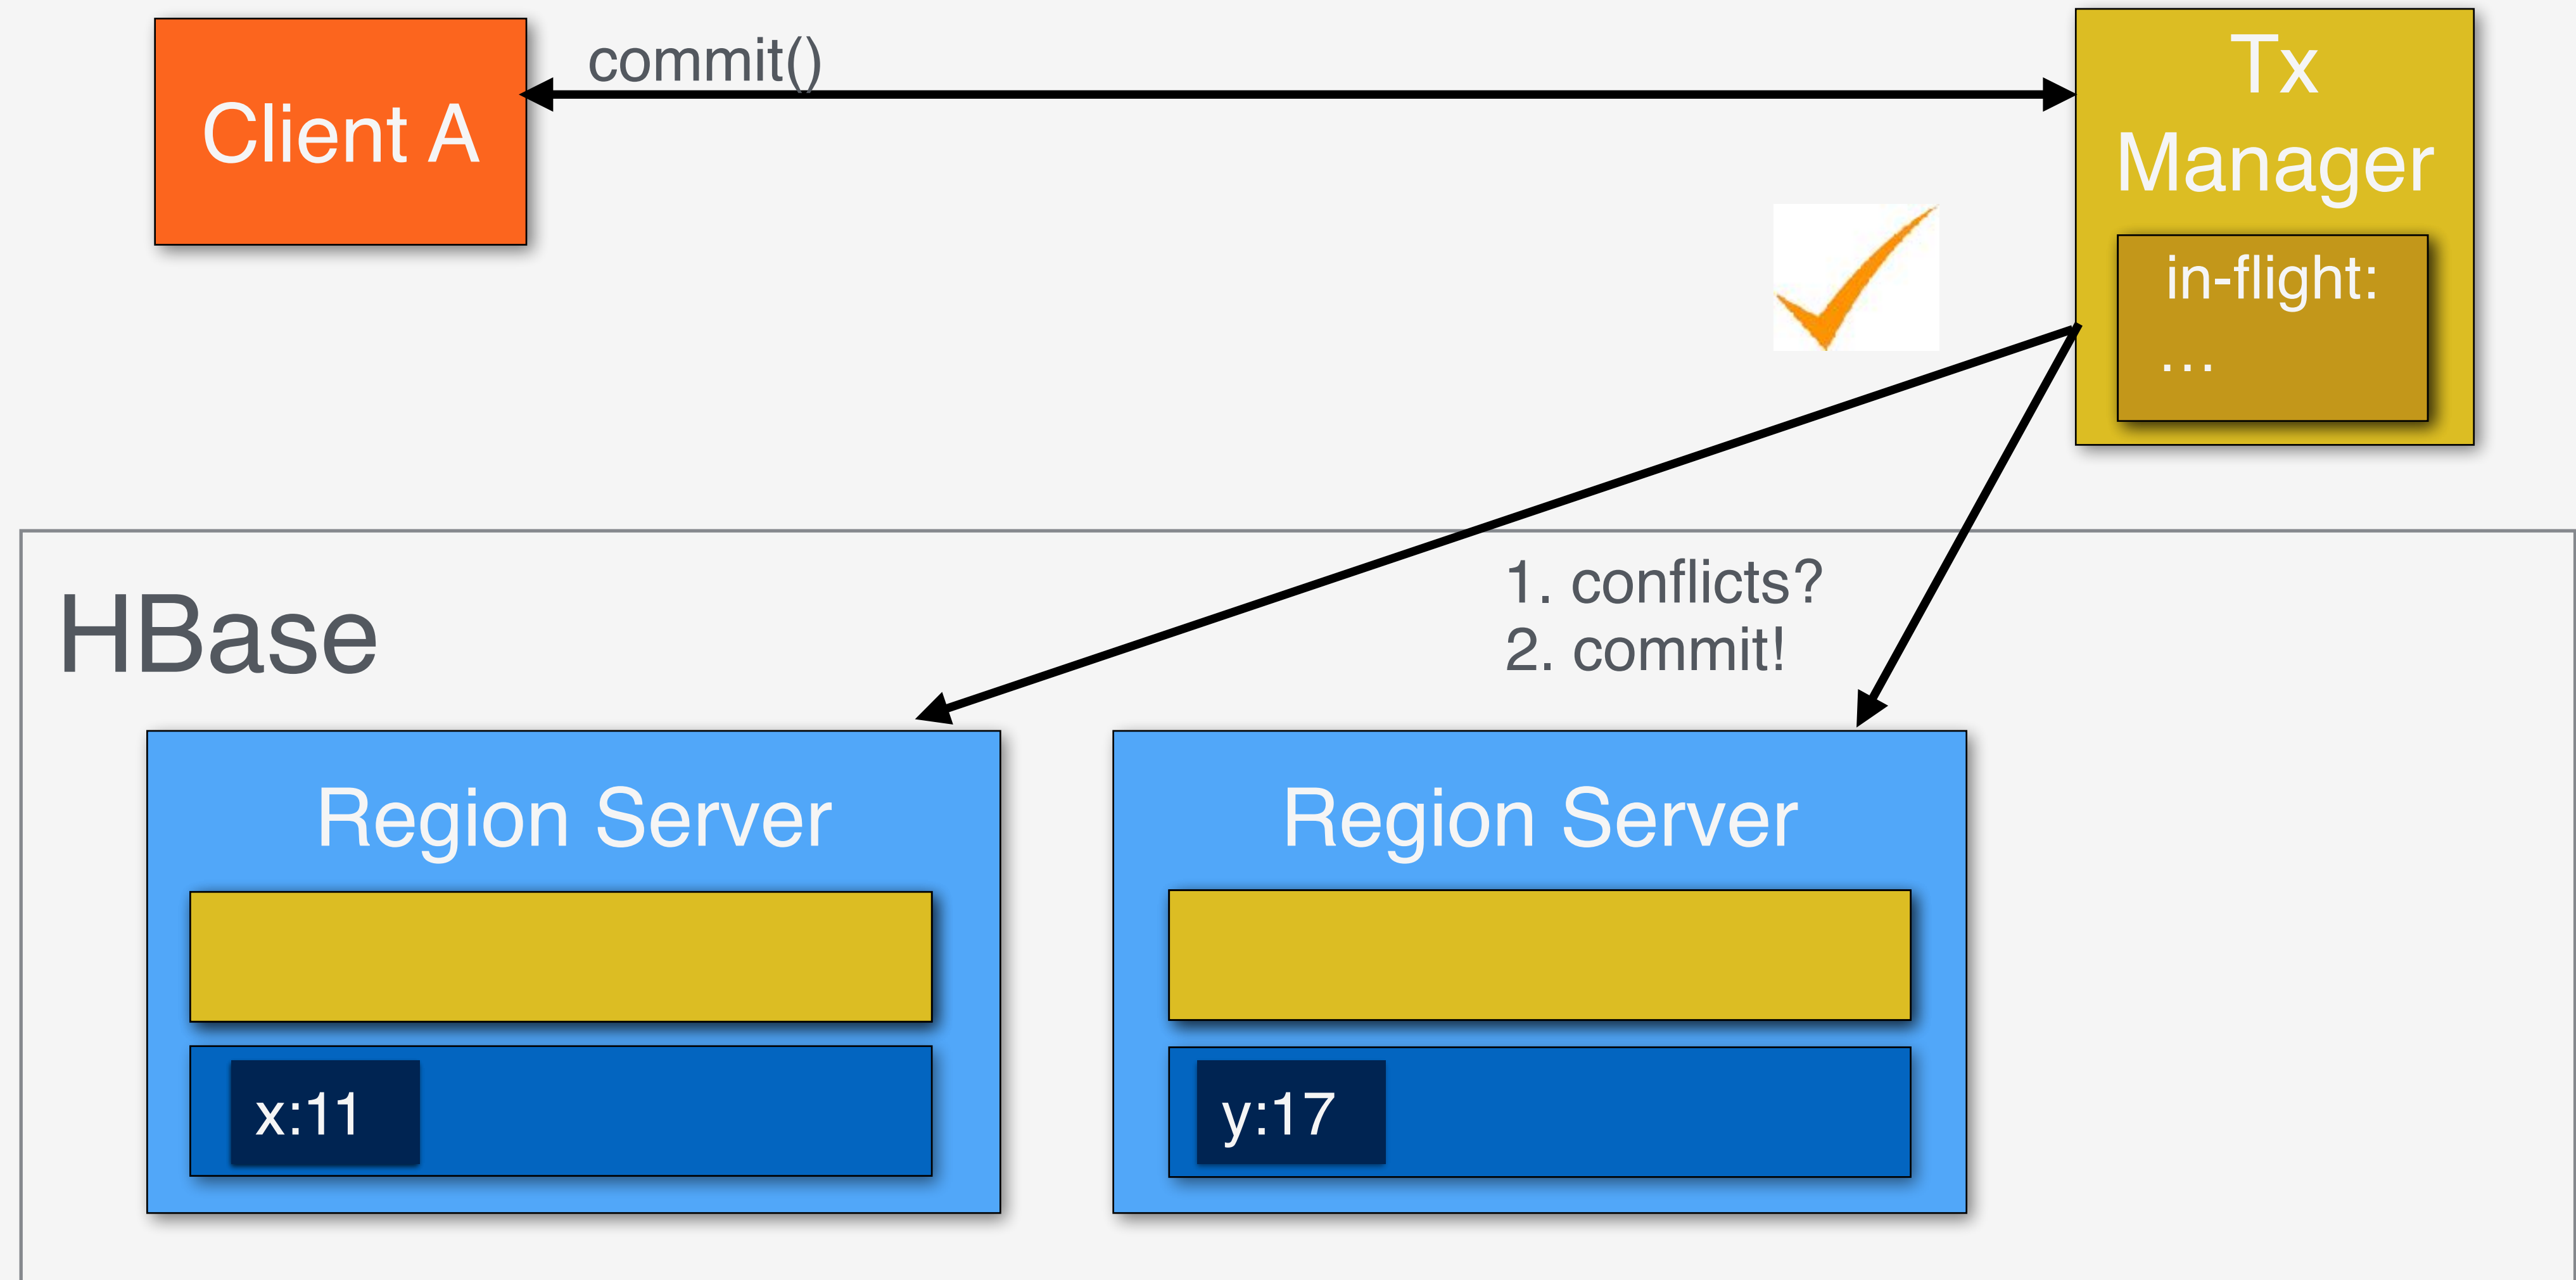
# Apache Trafodion



# Apache Trafodion

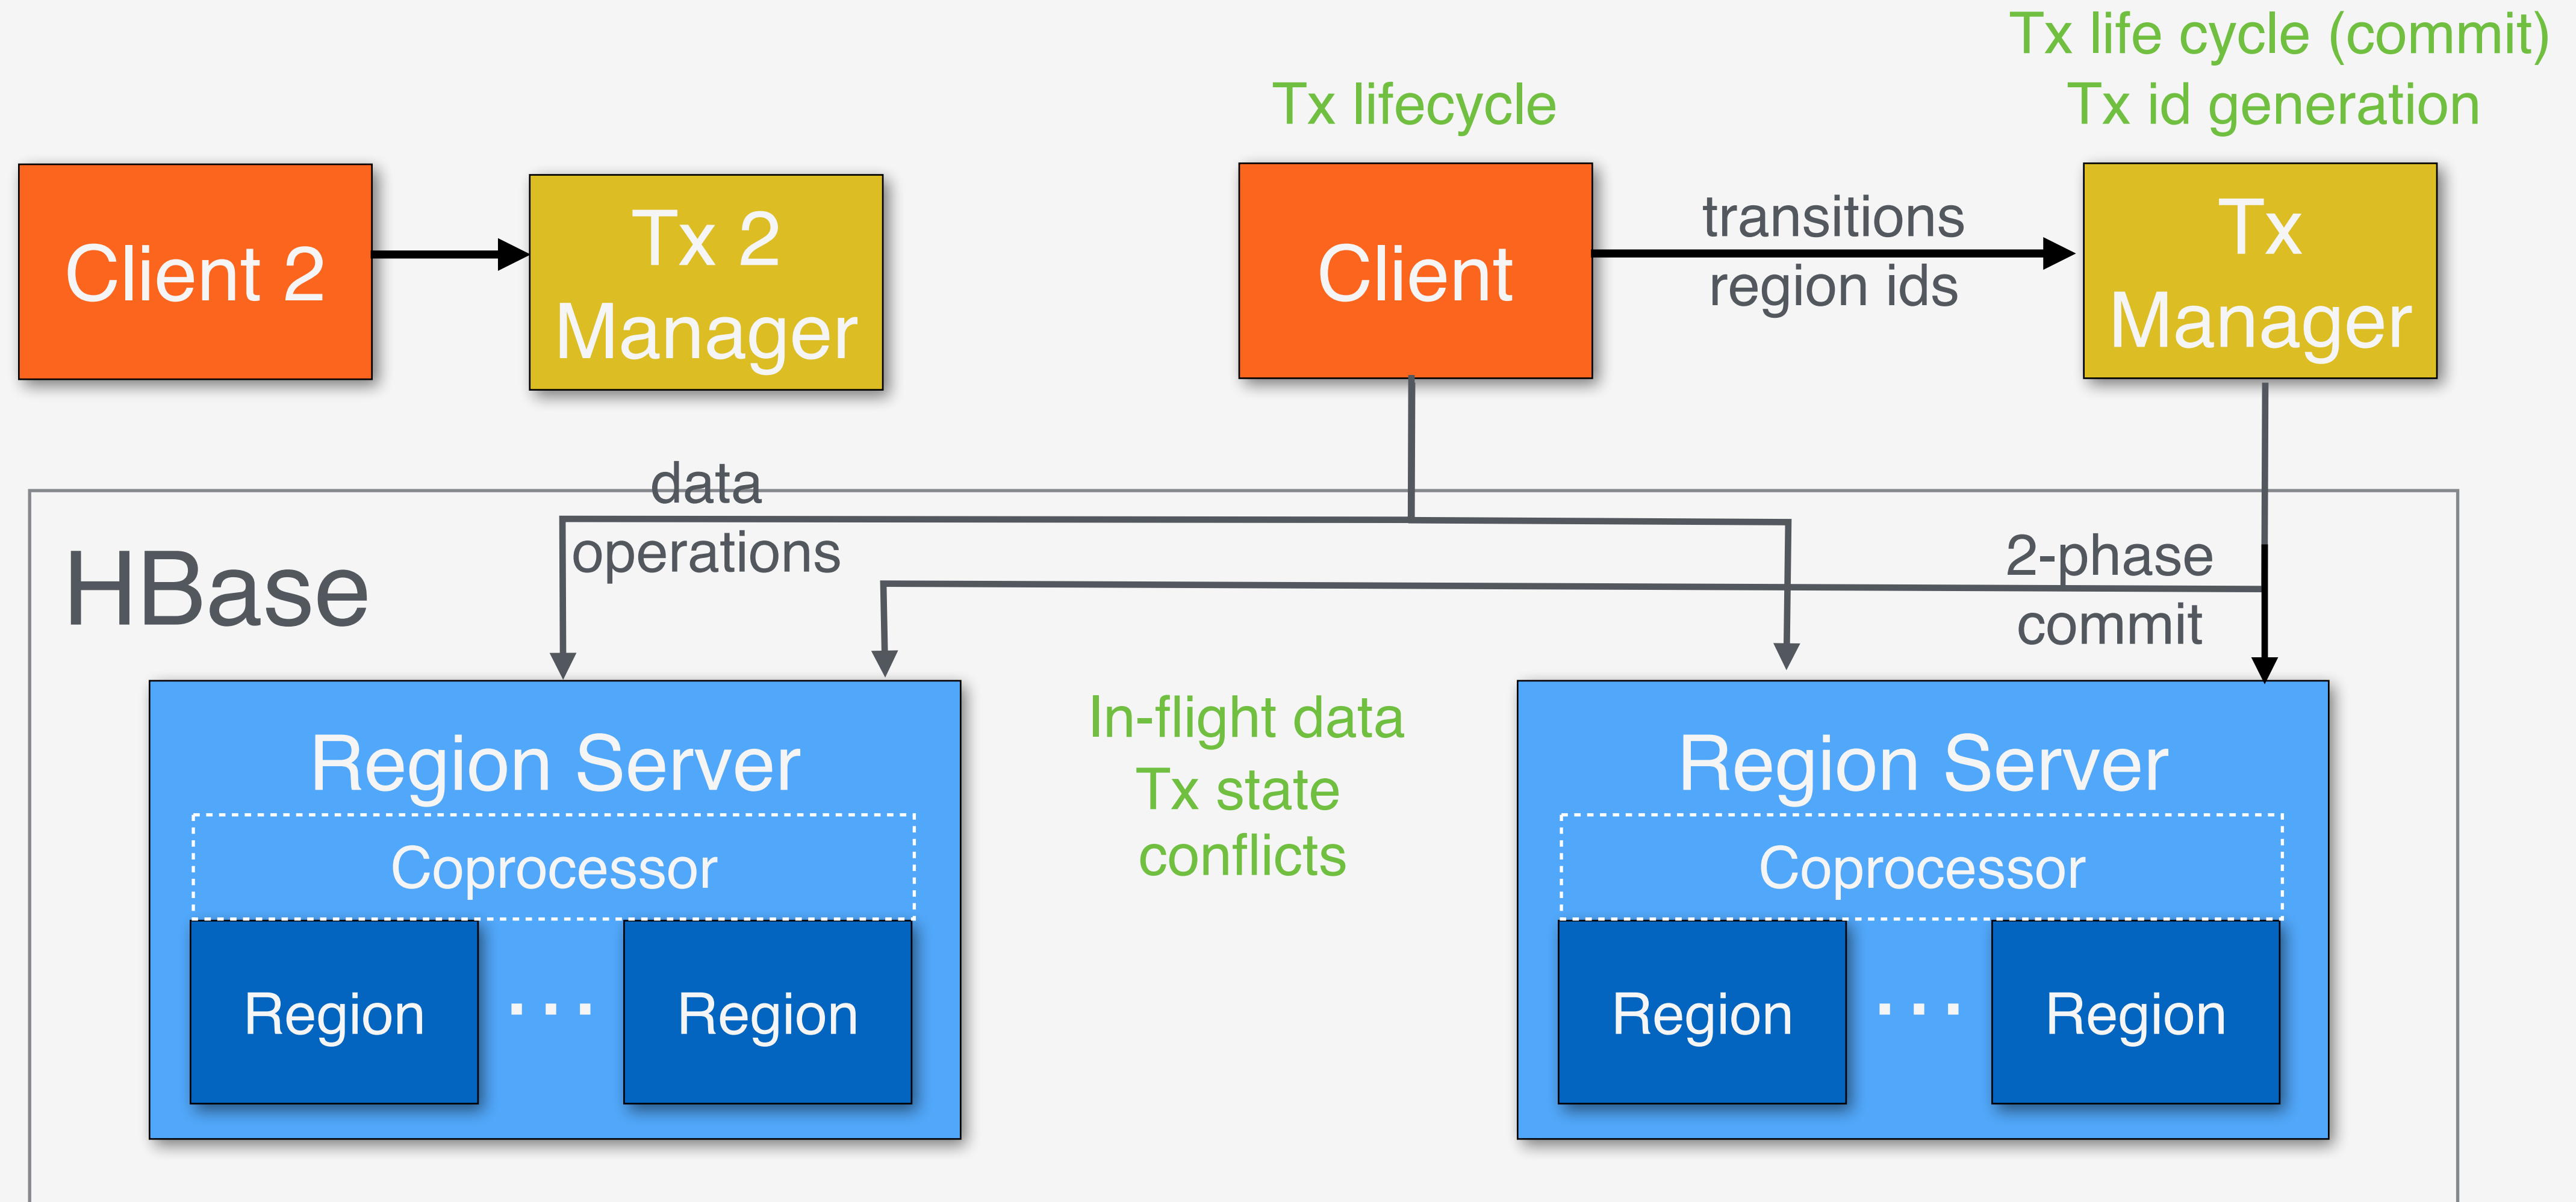


# Apache Trafodion





# Apache Trafodion



# Apache Trafodion

- Scales well:
  - Conflict detection is distributed: no single bottleneck
  - Commit coordination by multiple transaction managers
  - Optimization: bypass 2-phase commit if single region
- Coprocessors cache in-flight data in Memory
  - Flushed to HBase only on commit
  - Committed read (not snapshot, not repeatable read)
  - Option: cause conflicts for reads, too
- HA and Fault Tolerance
  - WAL for all state
  - All services are redundant and take over for each other
- Replication: Only in paid (non-Apache) add-on



# Apache Trafodion - Strengths

- Very good scalability
  - Scales almost linearly
  - Especially for very small transactions
- Familiar SQL/jdbc interface for RDB programmers
- Redundant and fault-tolerant
- Secure and multi-tenant:
  - Trafodion/SQL layer provides authn+authz



# Apache Trafodion - Not-So Strengths

- Monolithic, not available as standalone transaction system
- Heavy load on coprocessors
  - memory and compute
- Large transactions (e.g., MapReduce) will cause Out-of-memory
  - no special support for long-running transactions



# Apache Omid



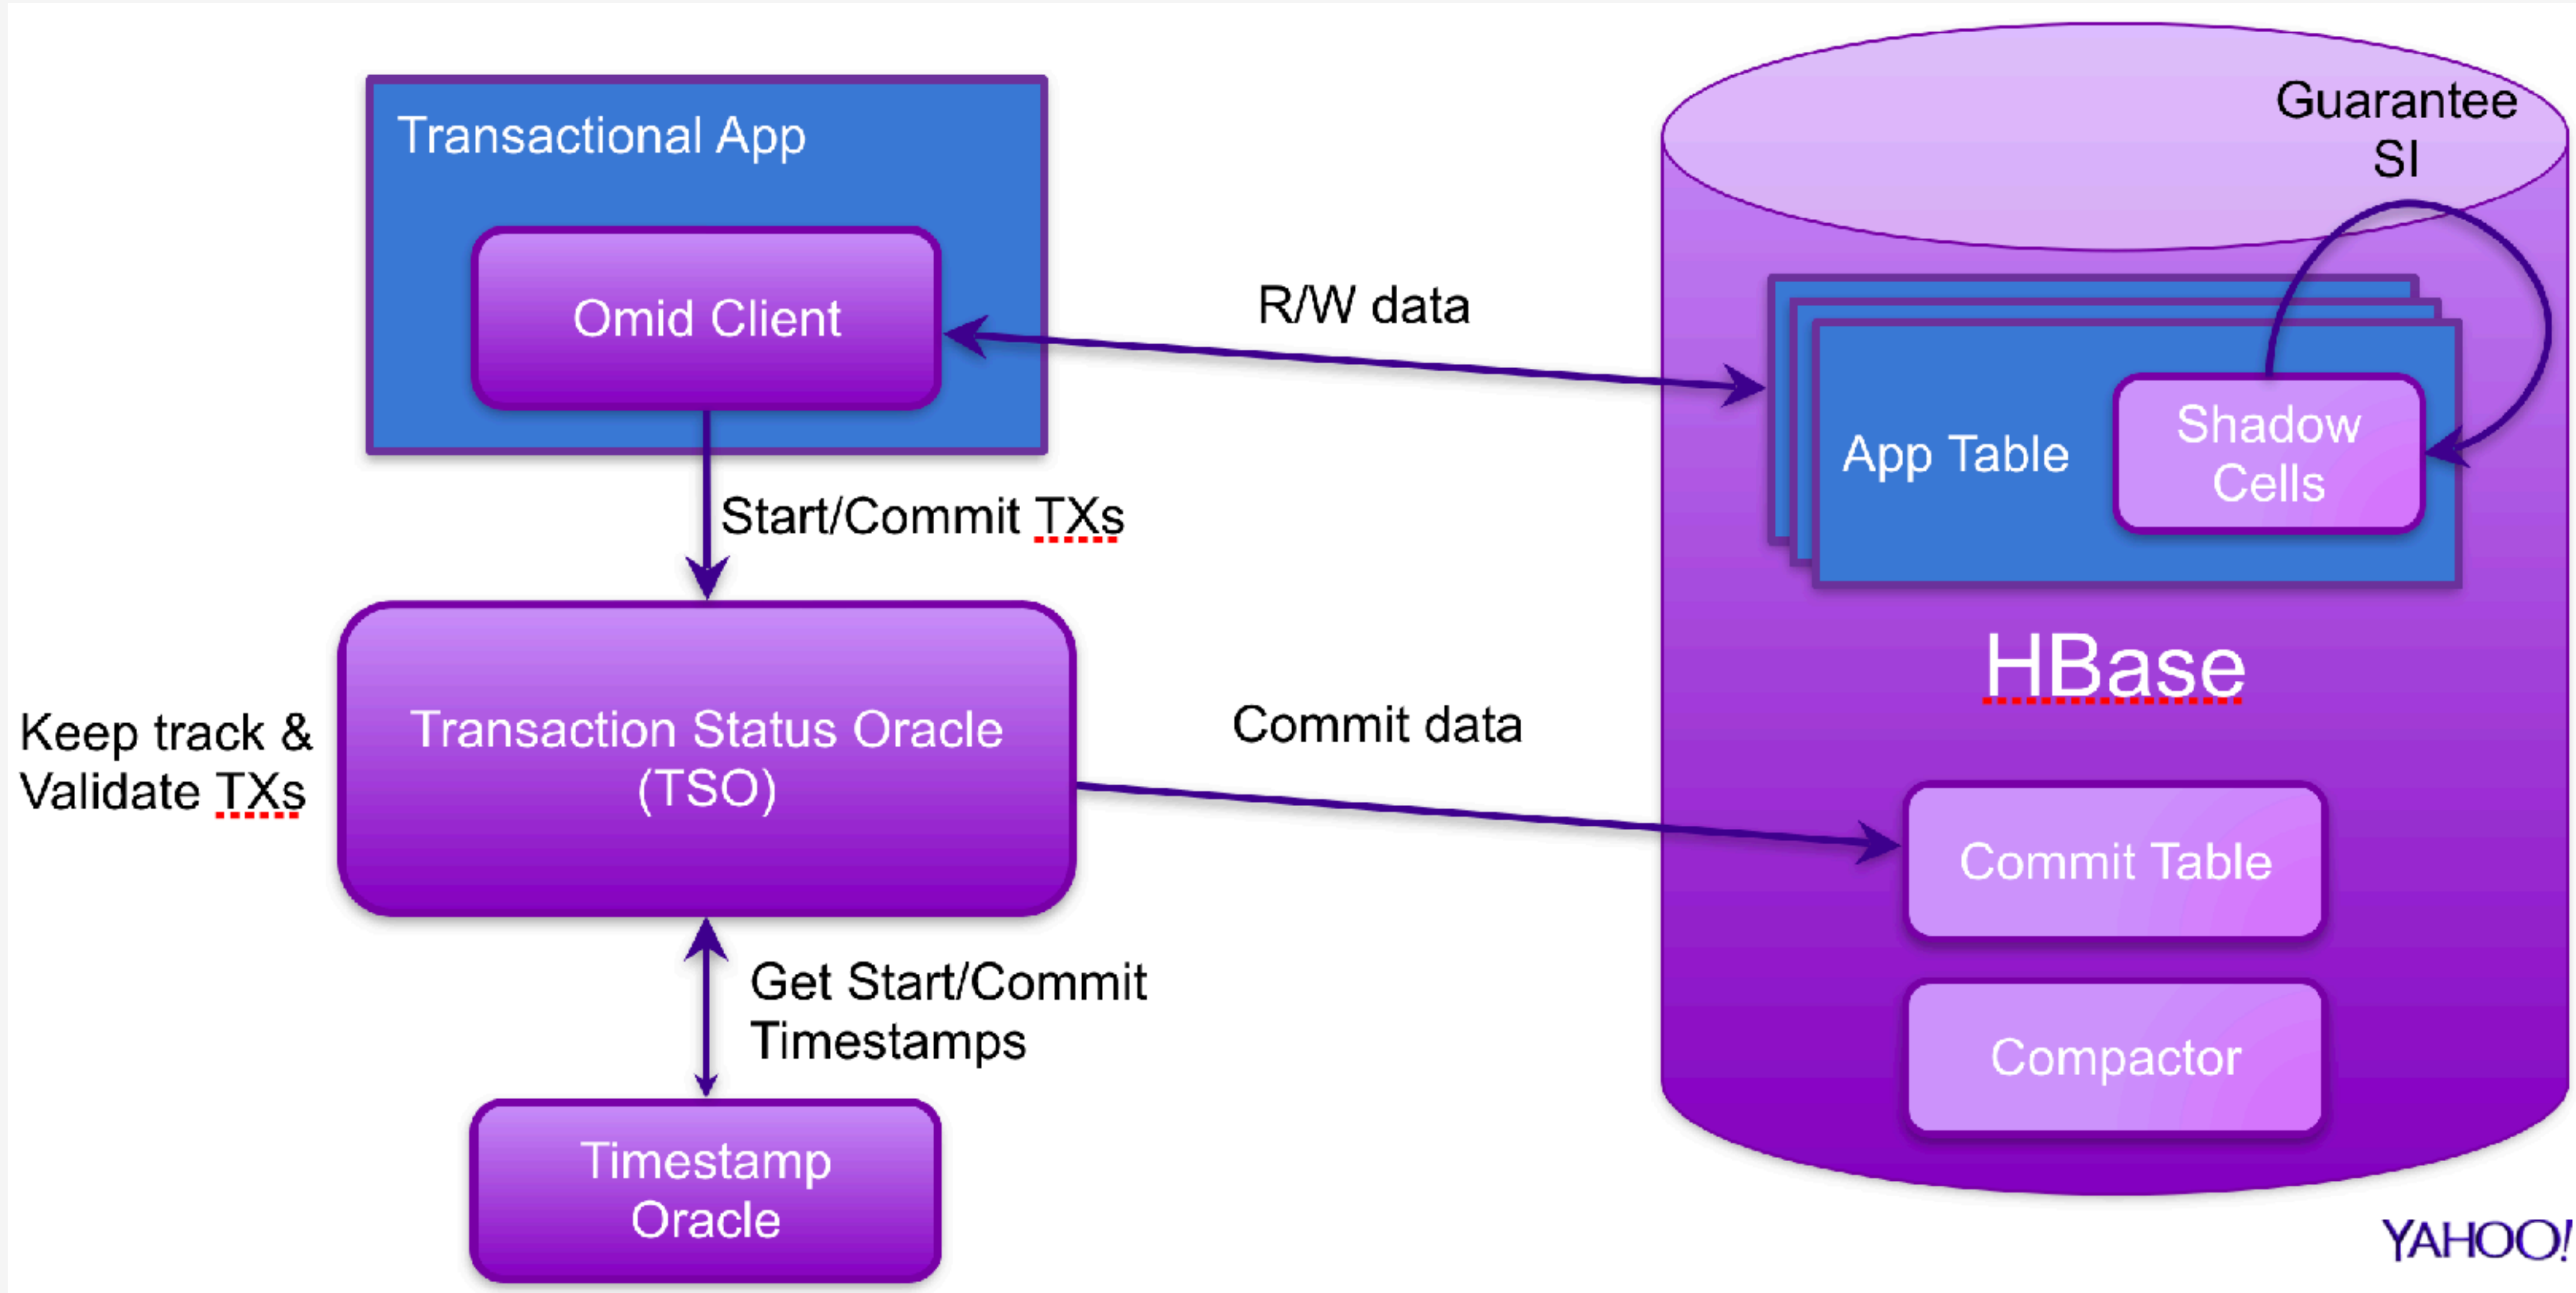
- Evolution of Omid based on the Google Percolator paper:

Daniel Peng, Frank Dabek: *Large-scale Incremental Processing Using Distributed Transactions and Notifications*, USENIX 2010.

- Idea: Move as much transaction state as possible into HBase
  - Shadow cells represent the state of a transaction
  - One shadow cell for every data cell written
  - Track committed transactions in an HBase table
  - Transaction Manager (TSO) has only 3 tasks
    - issue transaction IDs
    - conflict detection
    - write to commit table



# Apache Omid



# Apache Omid

Client A

Tx  
Manager

HBase

Region Server

x:10    37: commit.40

Region Server

Commits

37: 40

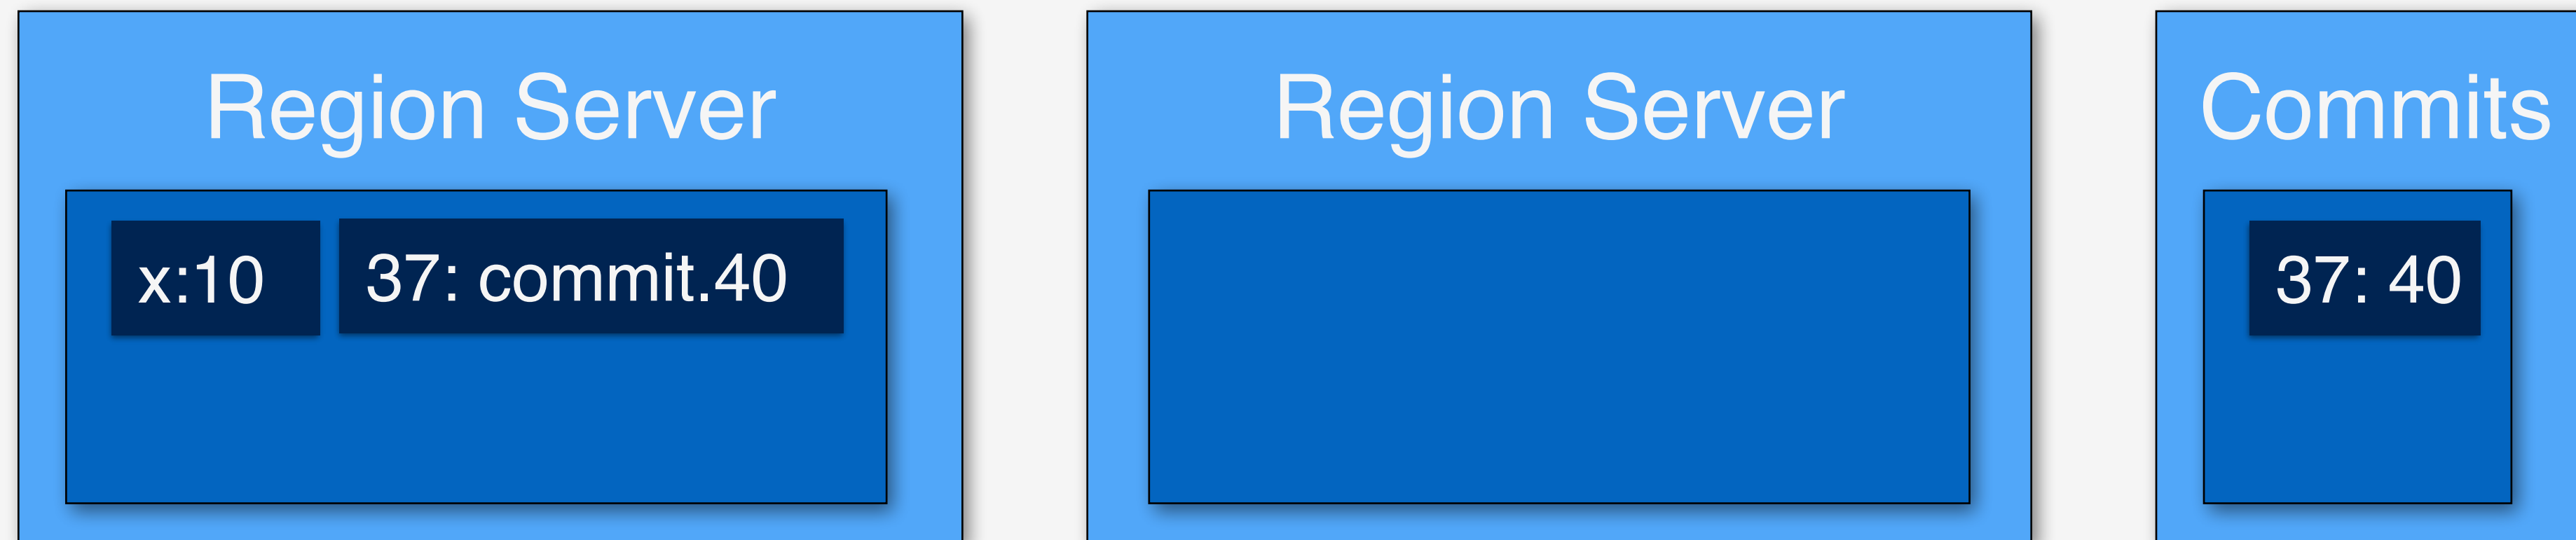




# Apache Omid

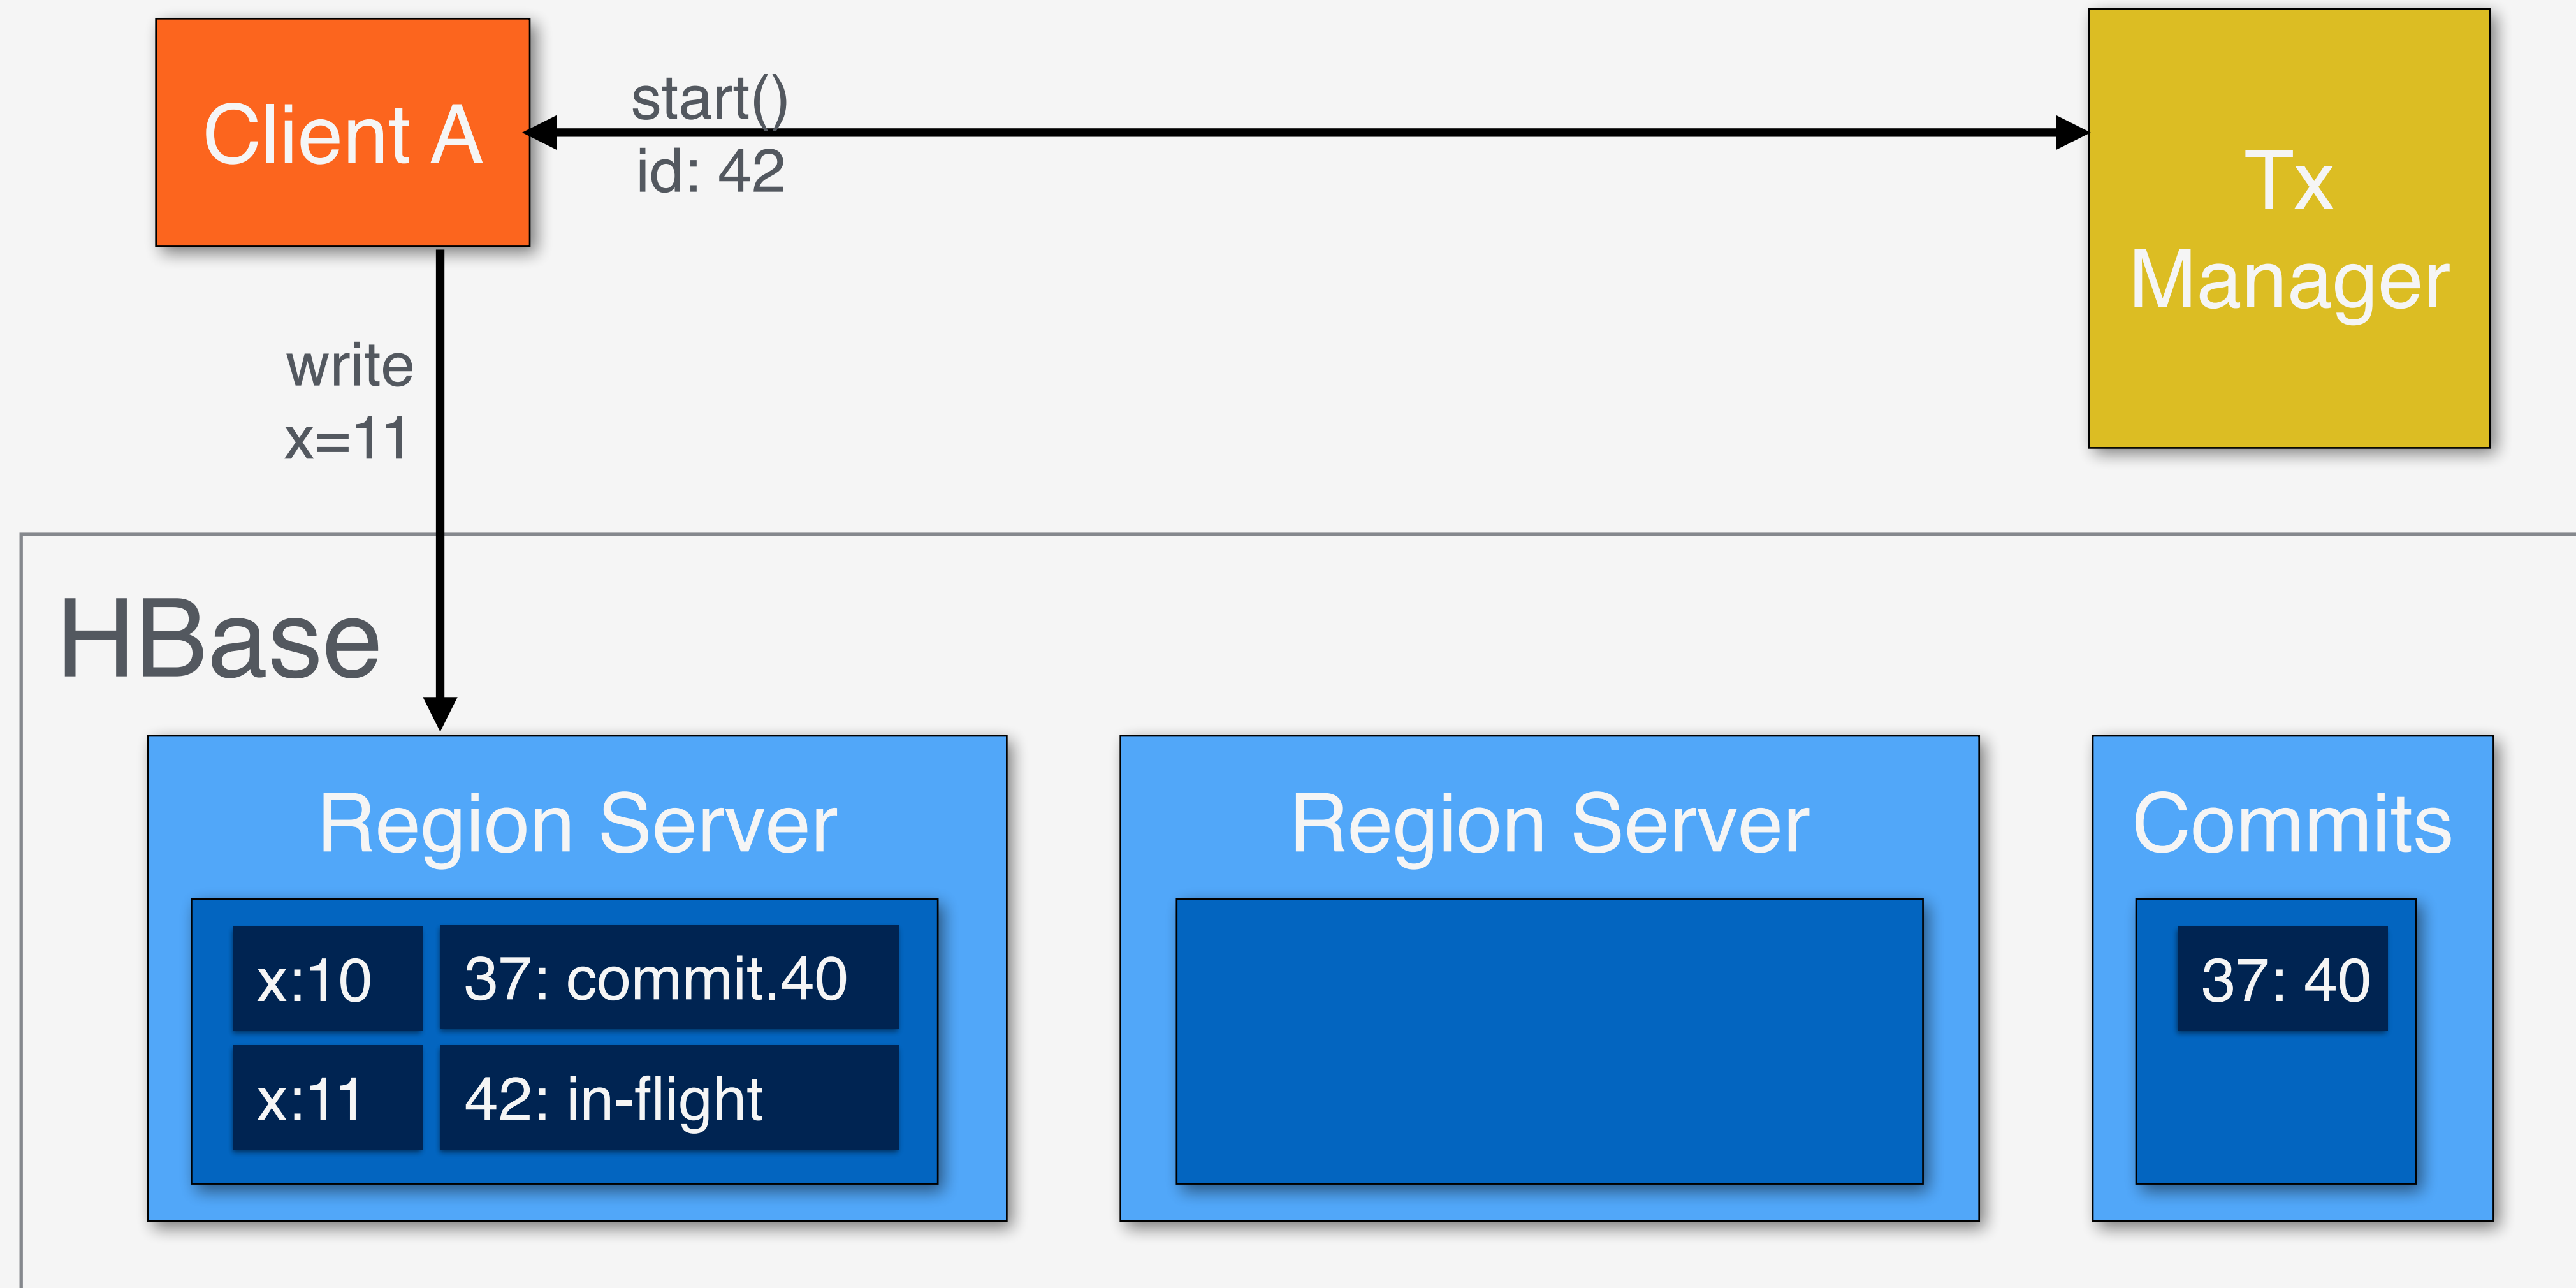


## HBase

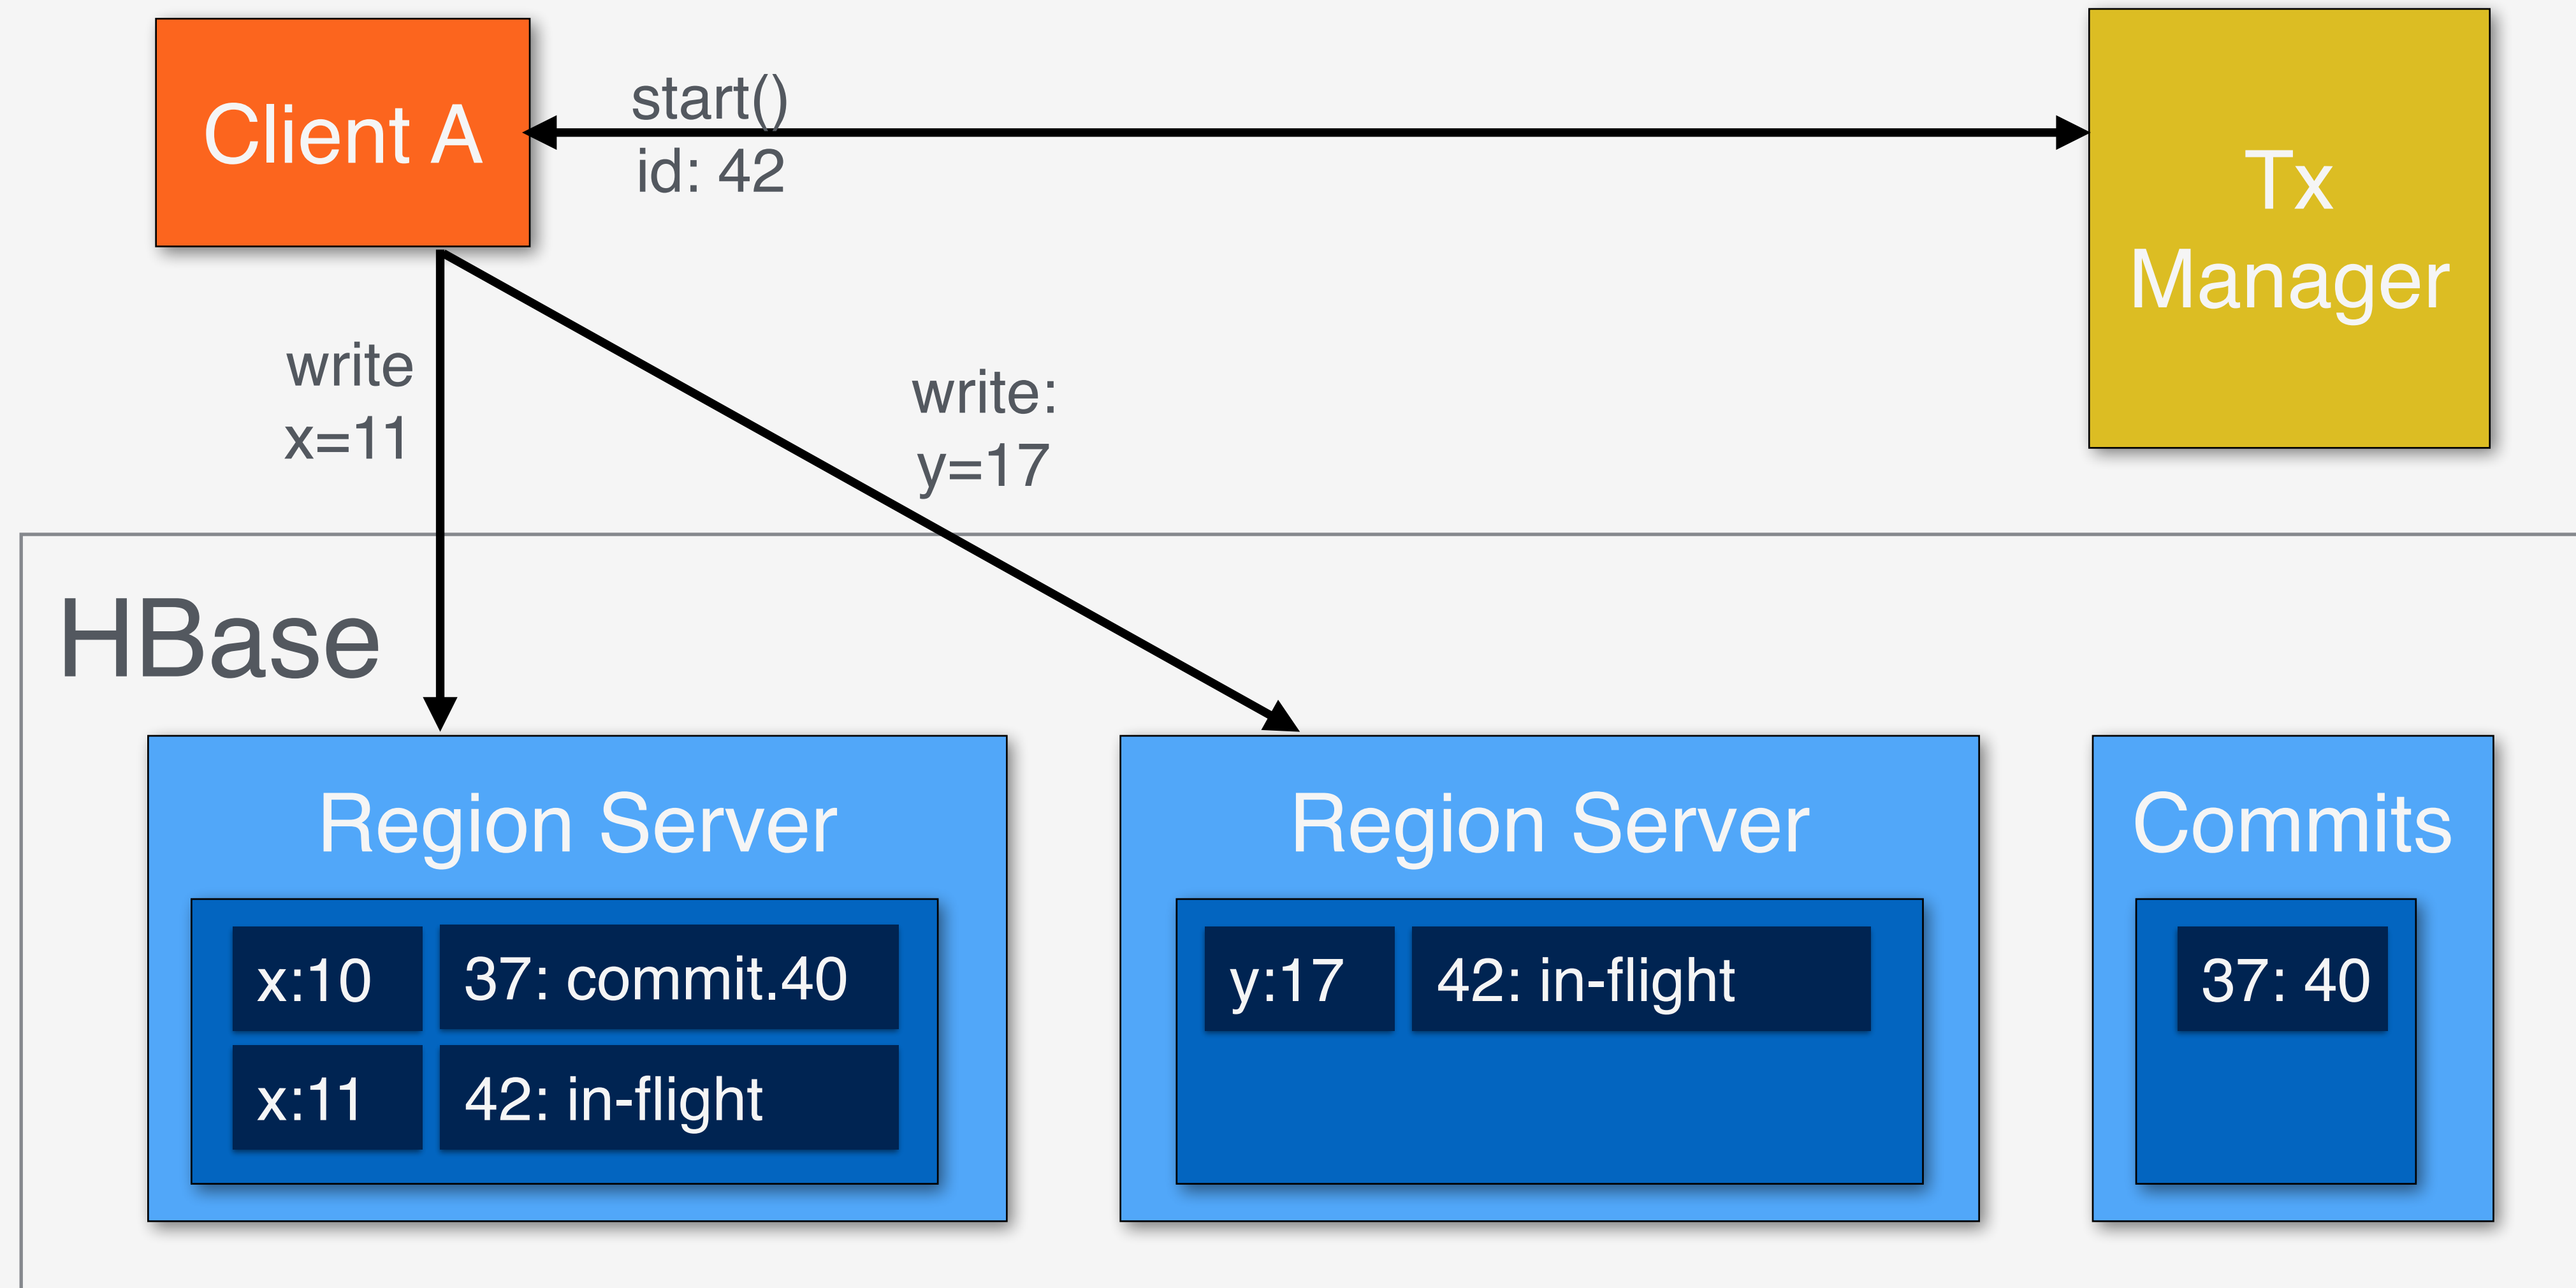




# Apache Omid



# Apache Omid



# Apache Omid

Client B

Tx  
Manager

## HBase

Region Server

x:10	37: commit.40
x:11	42: in-flight

Region Server


y:17	42: in-flight
------	---------------

Commits

37: 40
--------



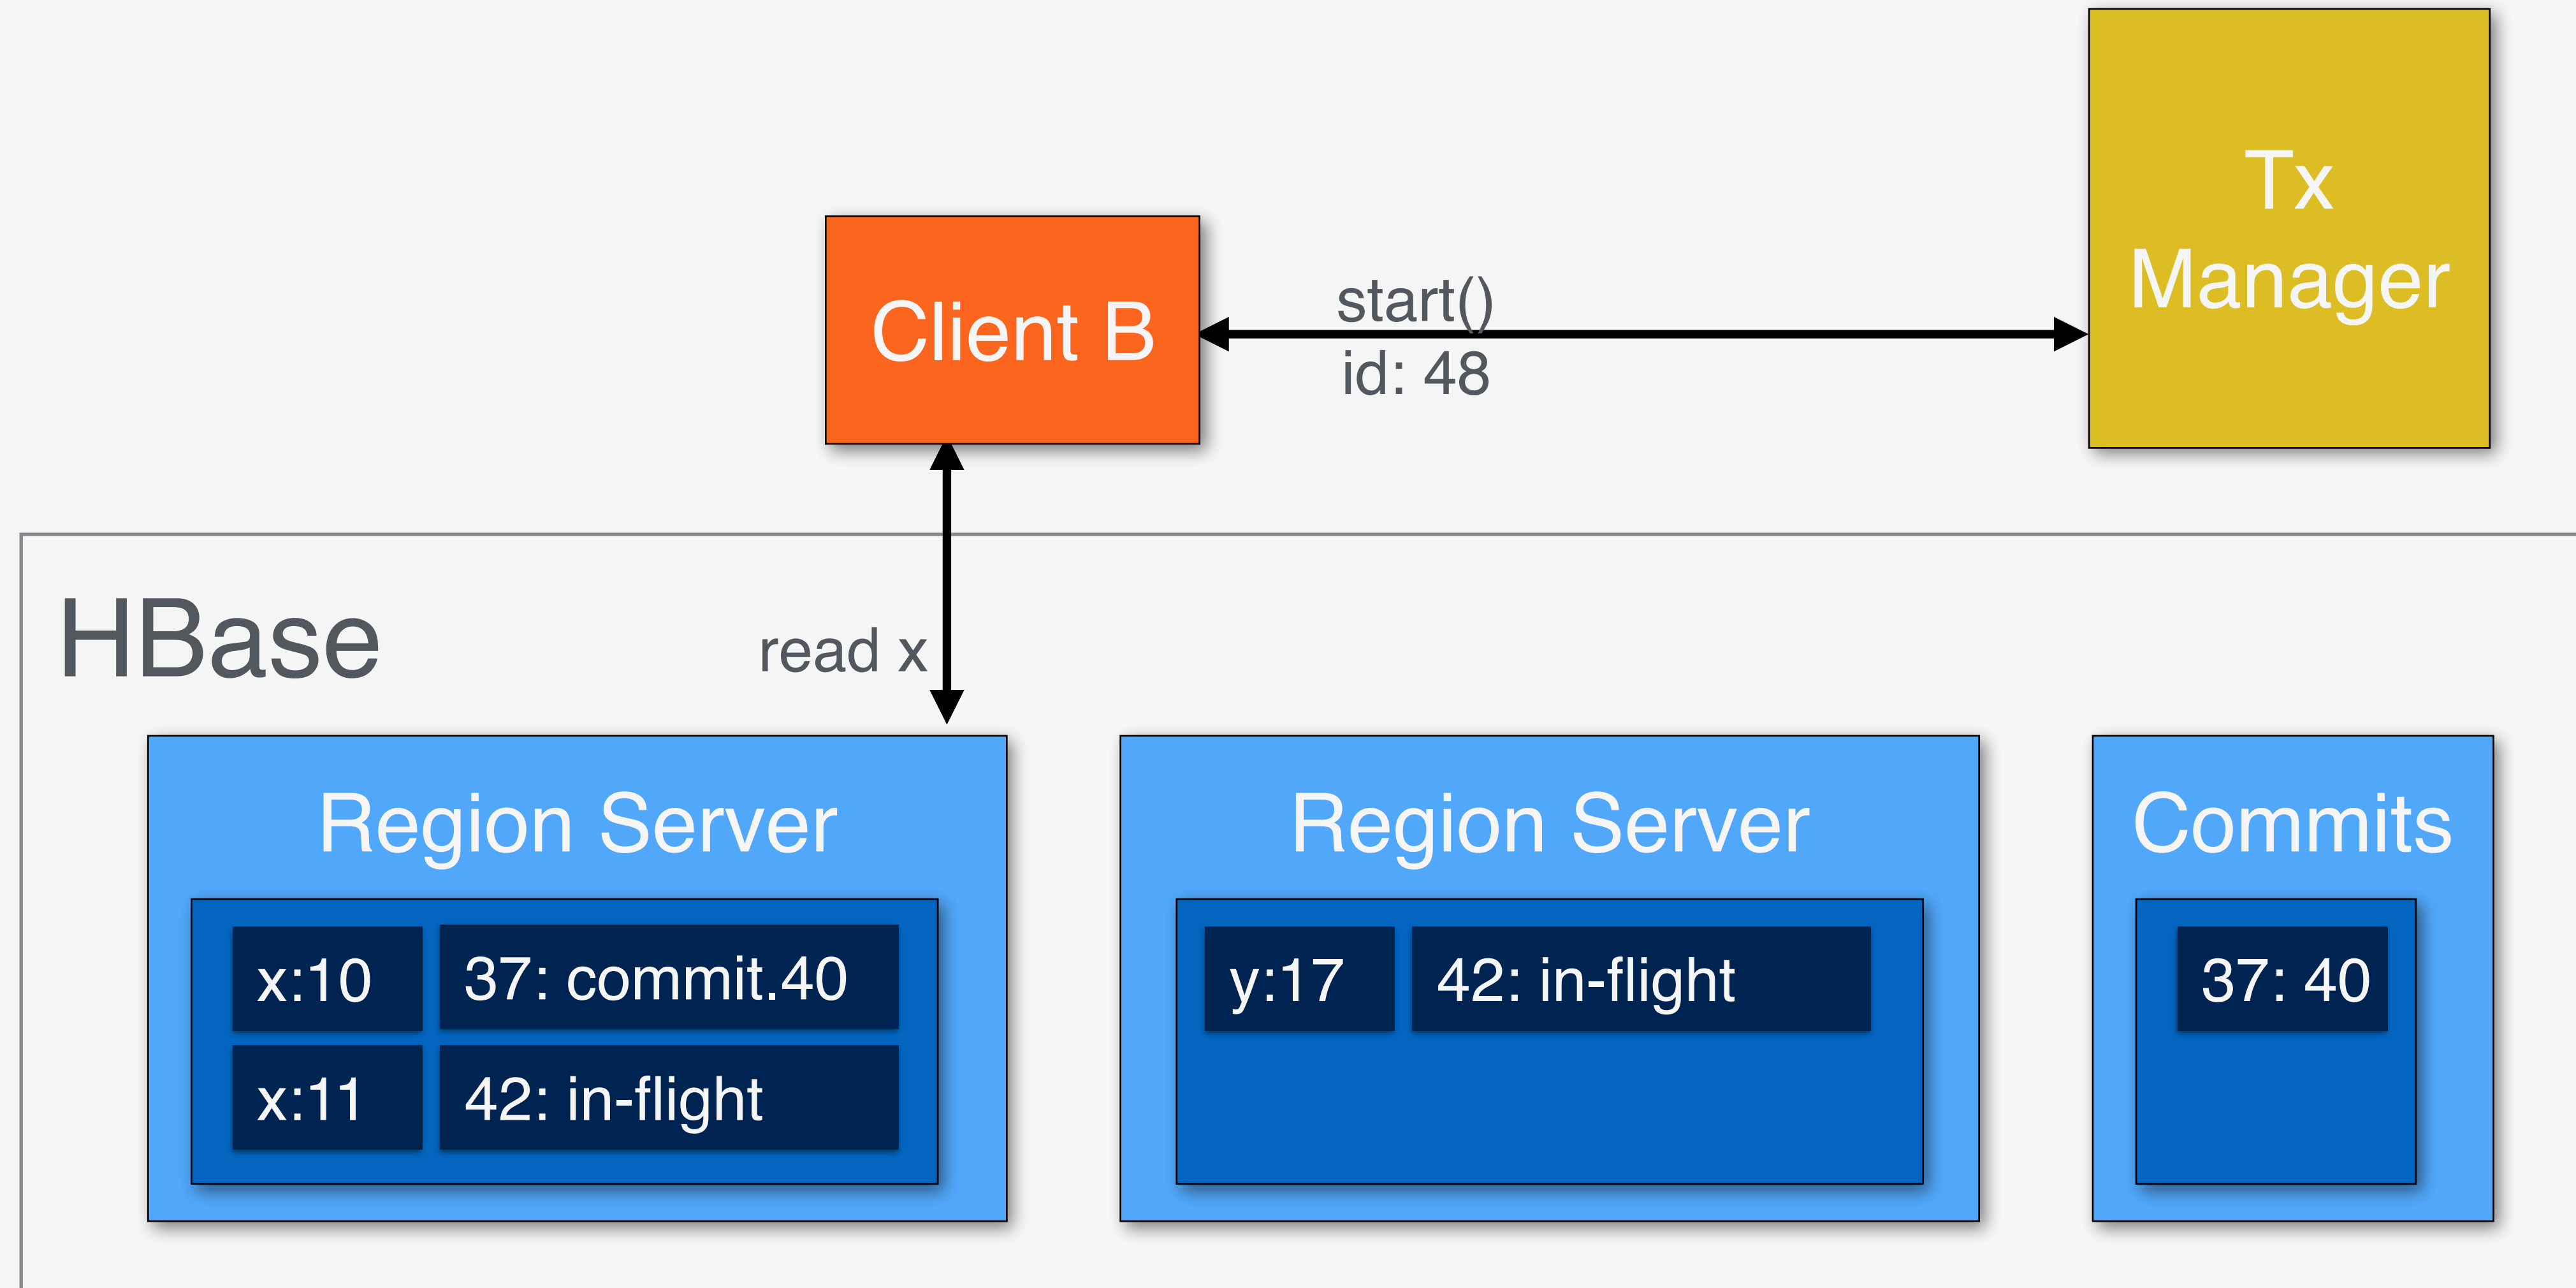
# Apache Omid




## HBase



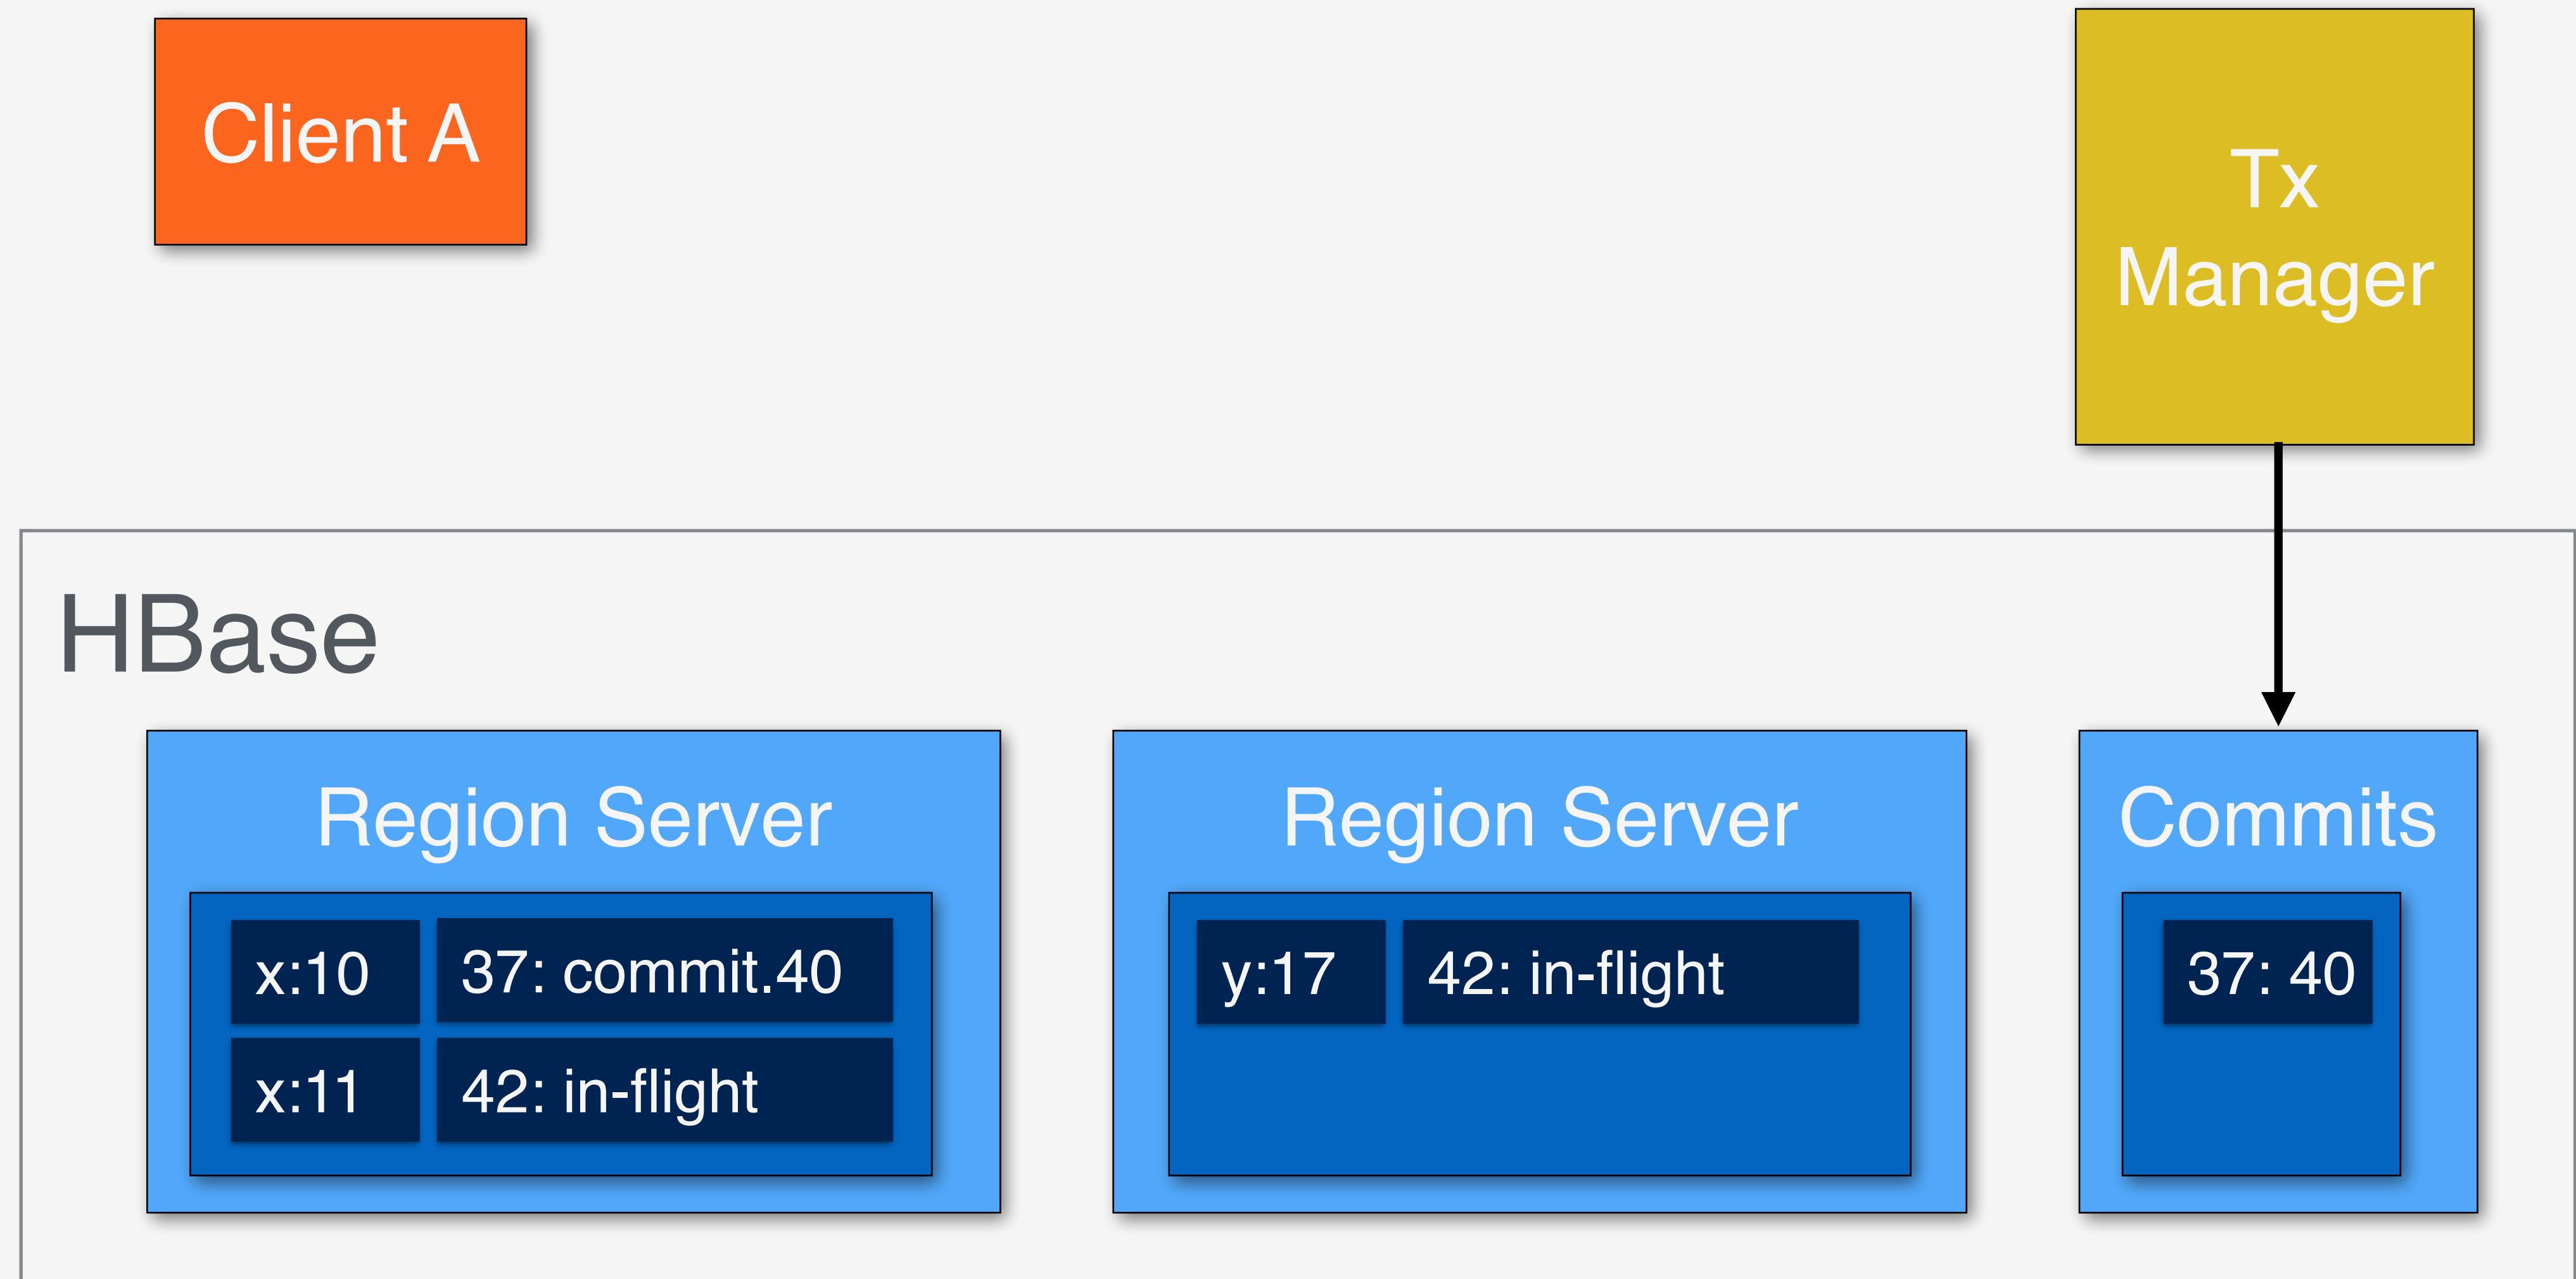
# Apache Omid



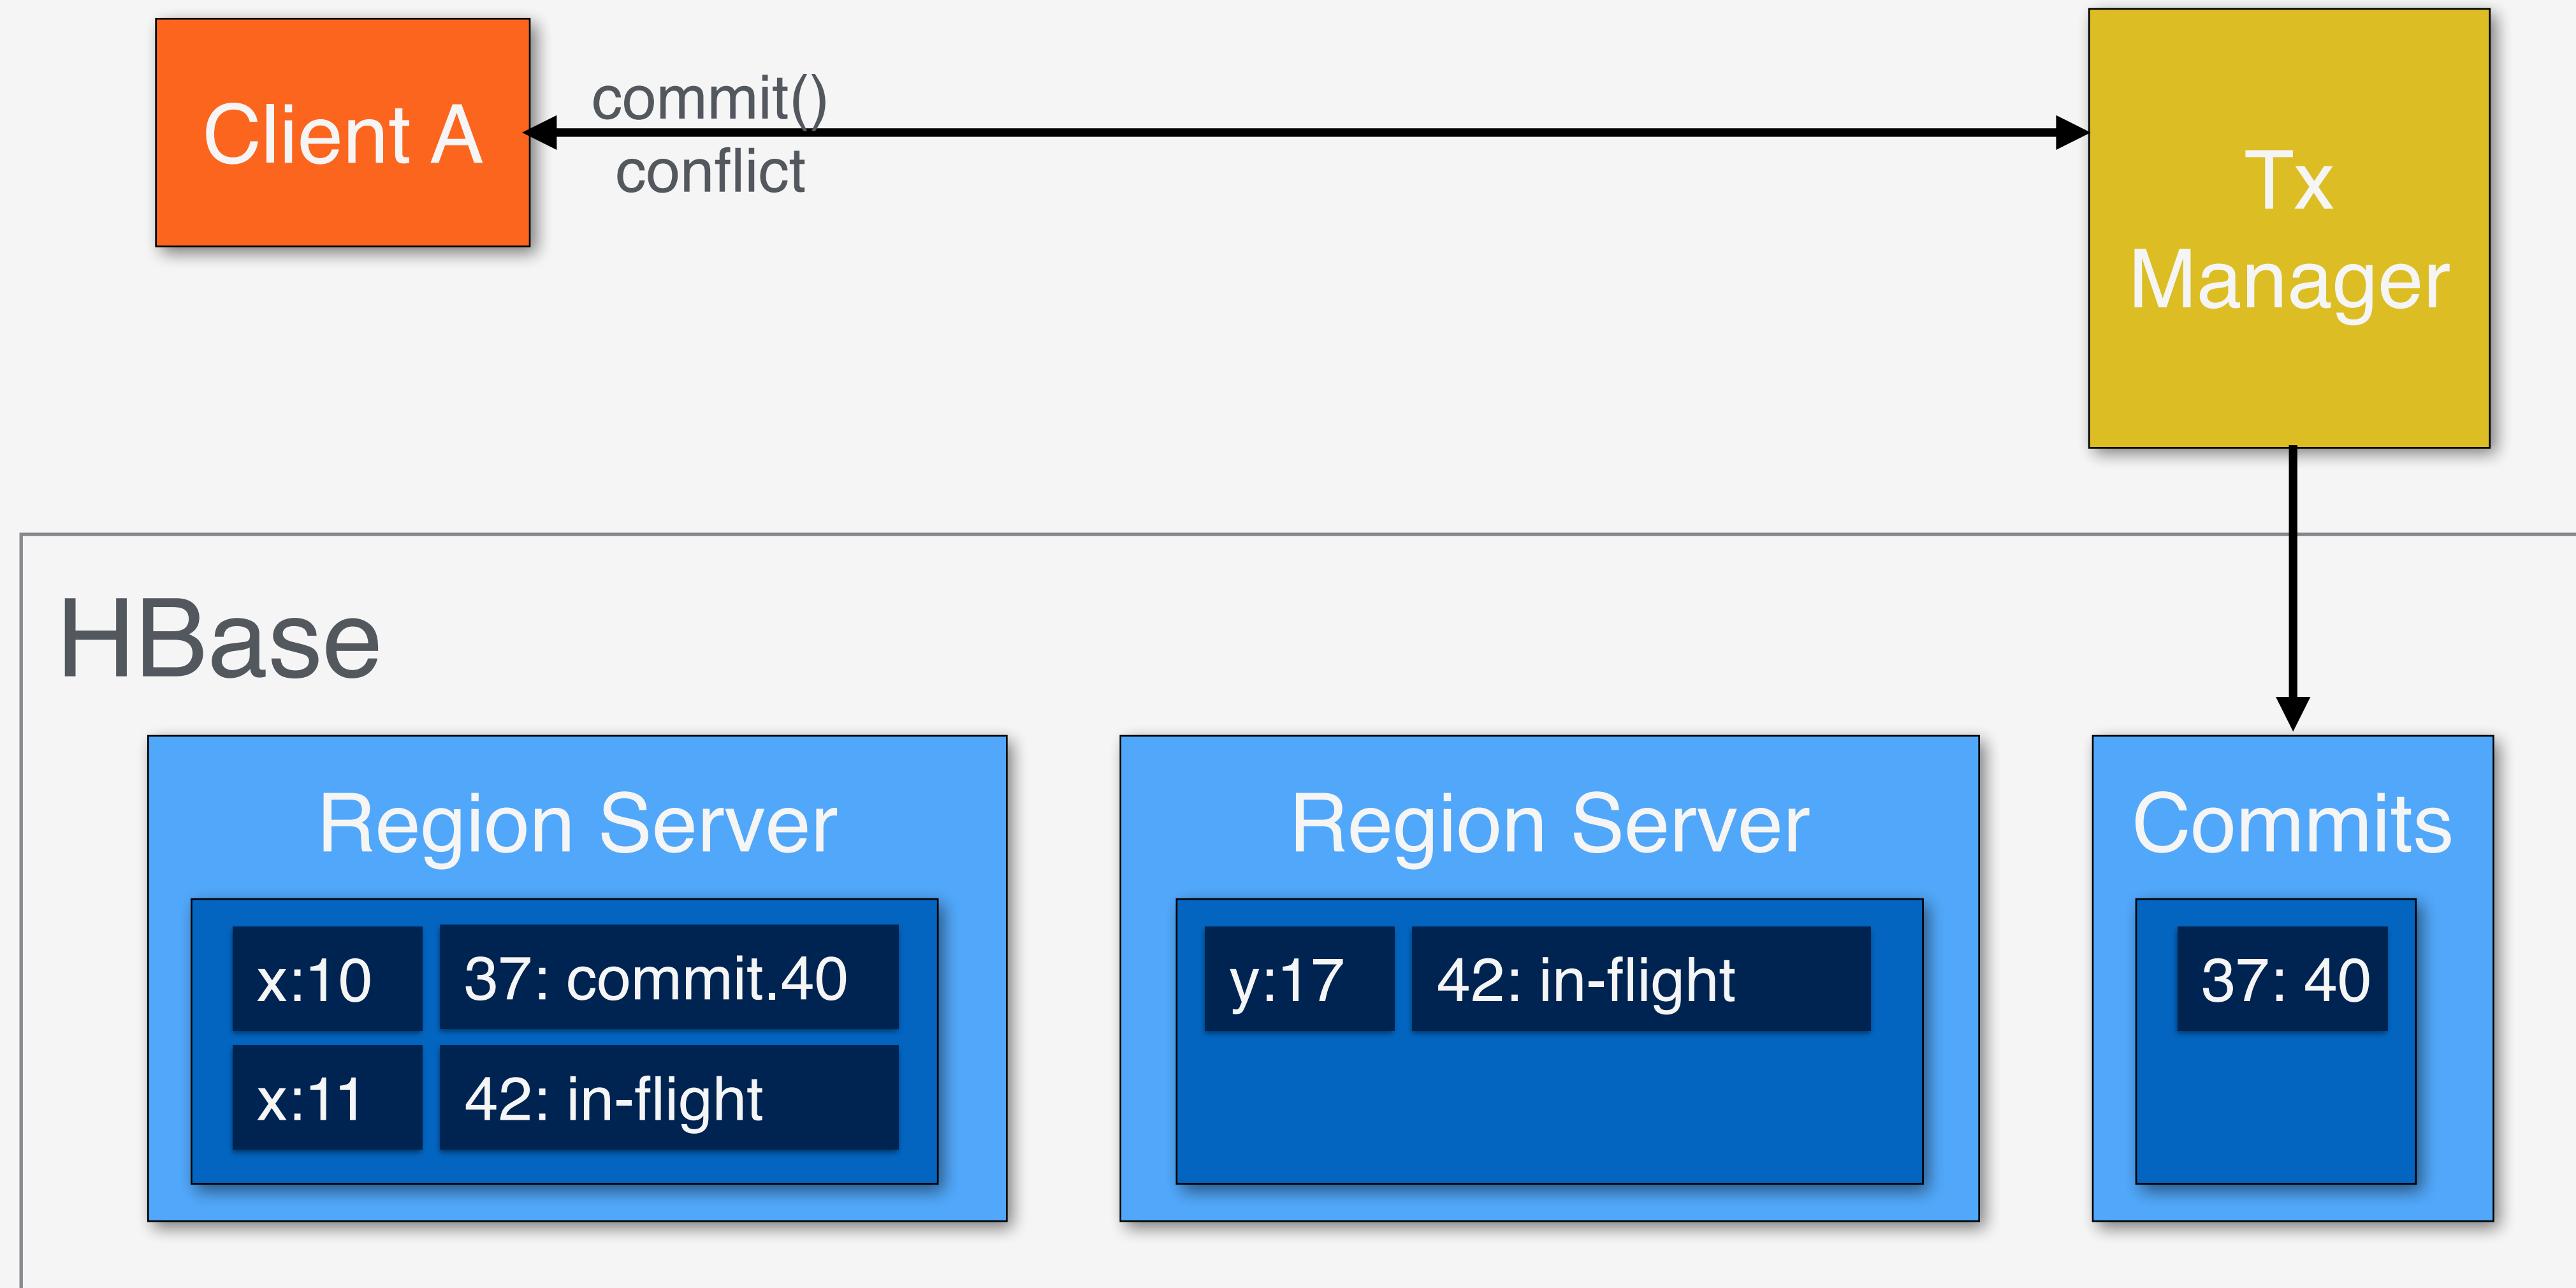
# Apache Omid



# Apache Omid

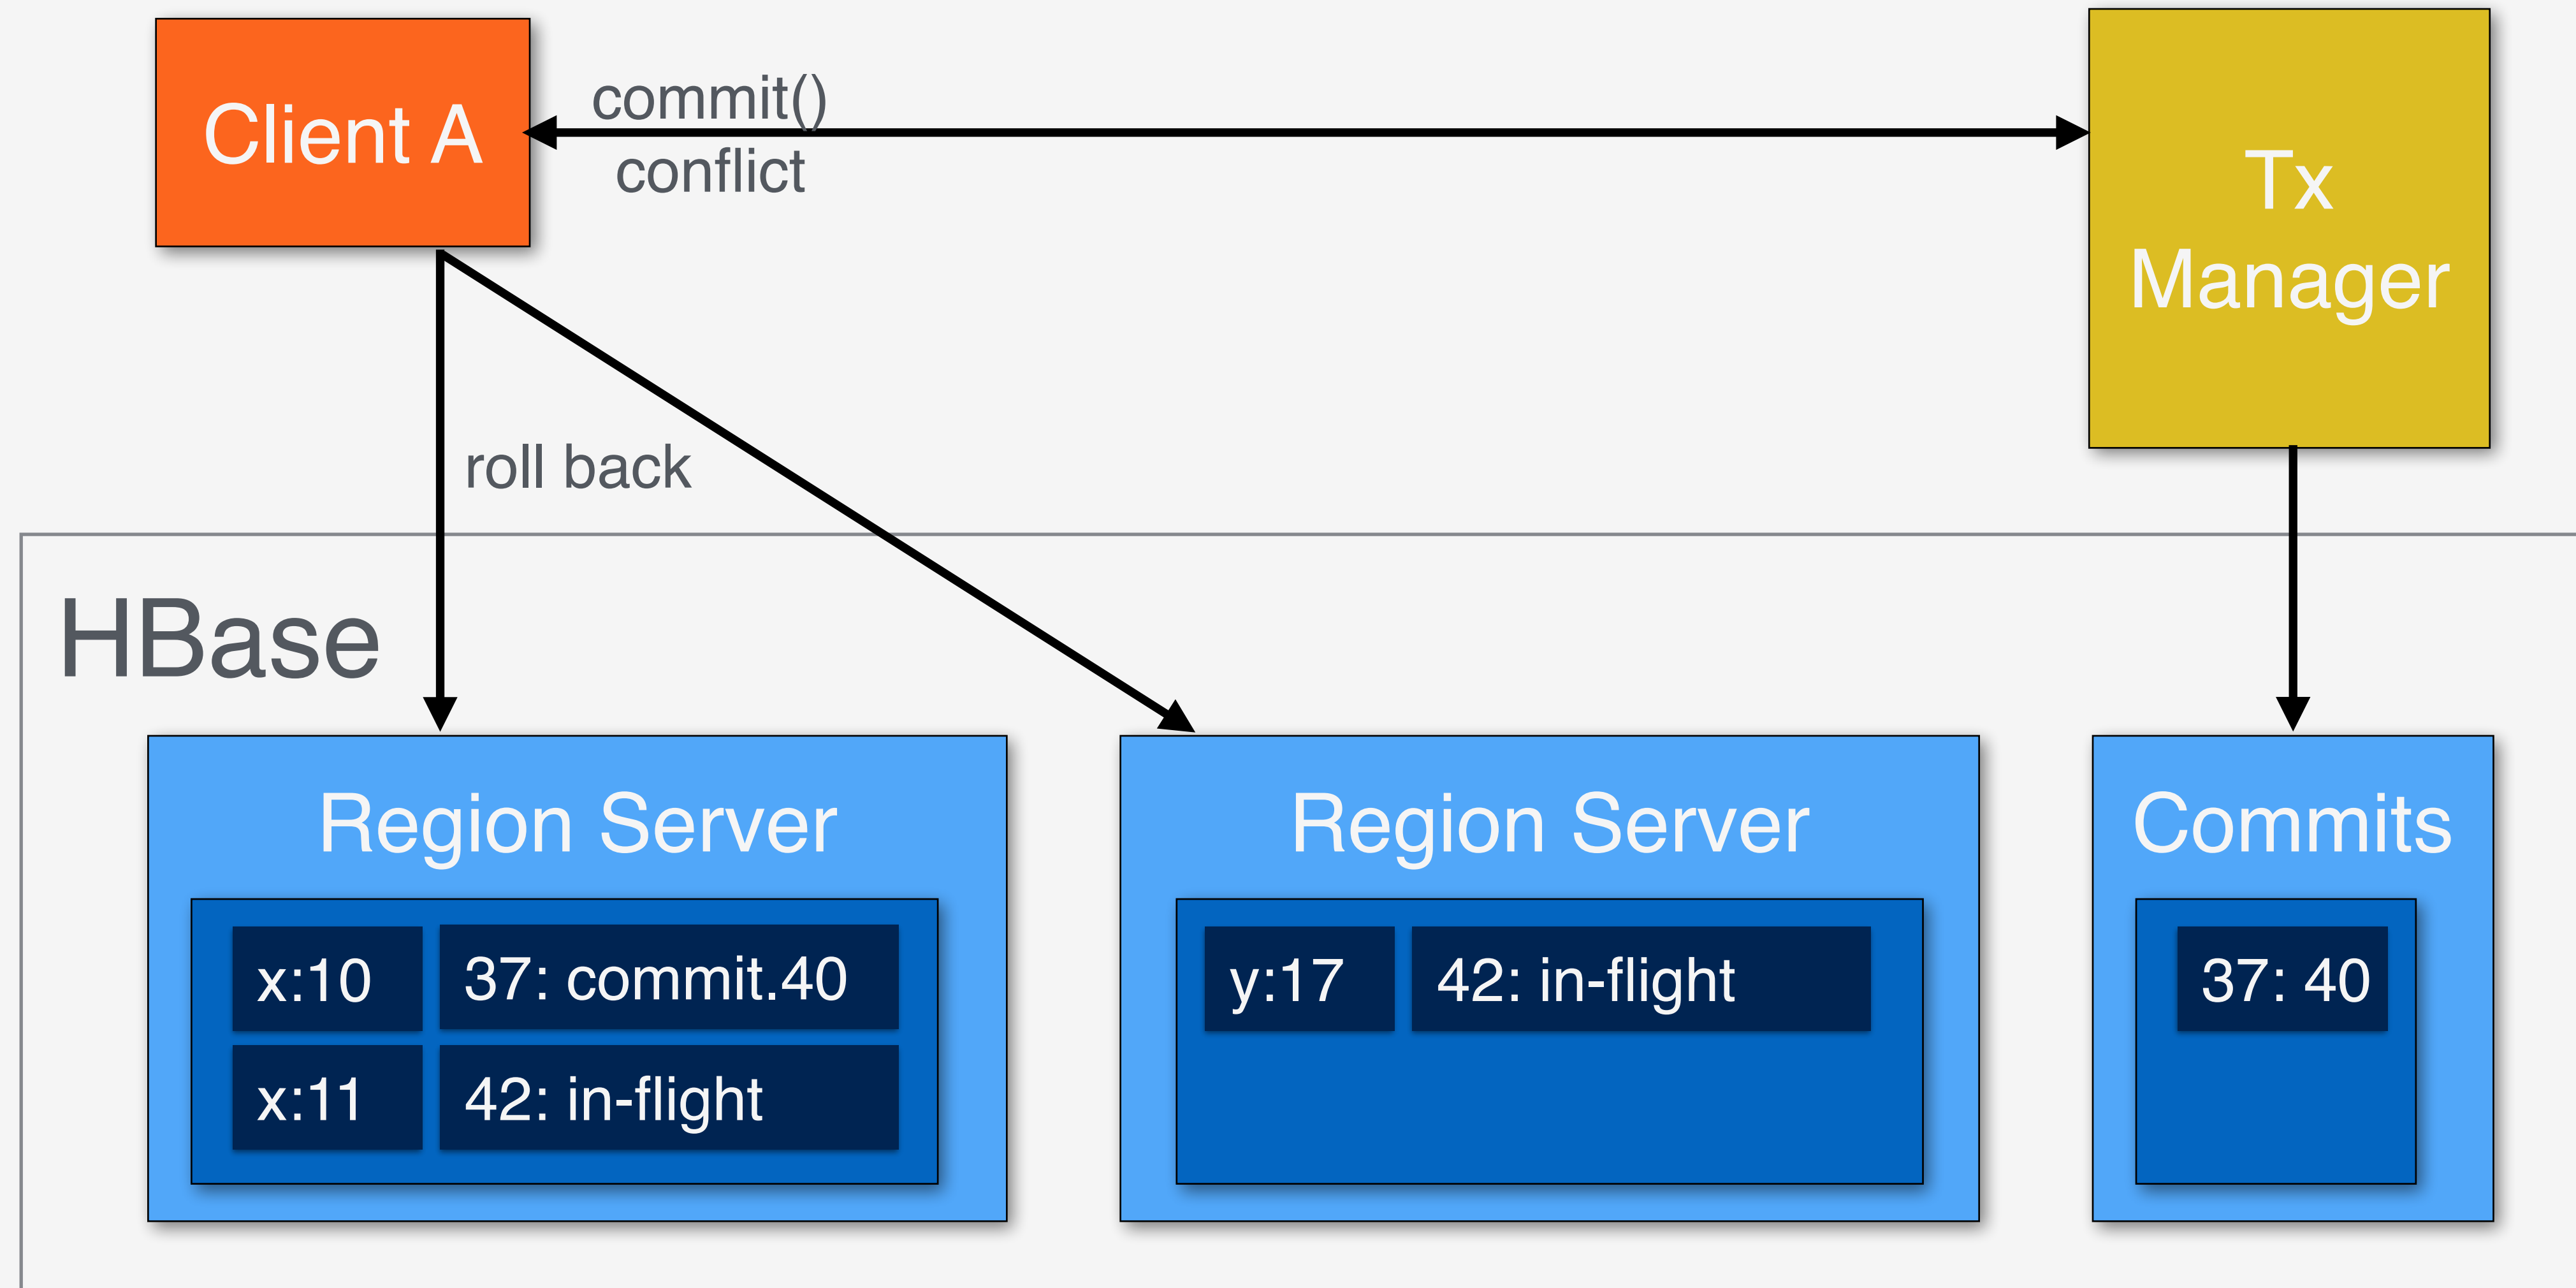


# Apache Omid

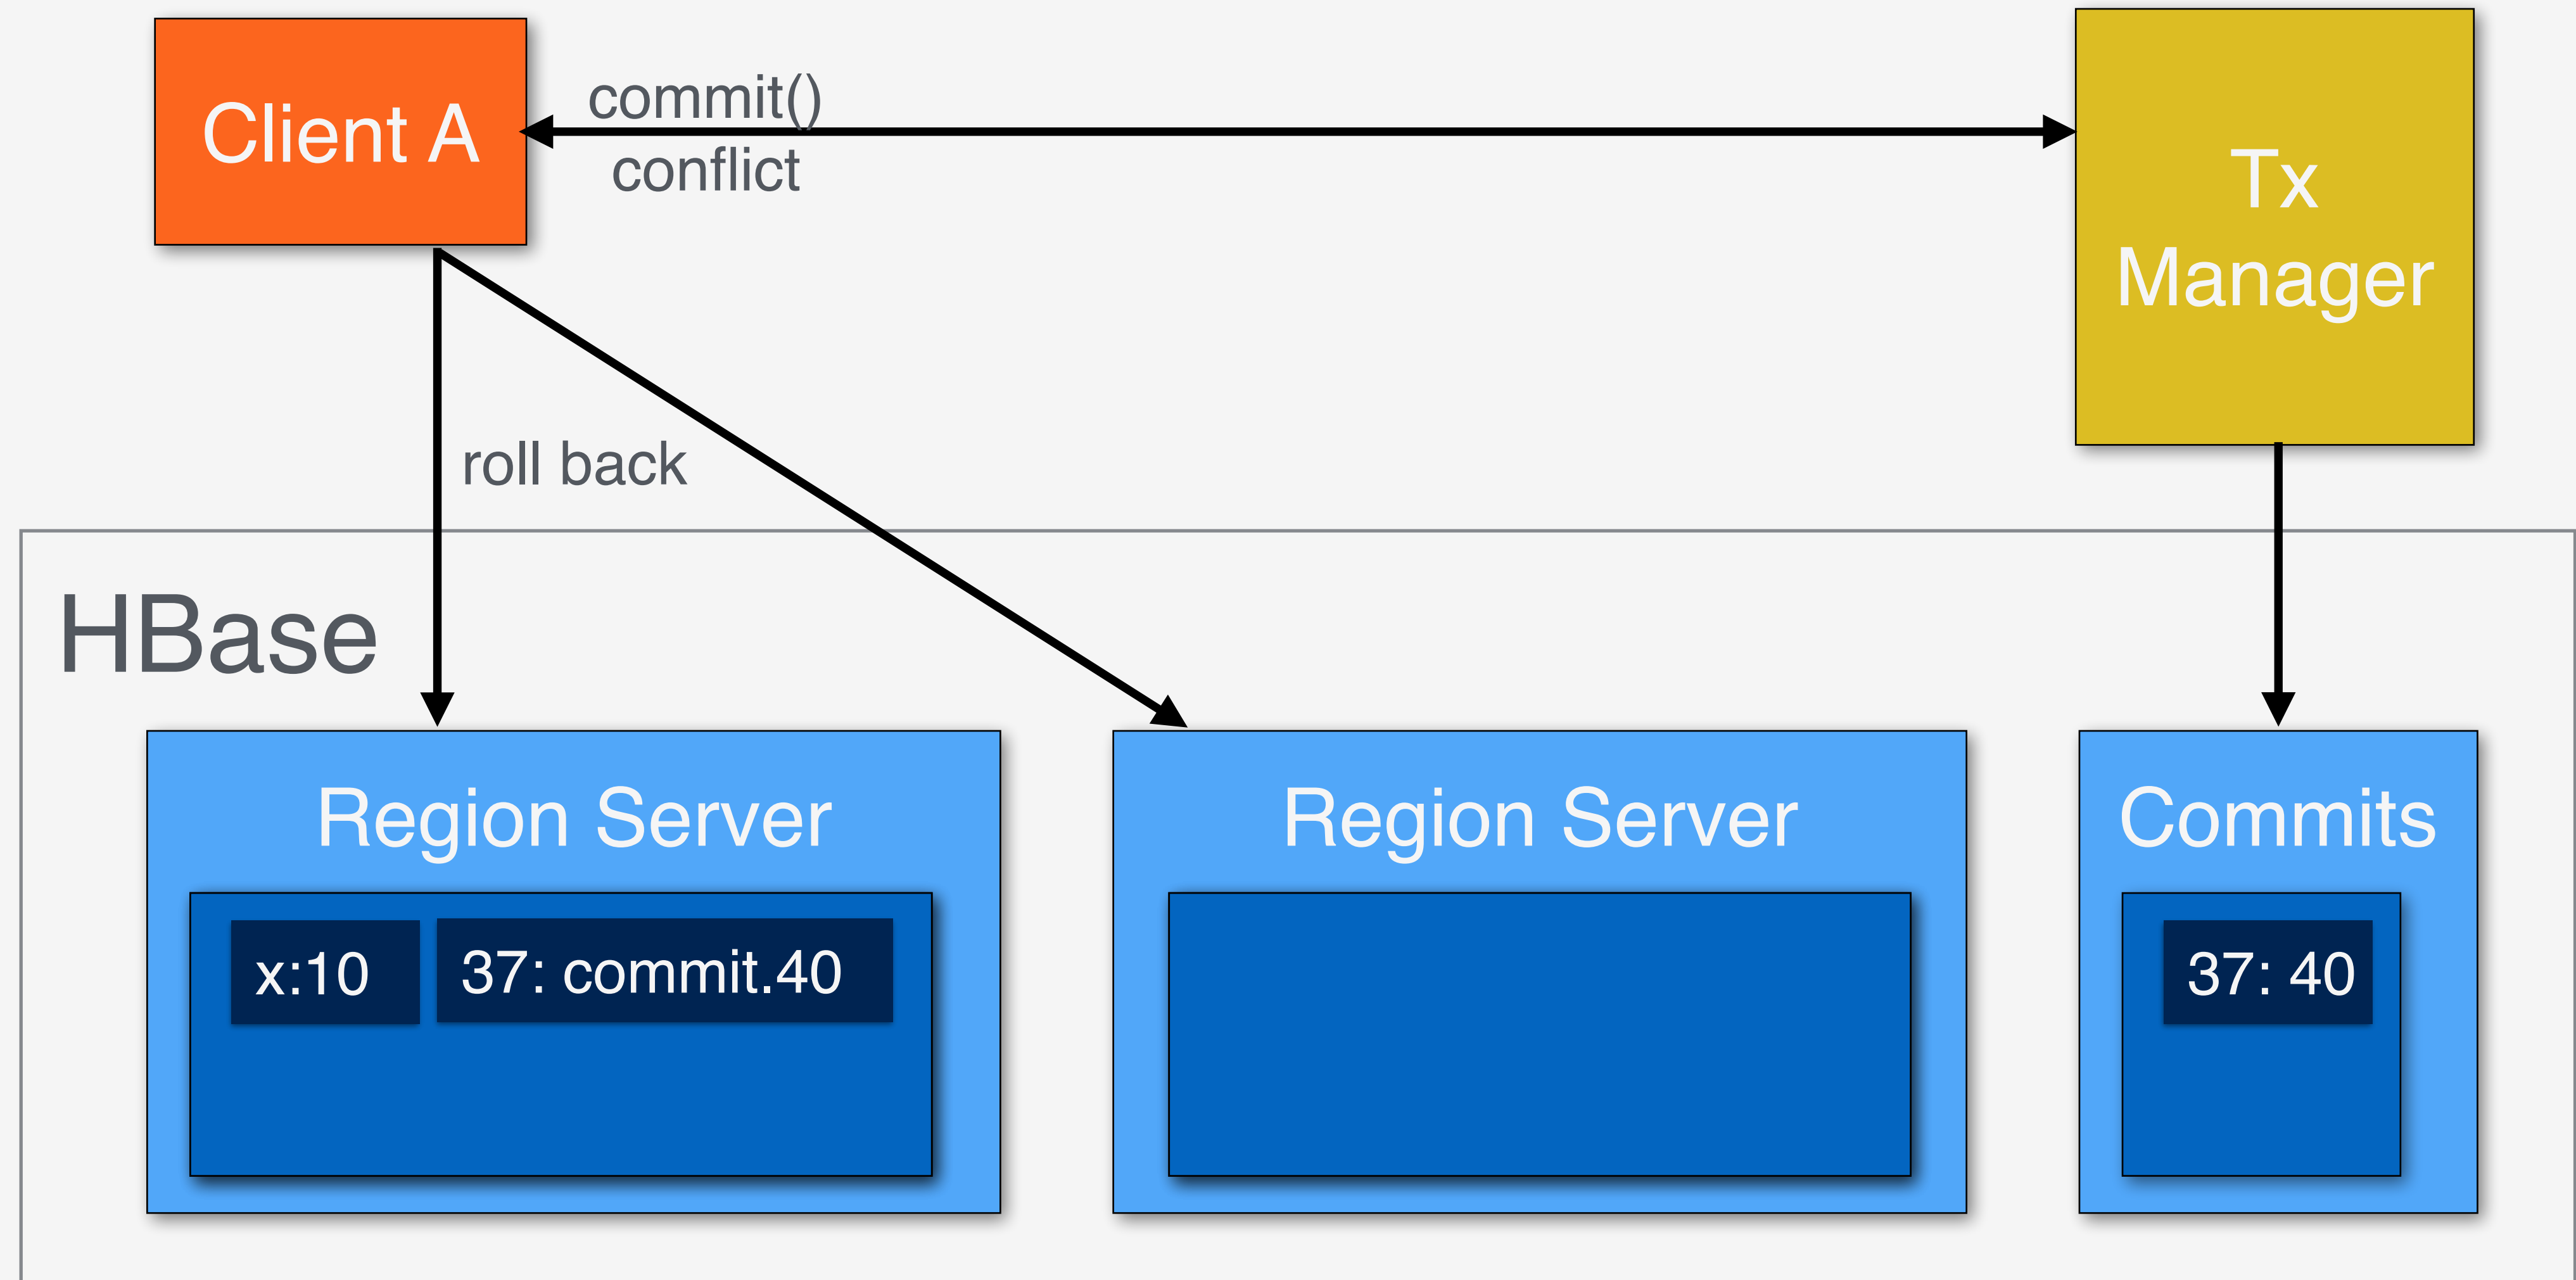




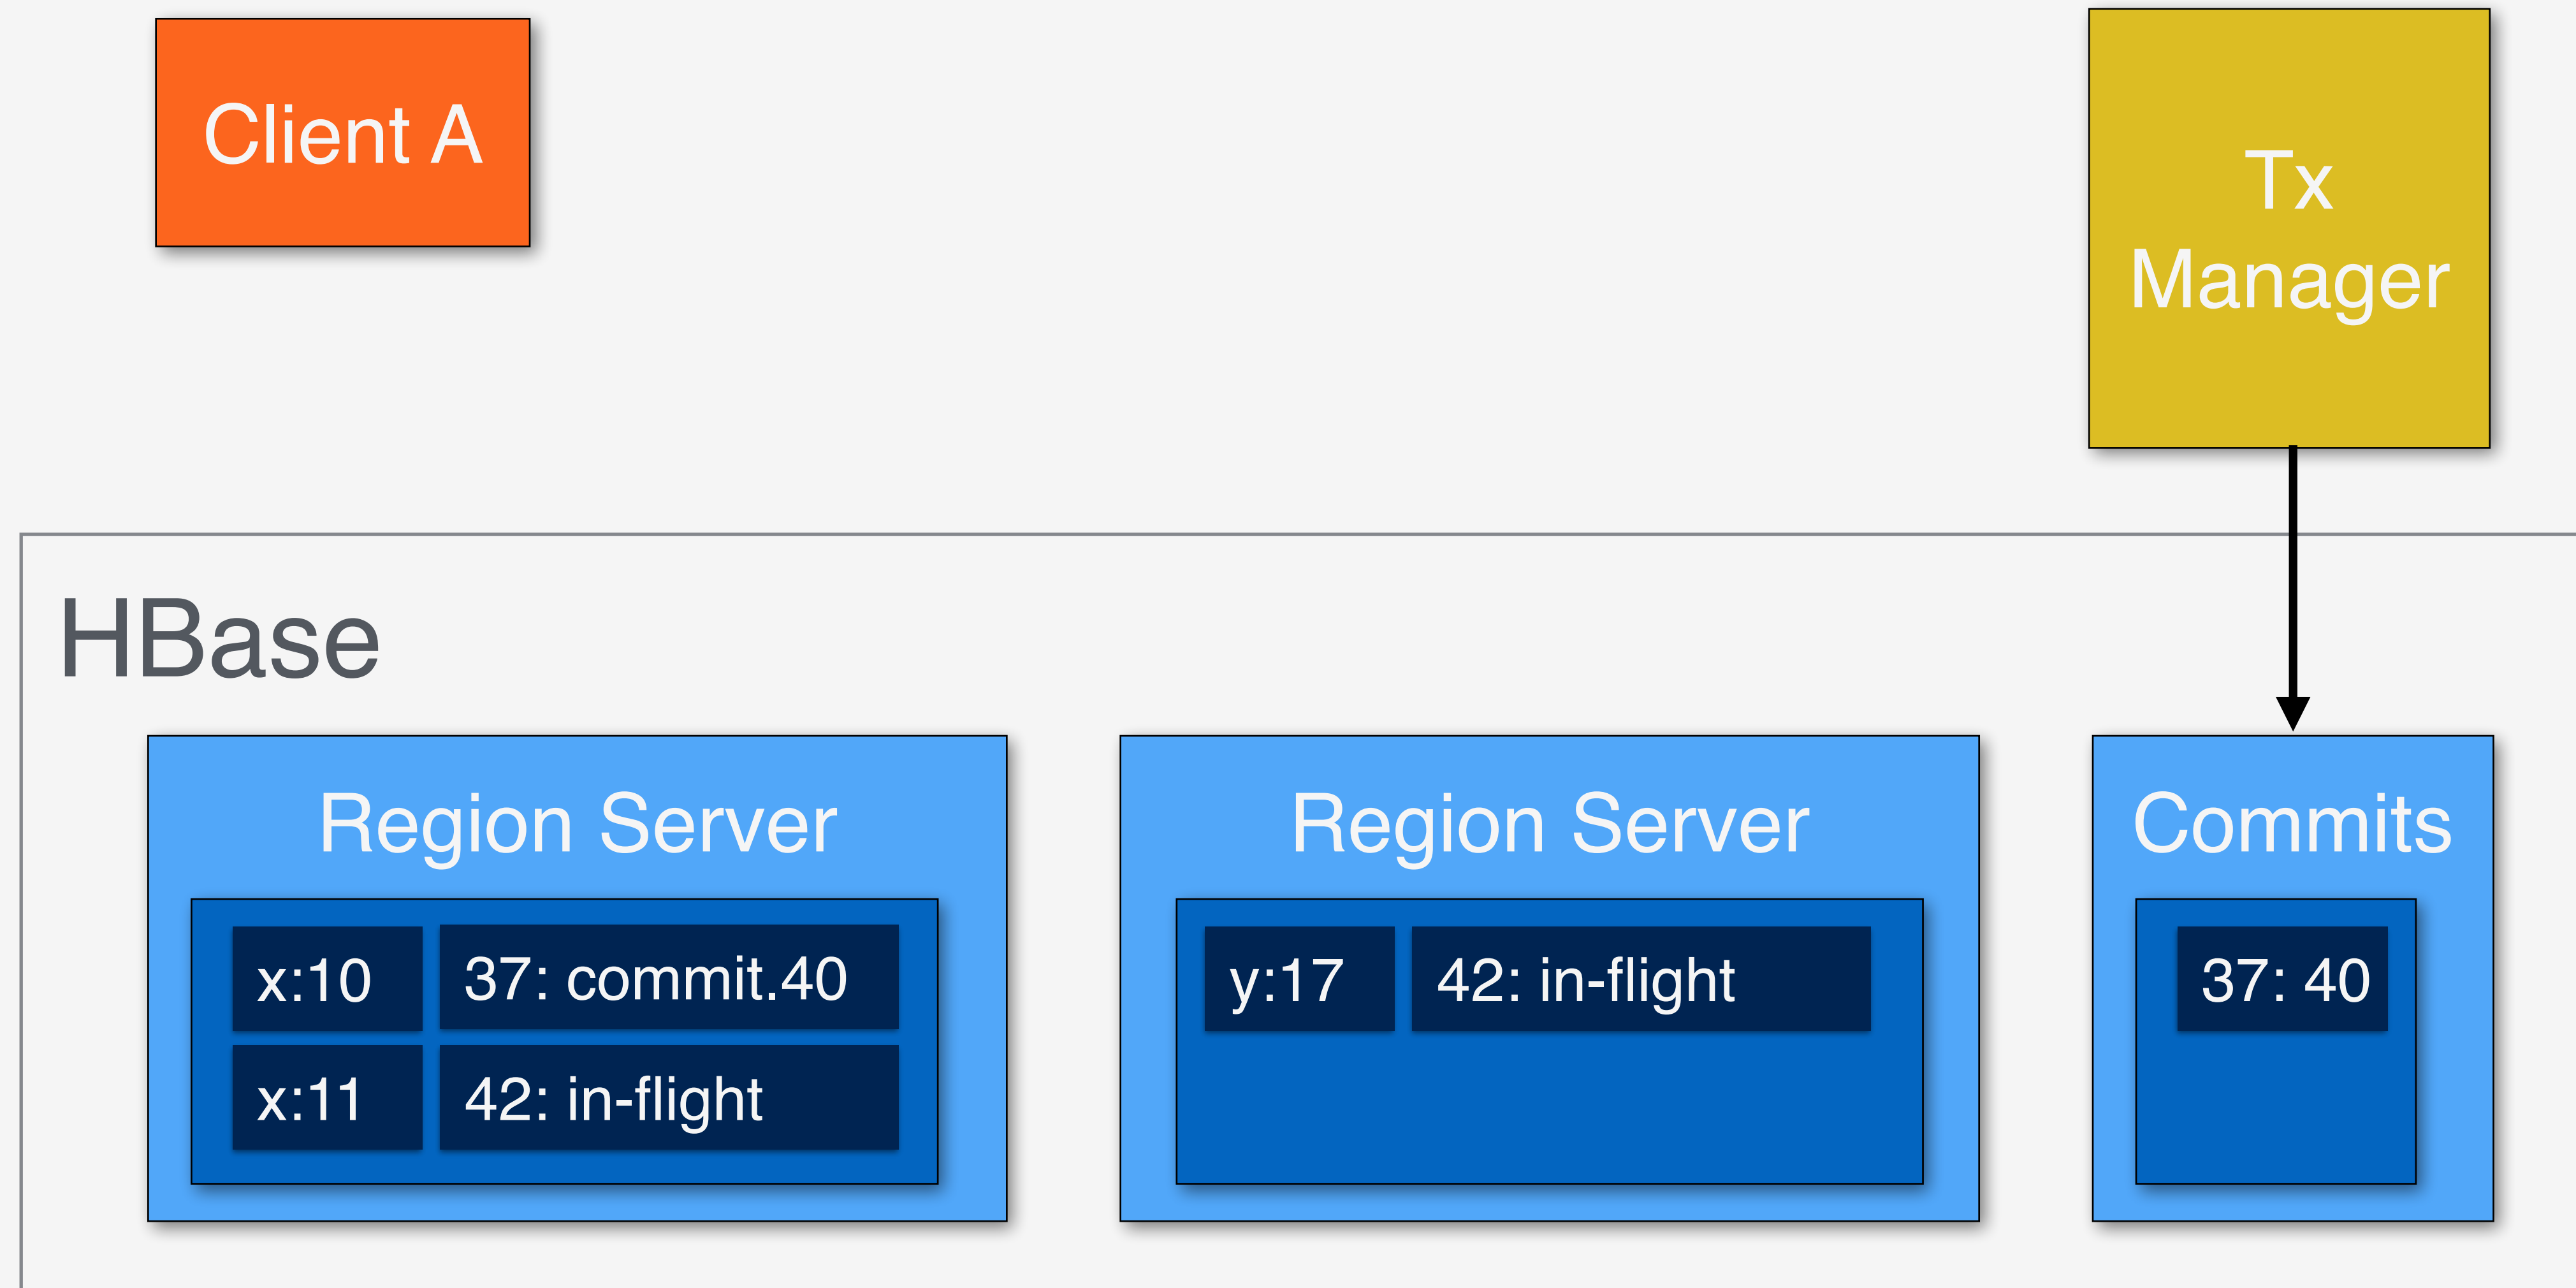
# Apache Omid



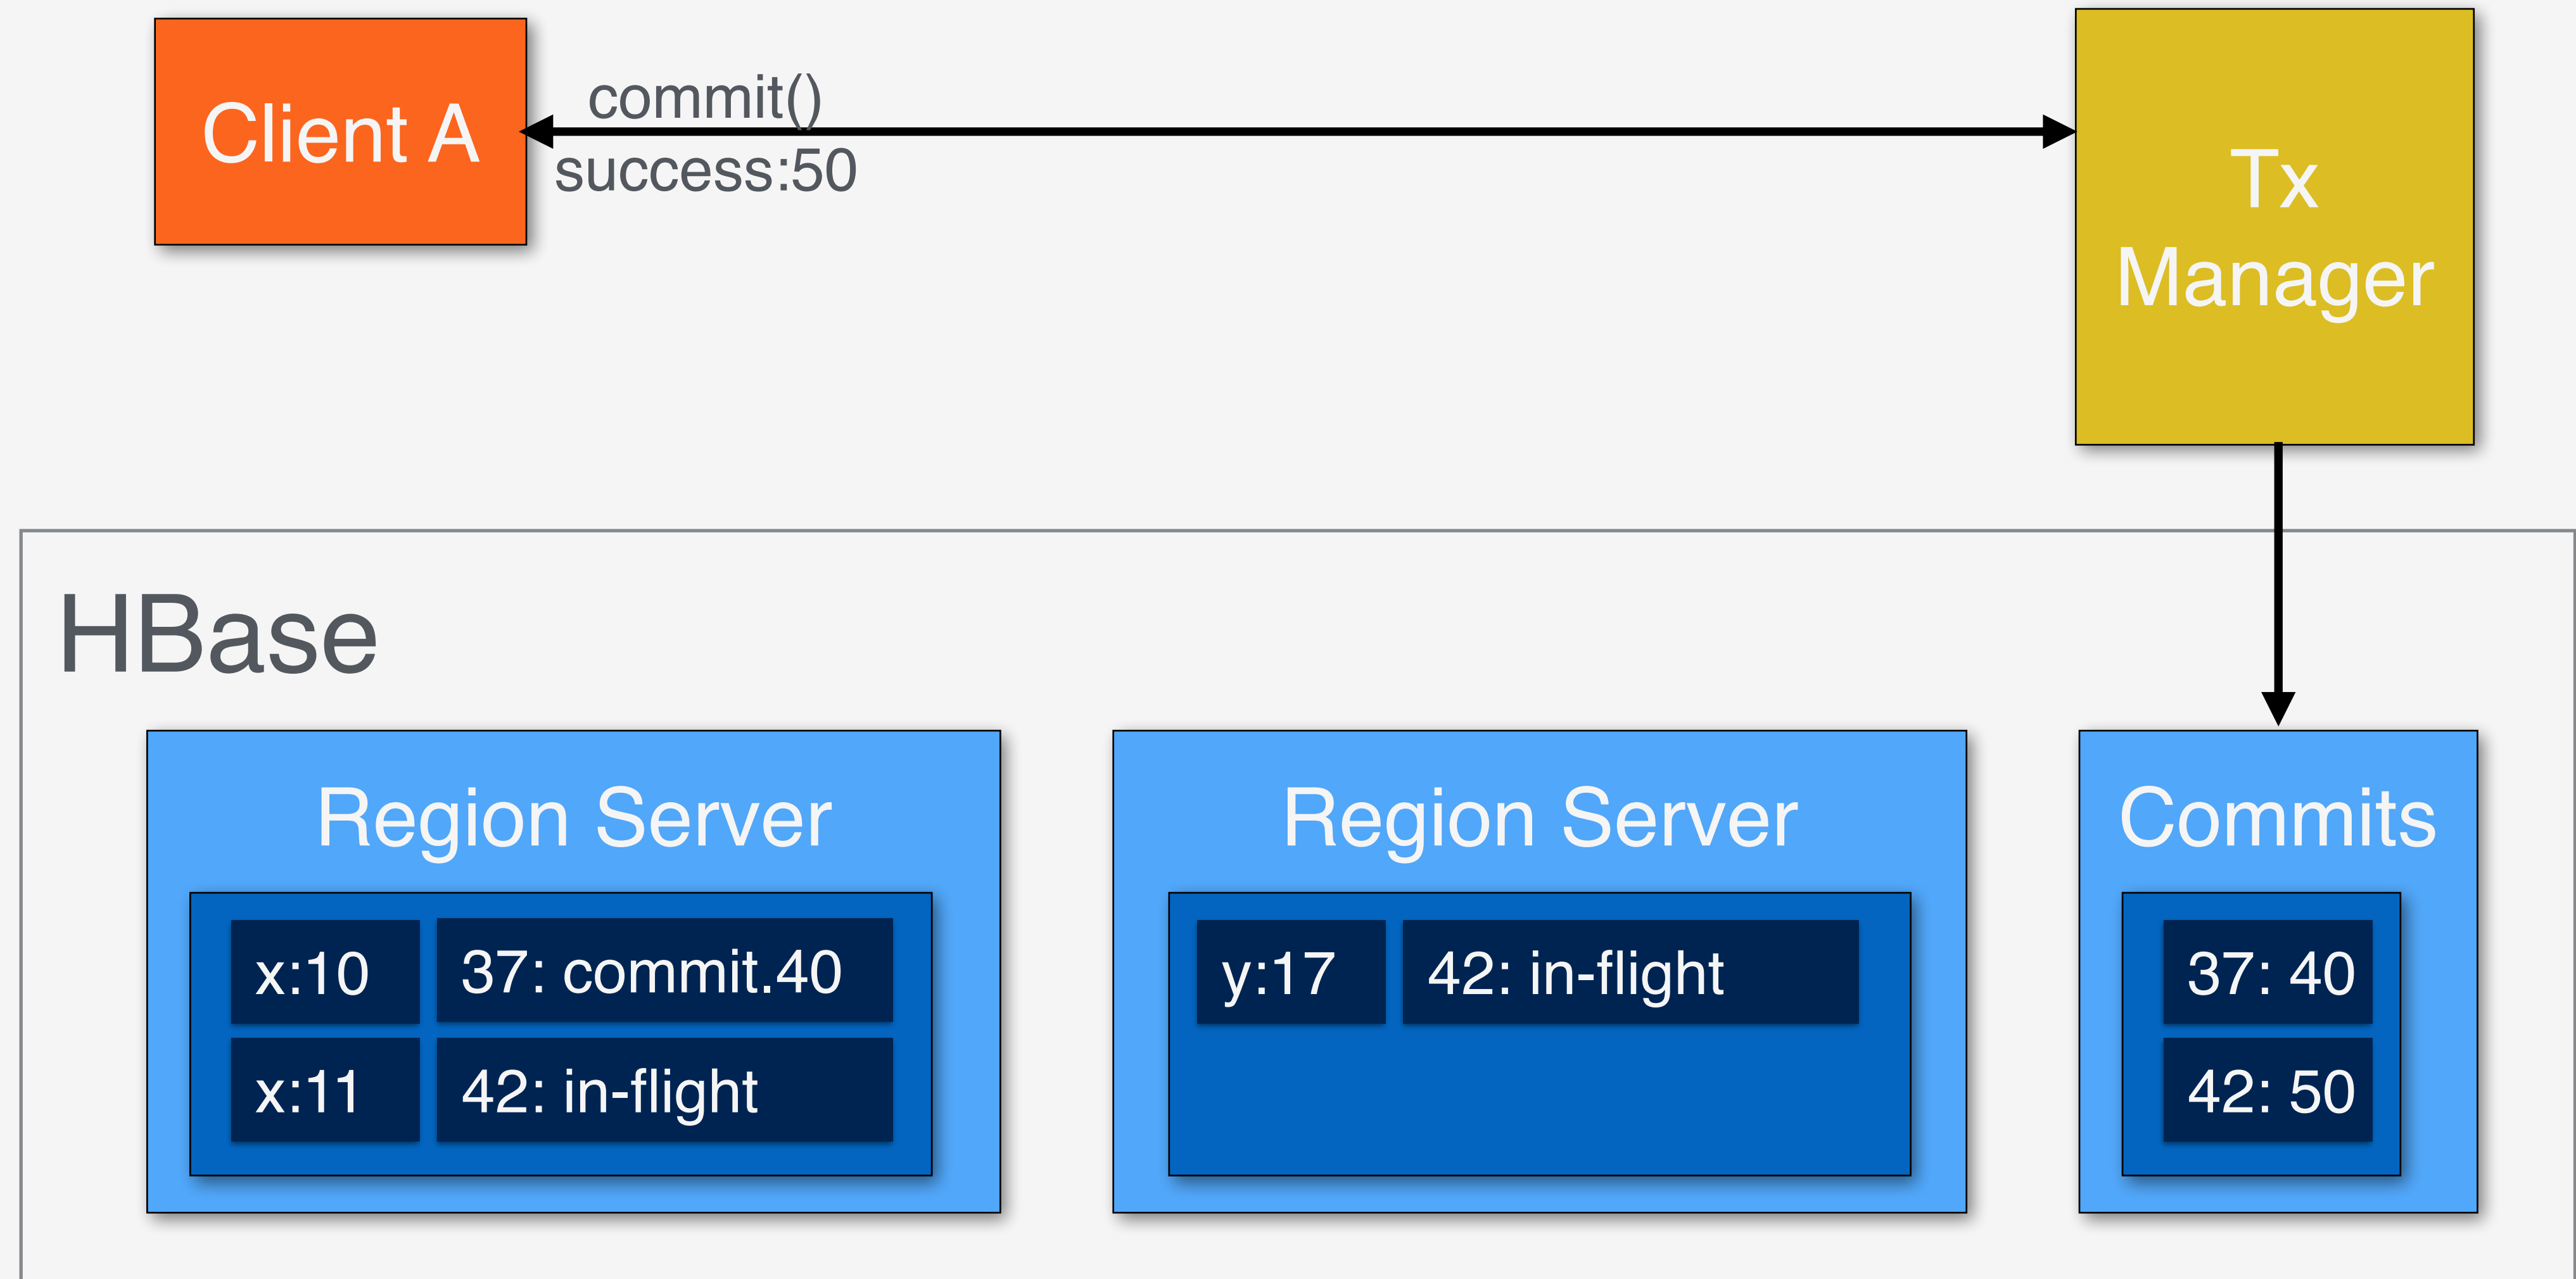
# Apache Omid



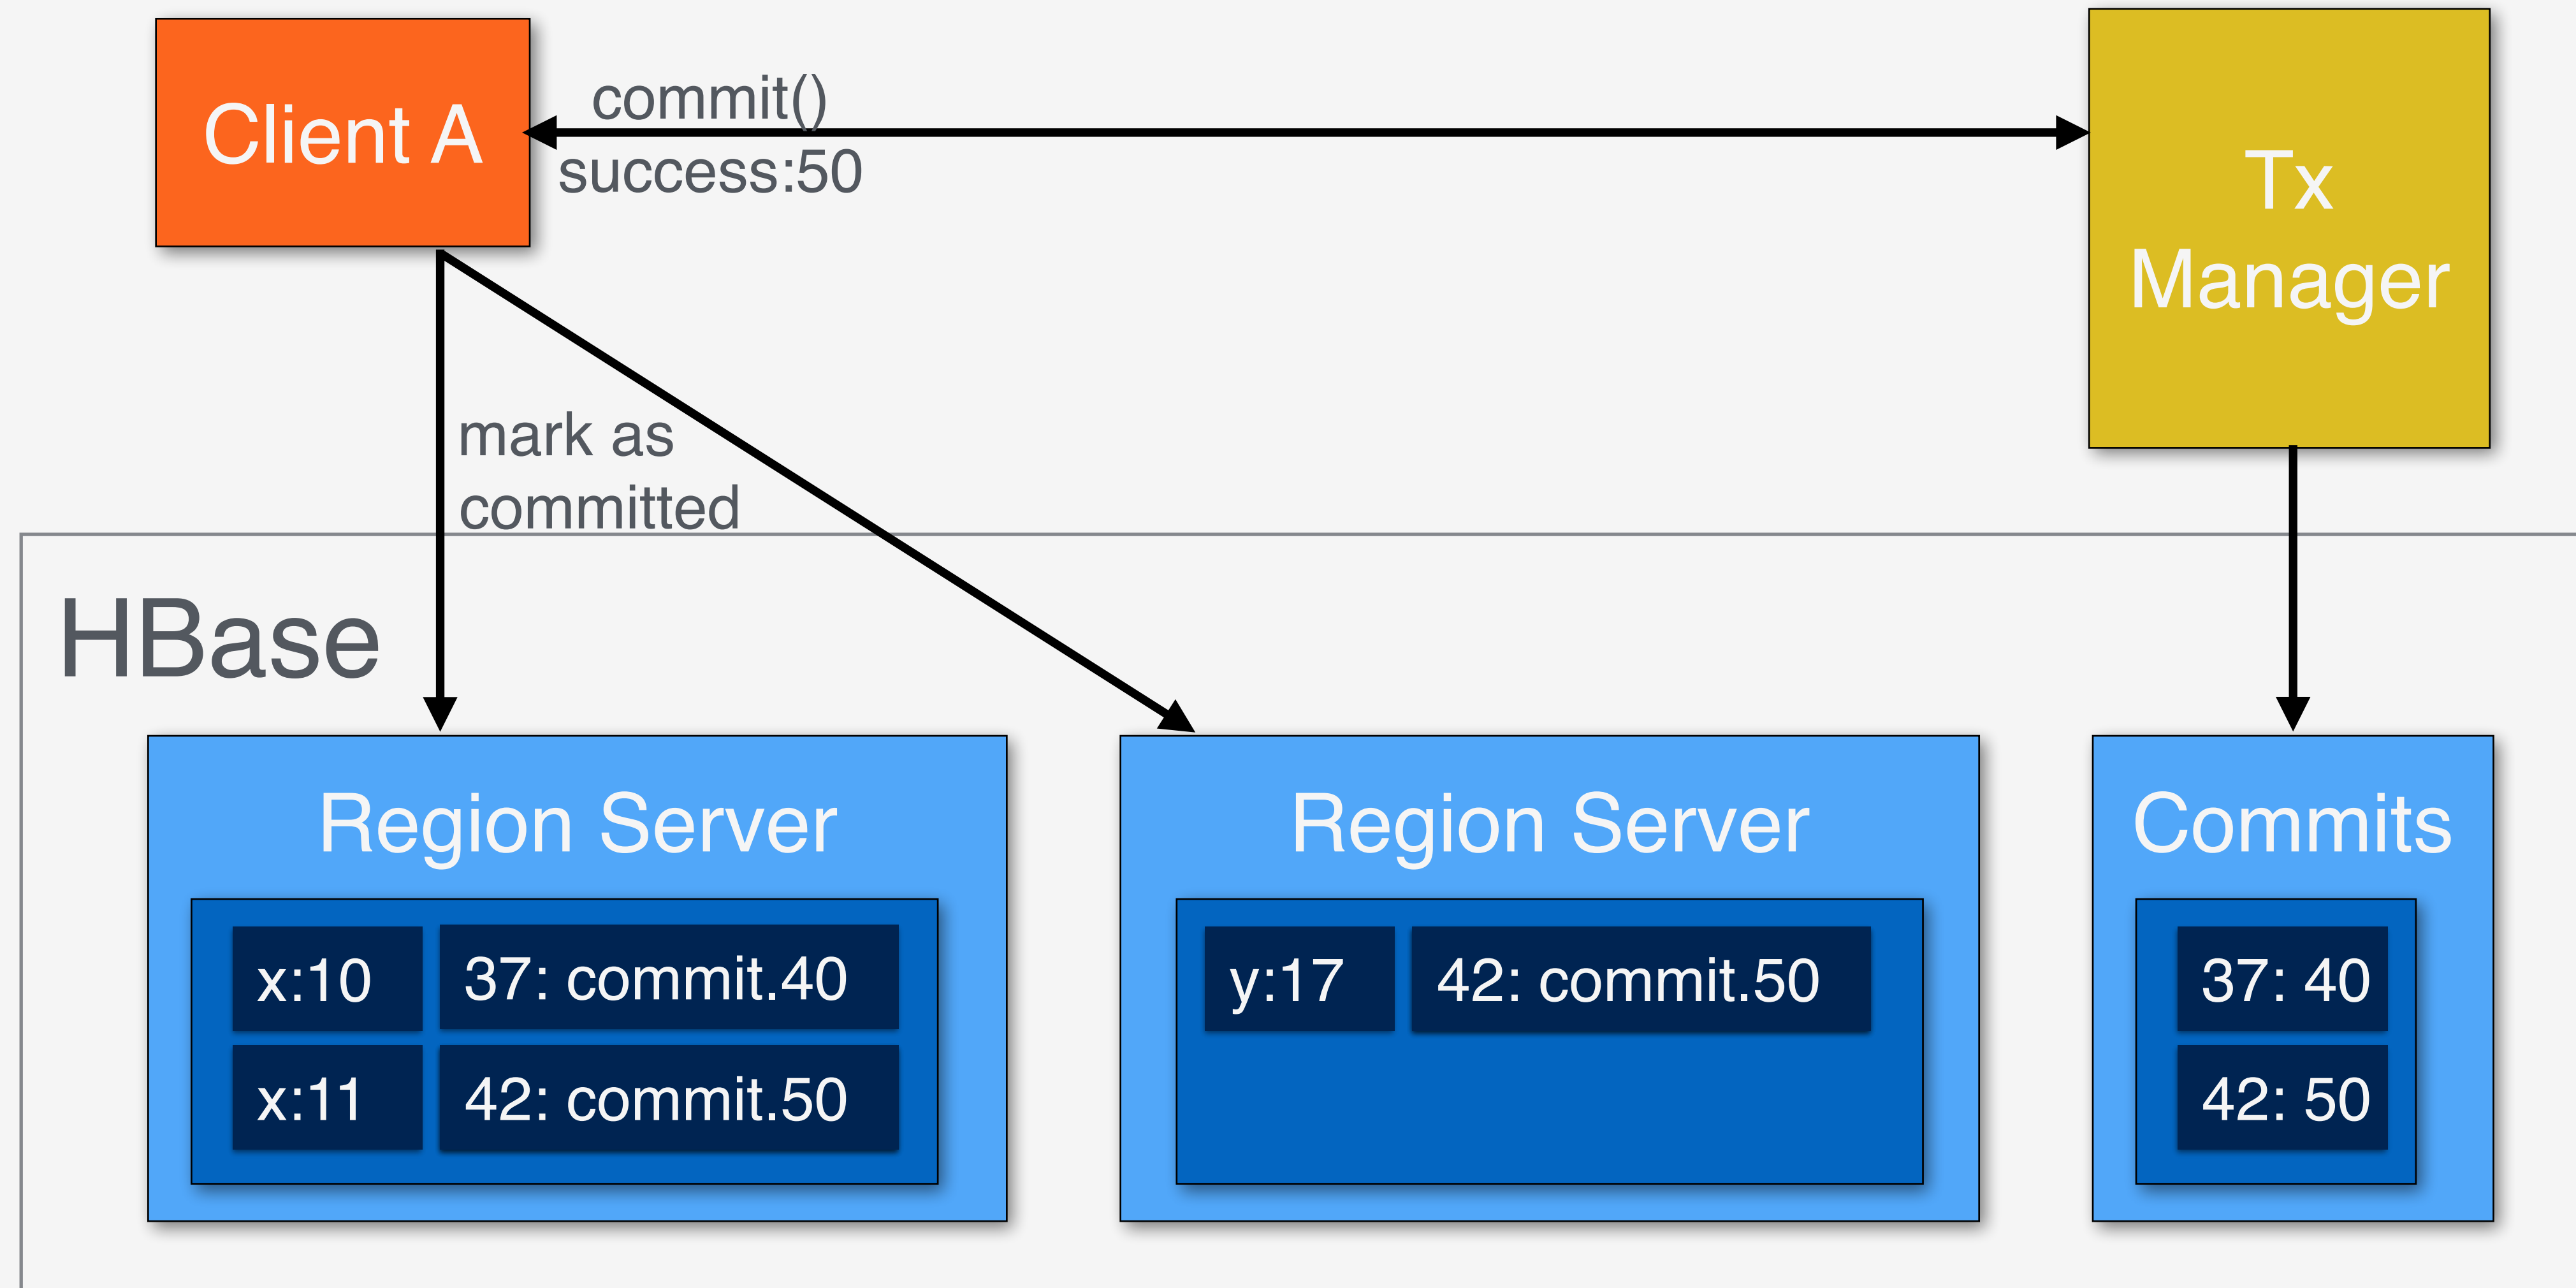
# Apache Omid



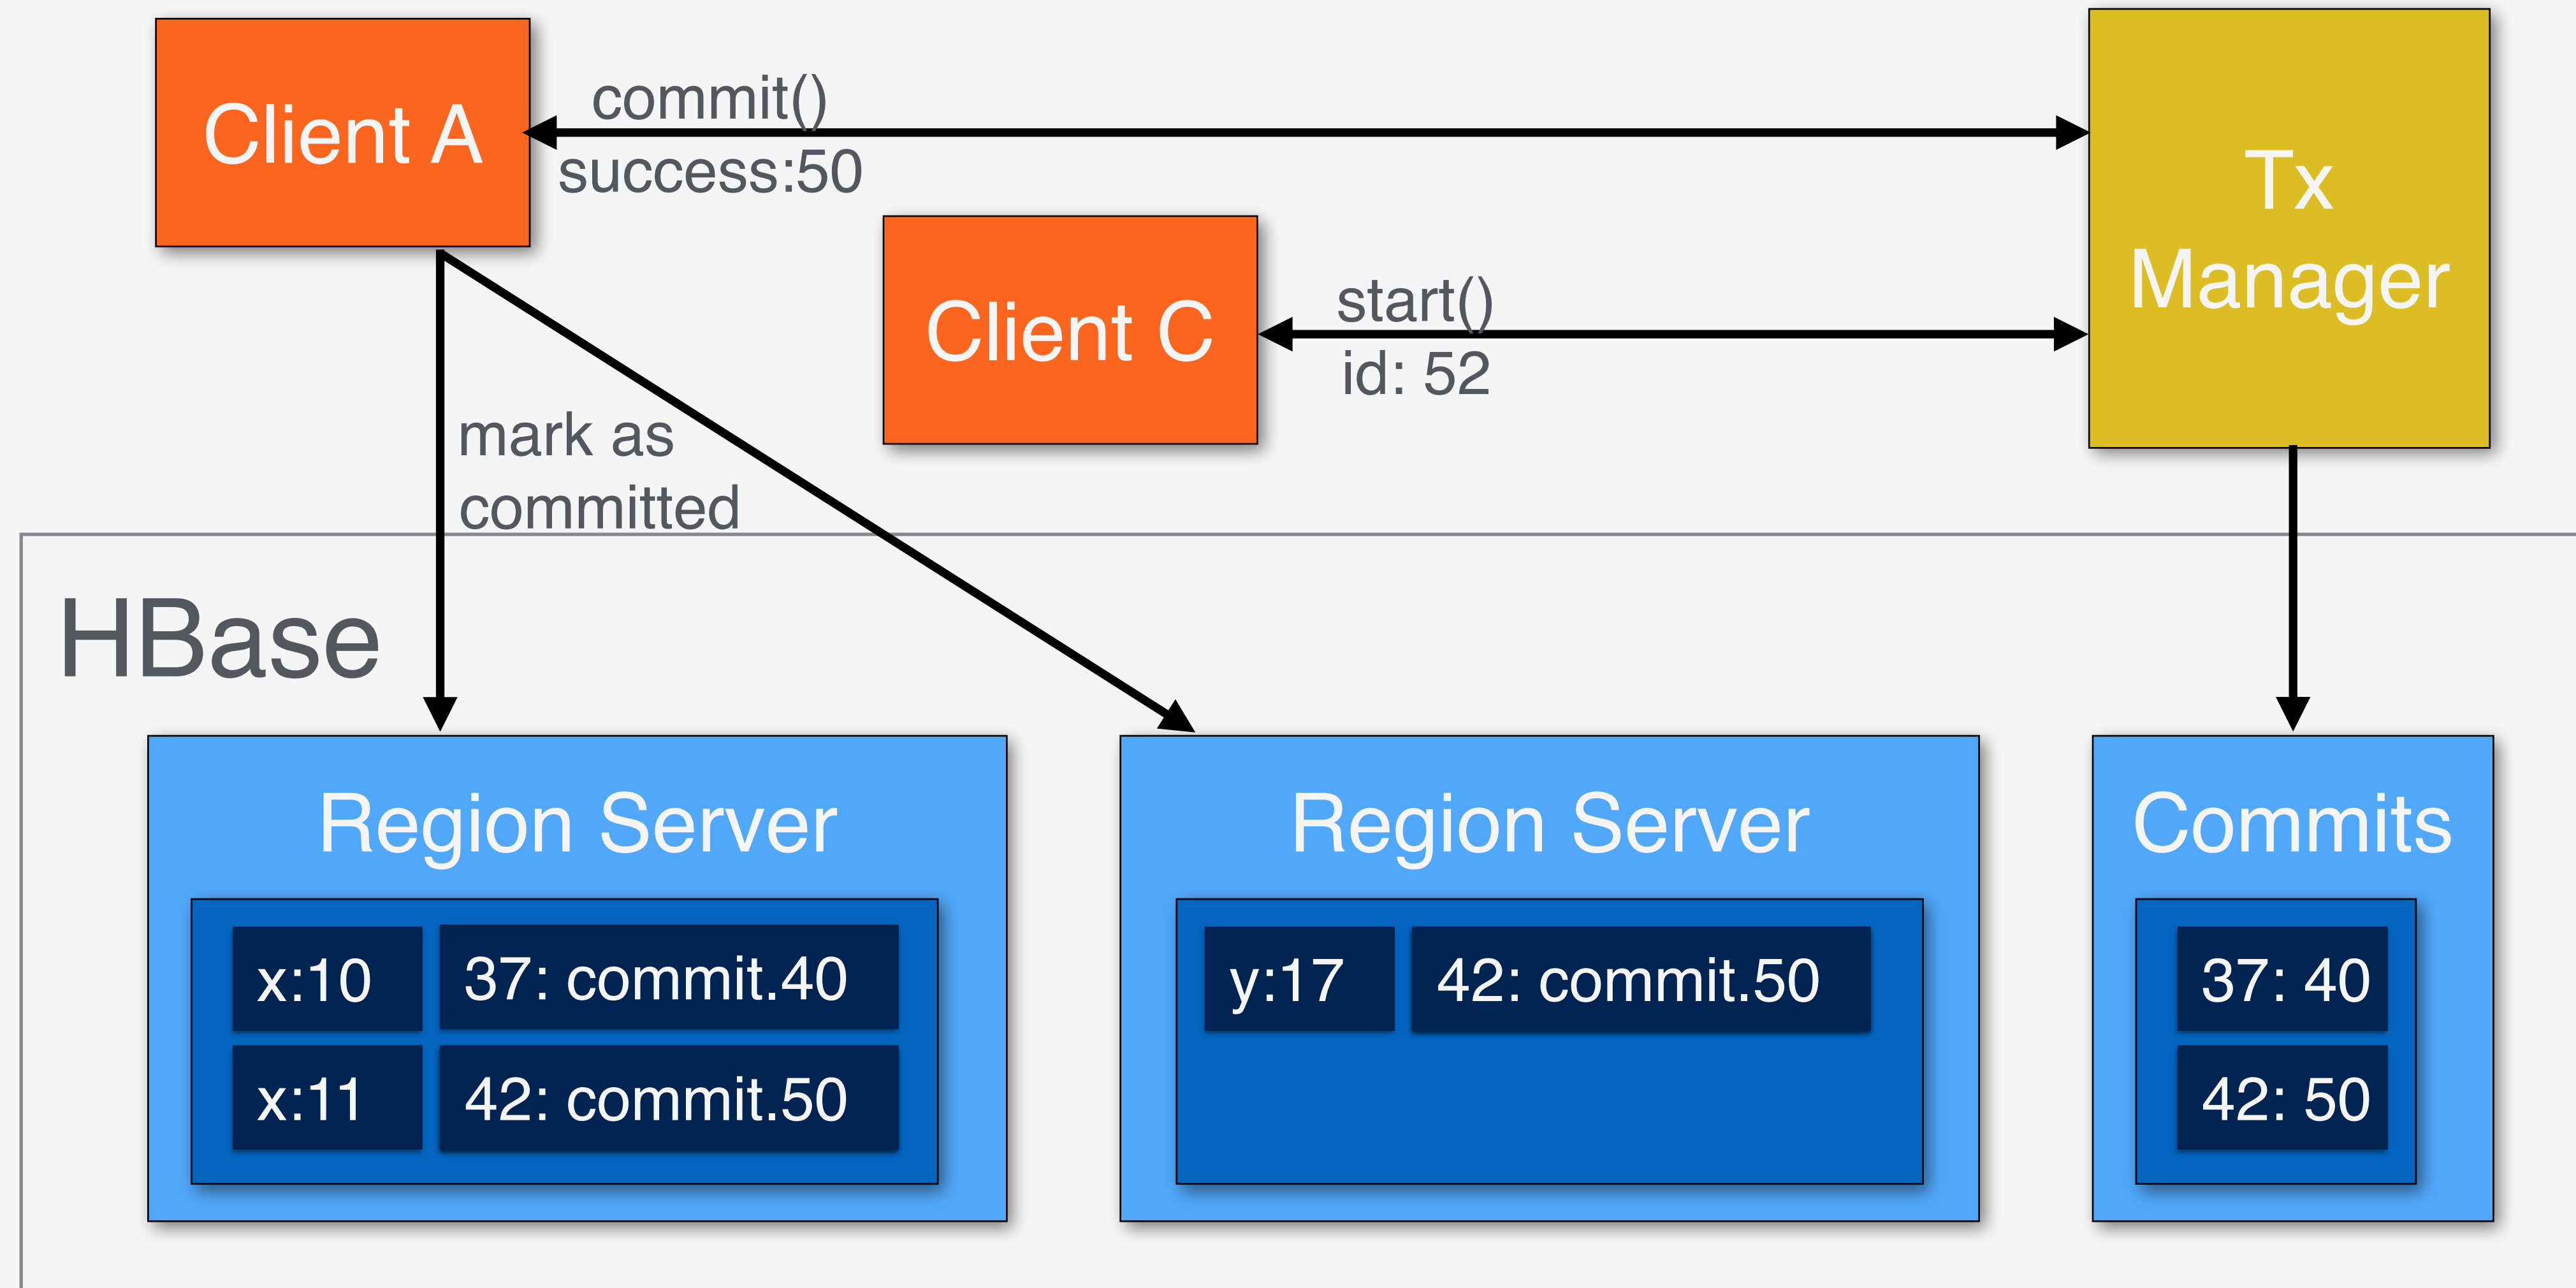
# Apache Omid



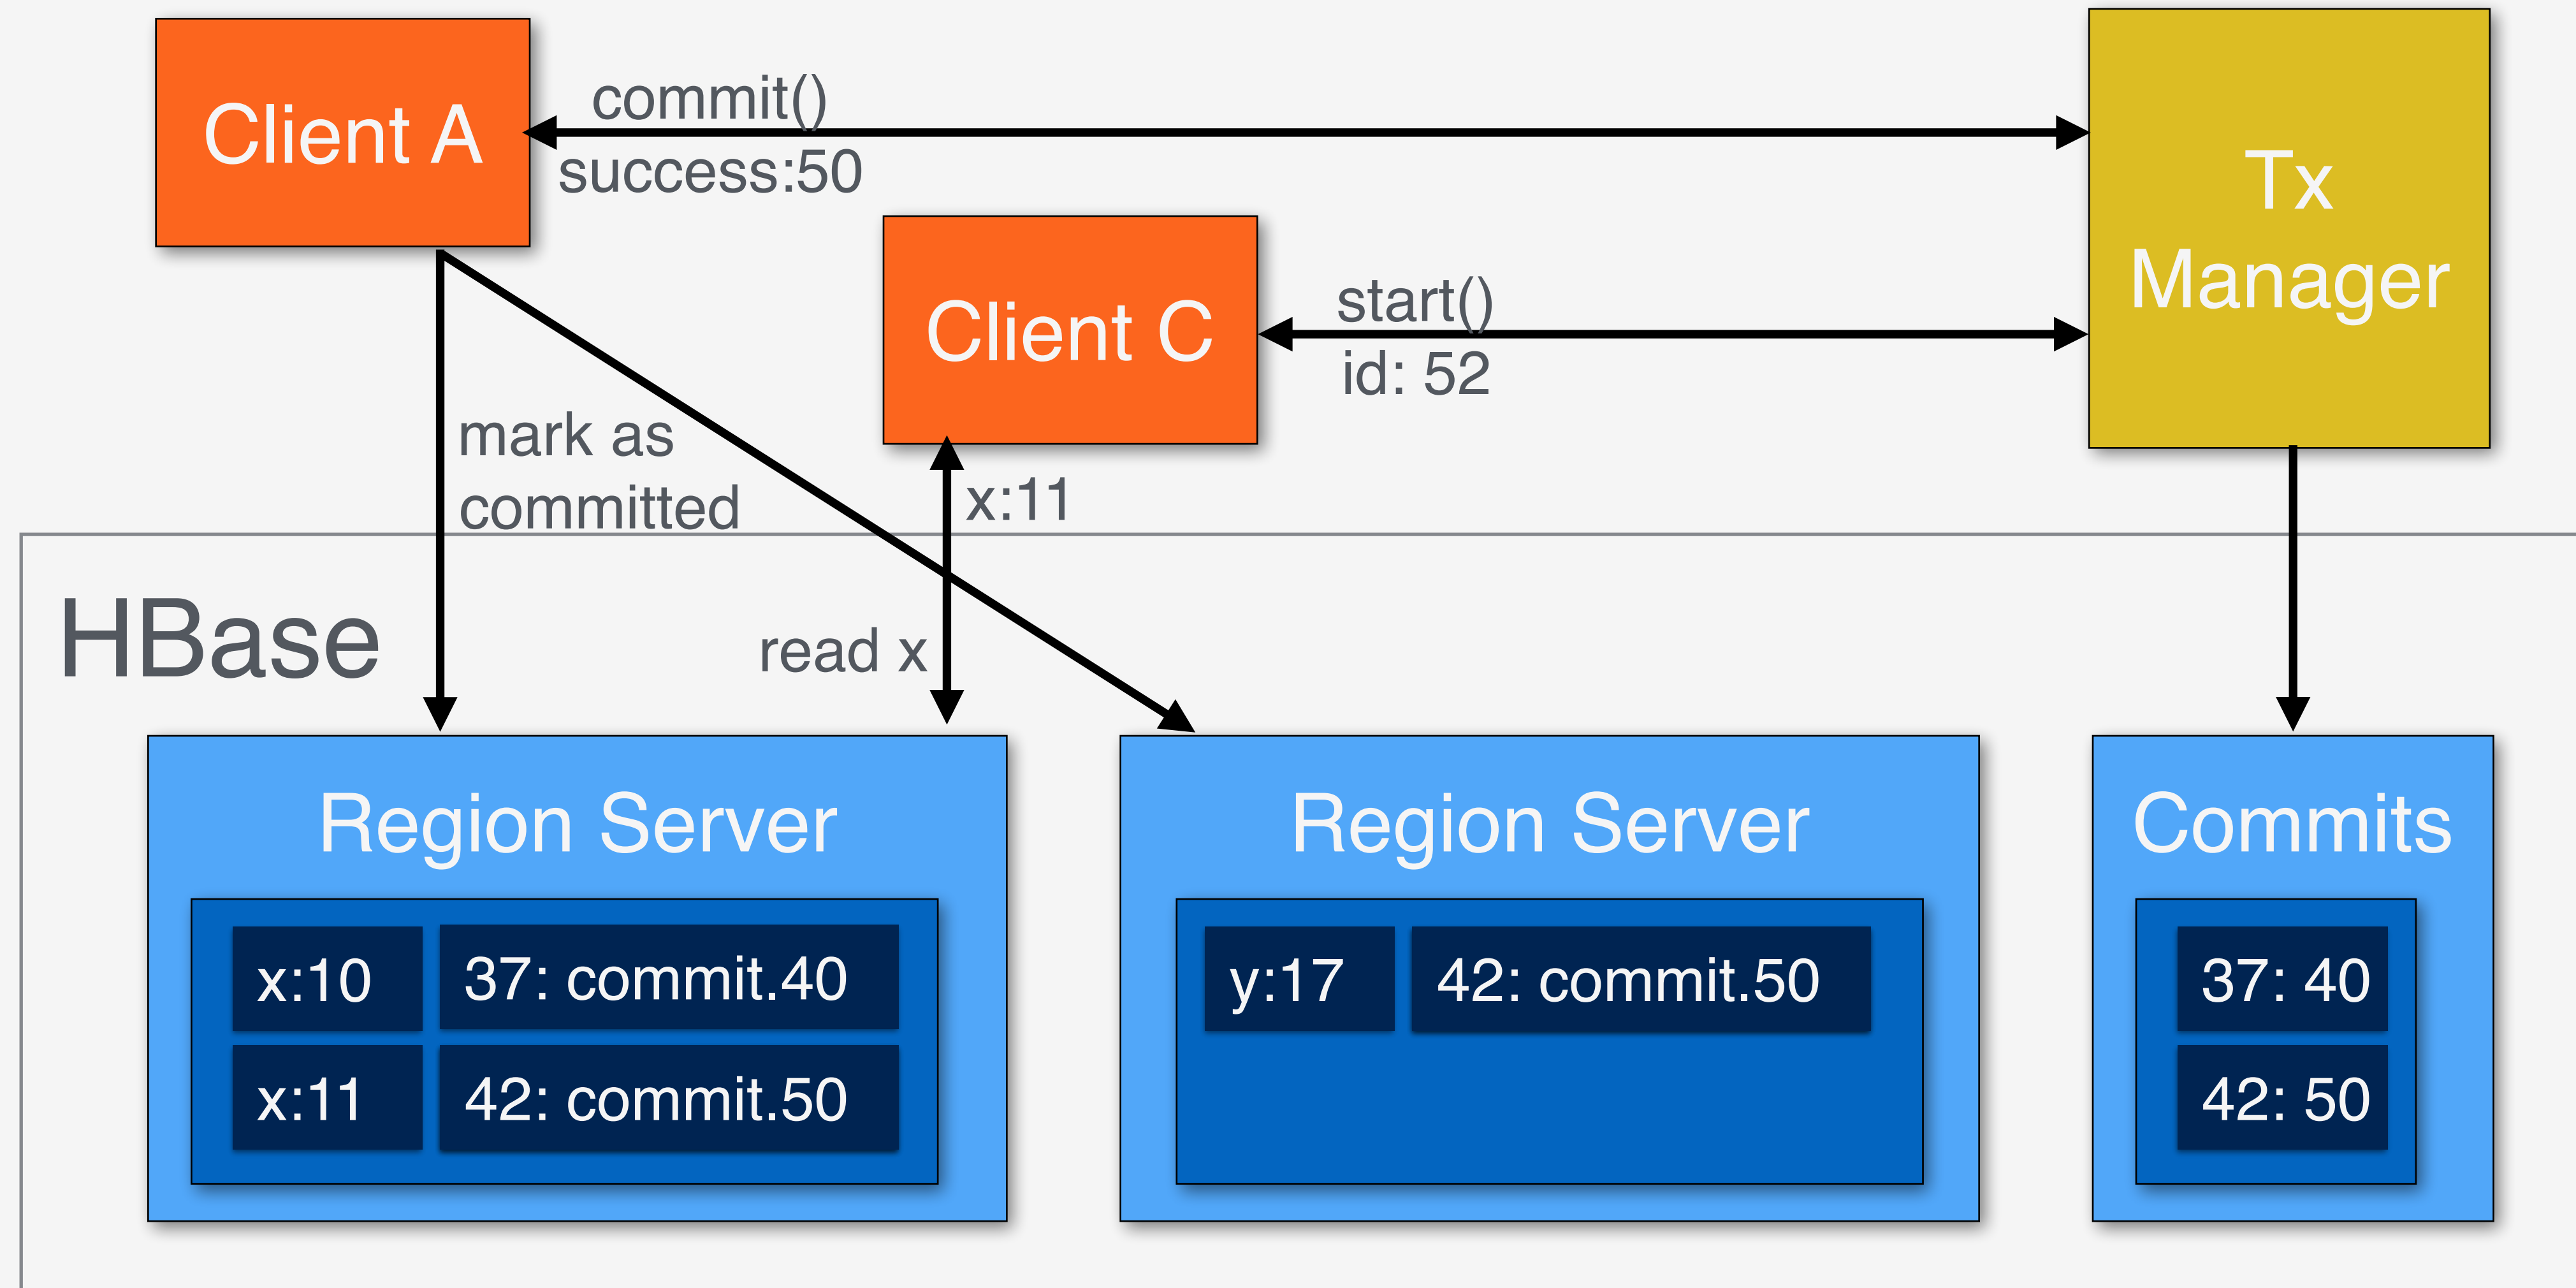
# Apache Omid



# Apache Omid




# Apache Omid



# Apache Omid - Future


- Atomic commit with linking?
  - Eliminate need for commit table

## HBase





# Apache Omid



# Apache Omid - Strengths

- Transaction state is in the database
  - Shadow cells plus commit table
  - Scales with the size of the cluster
- Transaction Manager is lightweight
  - Generation of tx IDs delegated to timestamp oracle
  - Conflict detection
  - Writing to commit table
- Fault Tolerance:
  - After failure, fail all existing transactions attempting to commit
  - Self-correcting: Read clients can delete invalid cells



# Apache Omid - Not So Strengths

- Storage intensive - shadow cells double the space
- I/O intensive - every cell requires two writes
  1. write data and shadow cell
  2. record commit in shadow cell
- Reads may also require two reads from HBase (commit table)
- Producer/Consumer: will often find the (uncommitted) shadow cell
  - Scans: high throughput sequential read disrupted by frequent lookups
- Security/Multi-tenancy:
  - All clients need access to commit table
  - Read clients need write access to repair invalid data
- Replication: Not implemented



# Summary

	Apache Tephra	Apache Trafodion	Apache Omid
<b>Tx State</b>	Tx Manager	Distributed to region servers	Tx Manager (changes) HBase (shadows/commits)
<b>Conflict detection</b>	Tx Manager	Distributed to regions, 2-phase commit	Tx Manager
<b>ID generation</b>	Tx Manager	Distributed to multiple Tx Managers	Tx Manager
<b>API</b>	HTable	SQL	Custom
<b>Multi-tenant</b>	Yes	Yes	No
<b>Strength</b>	Scans, Large Tx, API	Scalable, full SQL	Scale, throughput
<b>So so</b>	Scale, Throughput	API not Hbase, Large Tx	Scans, Producer/Consumer



# Links

Join the community:



(incubating)

<http://tephra.apache.org/>



(incubating)

<http://trafodion.apache.org/>



Apache Omid (incubating)

<http://omid.apache.org/>



# Thank you

... for listening to my talk.

Credits:

- Sean Broeder, Narendra Goyal (Trafodion)
- Francisco Perez-Sorrosal (Omid)



# Thank you

... for listening to my talk.

Credits:

- Sean Broeder, Narendra Goyal (Trafodion)
- Francisco Perez-Sorrosal (Omid)

## Questions?

